

# KUBERNETES FROM SCRATCH

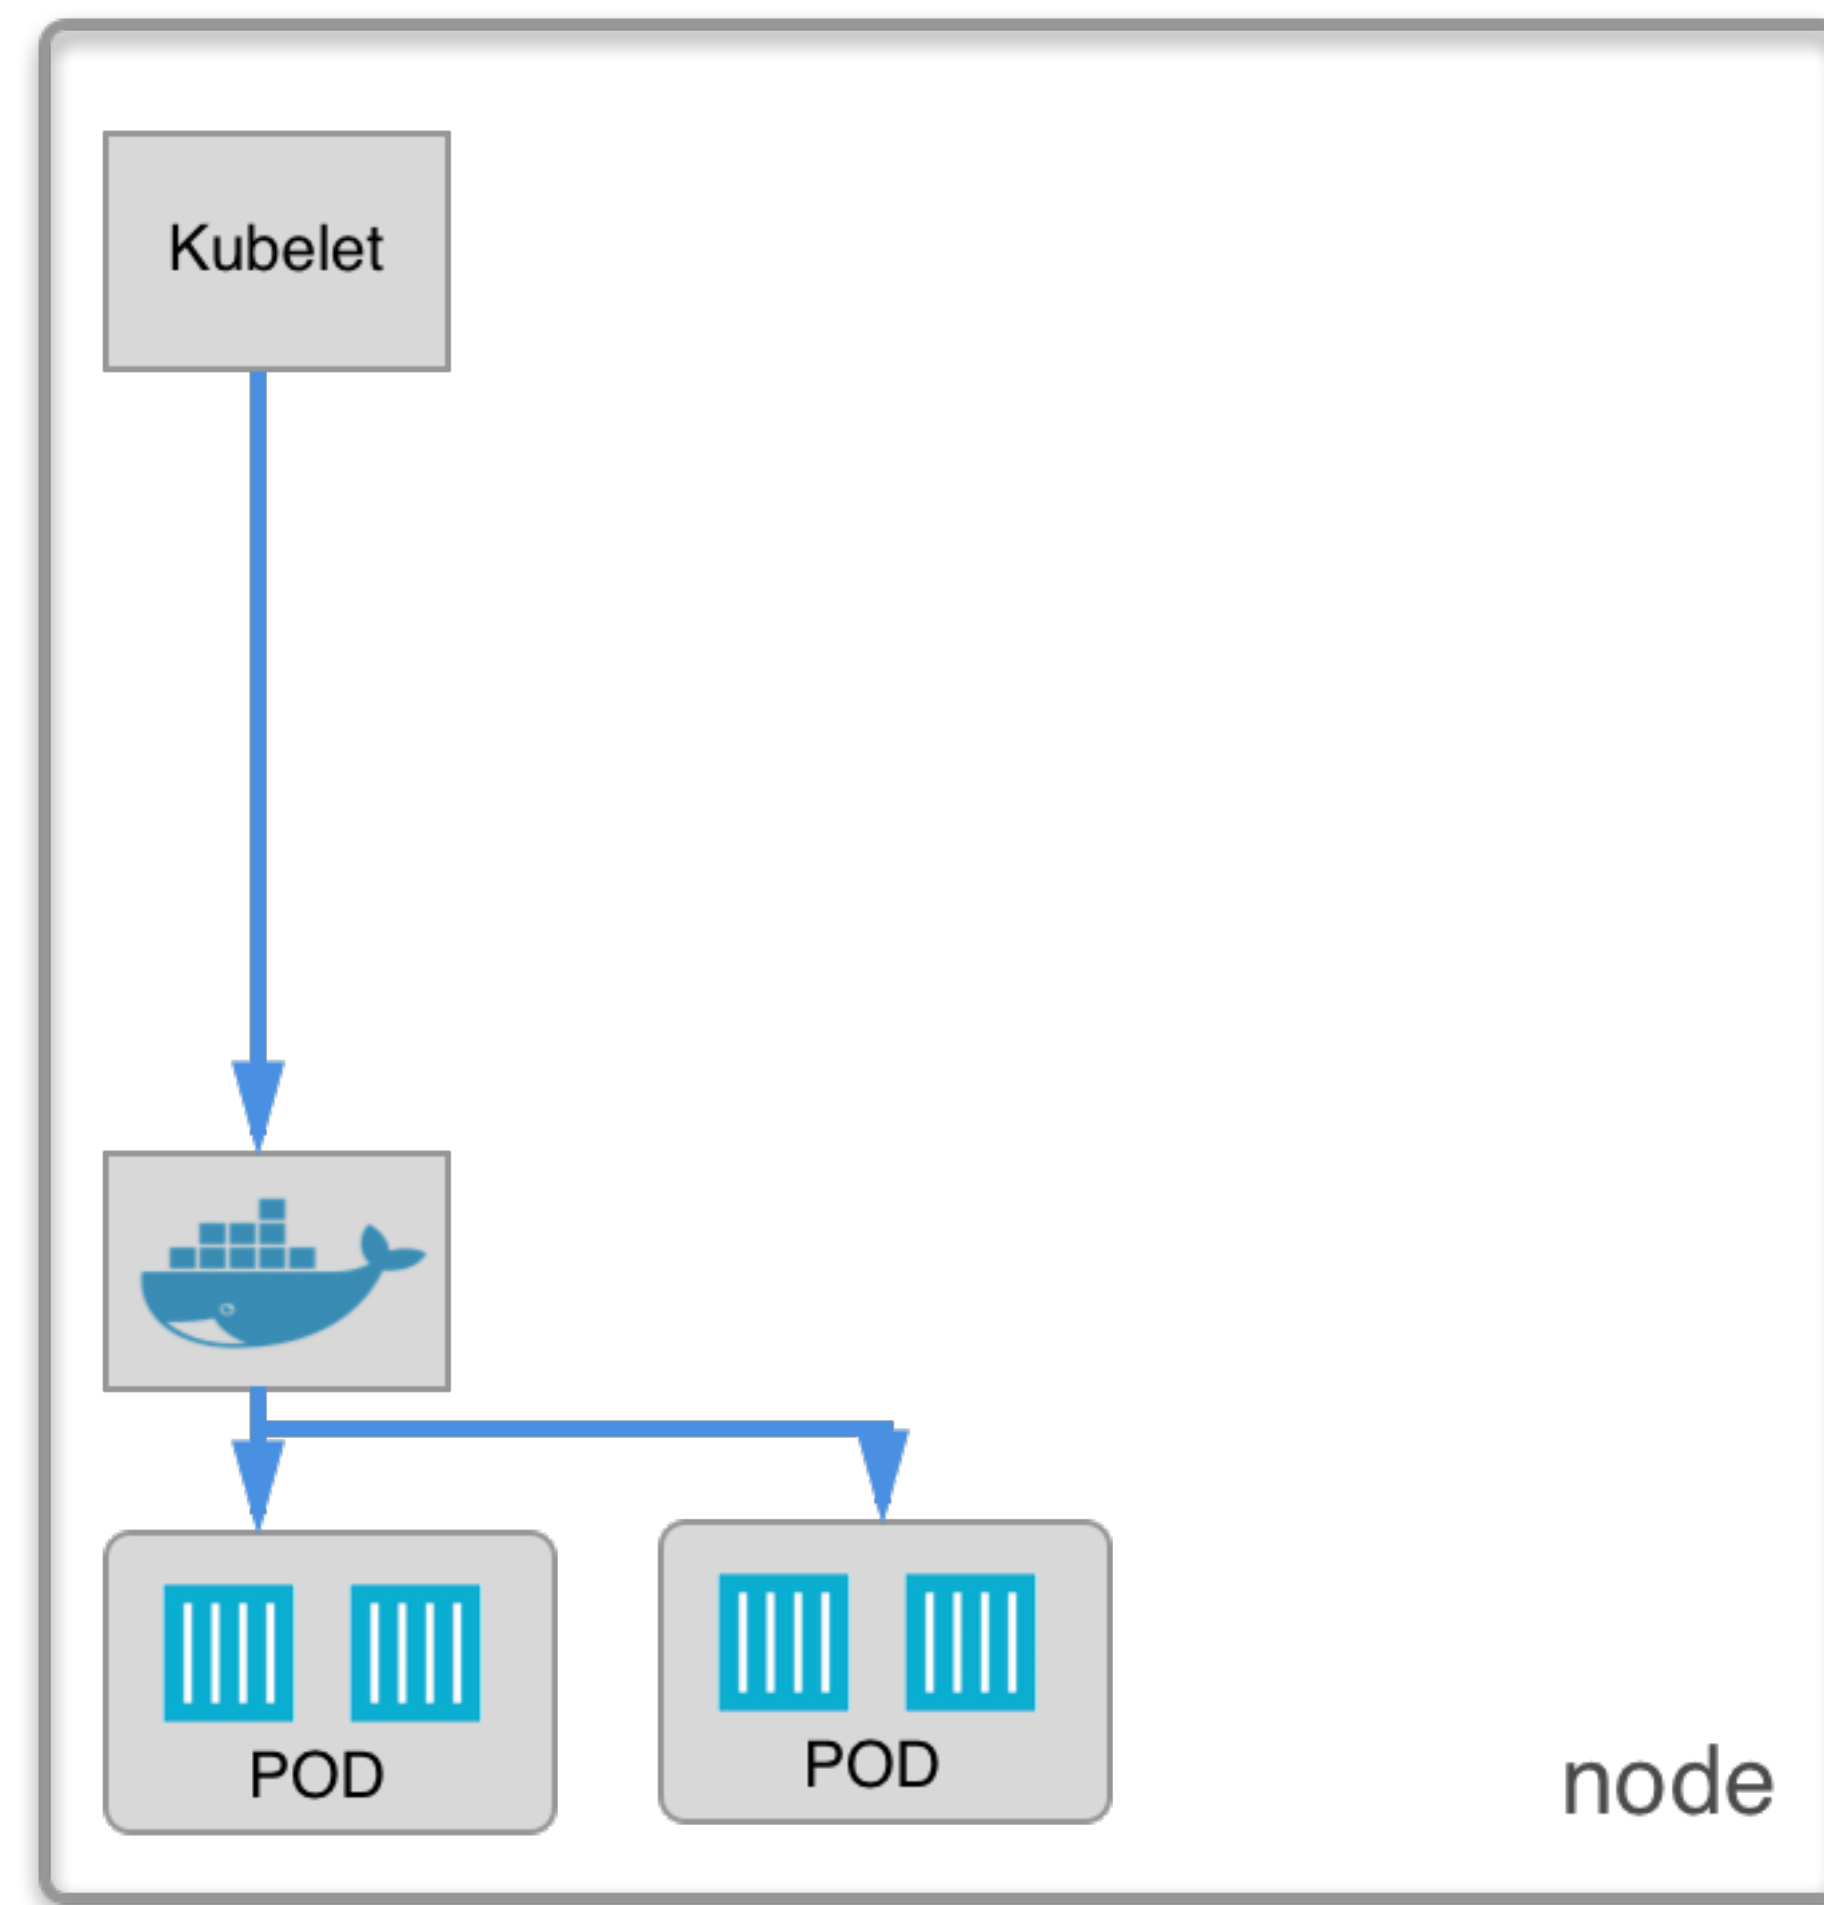
---

Konrad Rotkiewicz  
Owner @ Ulam Labs  
[konrad@ulam.io](mailto:konrad@ulam.io)

# AGENDA

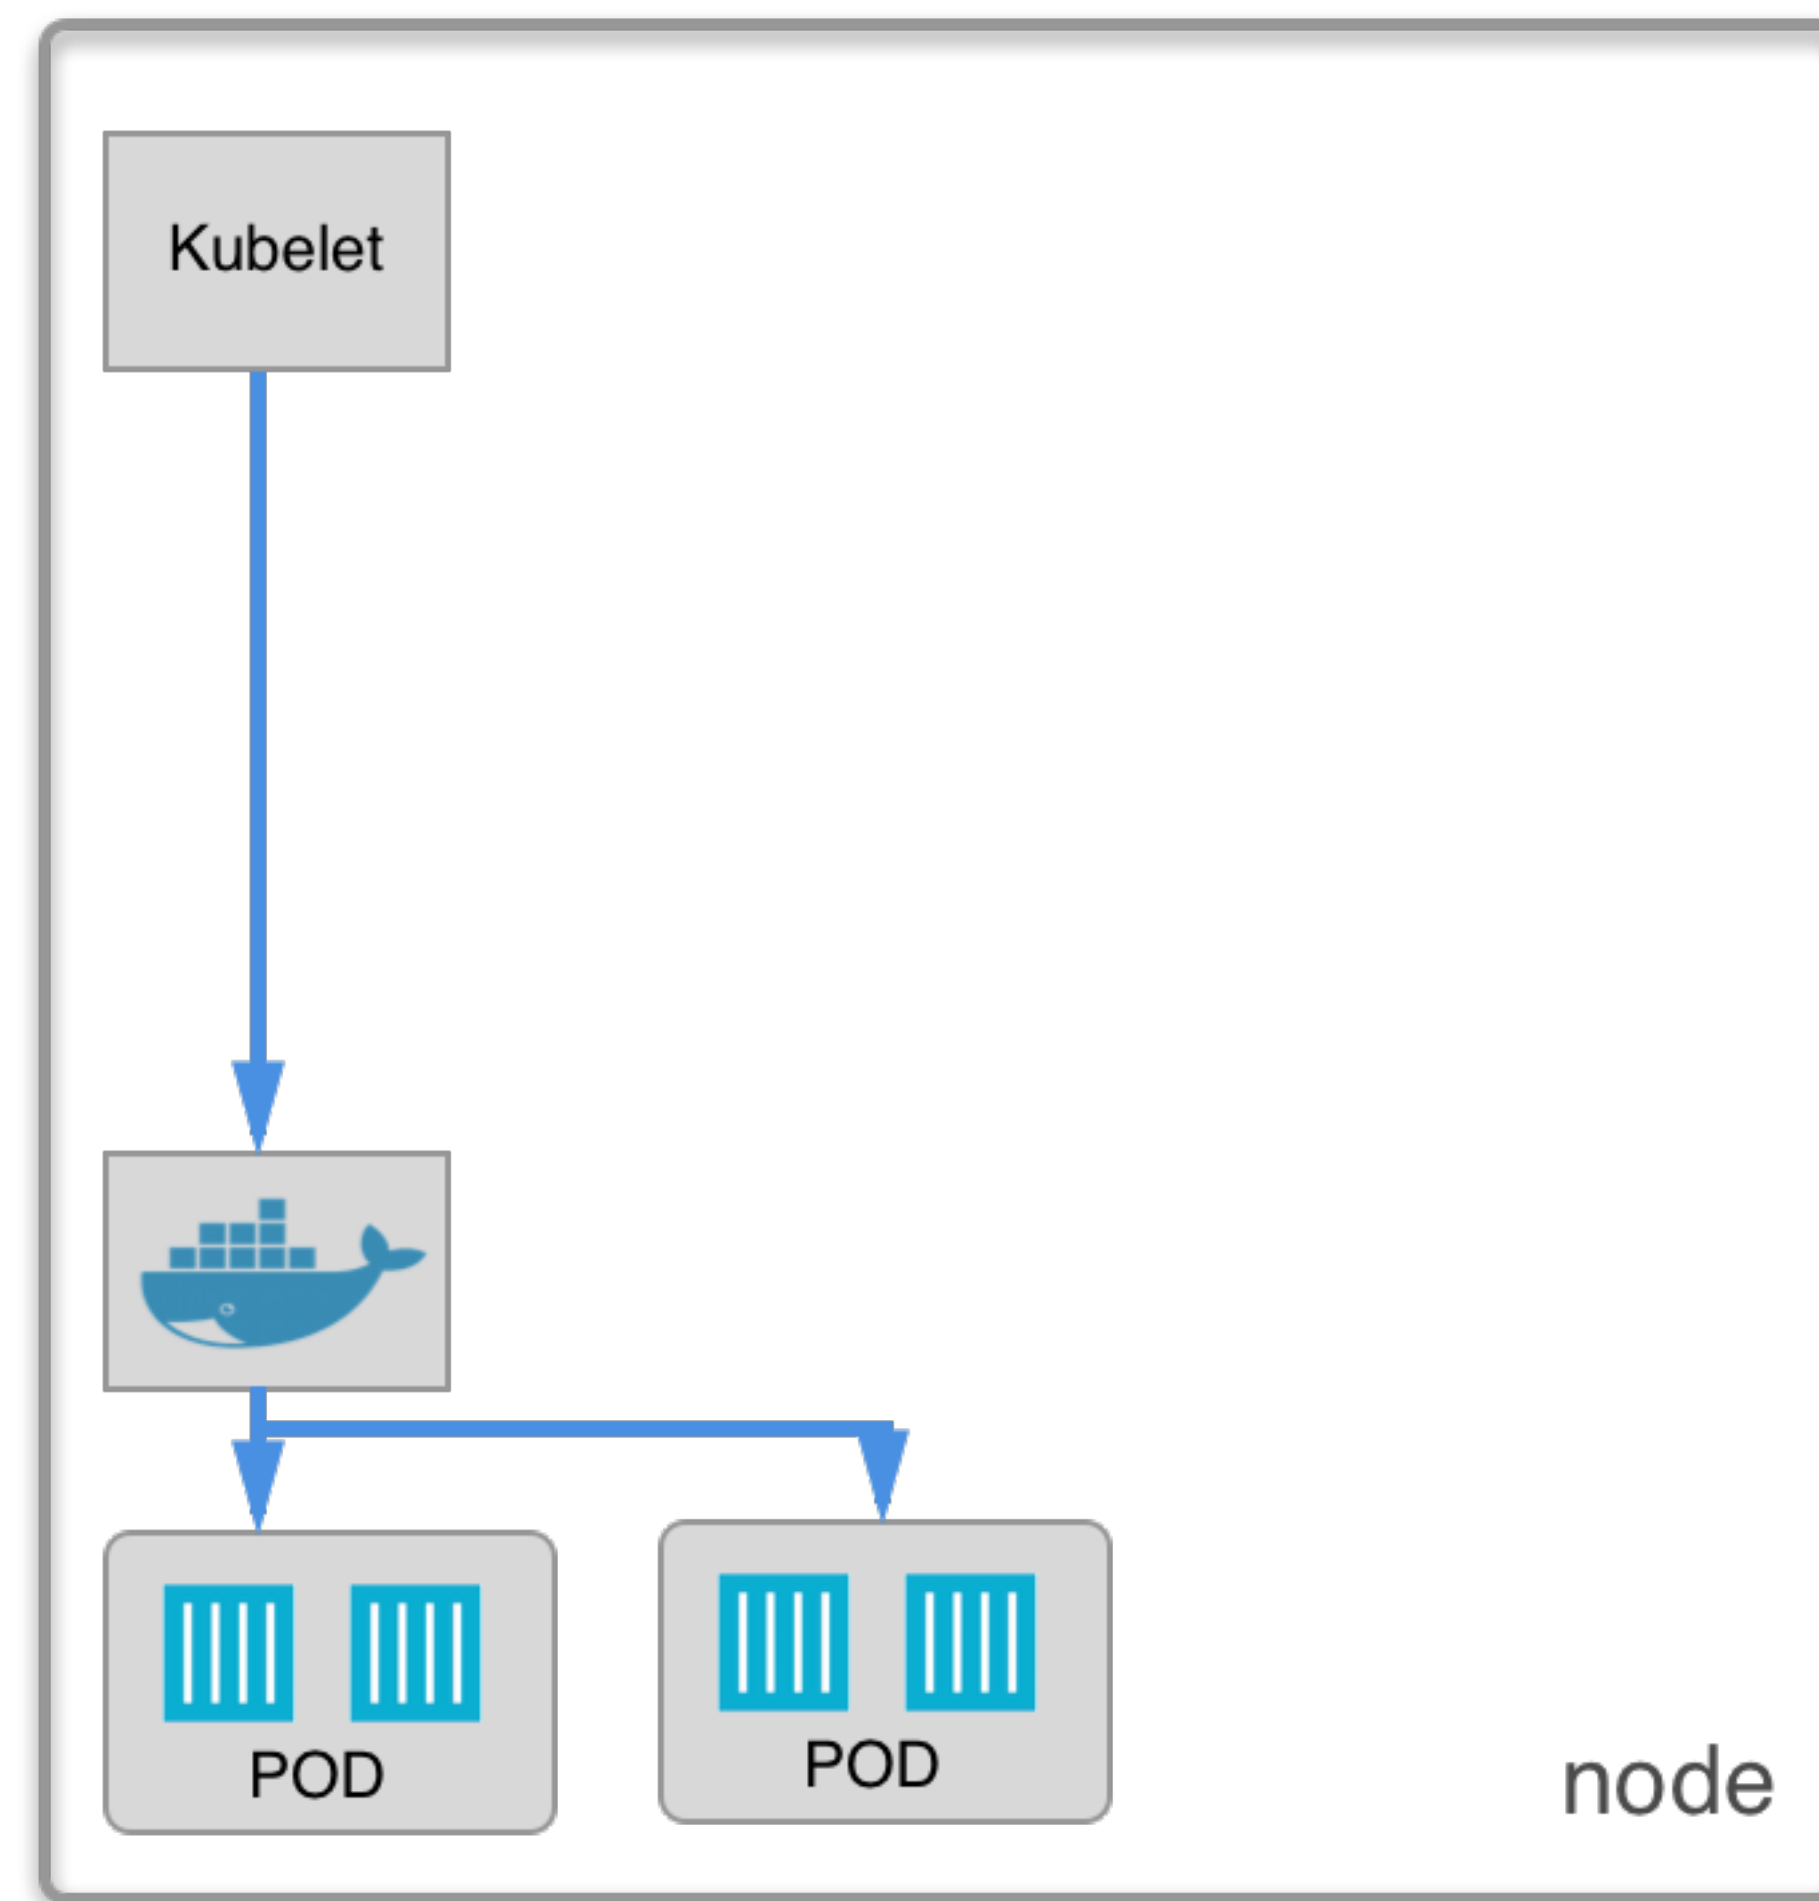
- kubelet
- api server
- etcd
- scheduler
- controller manager
- proxy
- networking
- IaaS integration
- Q&A

# KUBELET



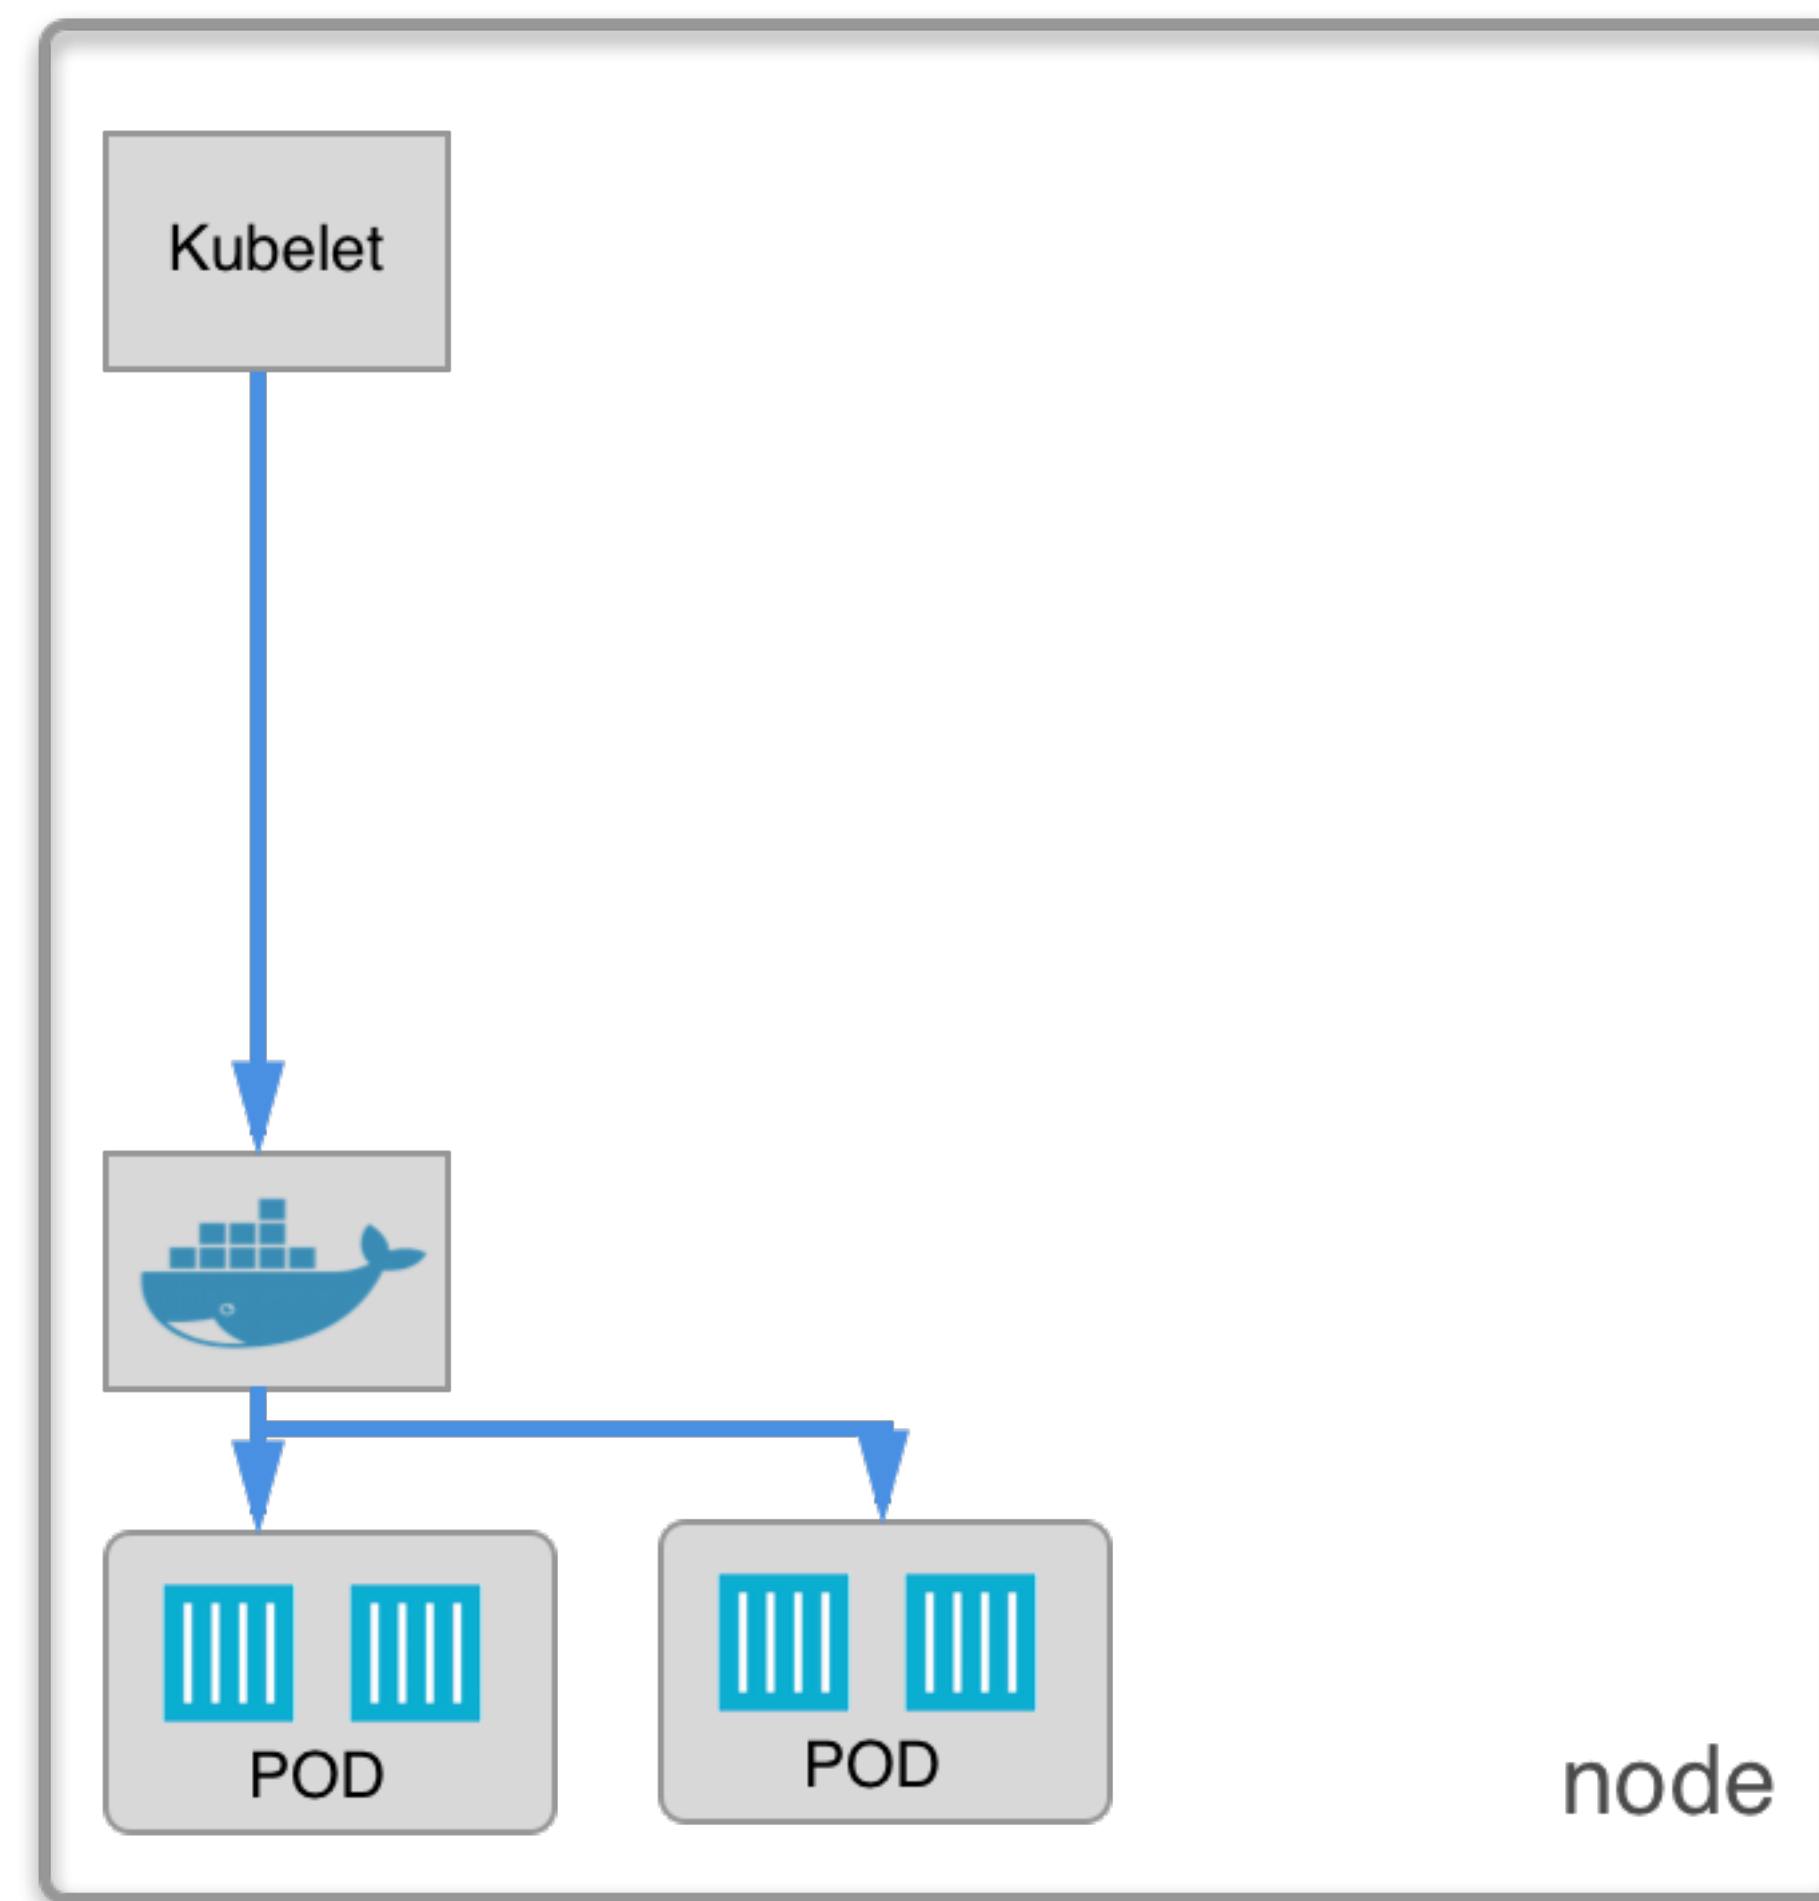
# KUBELET

- the “worker”



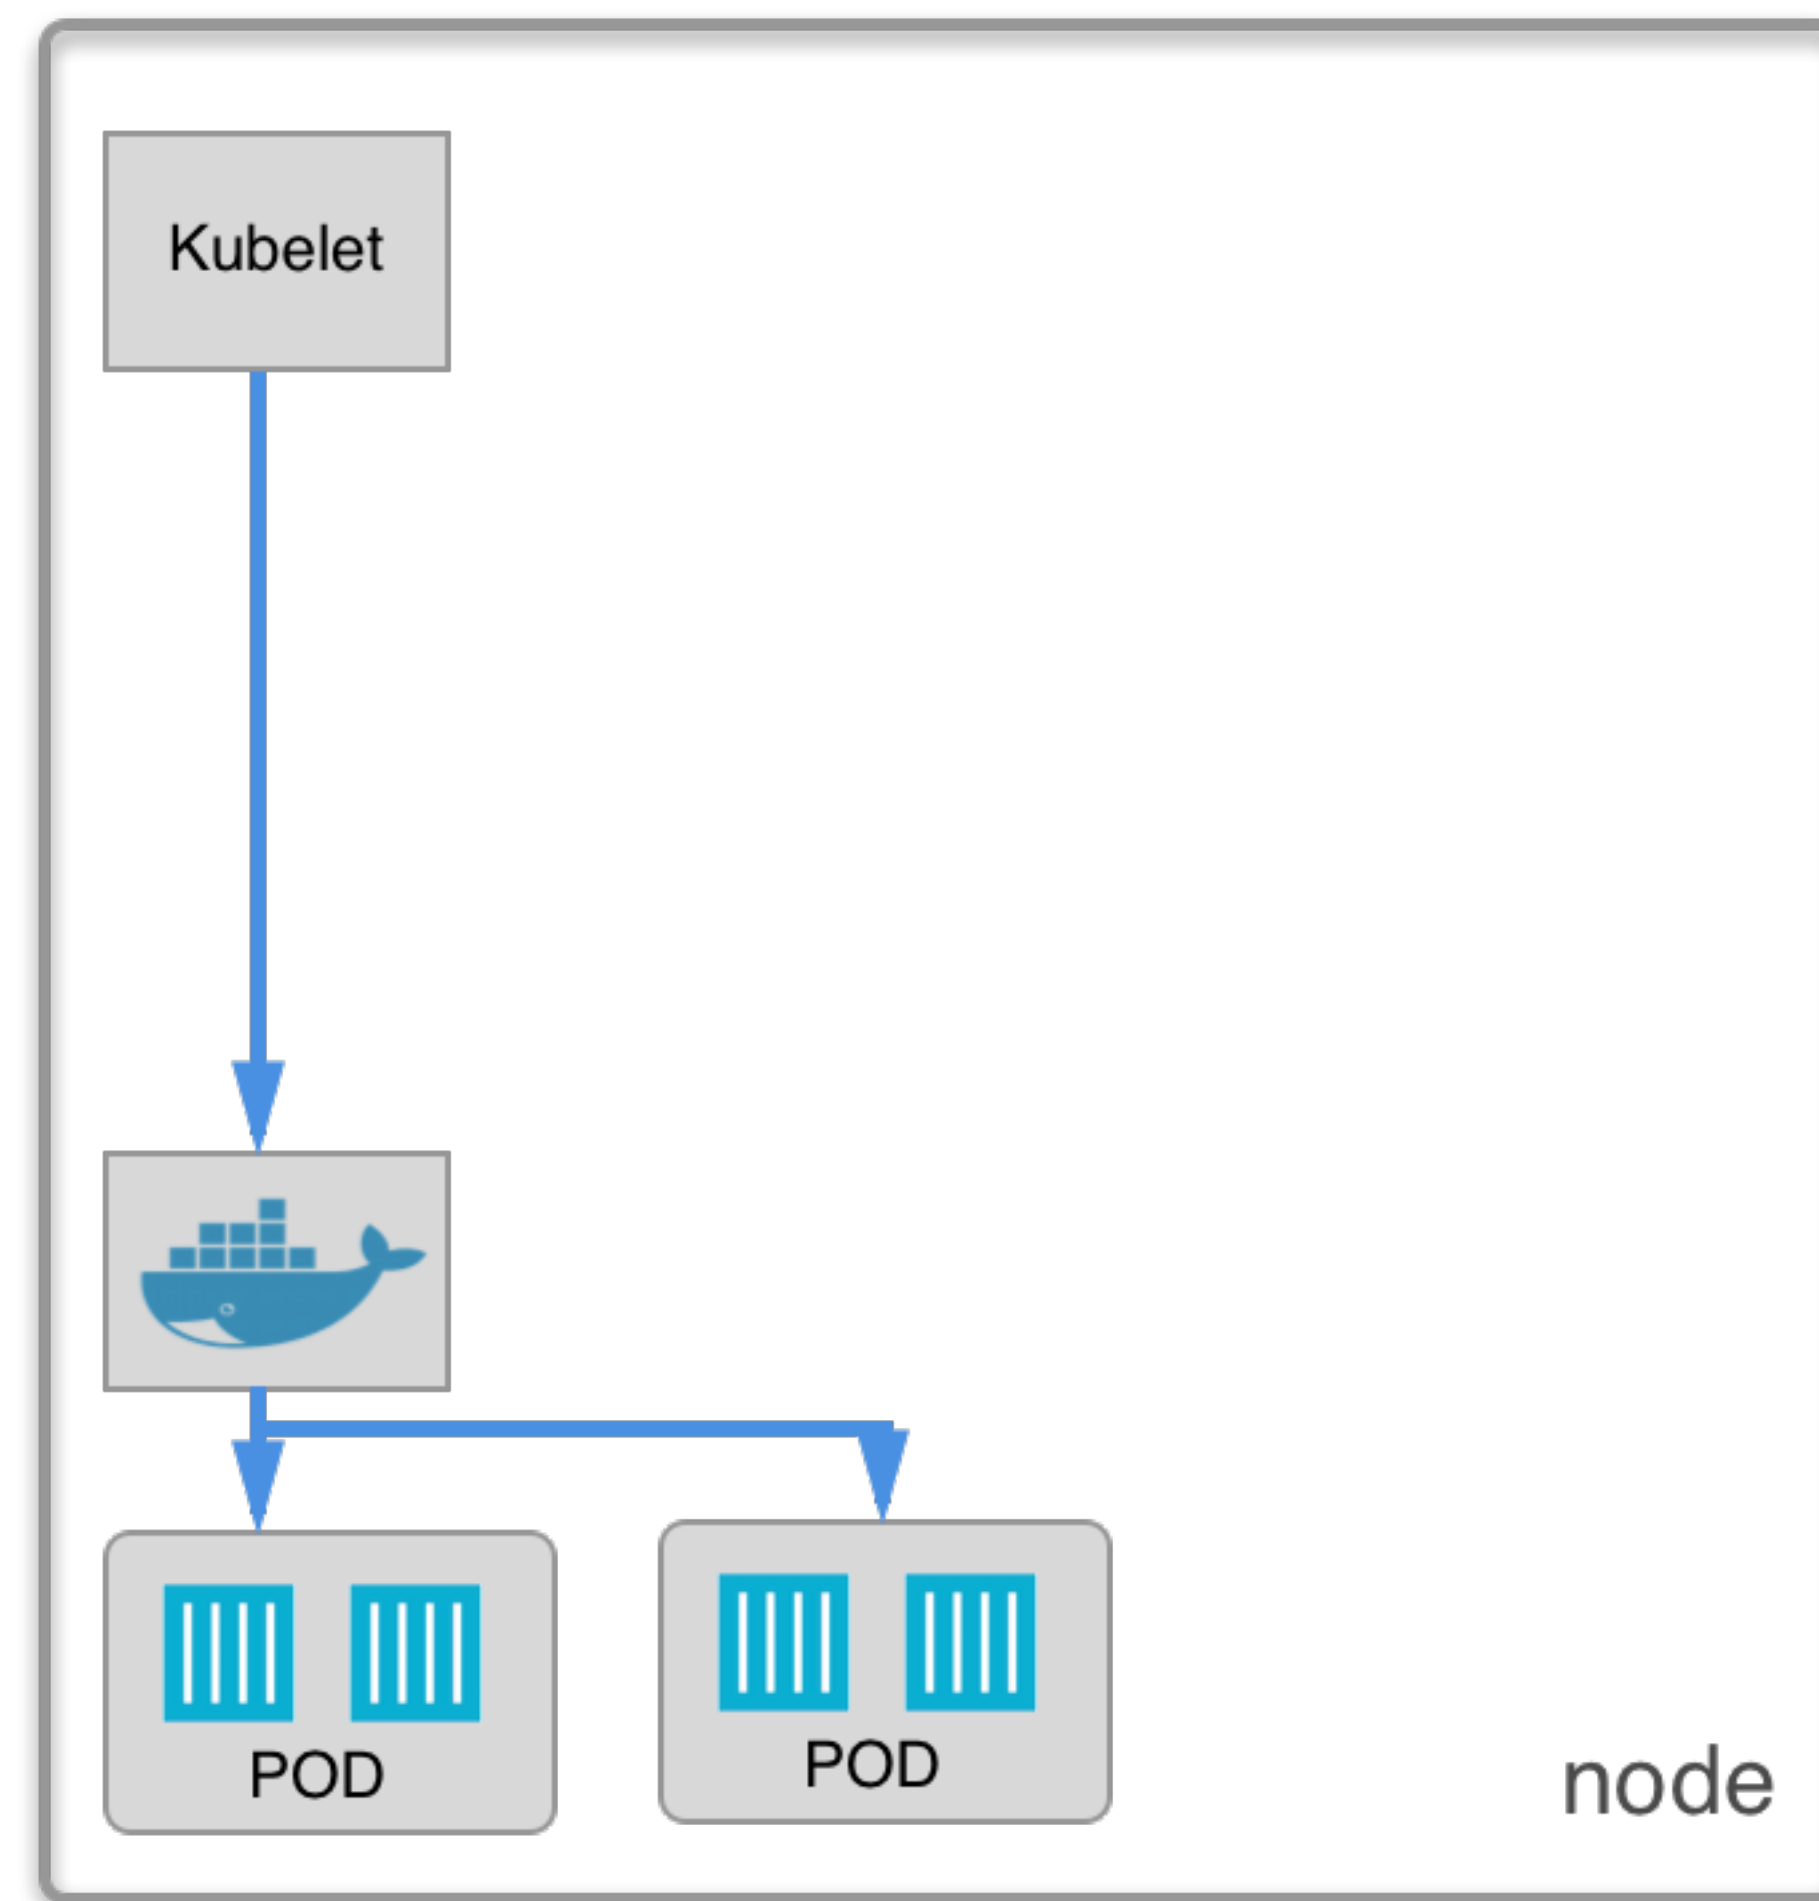
# KUBELET

- the “worker”
- reads pod's specification in YAML/JSON (from directory, own api or etcd)



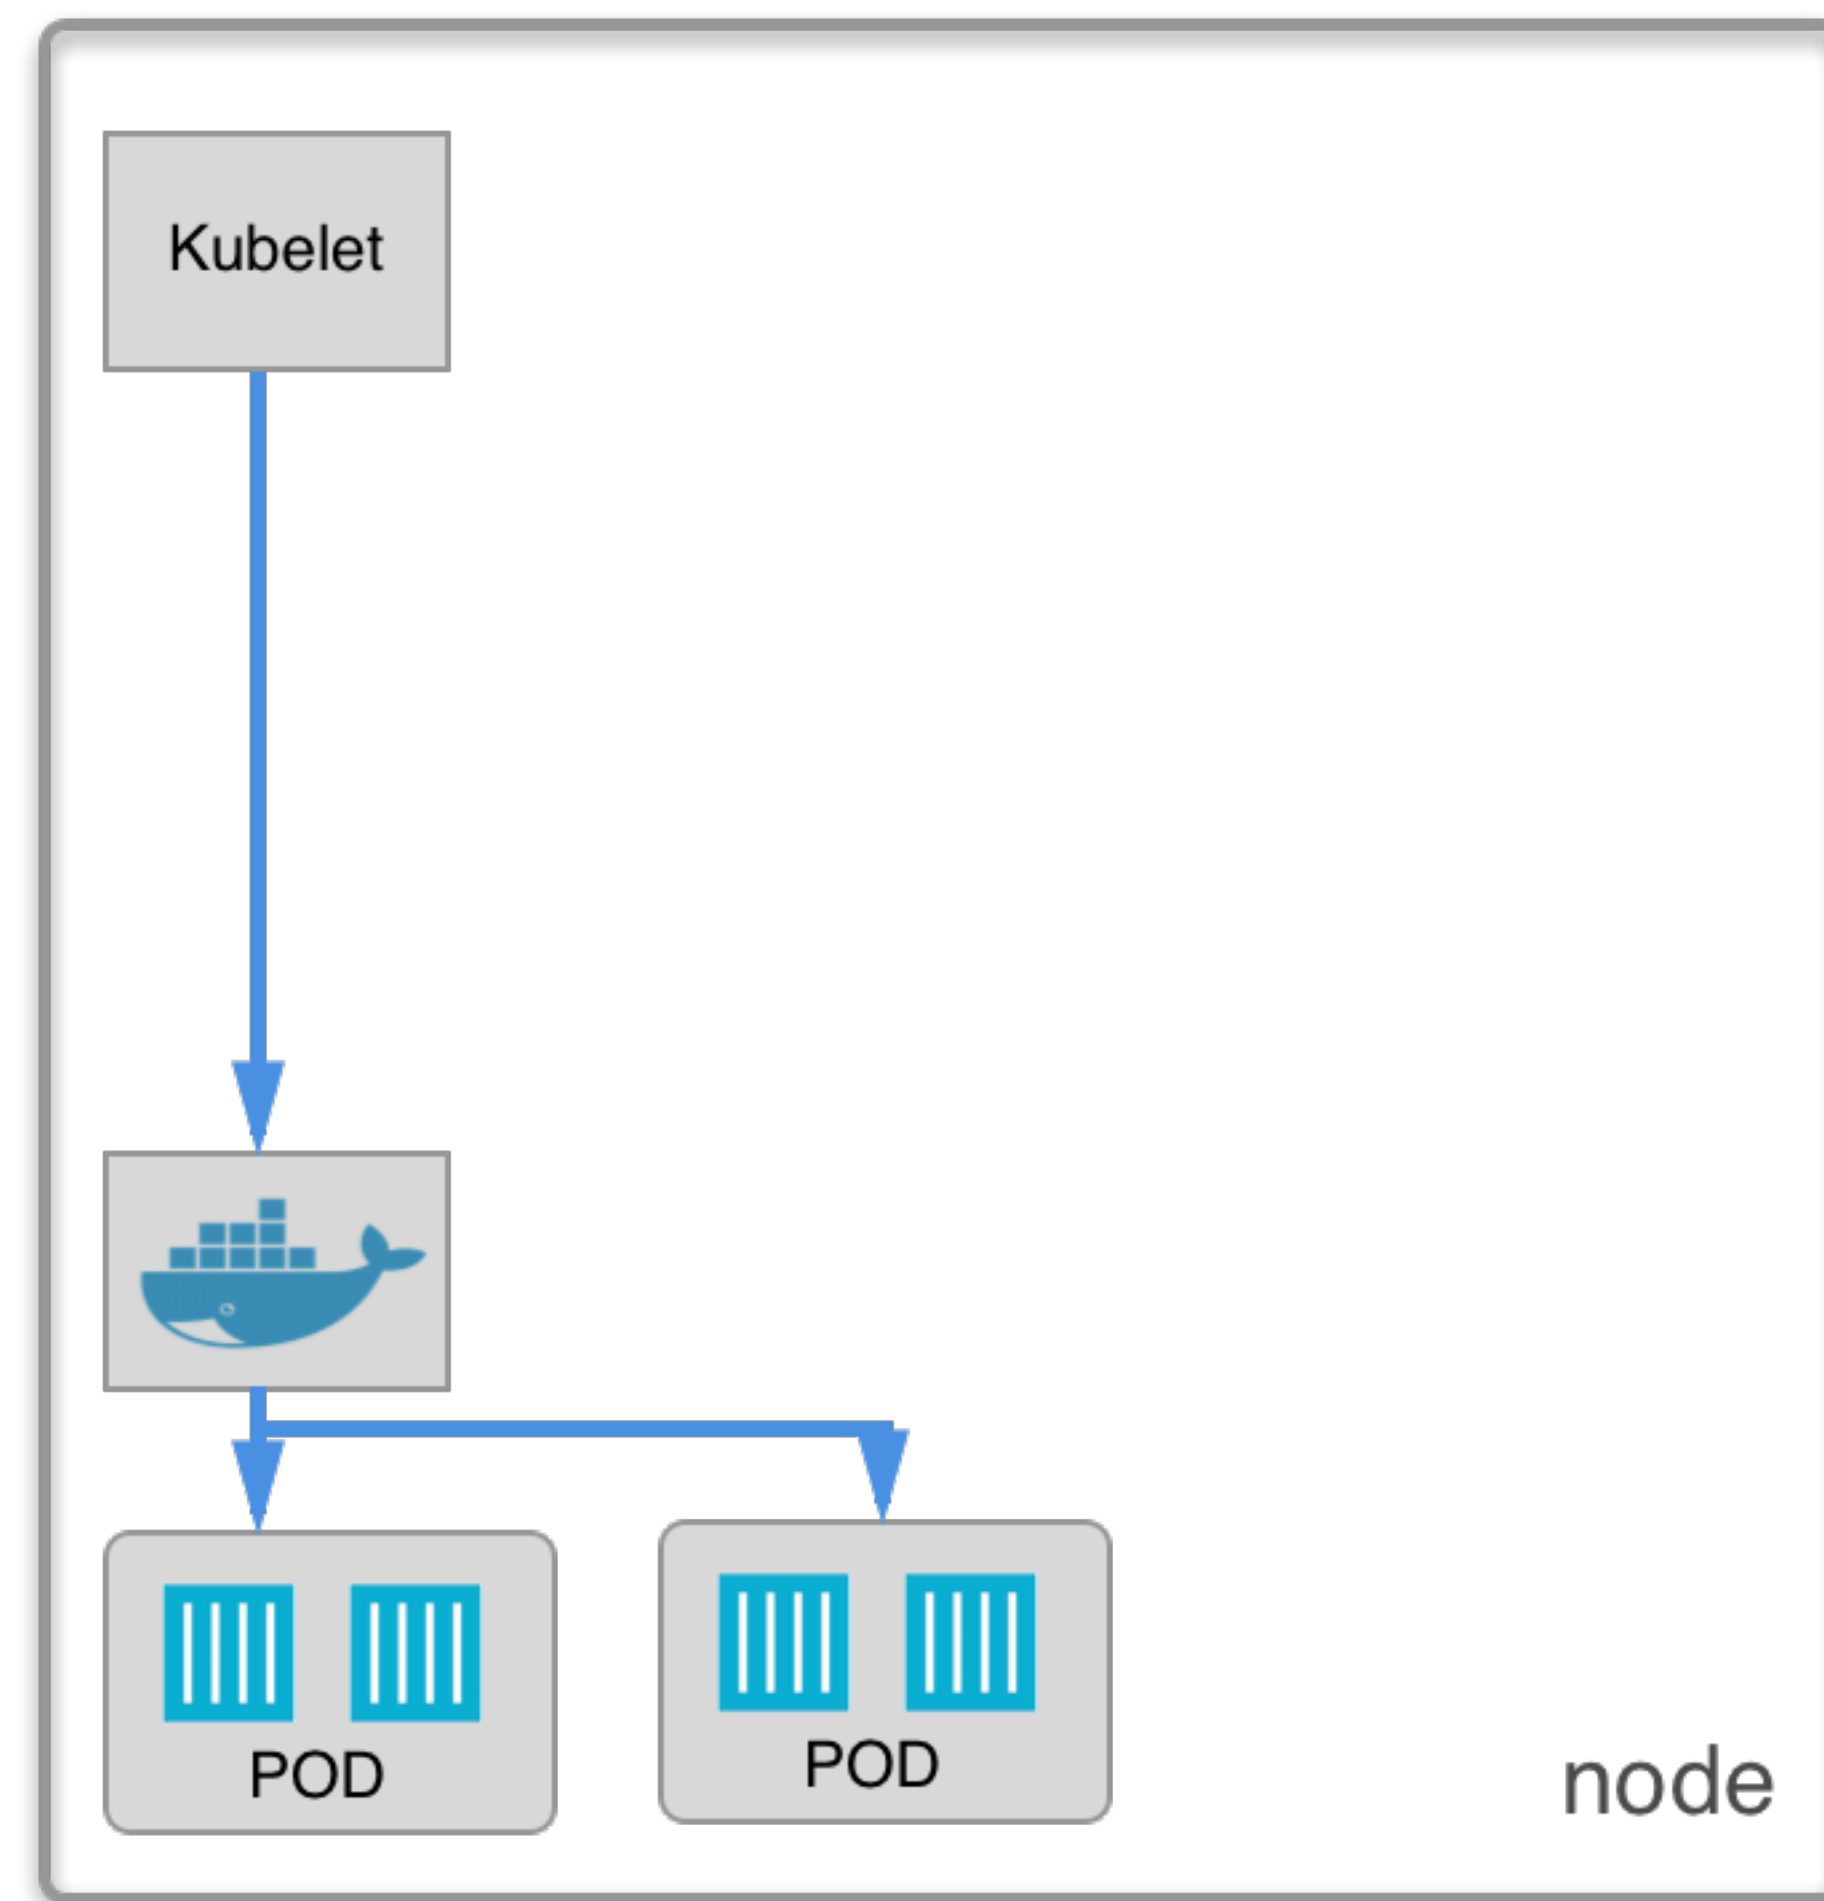
# KUBELET

- the “worker”
- reads pod's specification in YAML/JSON (from directory, own api or etcd)
- talks to the Docker



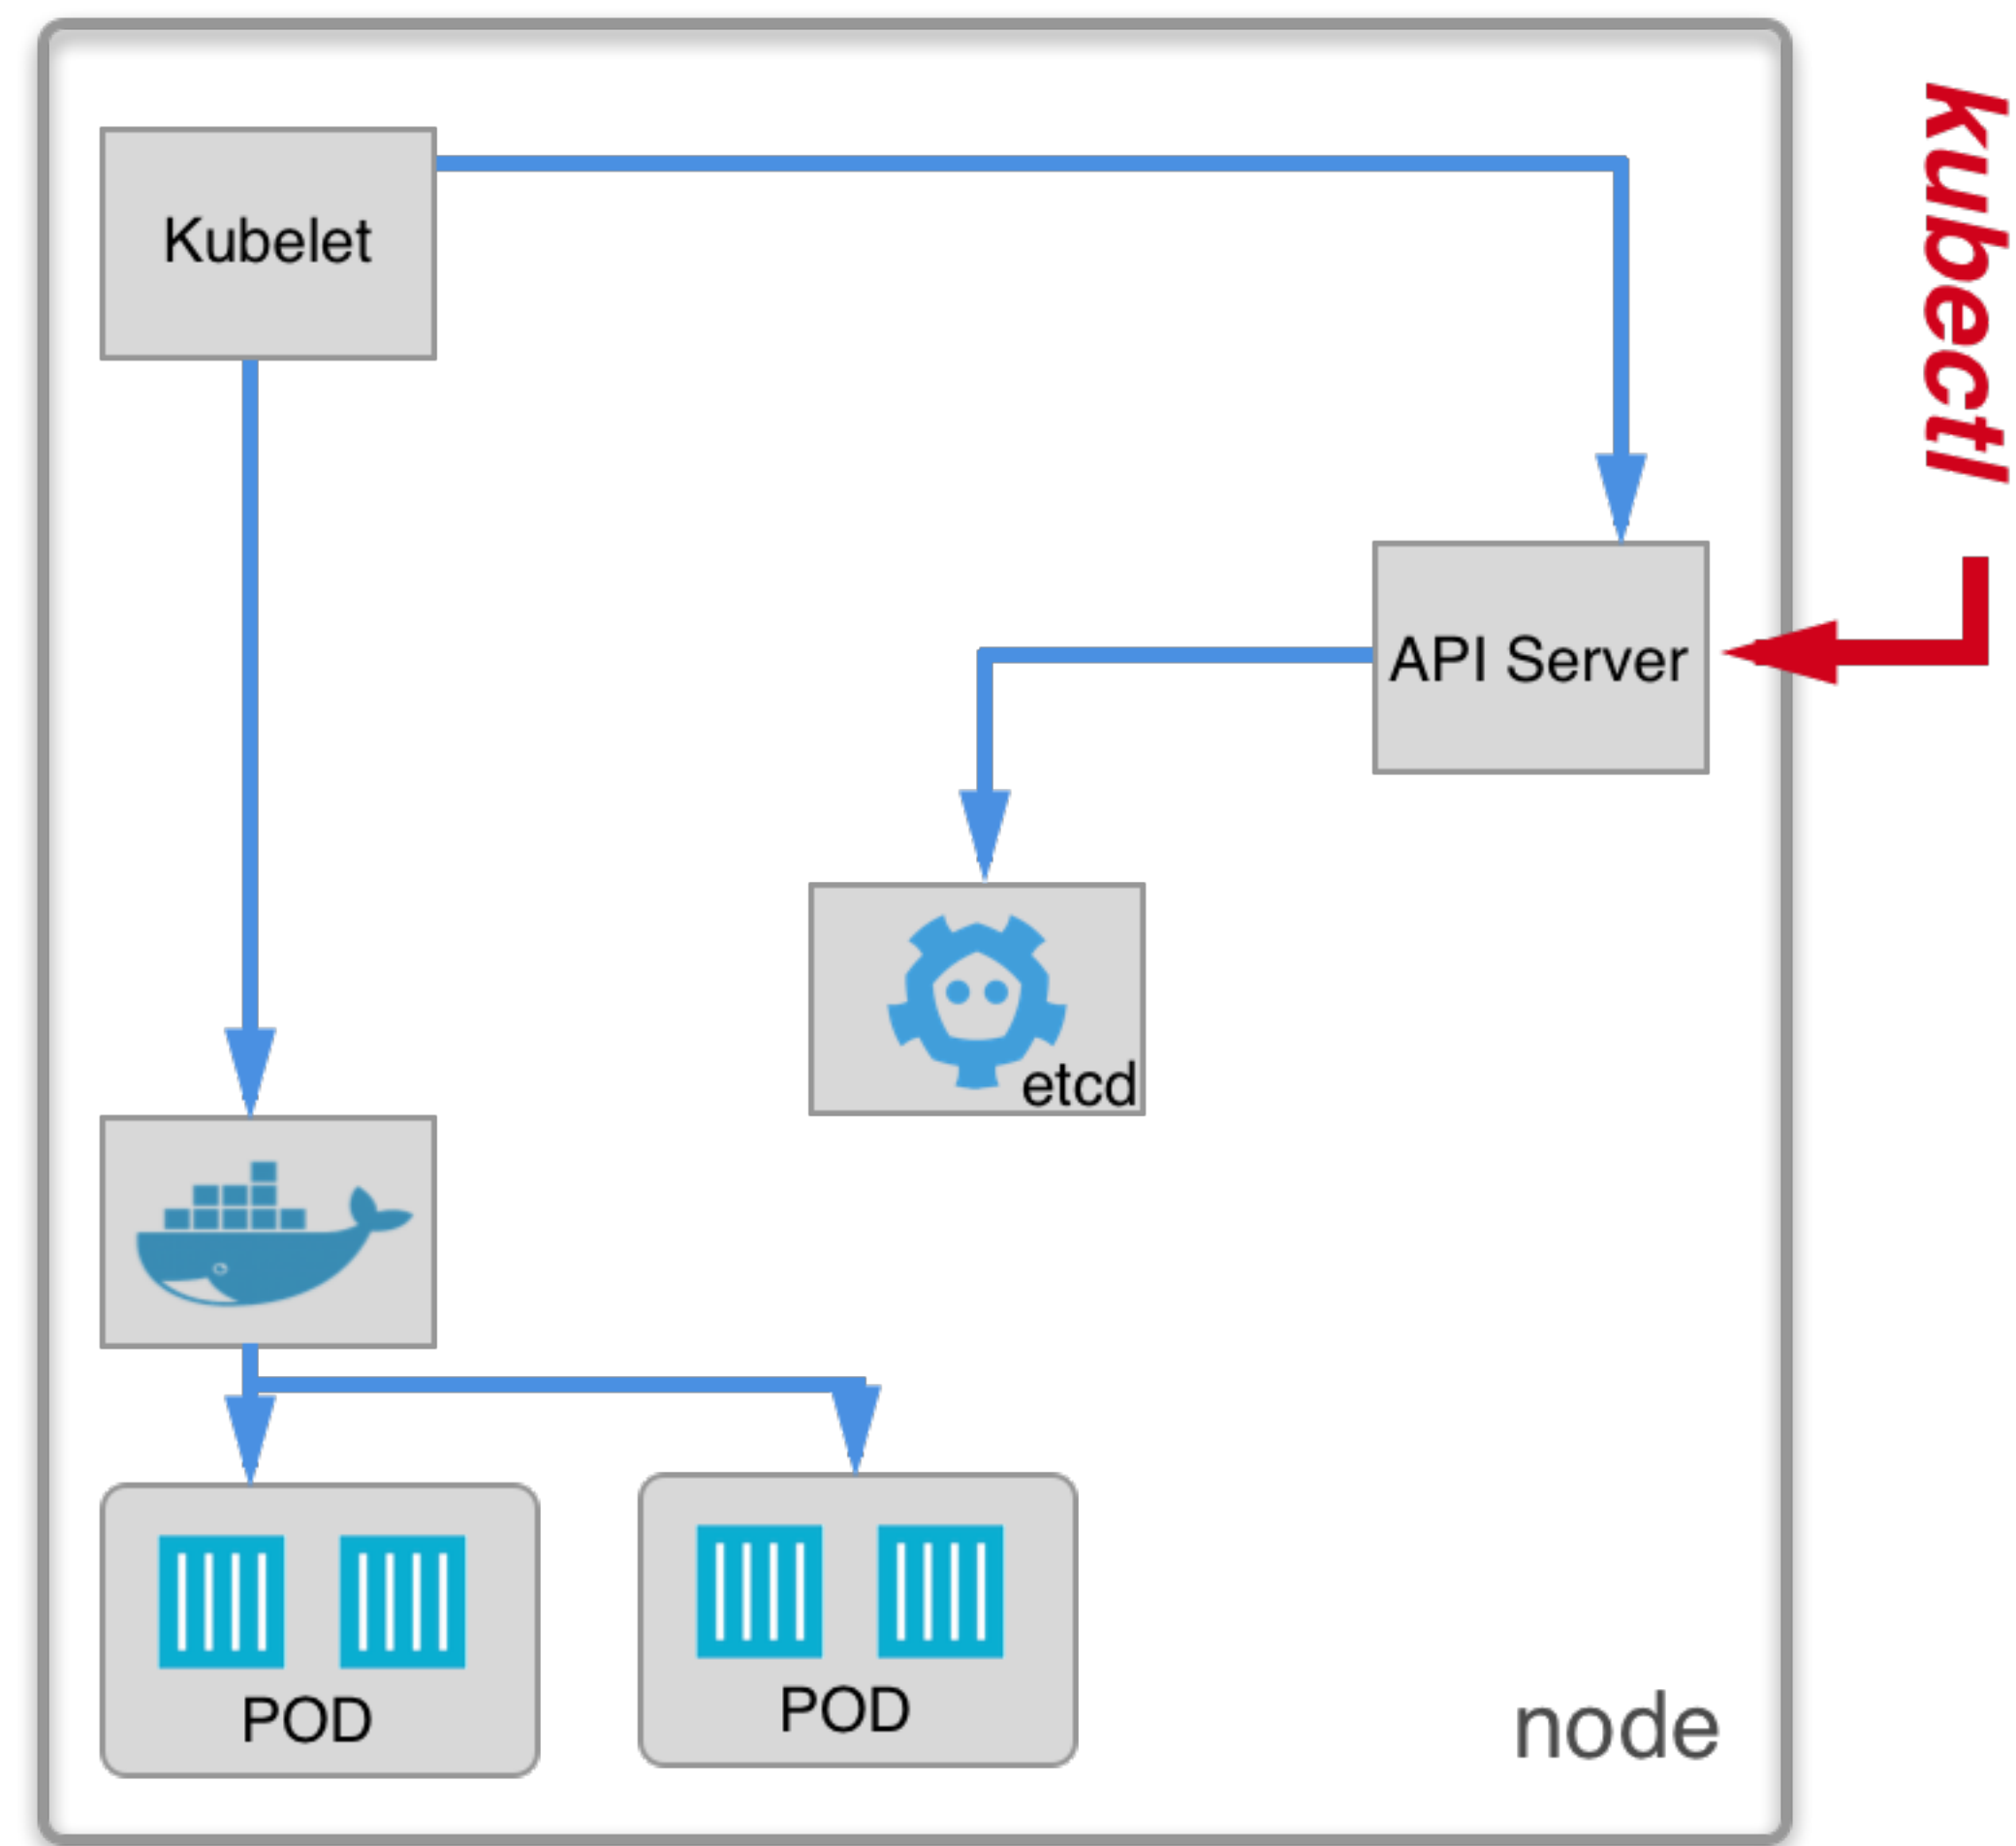
# KUBELET

- the “worker”
- reads pod’s specification in YAML/JSON (from directory, own api or etcd)
- talks to the Docker
- creates pods / containers



# ETCD + API SERVER

etcd

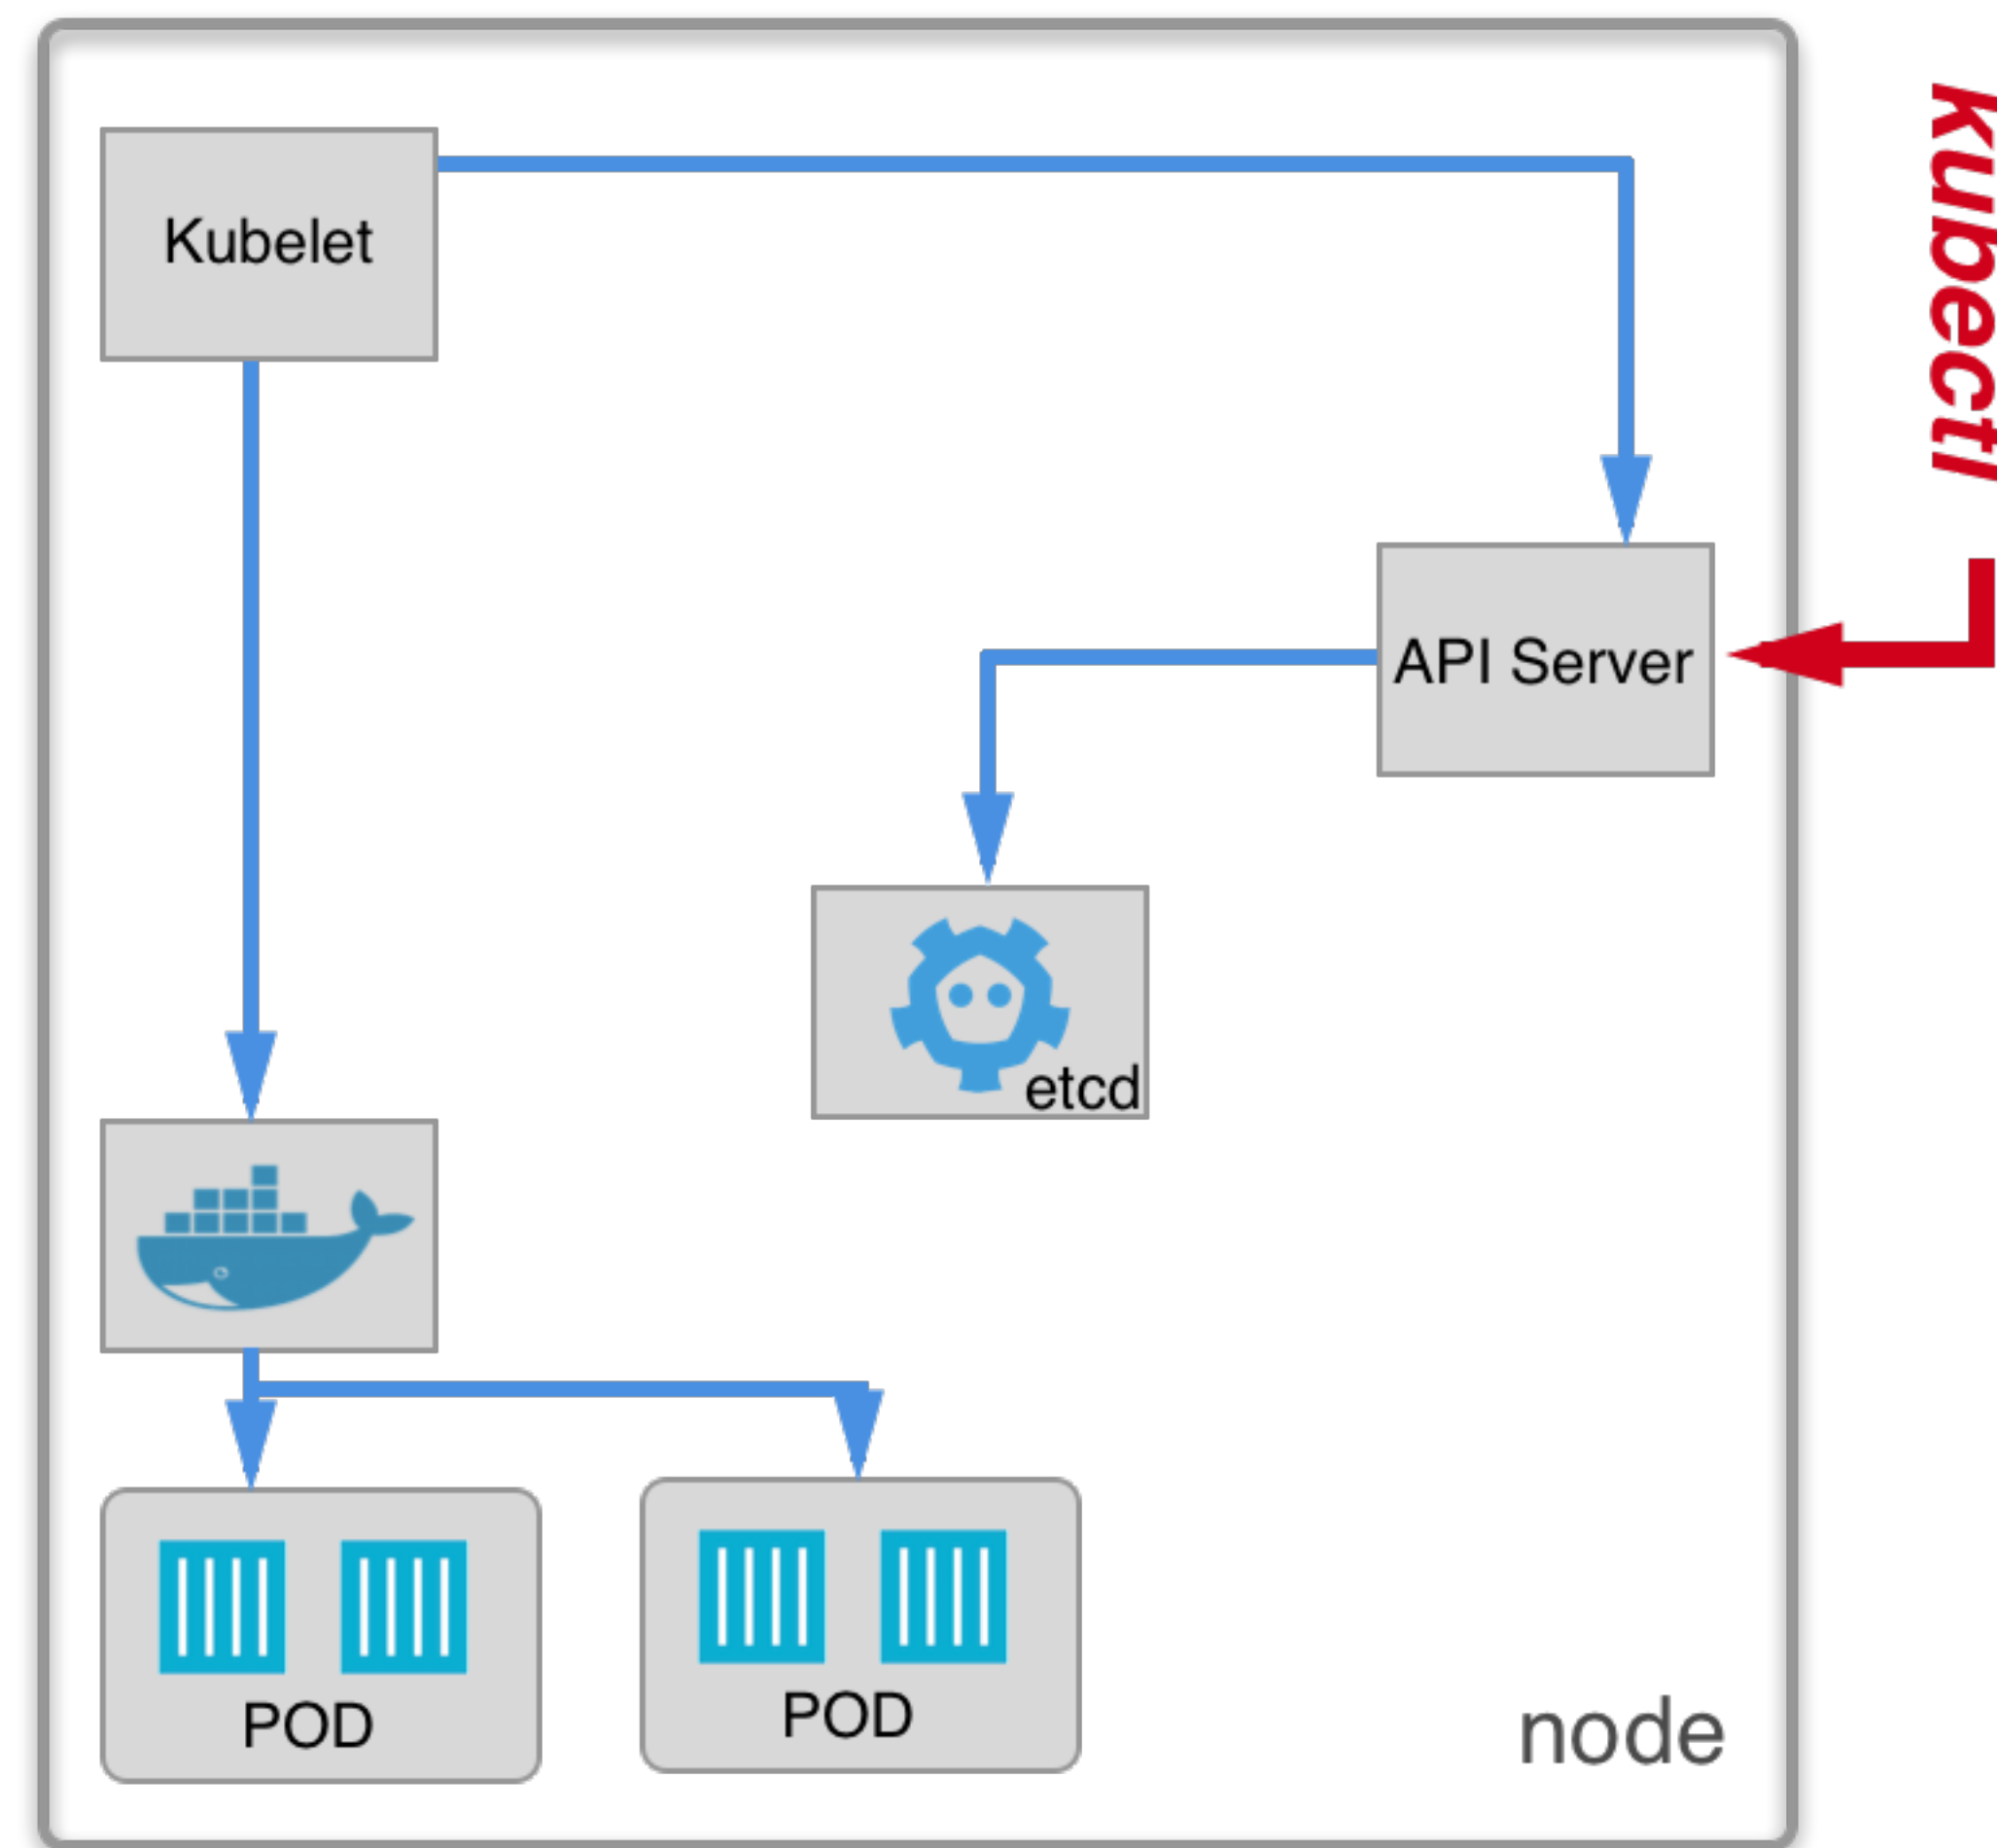




# ETCD + API SERVER

etcd

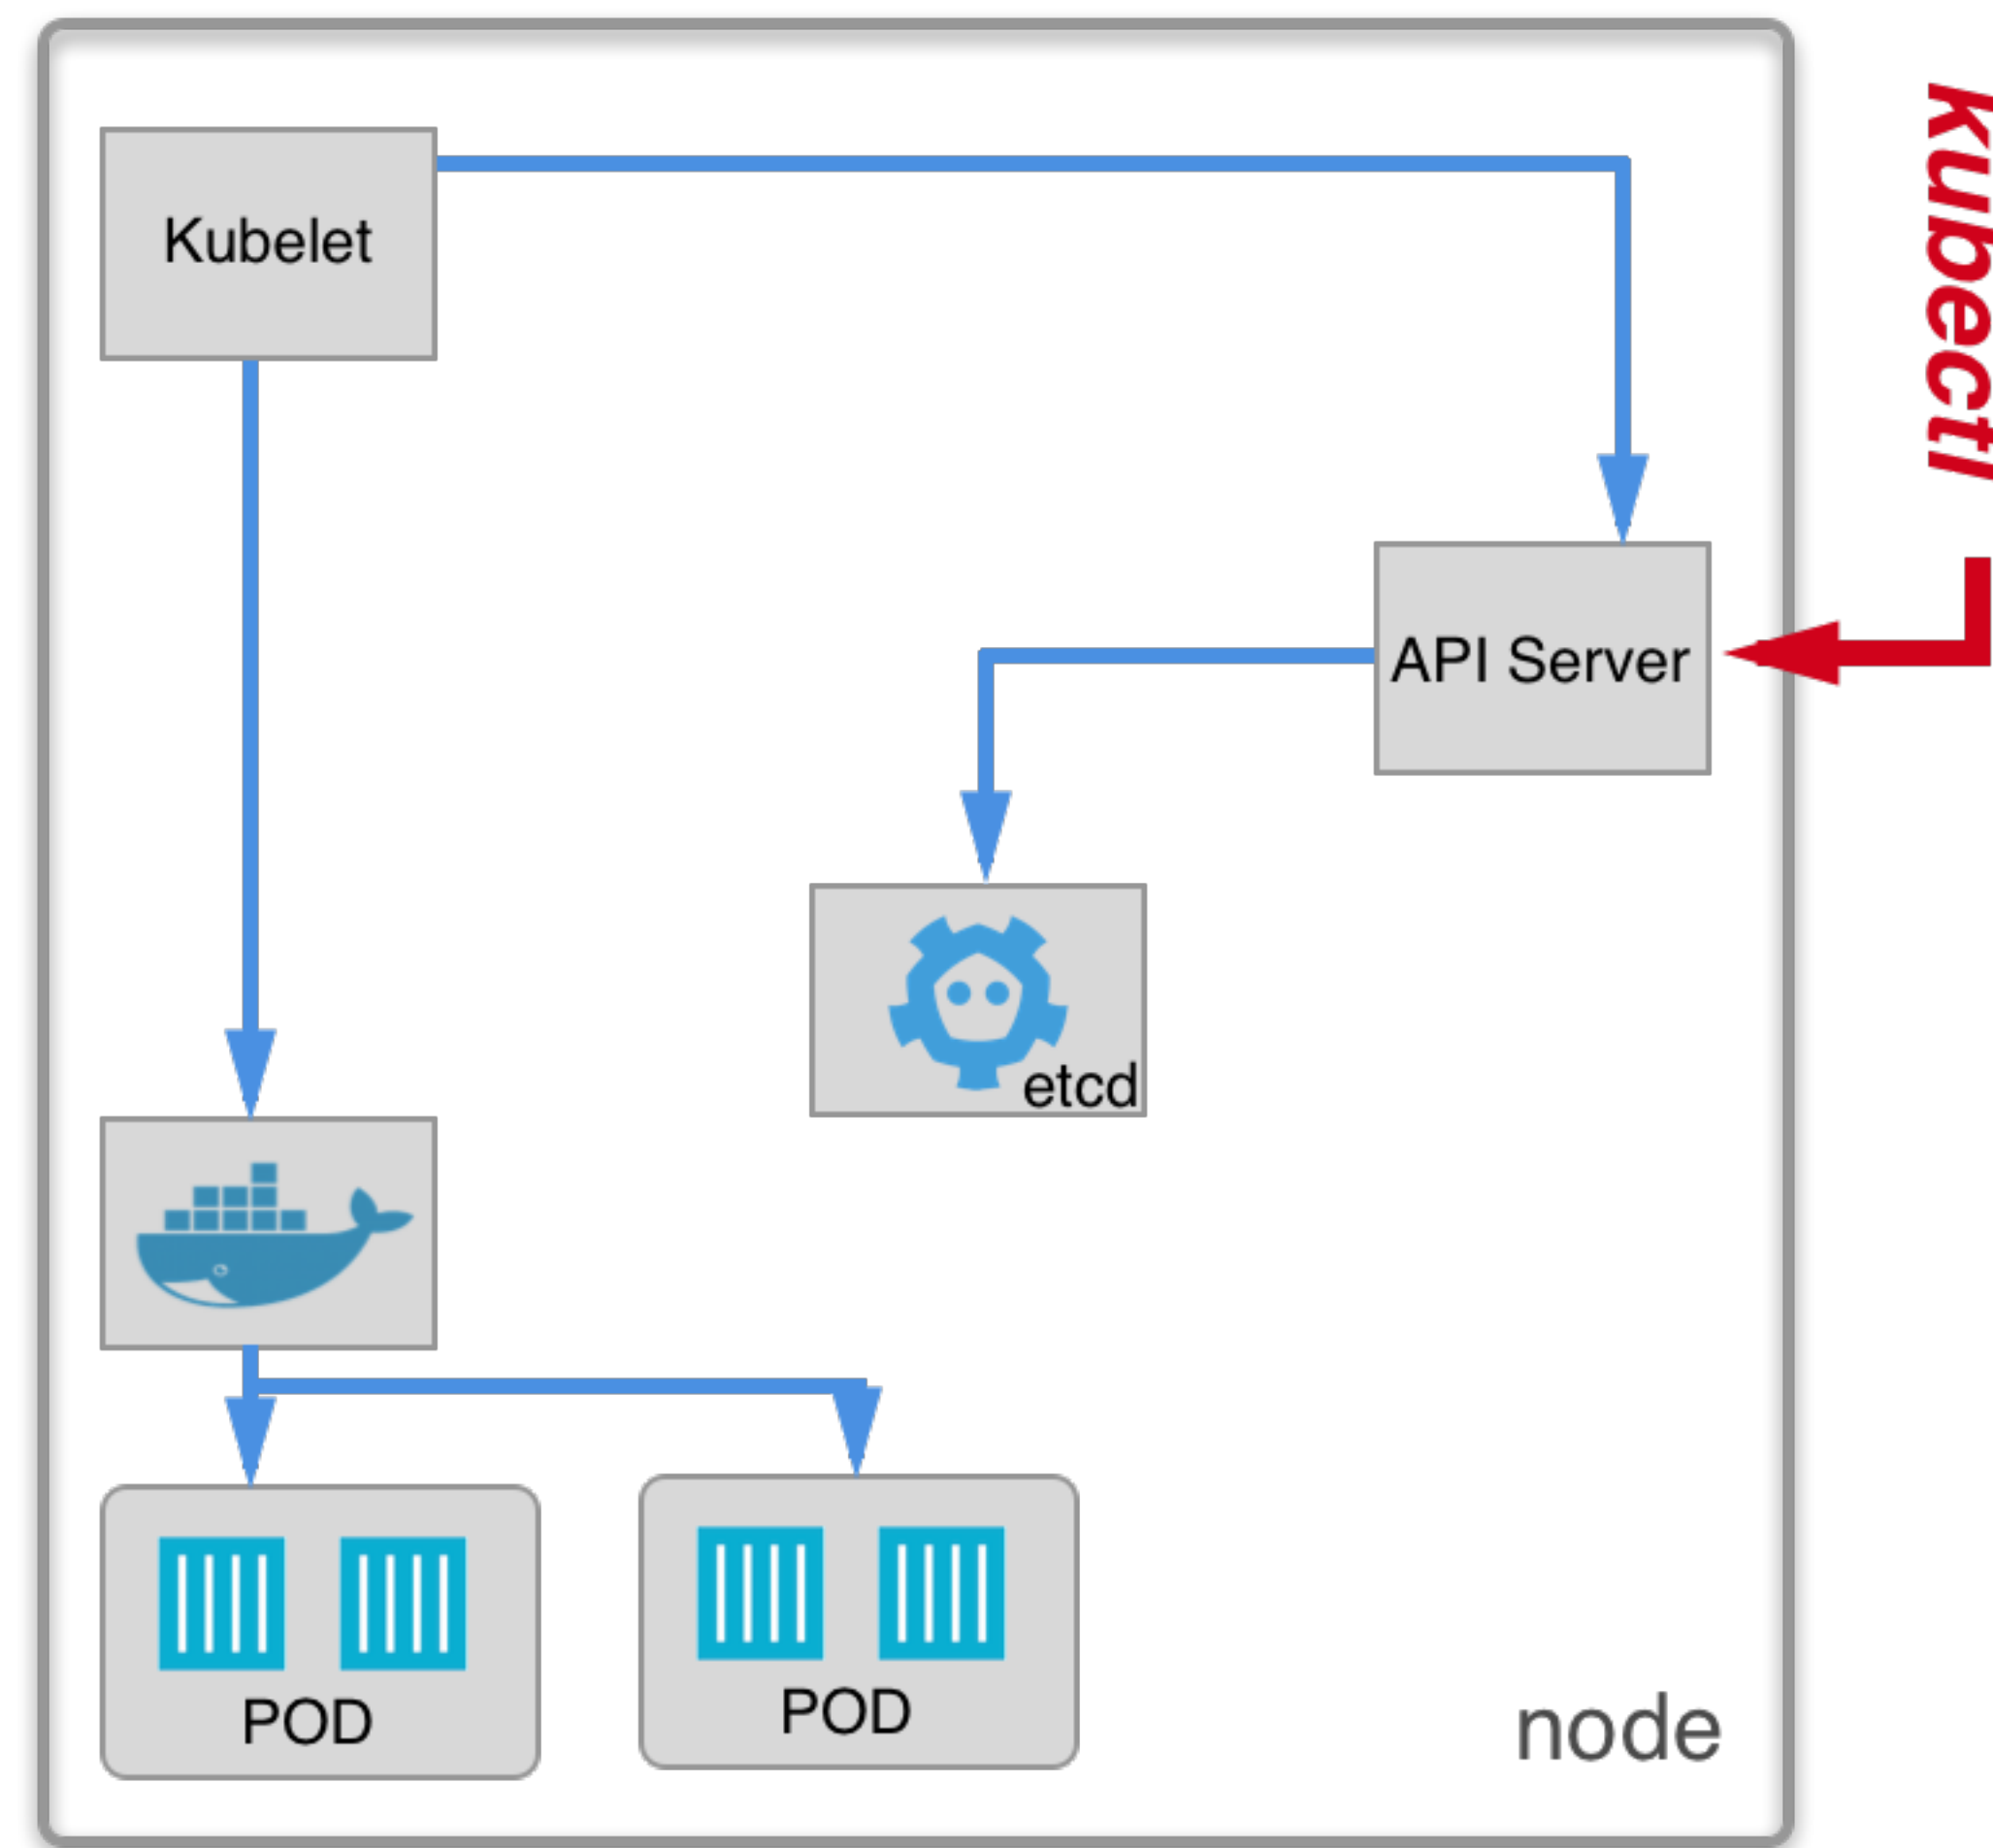
- key-value distributed database



# ETCD + API SERVER

## etcd

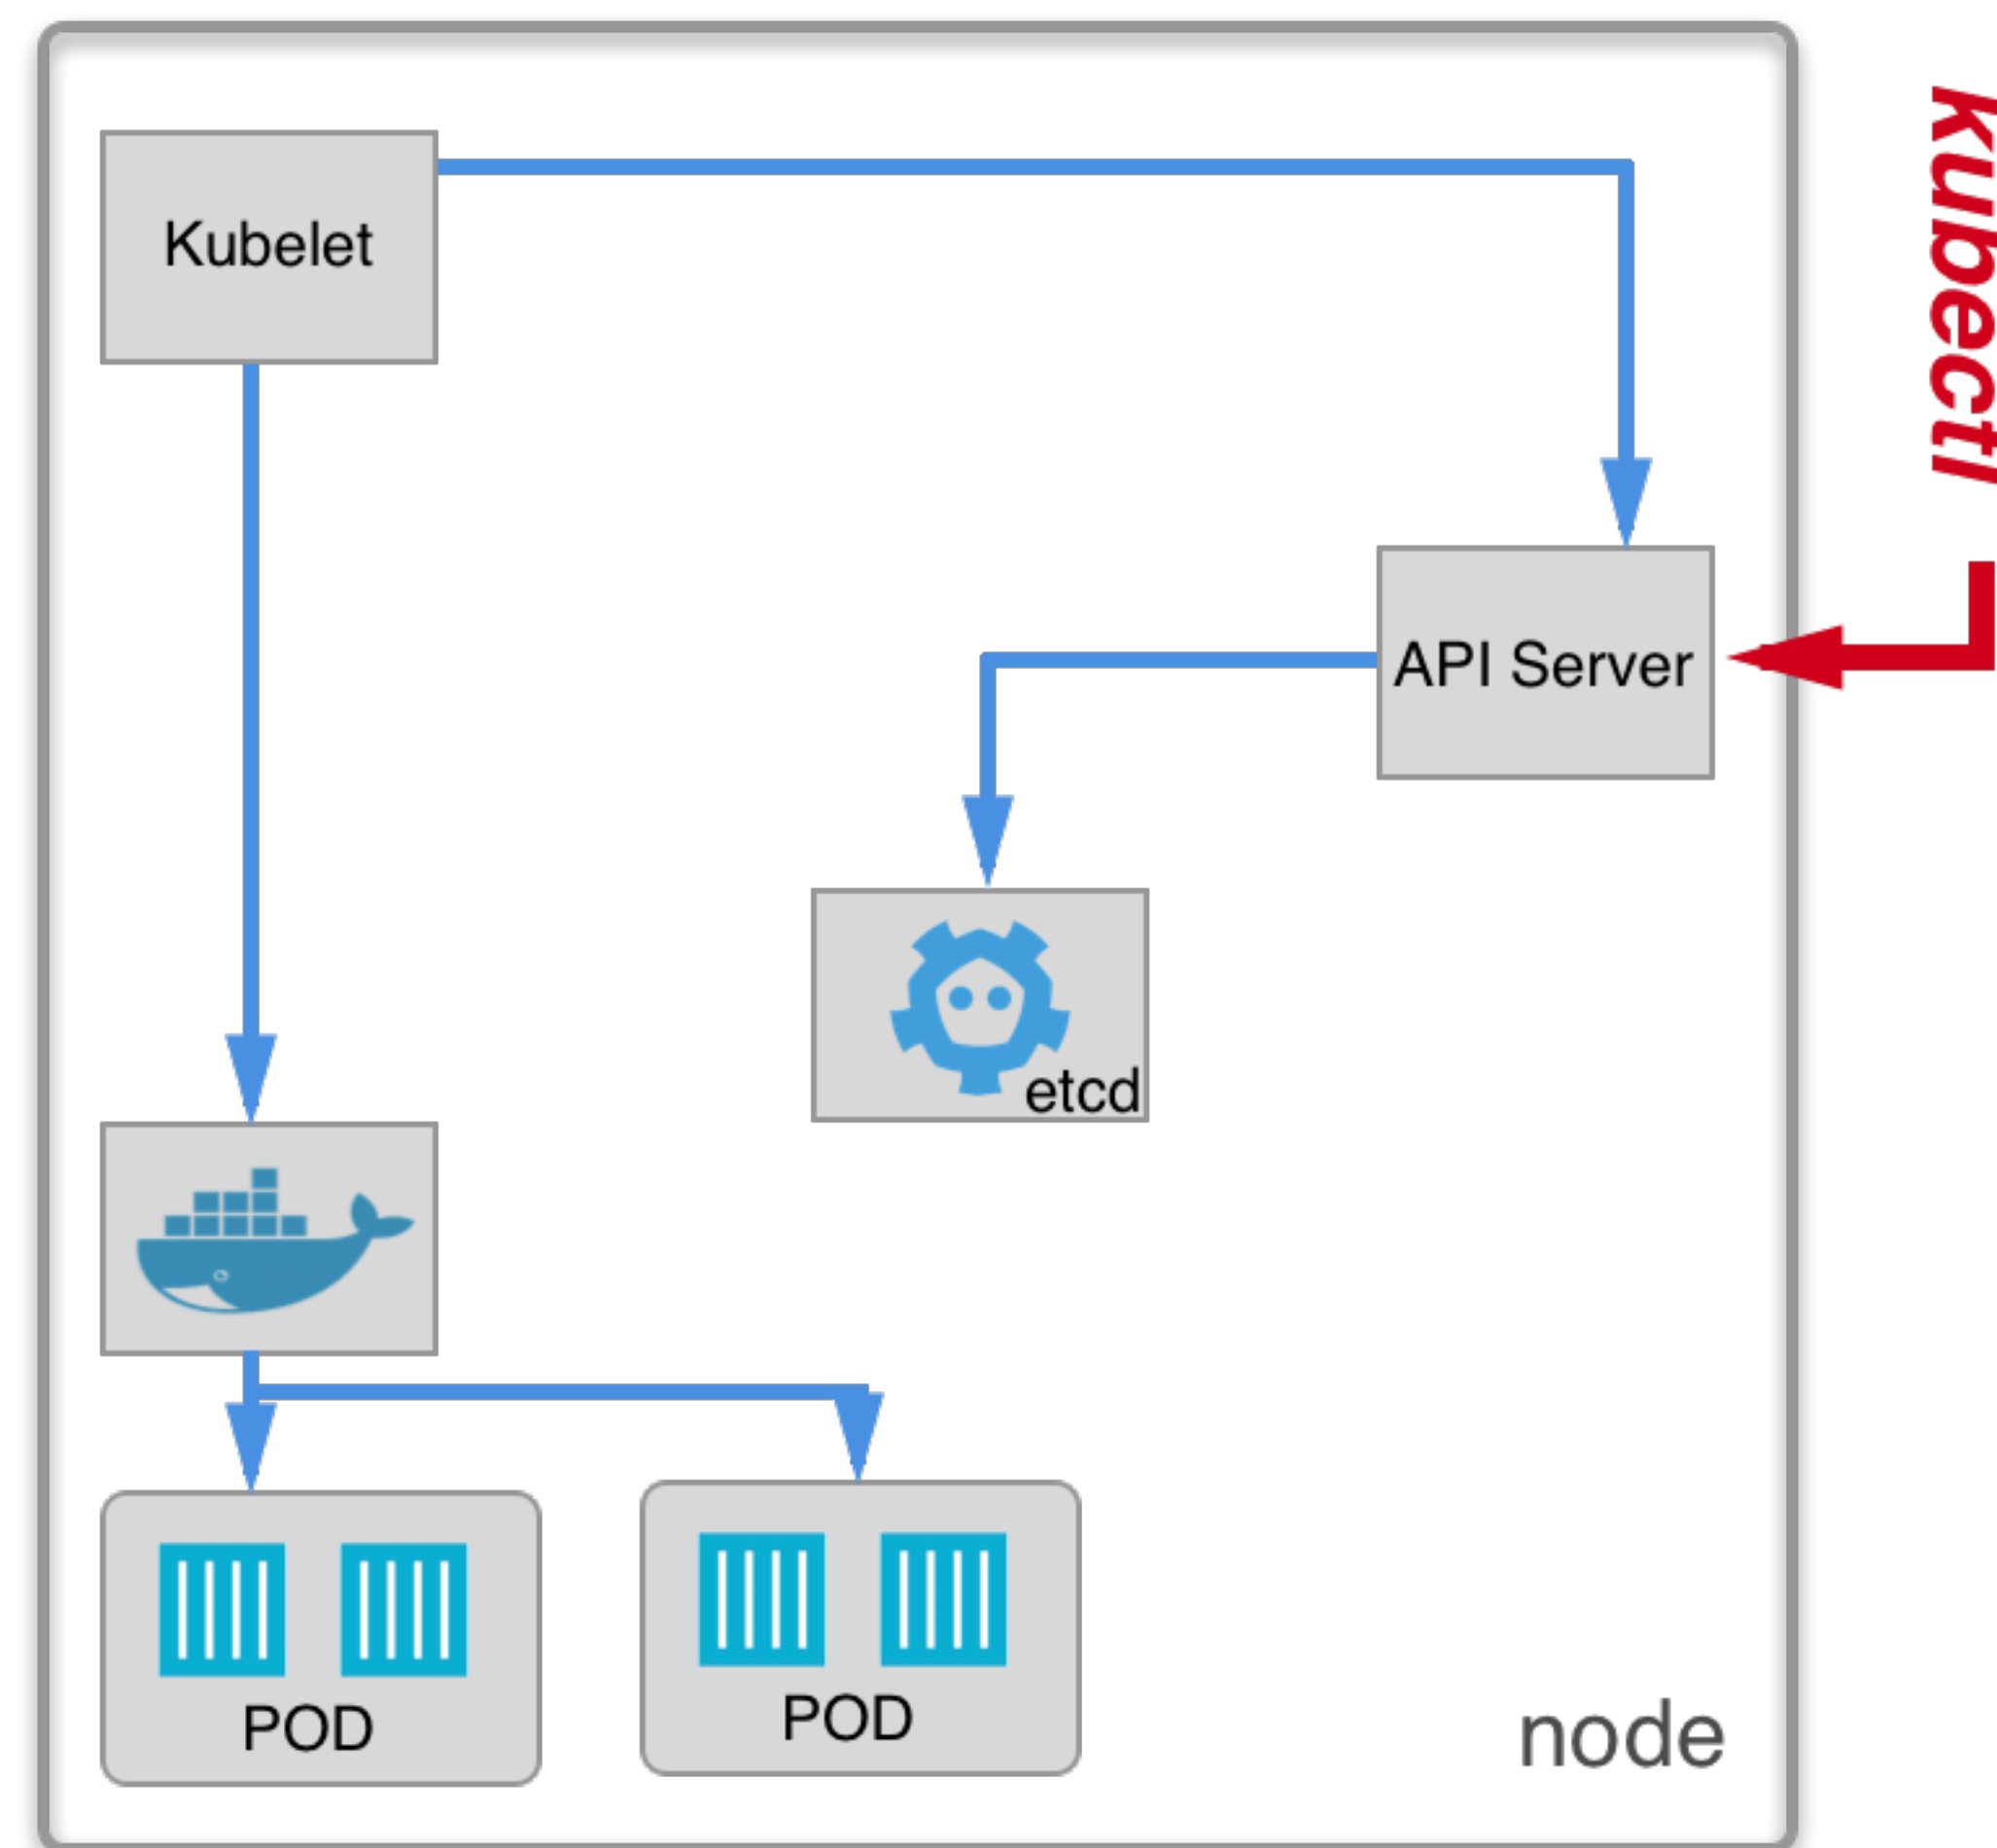
- key-value distributed database
- keys can be watchable



# ETCD + API SERVER

## etcd

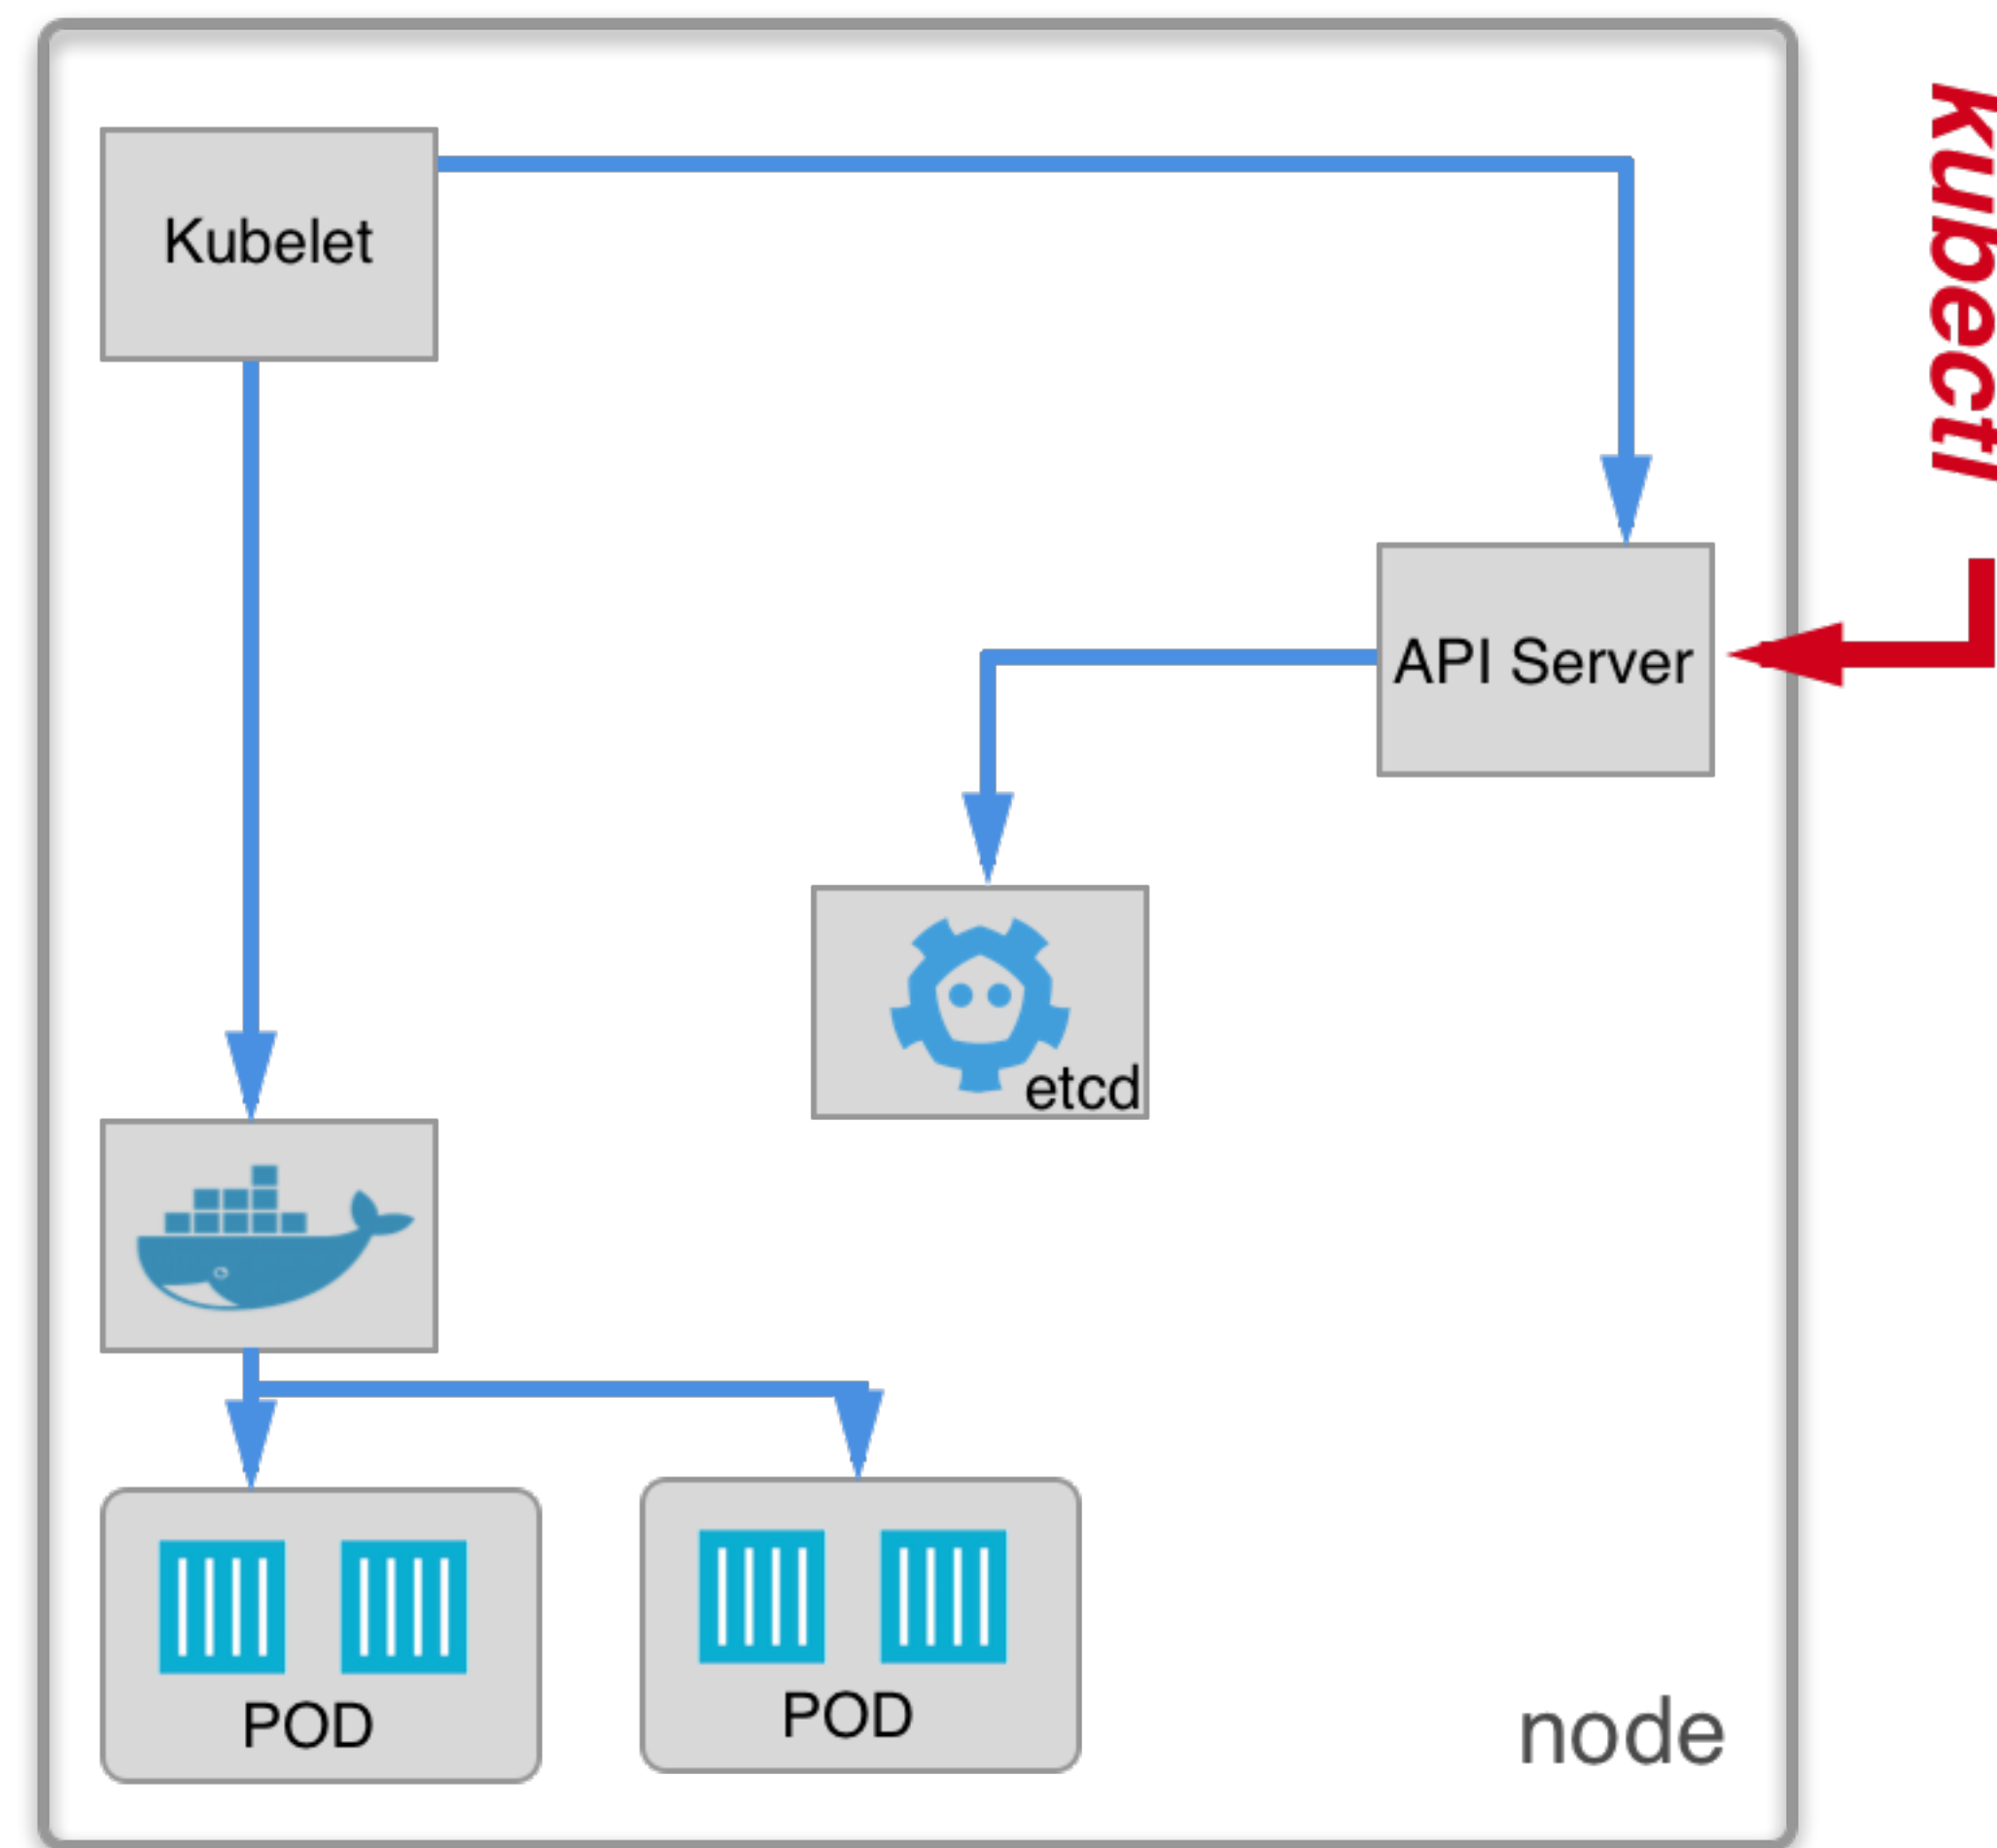
- key-value distributed database
- keys can be watchable
- “similar” to zookeeper or consul



# ETCD + API SERVER

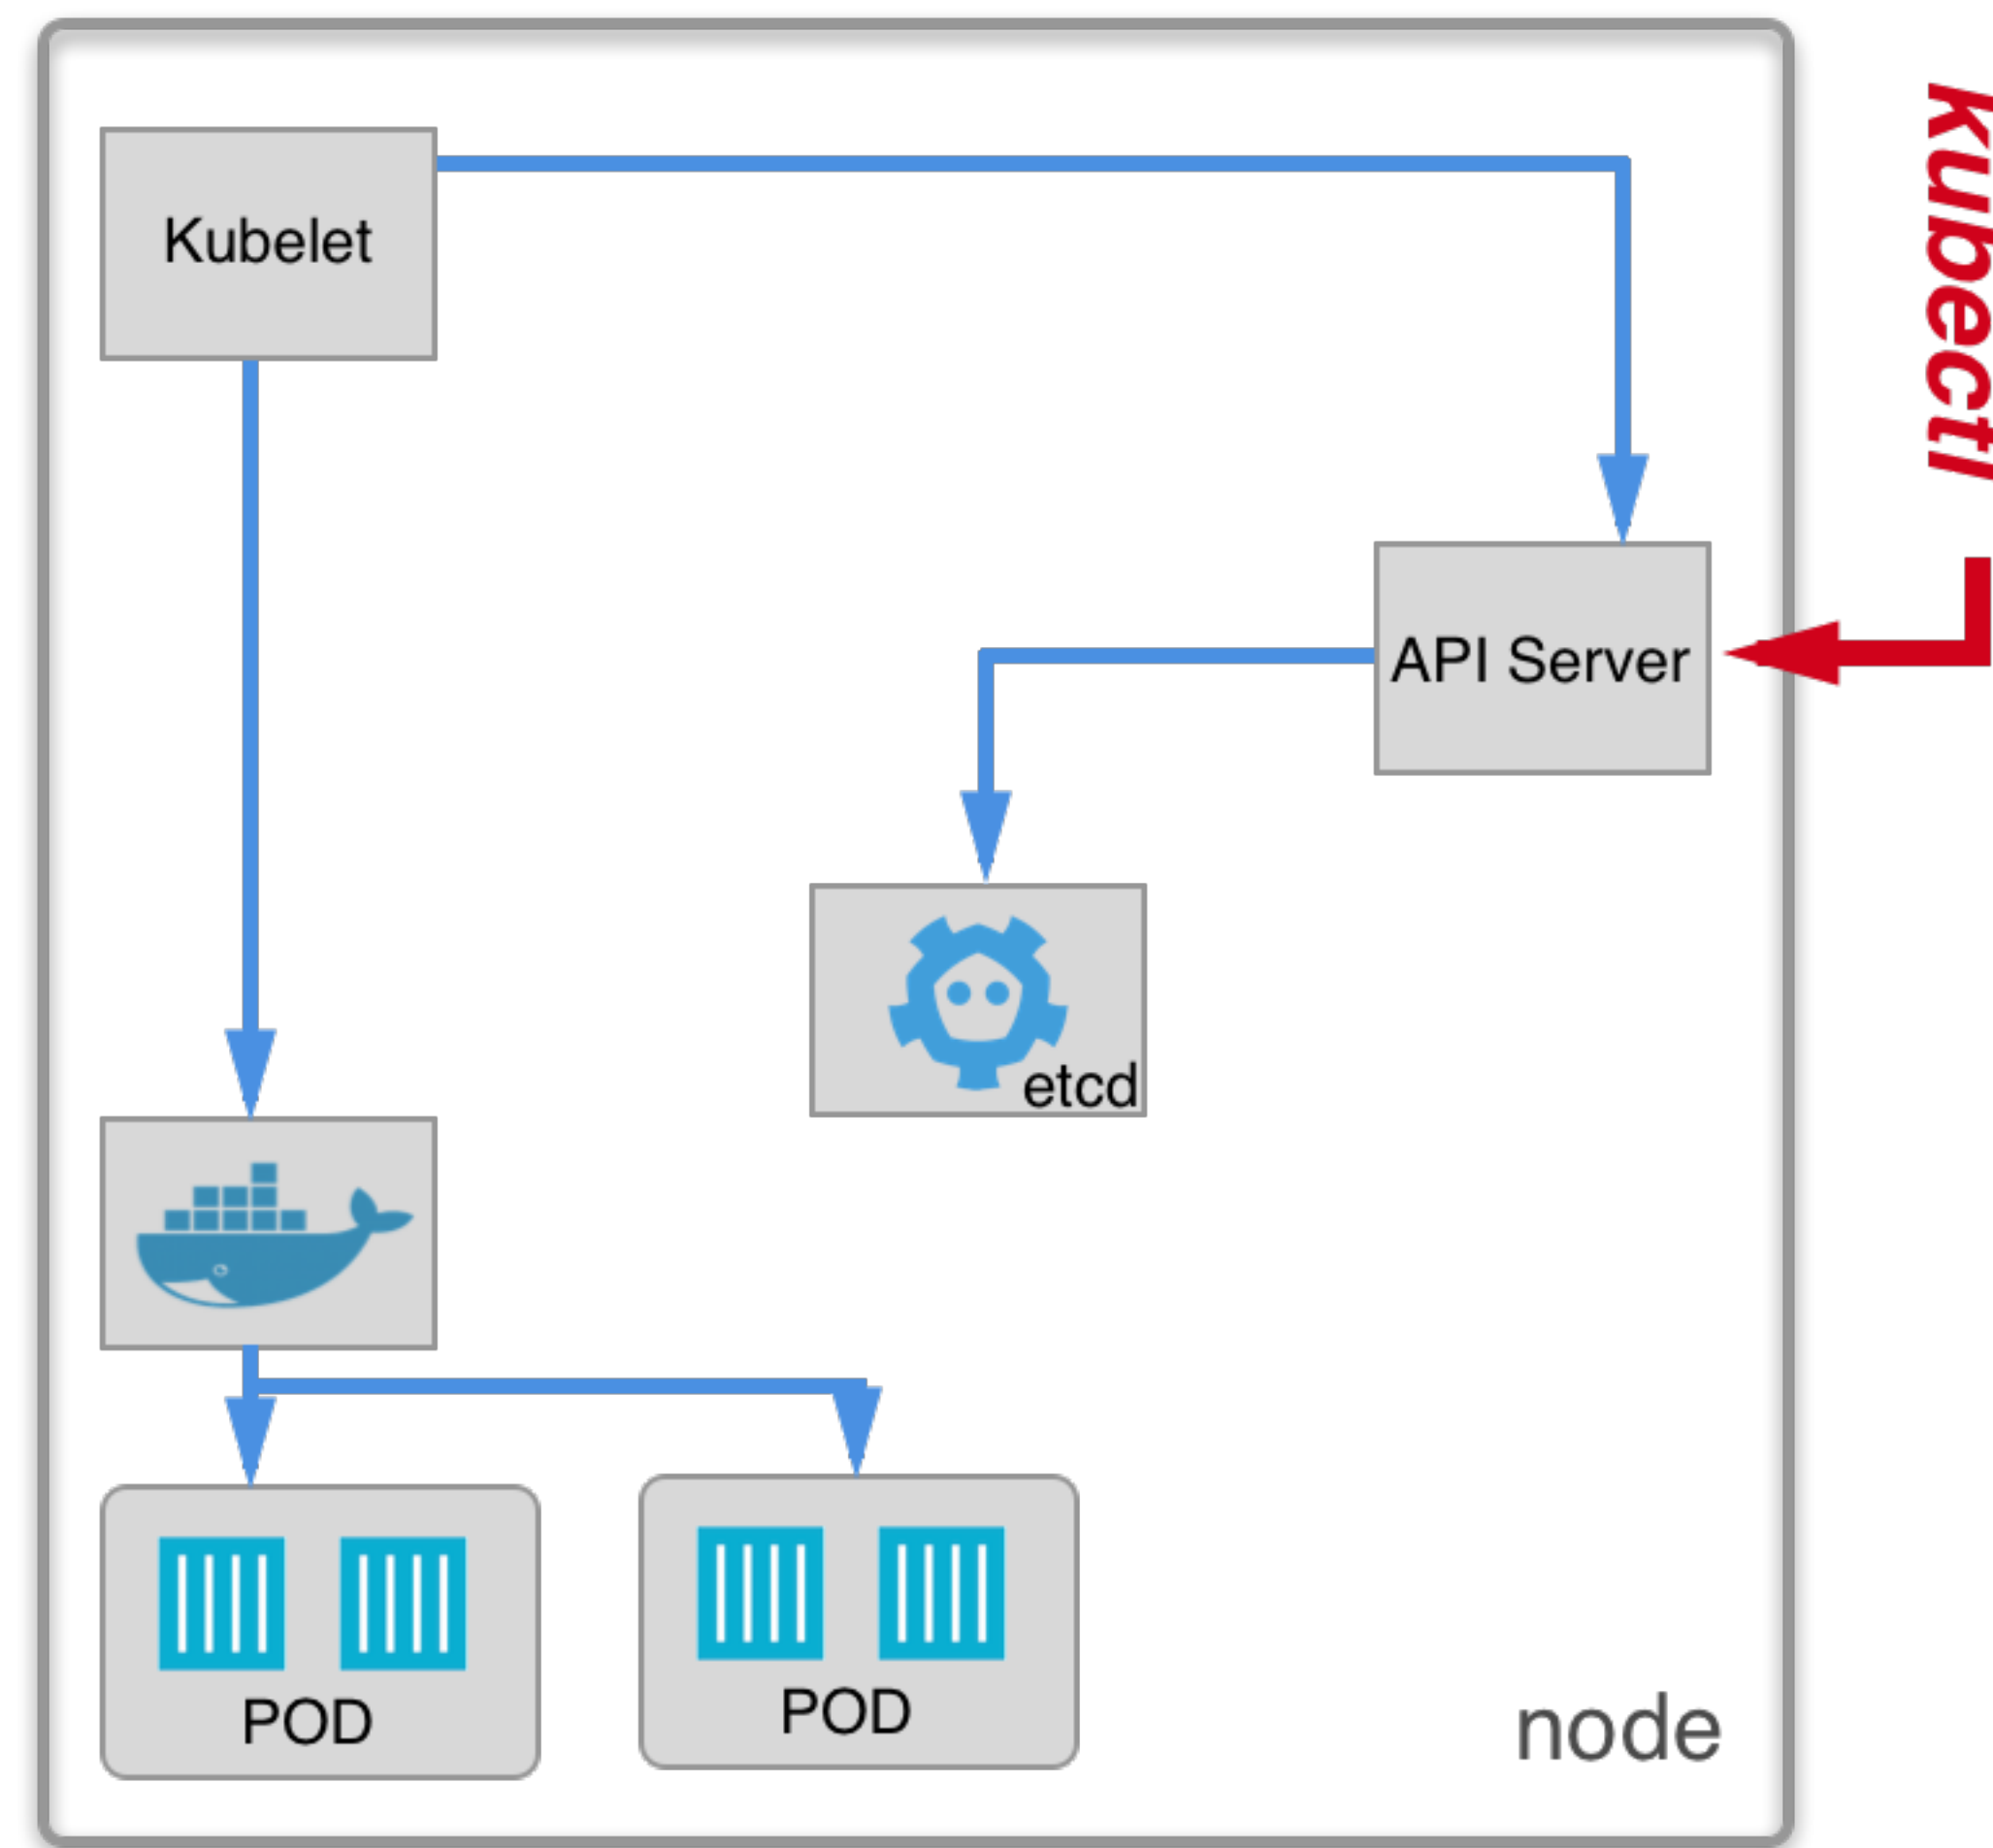
## etcd

- key-value distributed database
- keys can be watchable
- “similar” to zookeeper or consul
- stores state of the cluster



# ETCD + API SERVER

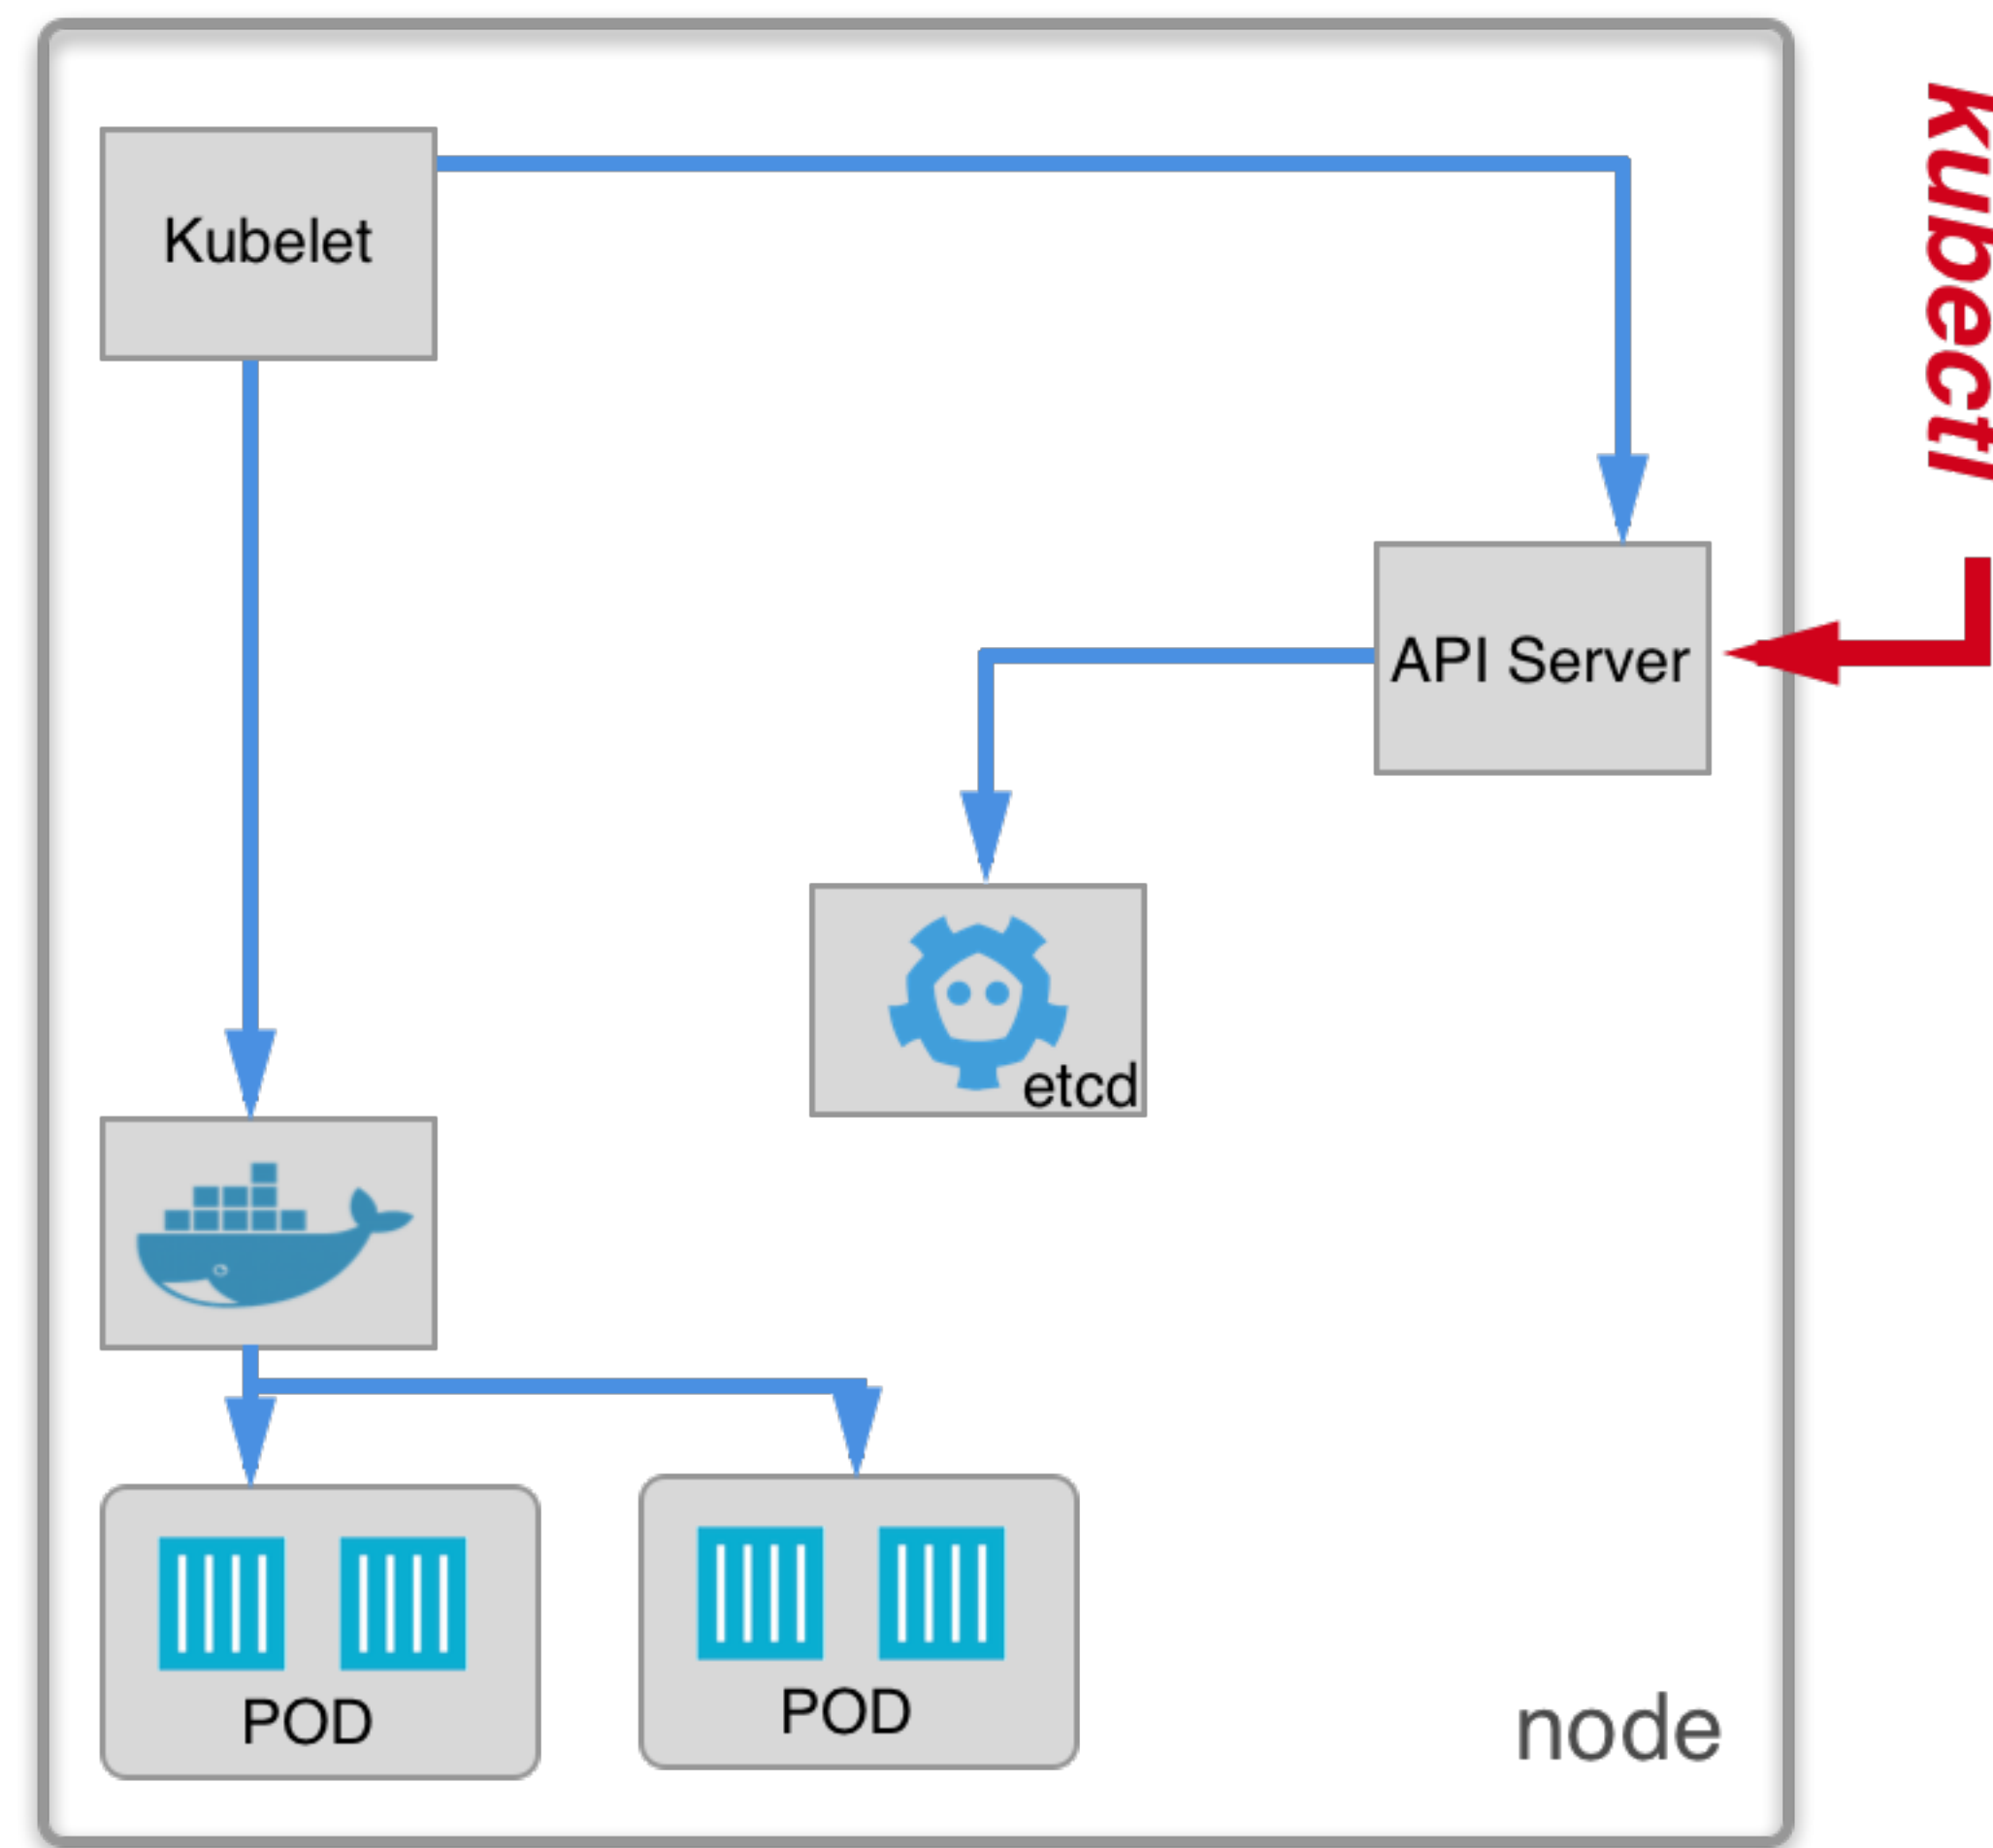
API Server



# ETCD + API SERVER

## API Server

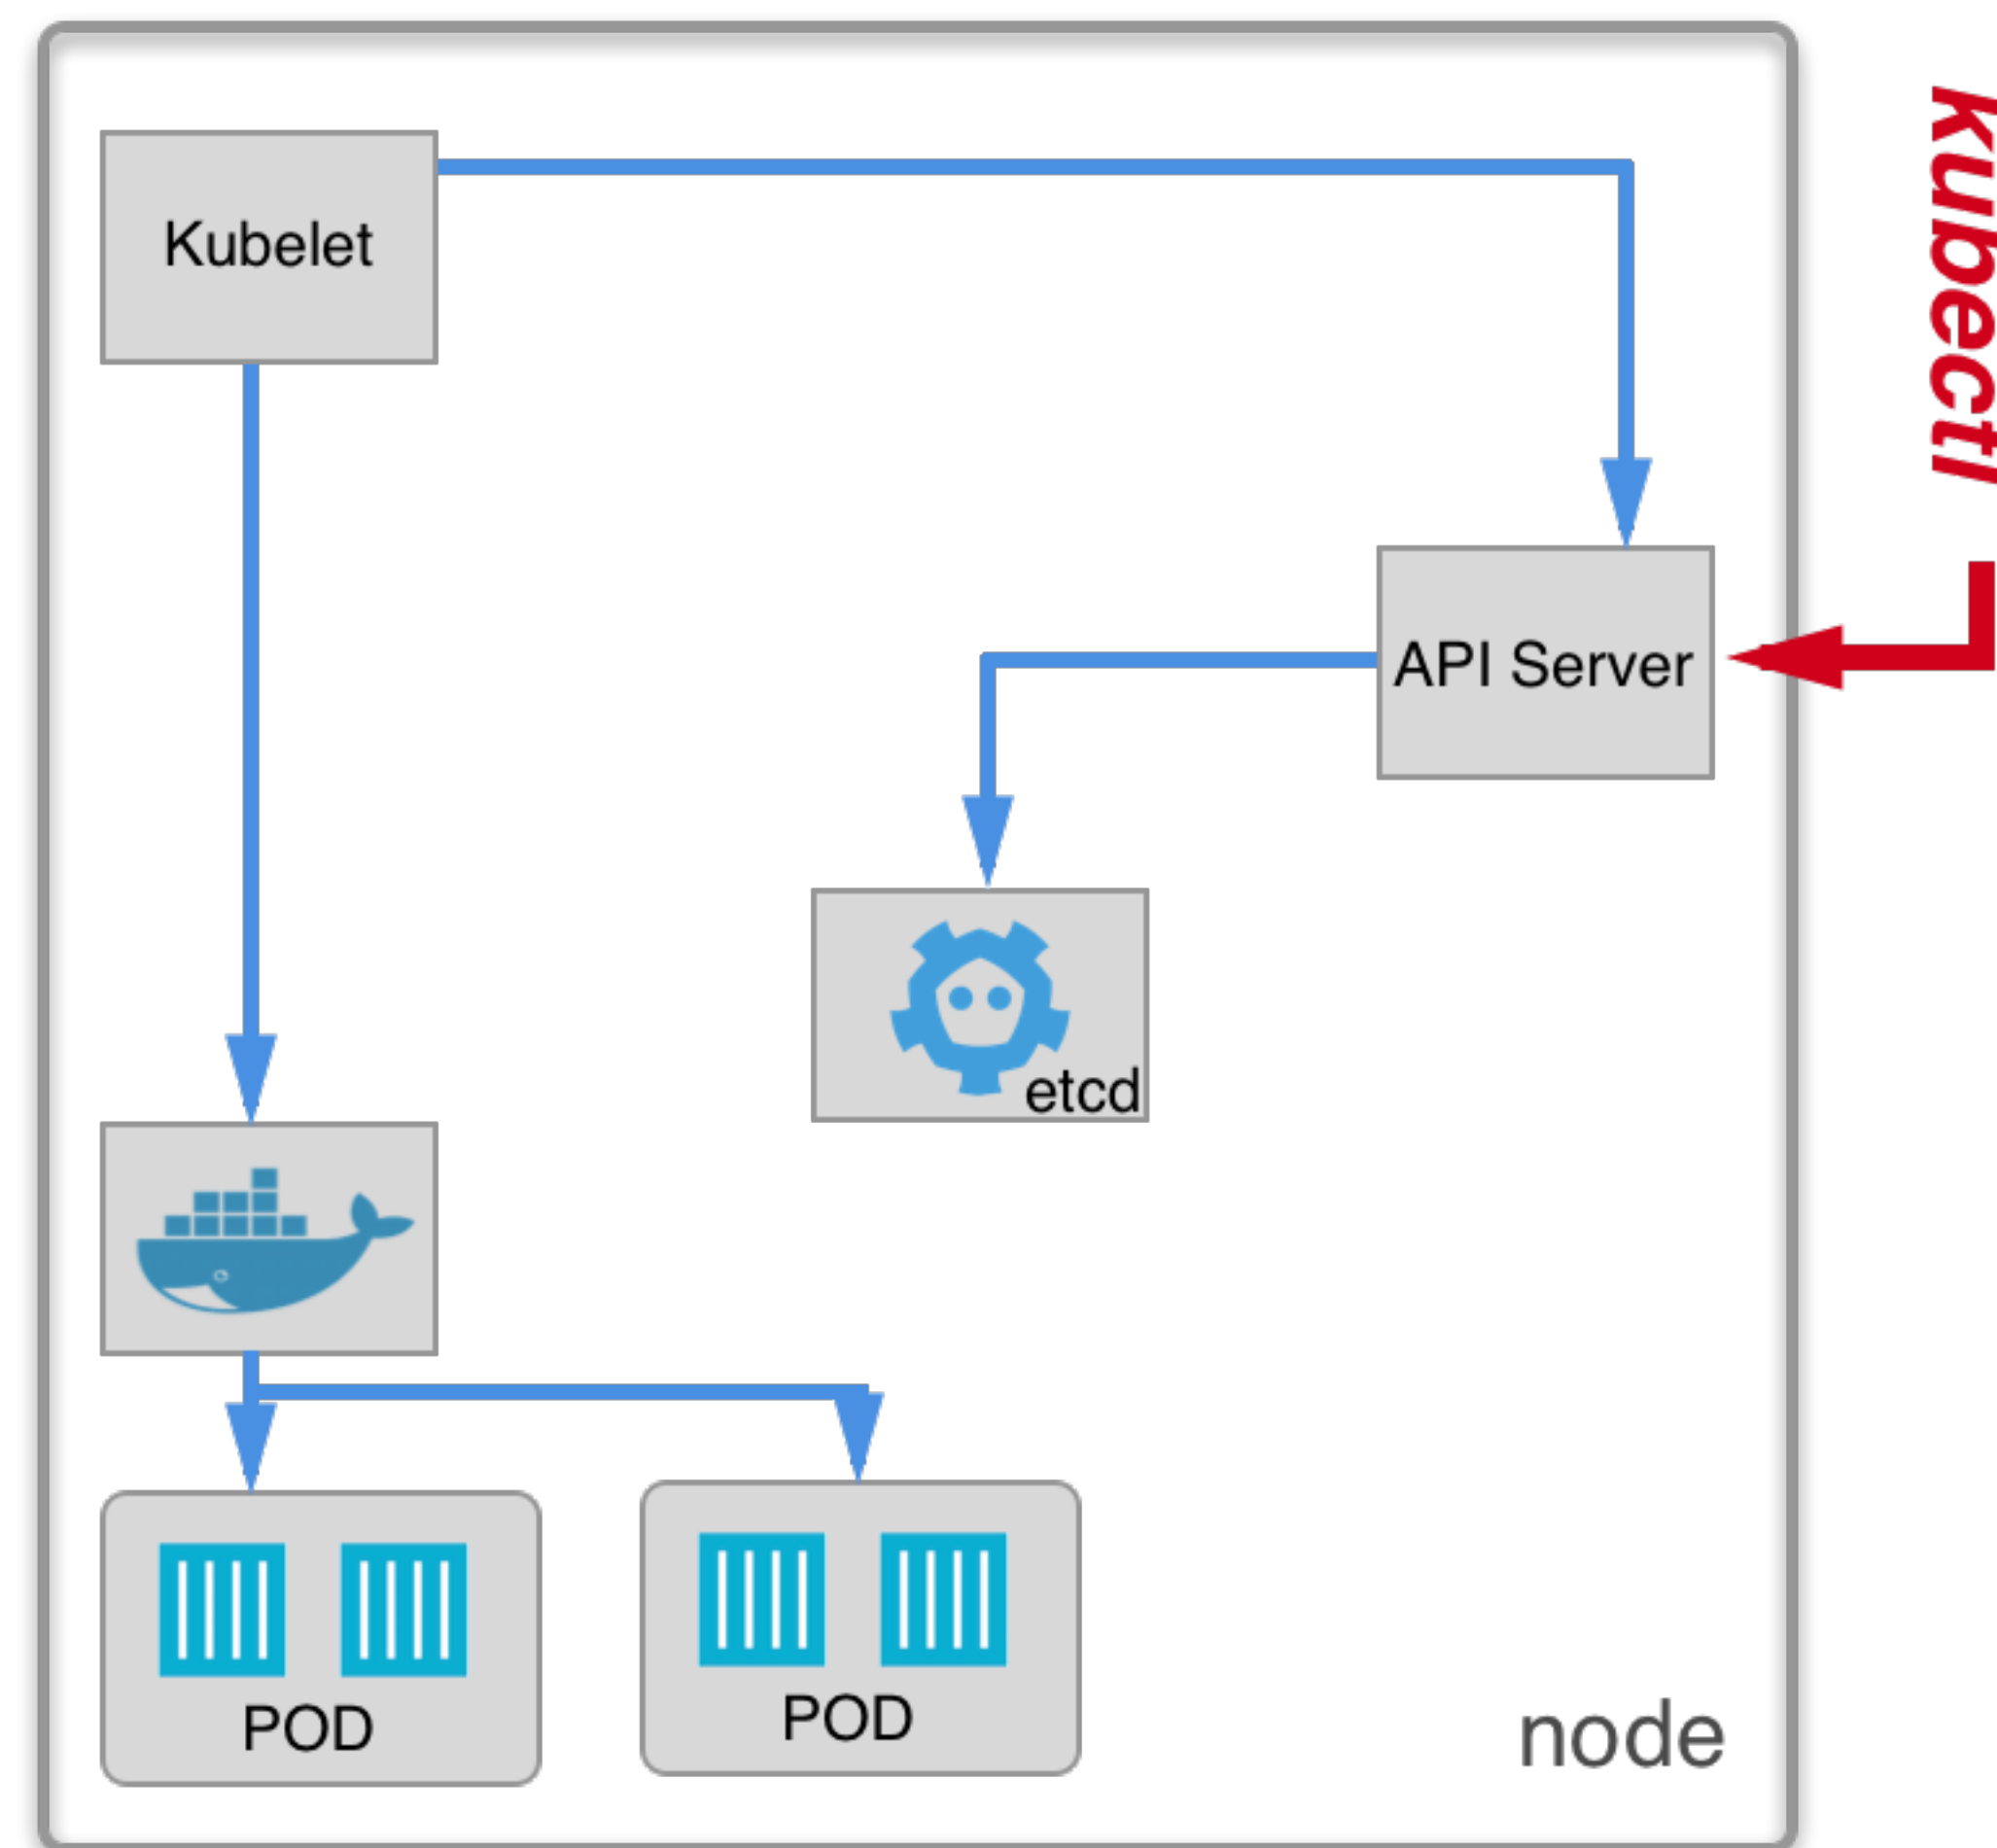
- "gateway" to etcd



# ETCD + API SERVER

## API Server

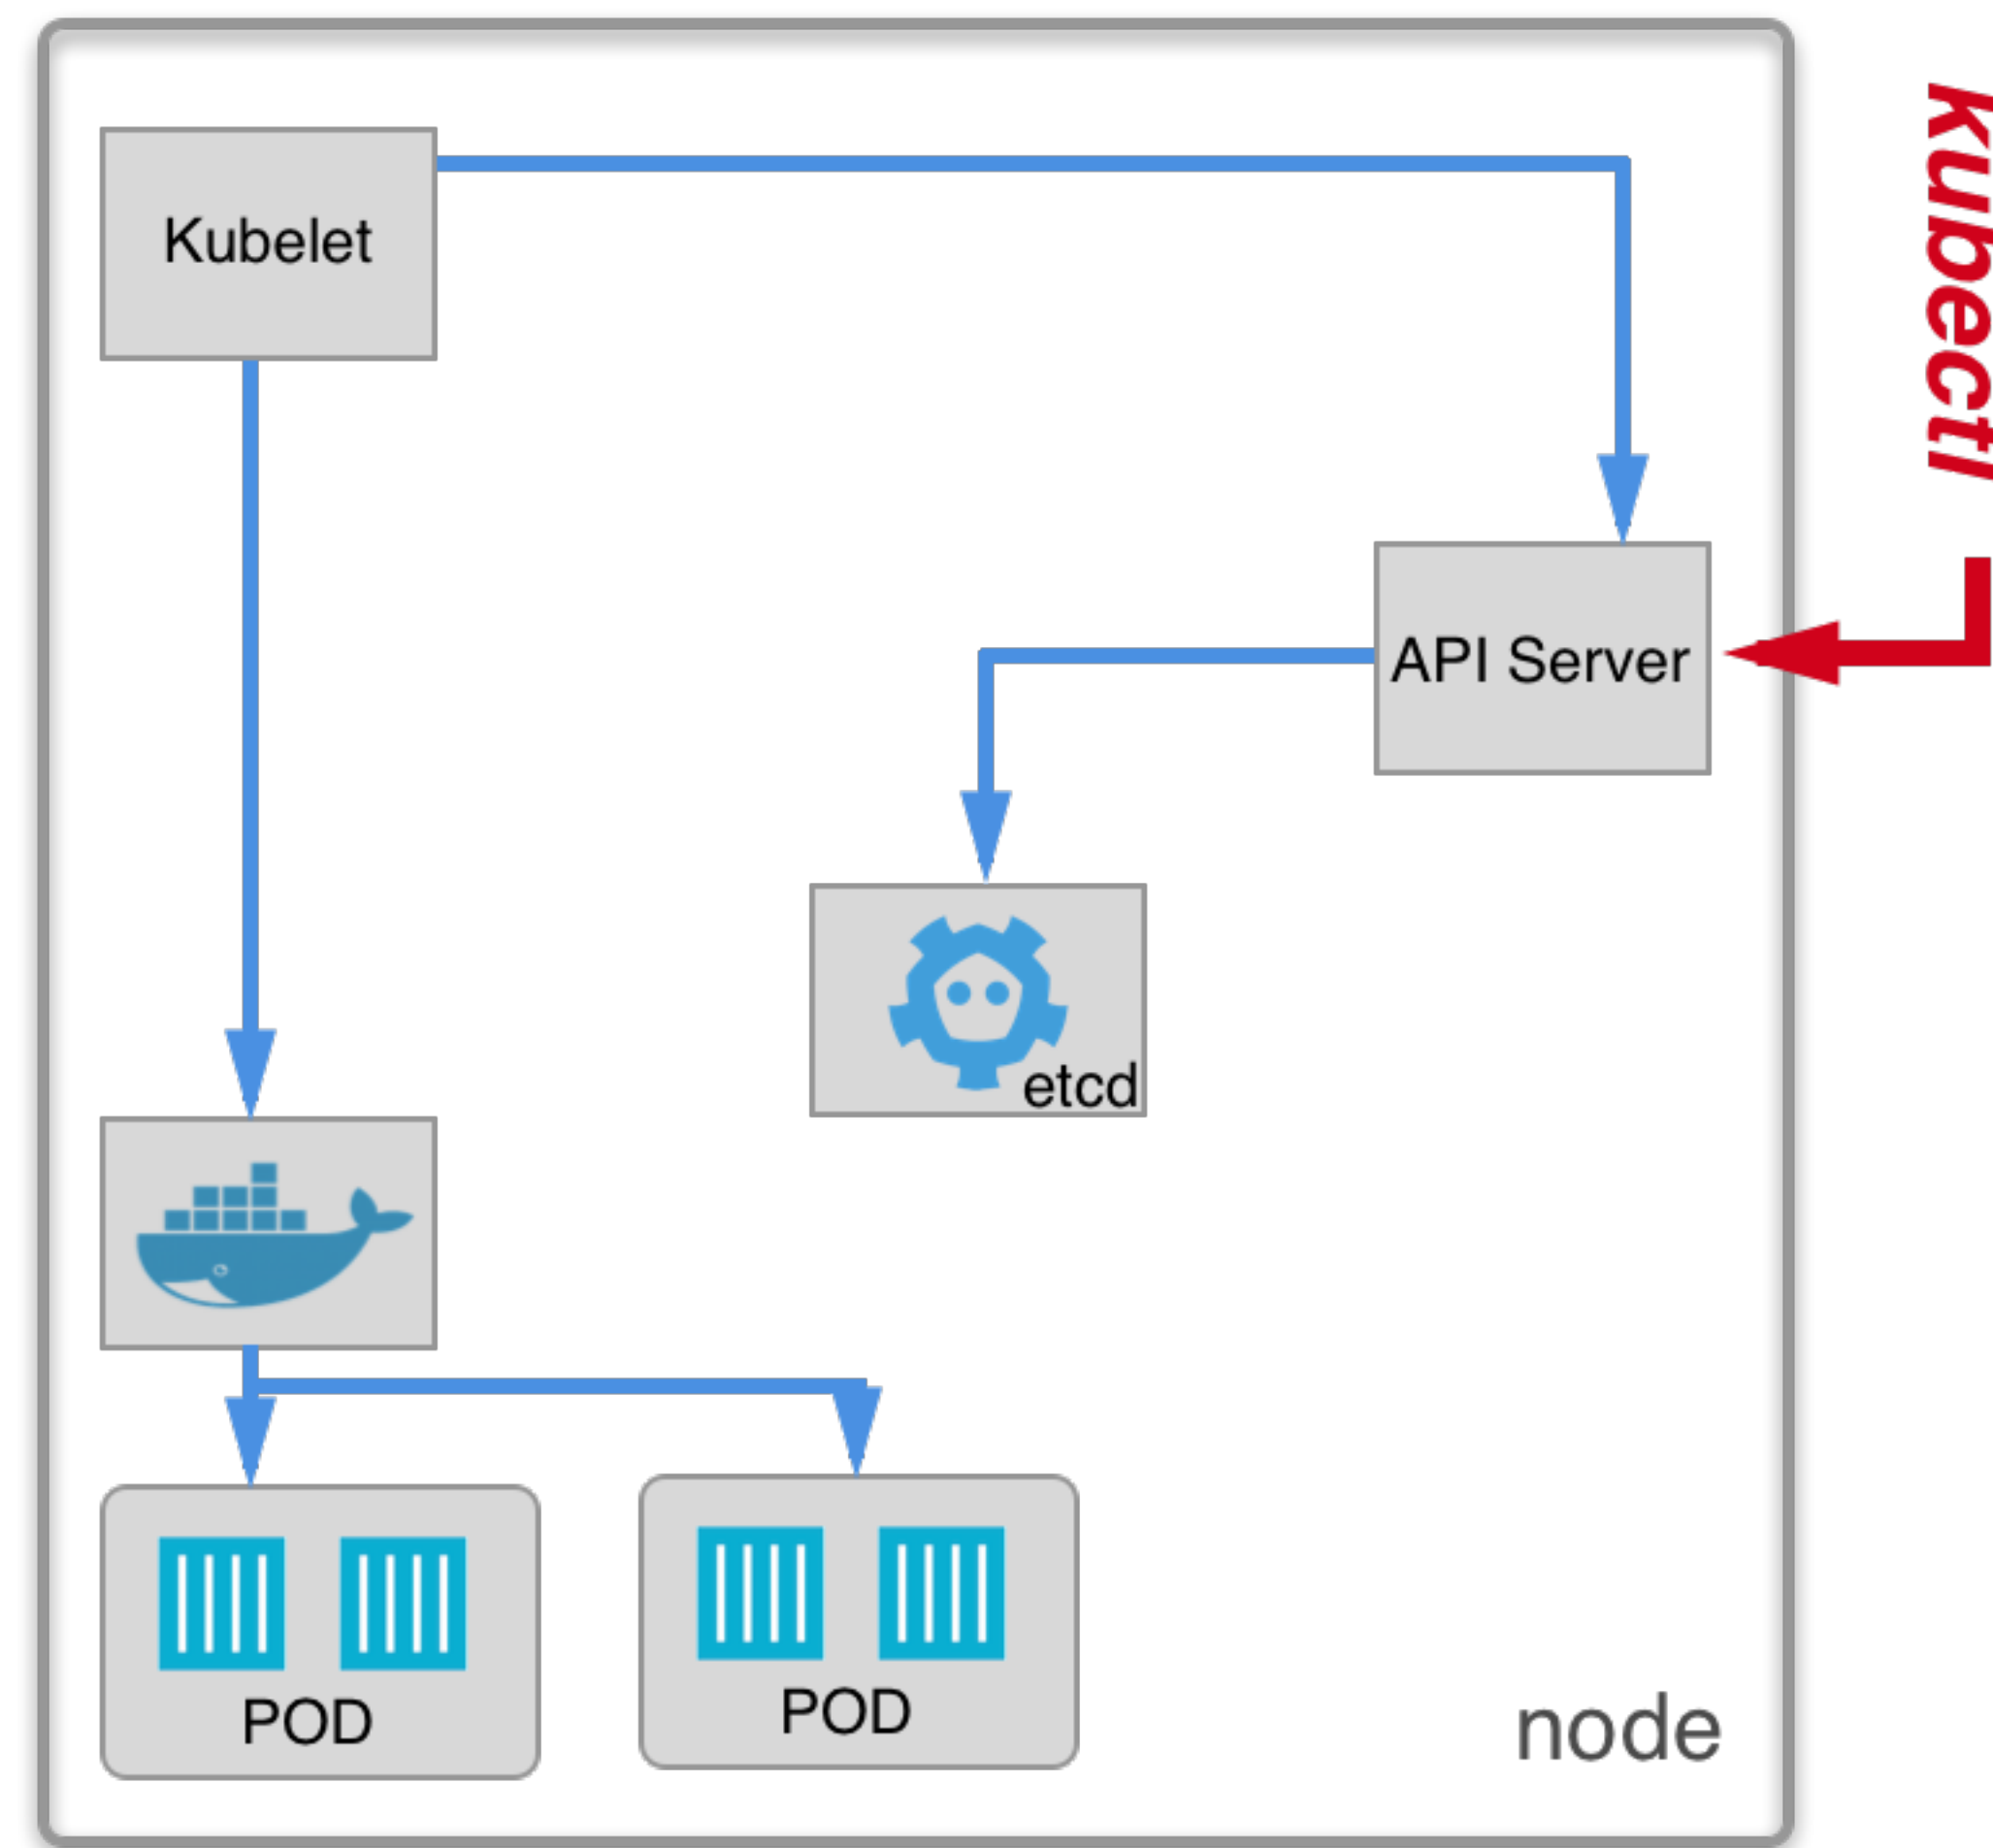
- “gateway” to etcd
- authenticates and authorizes requests



# ETCD + API SERVER

## API Server

- “gateway” to etcd
- authenticates and authorizes requests
- all k8s components talk to it

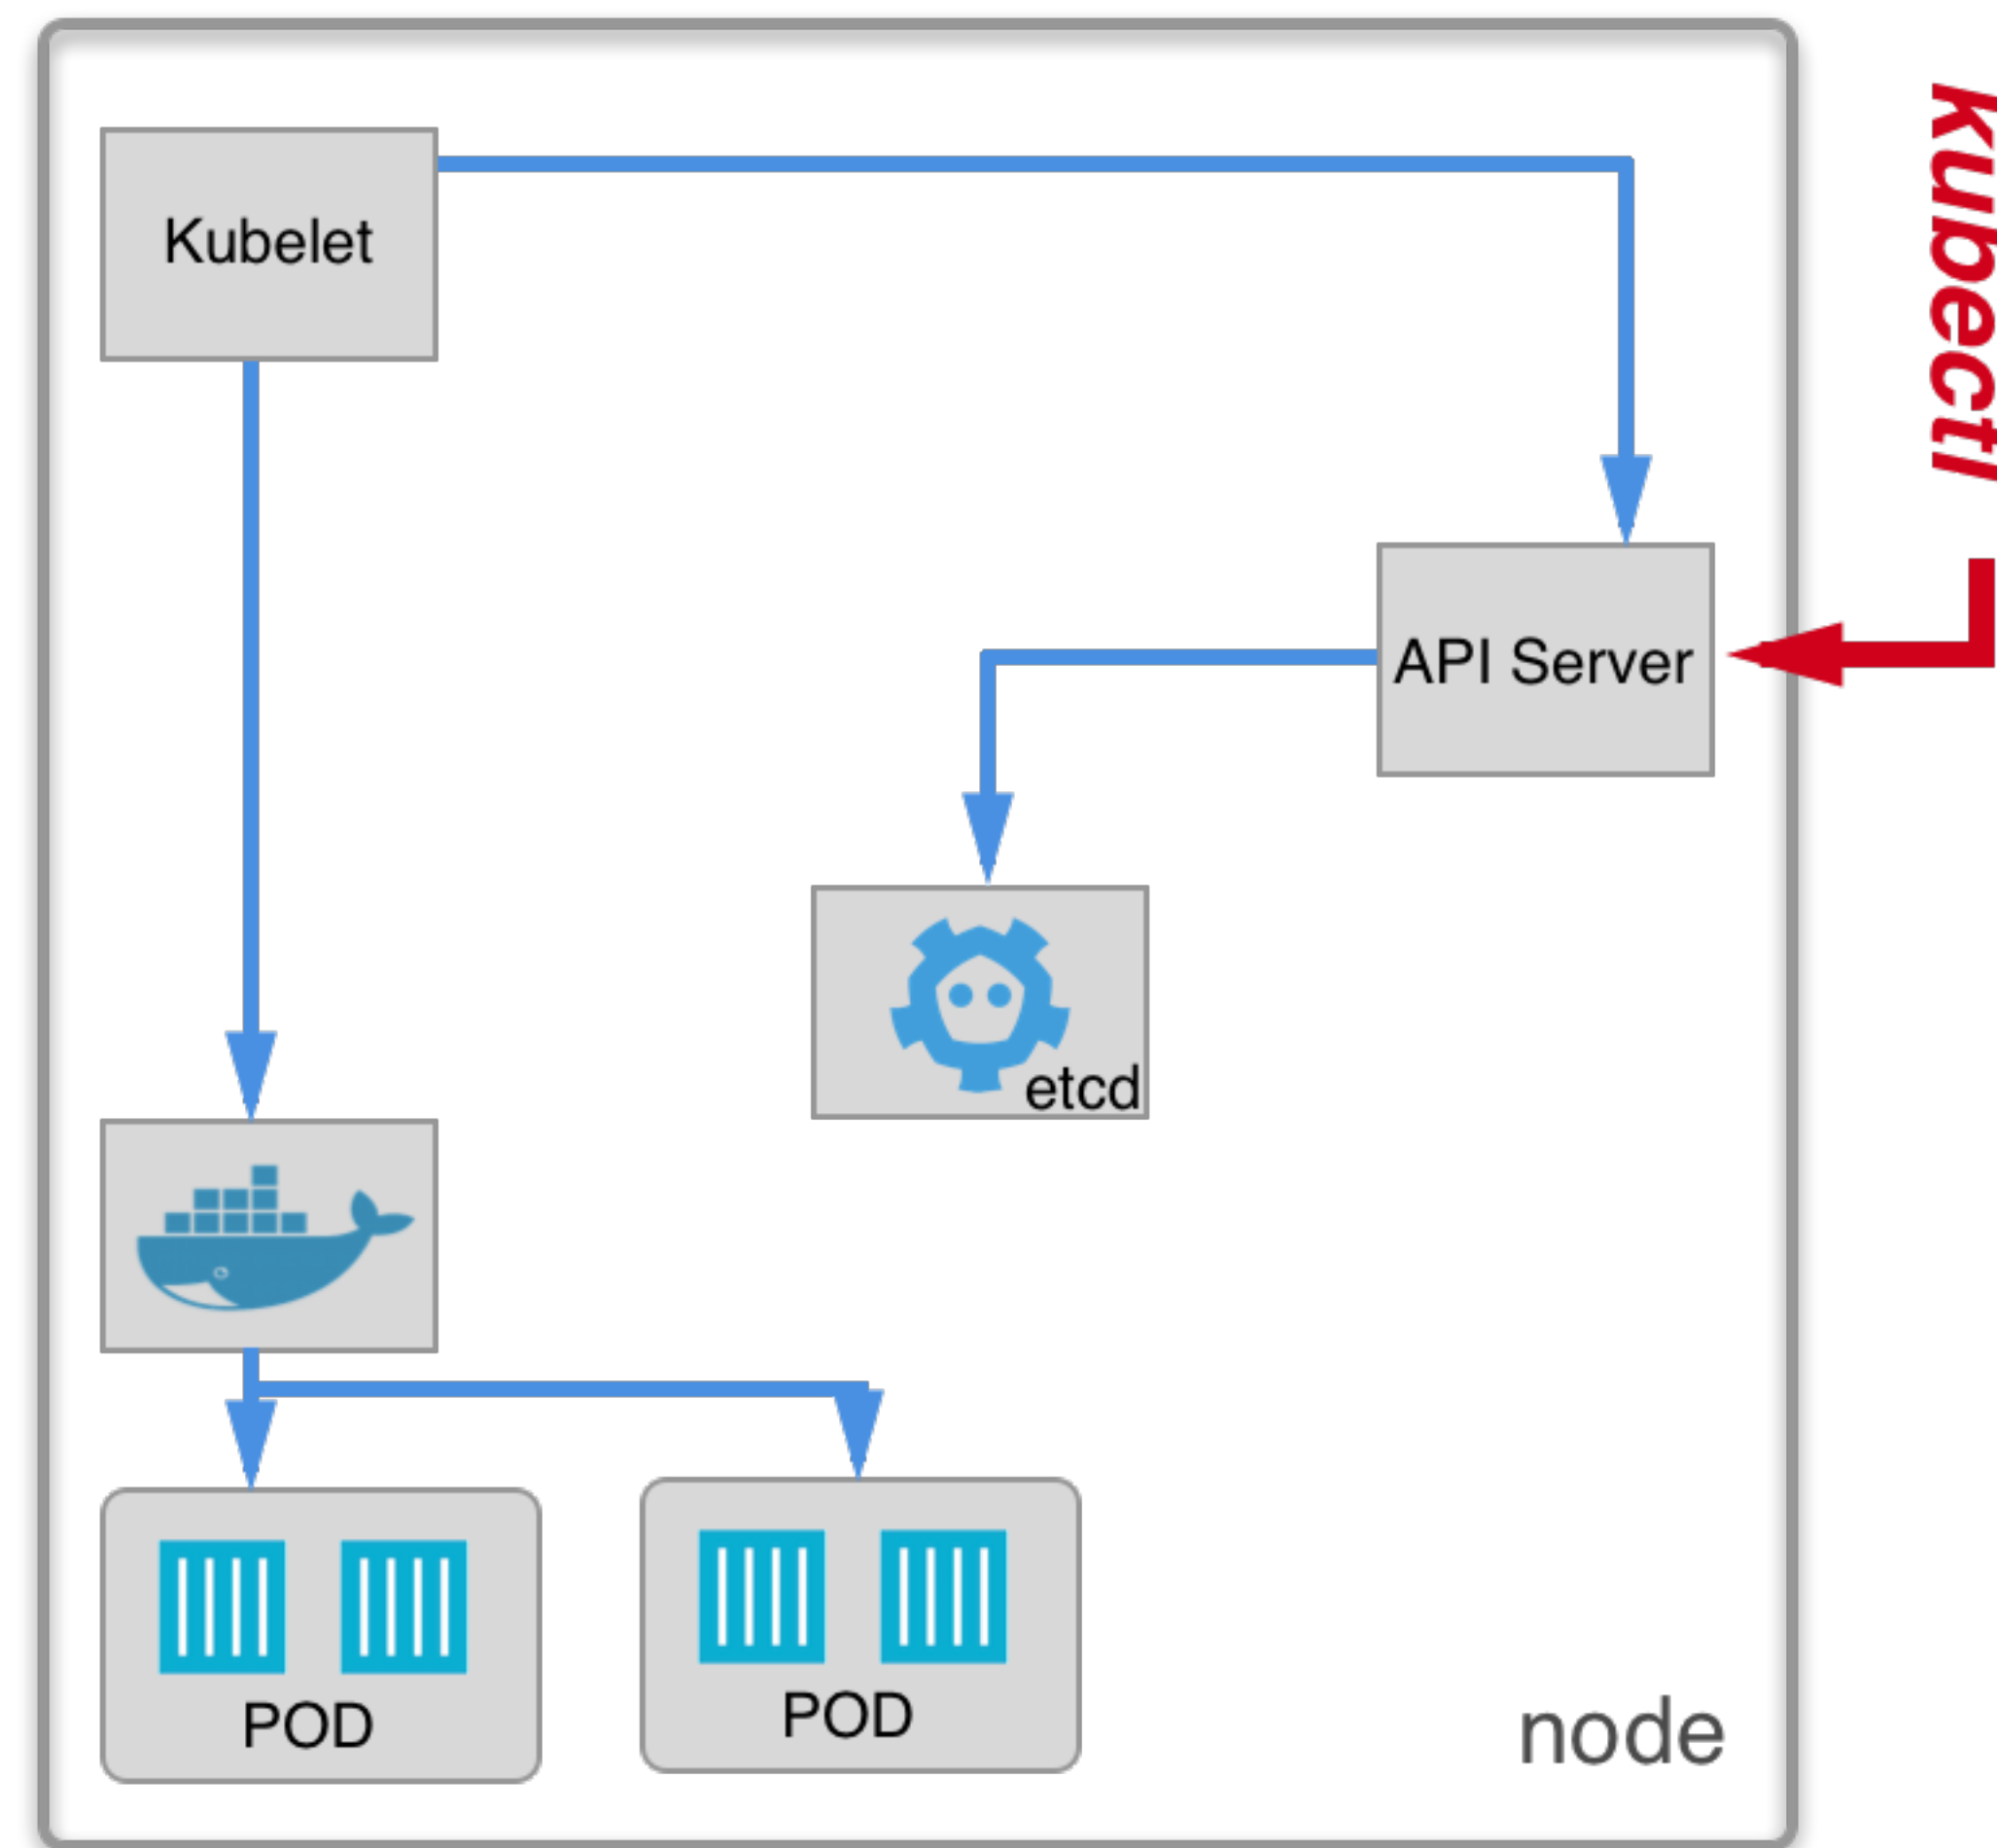




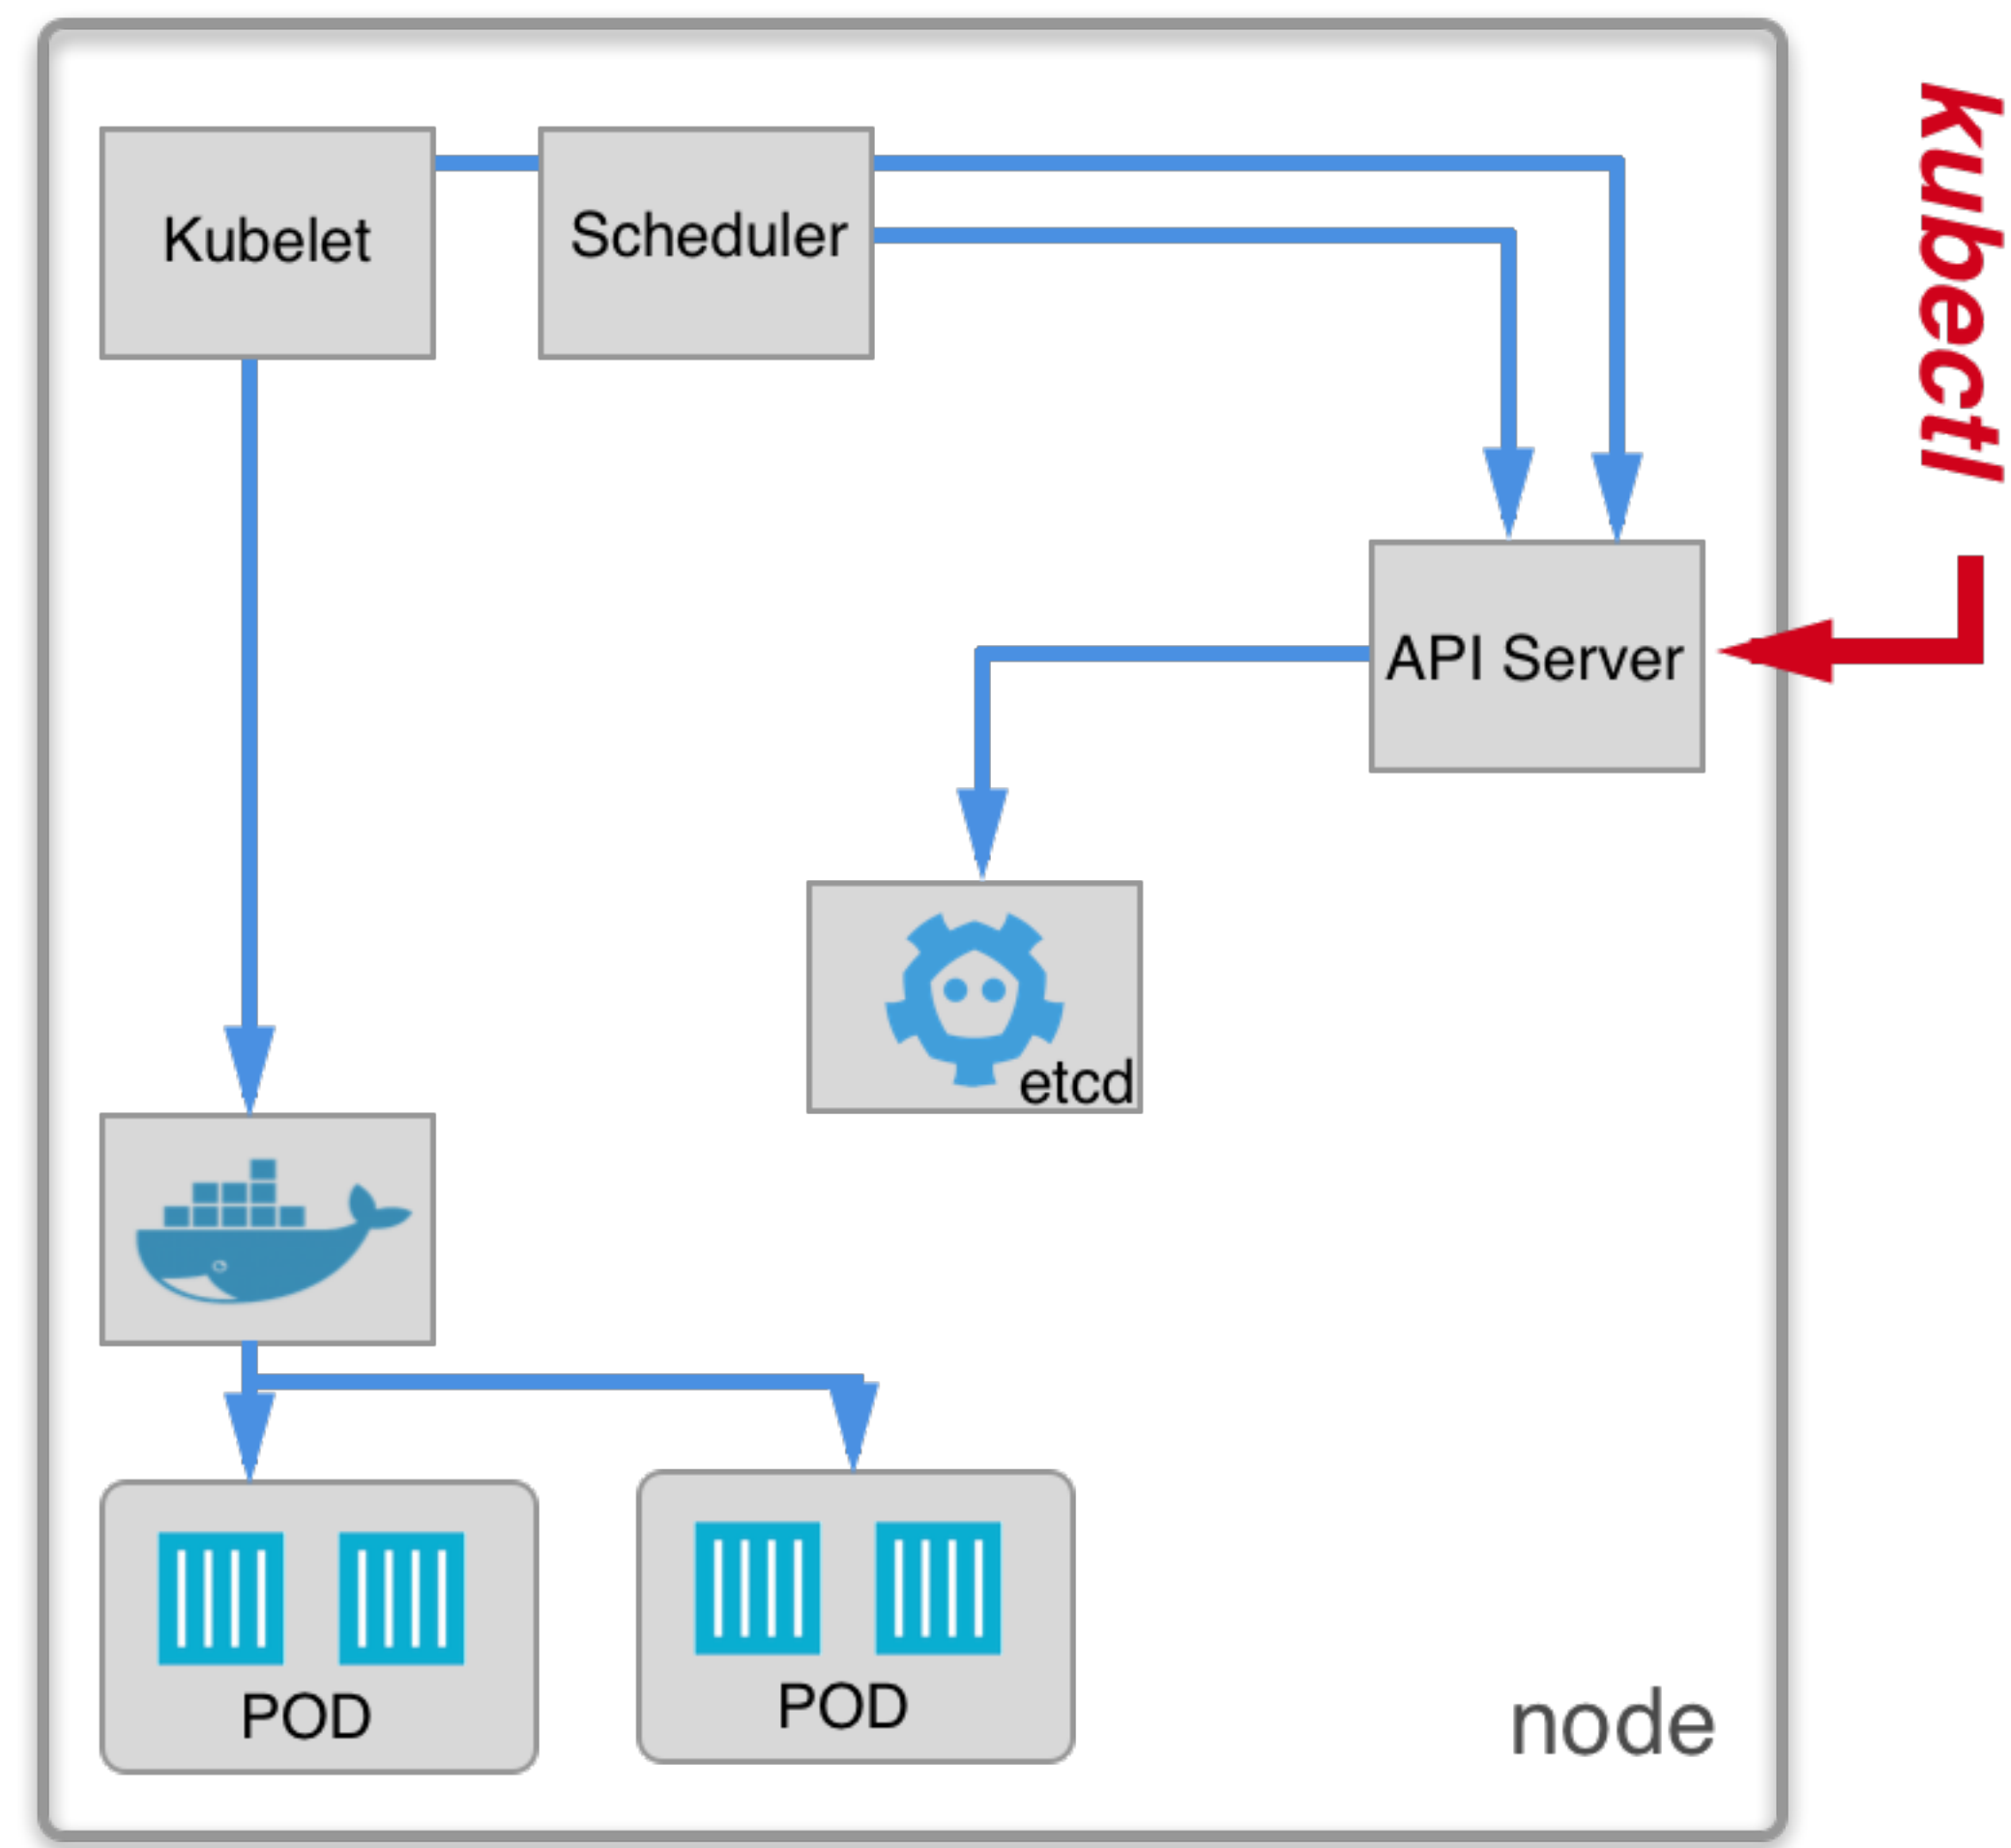
# ETCD + API SERVER

## API Server

- “gateway” to etcd
- authenticates and authorizes requests
- all k8s components talk to it
- kubectl talks to it

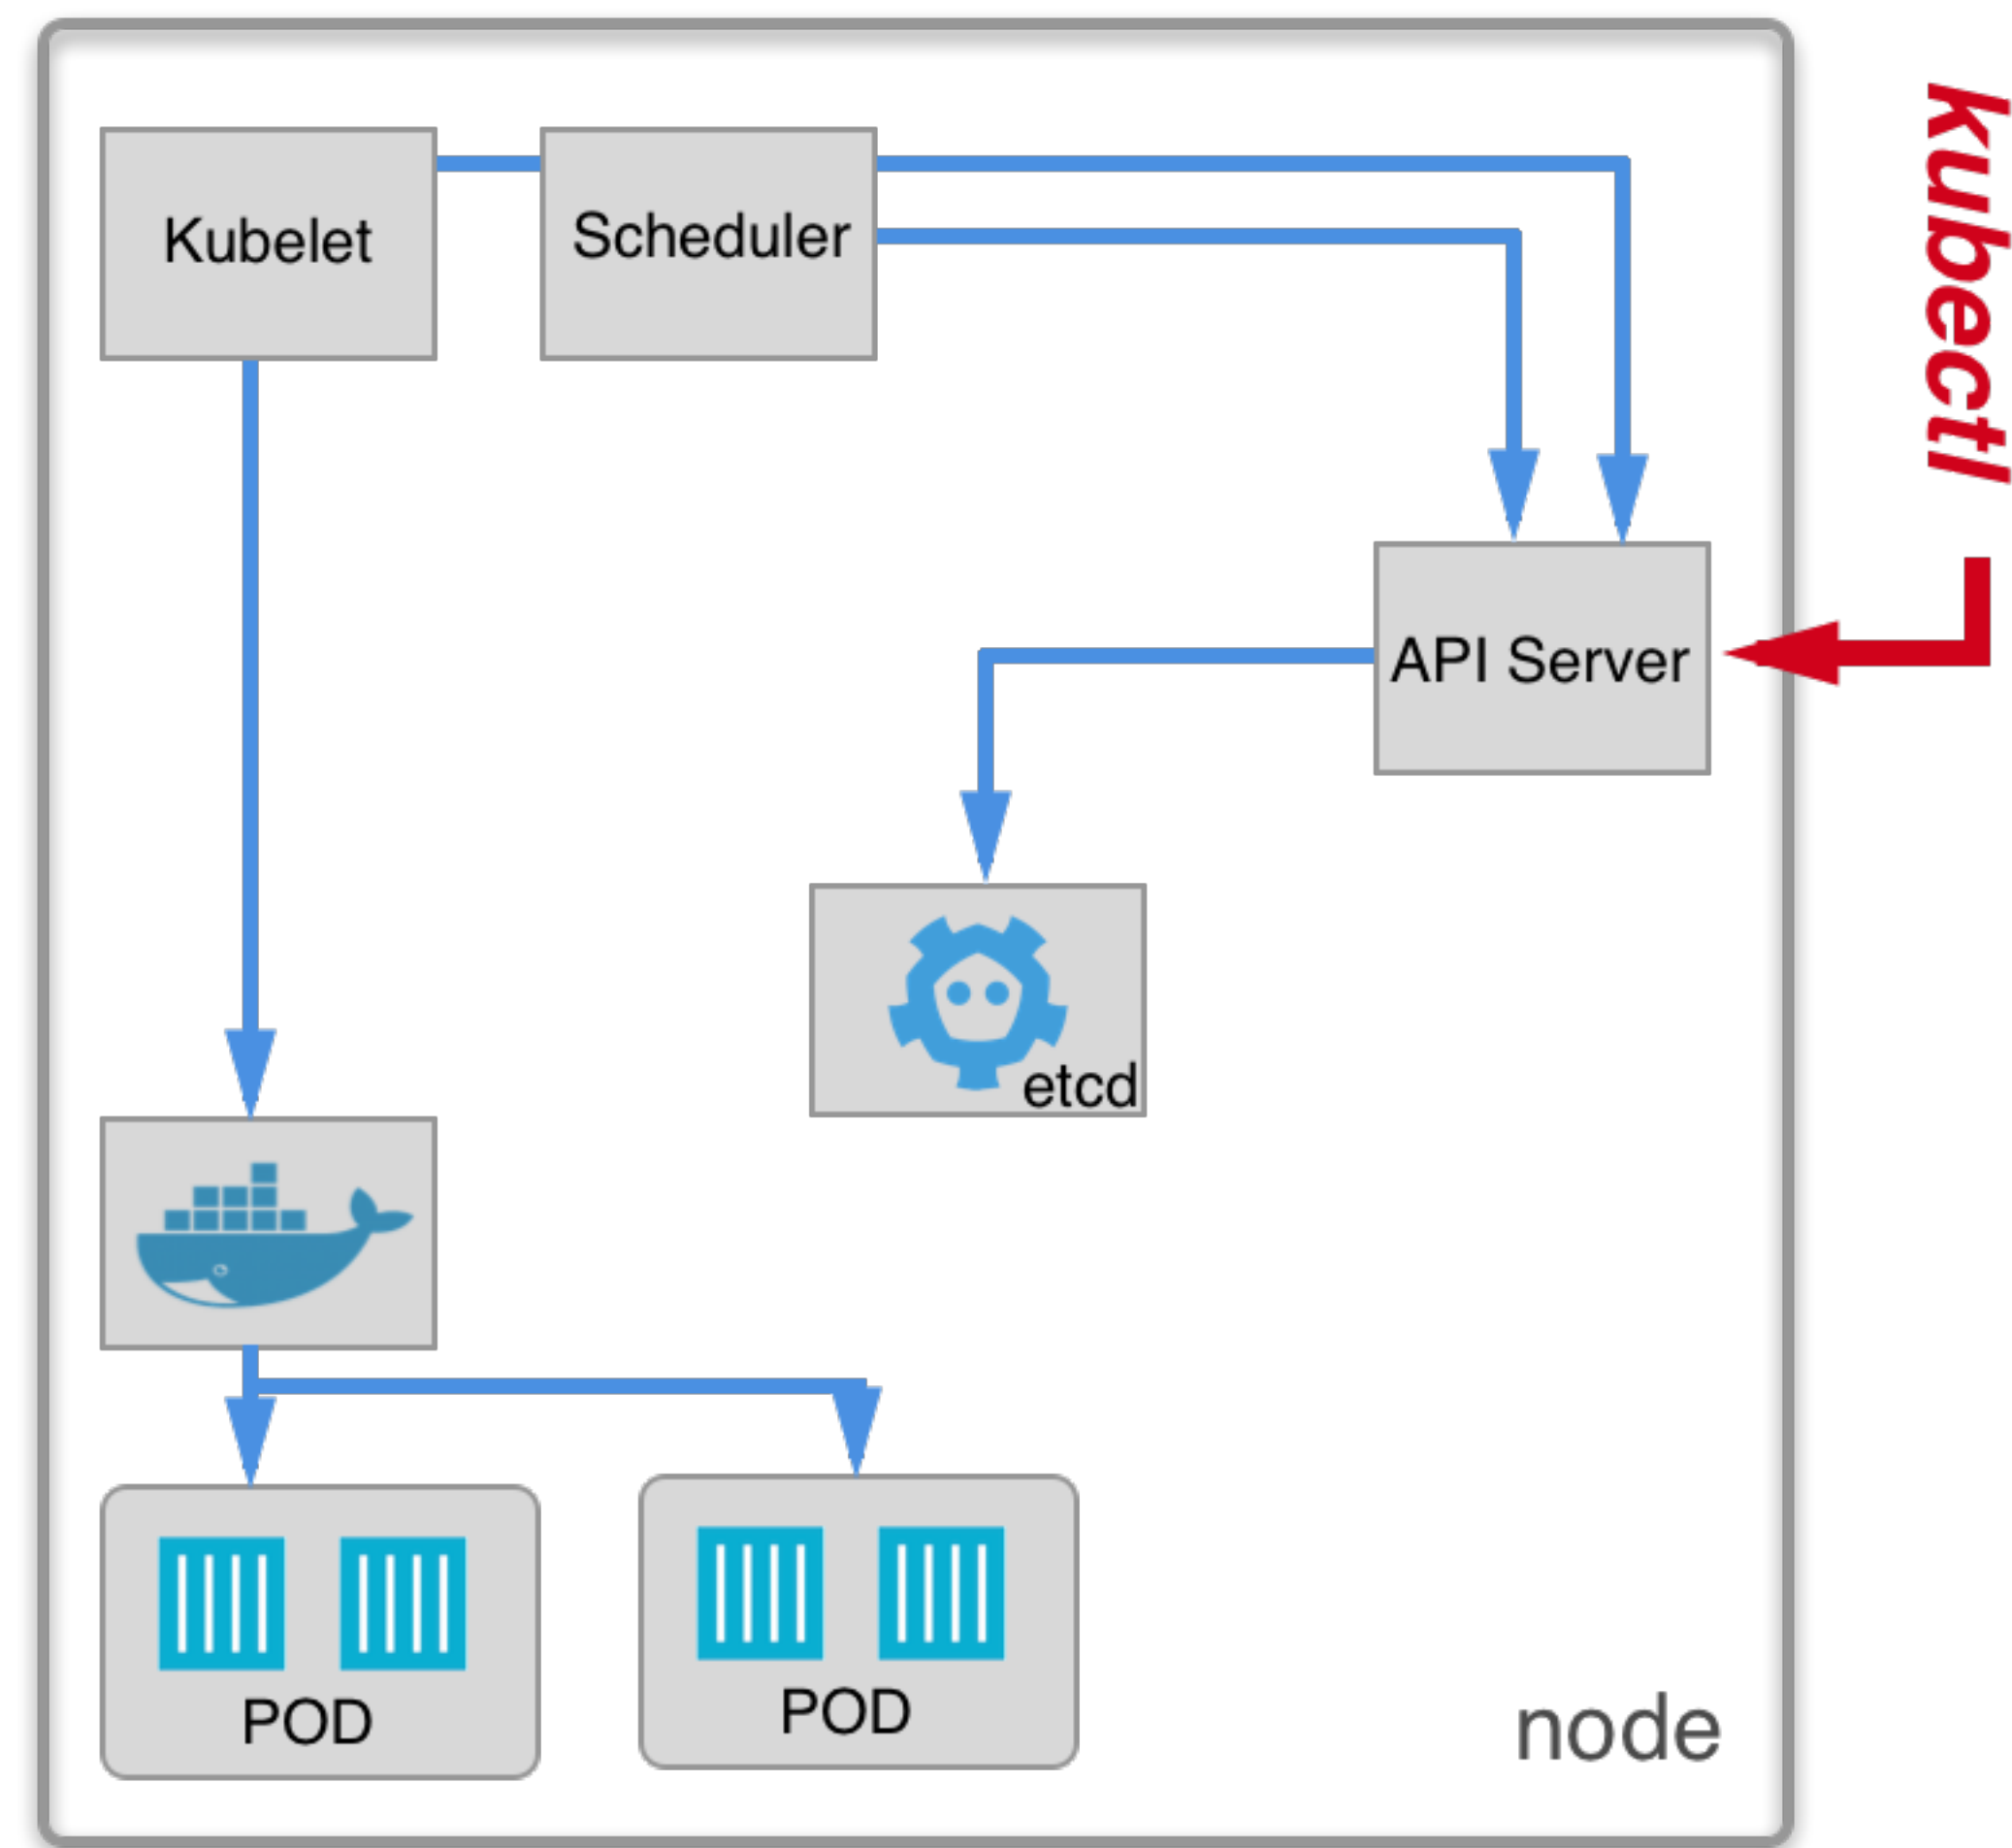


# SCHEDULER



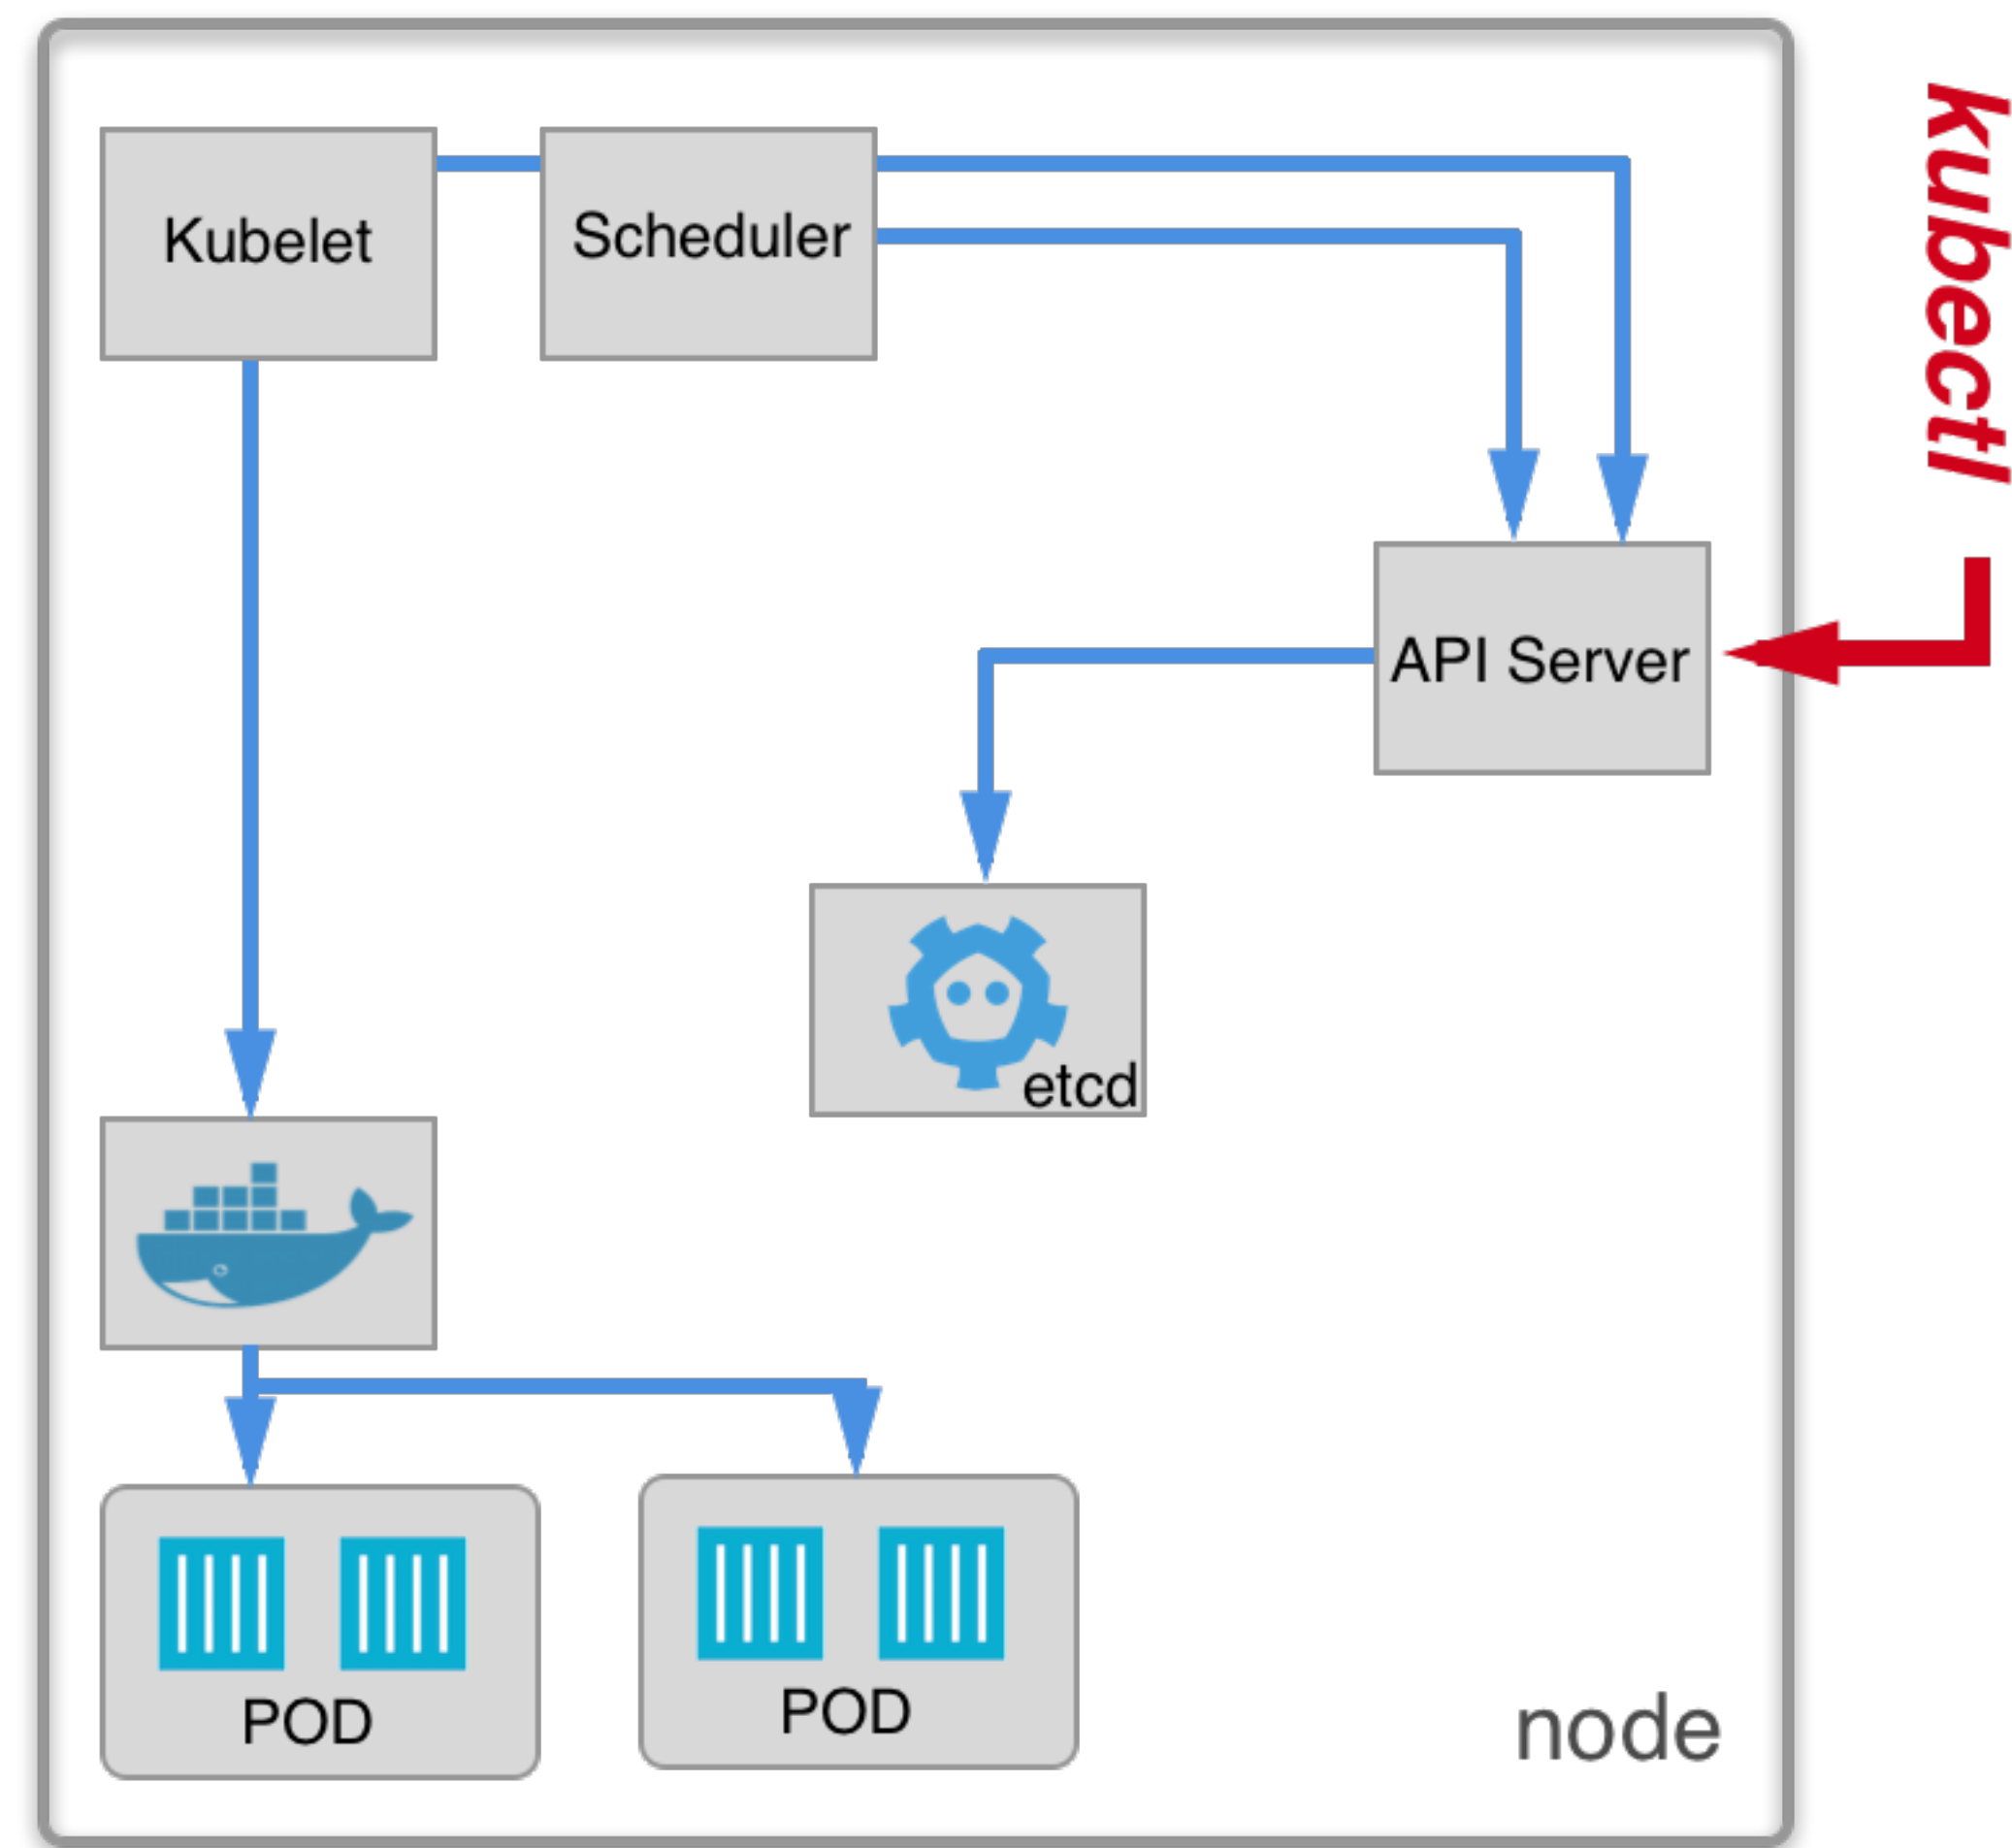
# SCHEDULER

- responsible for scheduling pods on nodes



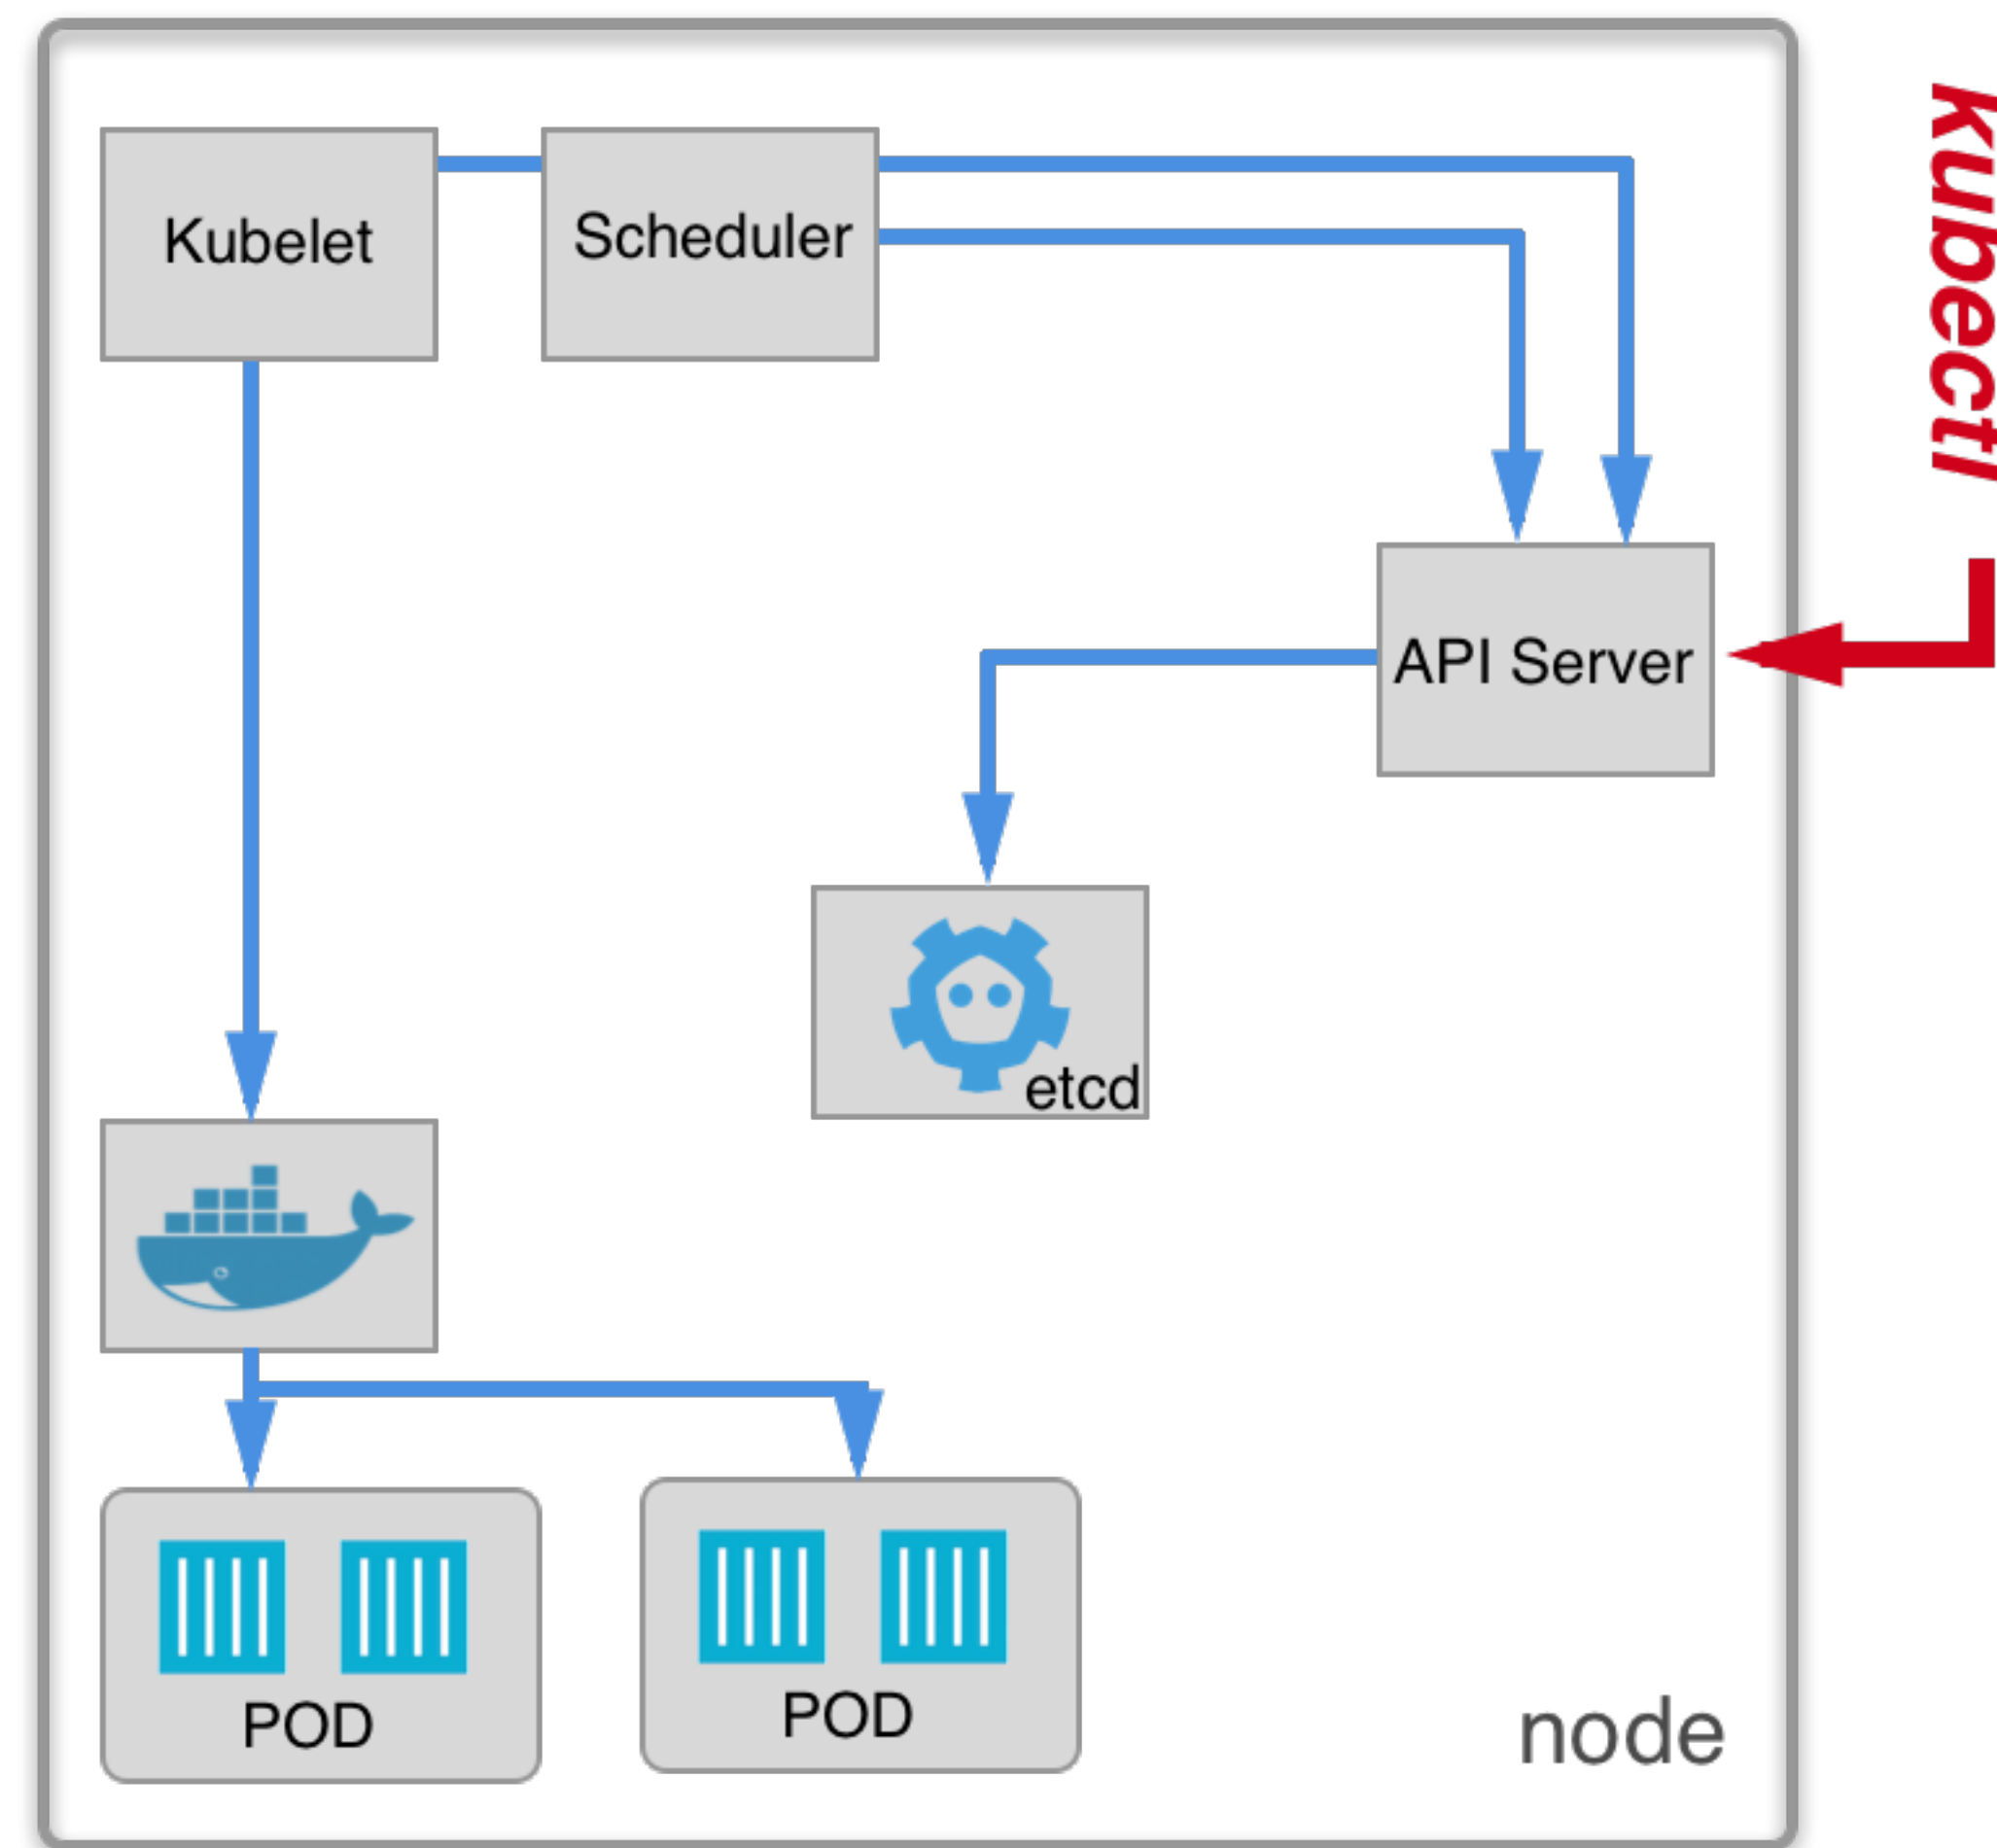
# SCHEDULER

- responsible for scheduling pods on nodes
- connects to API Server

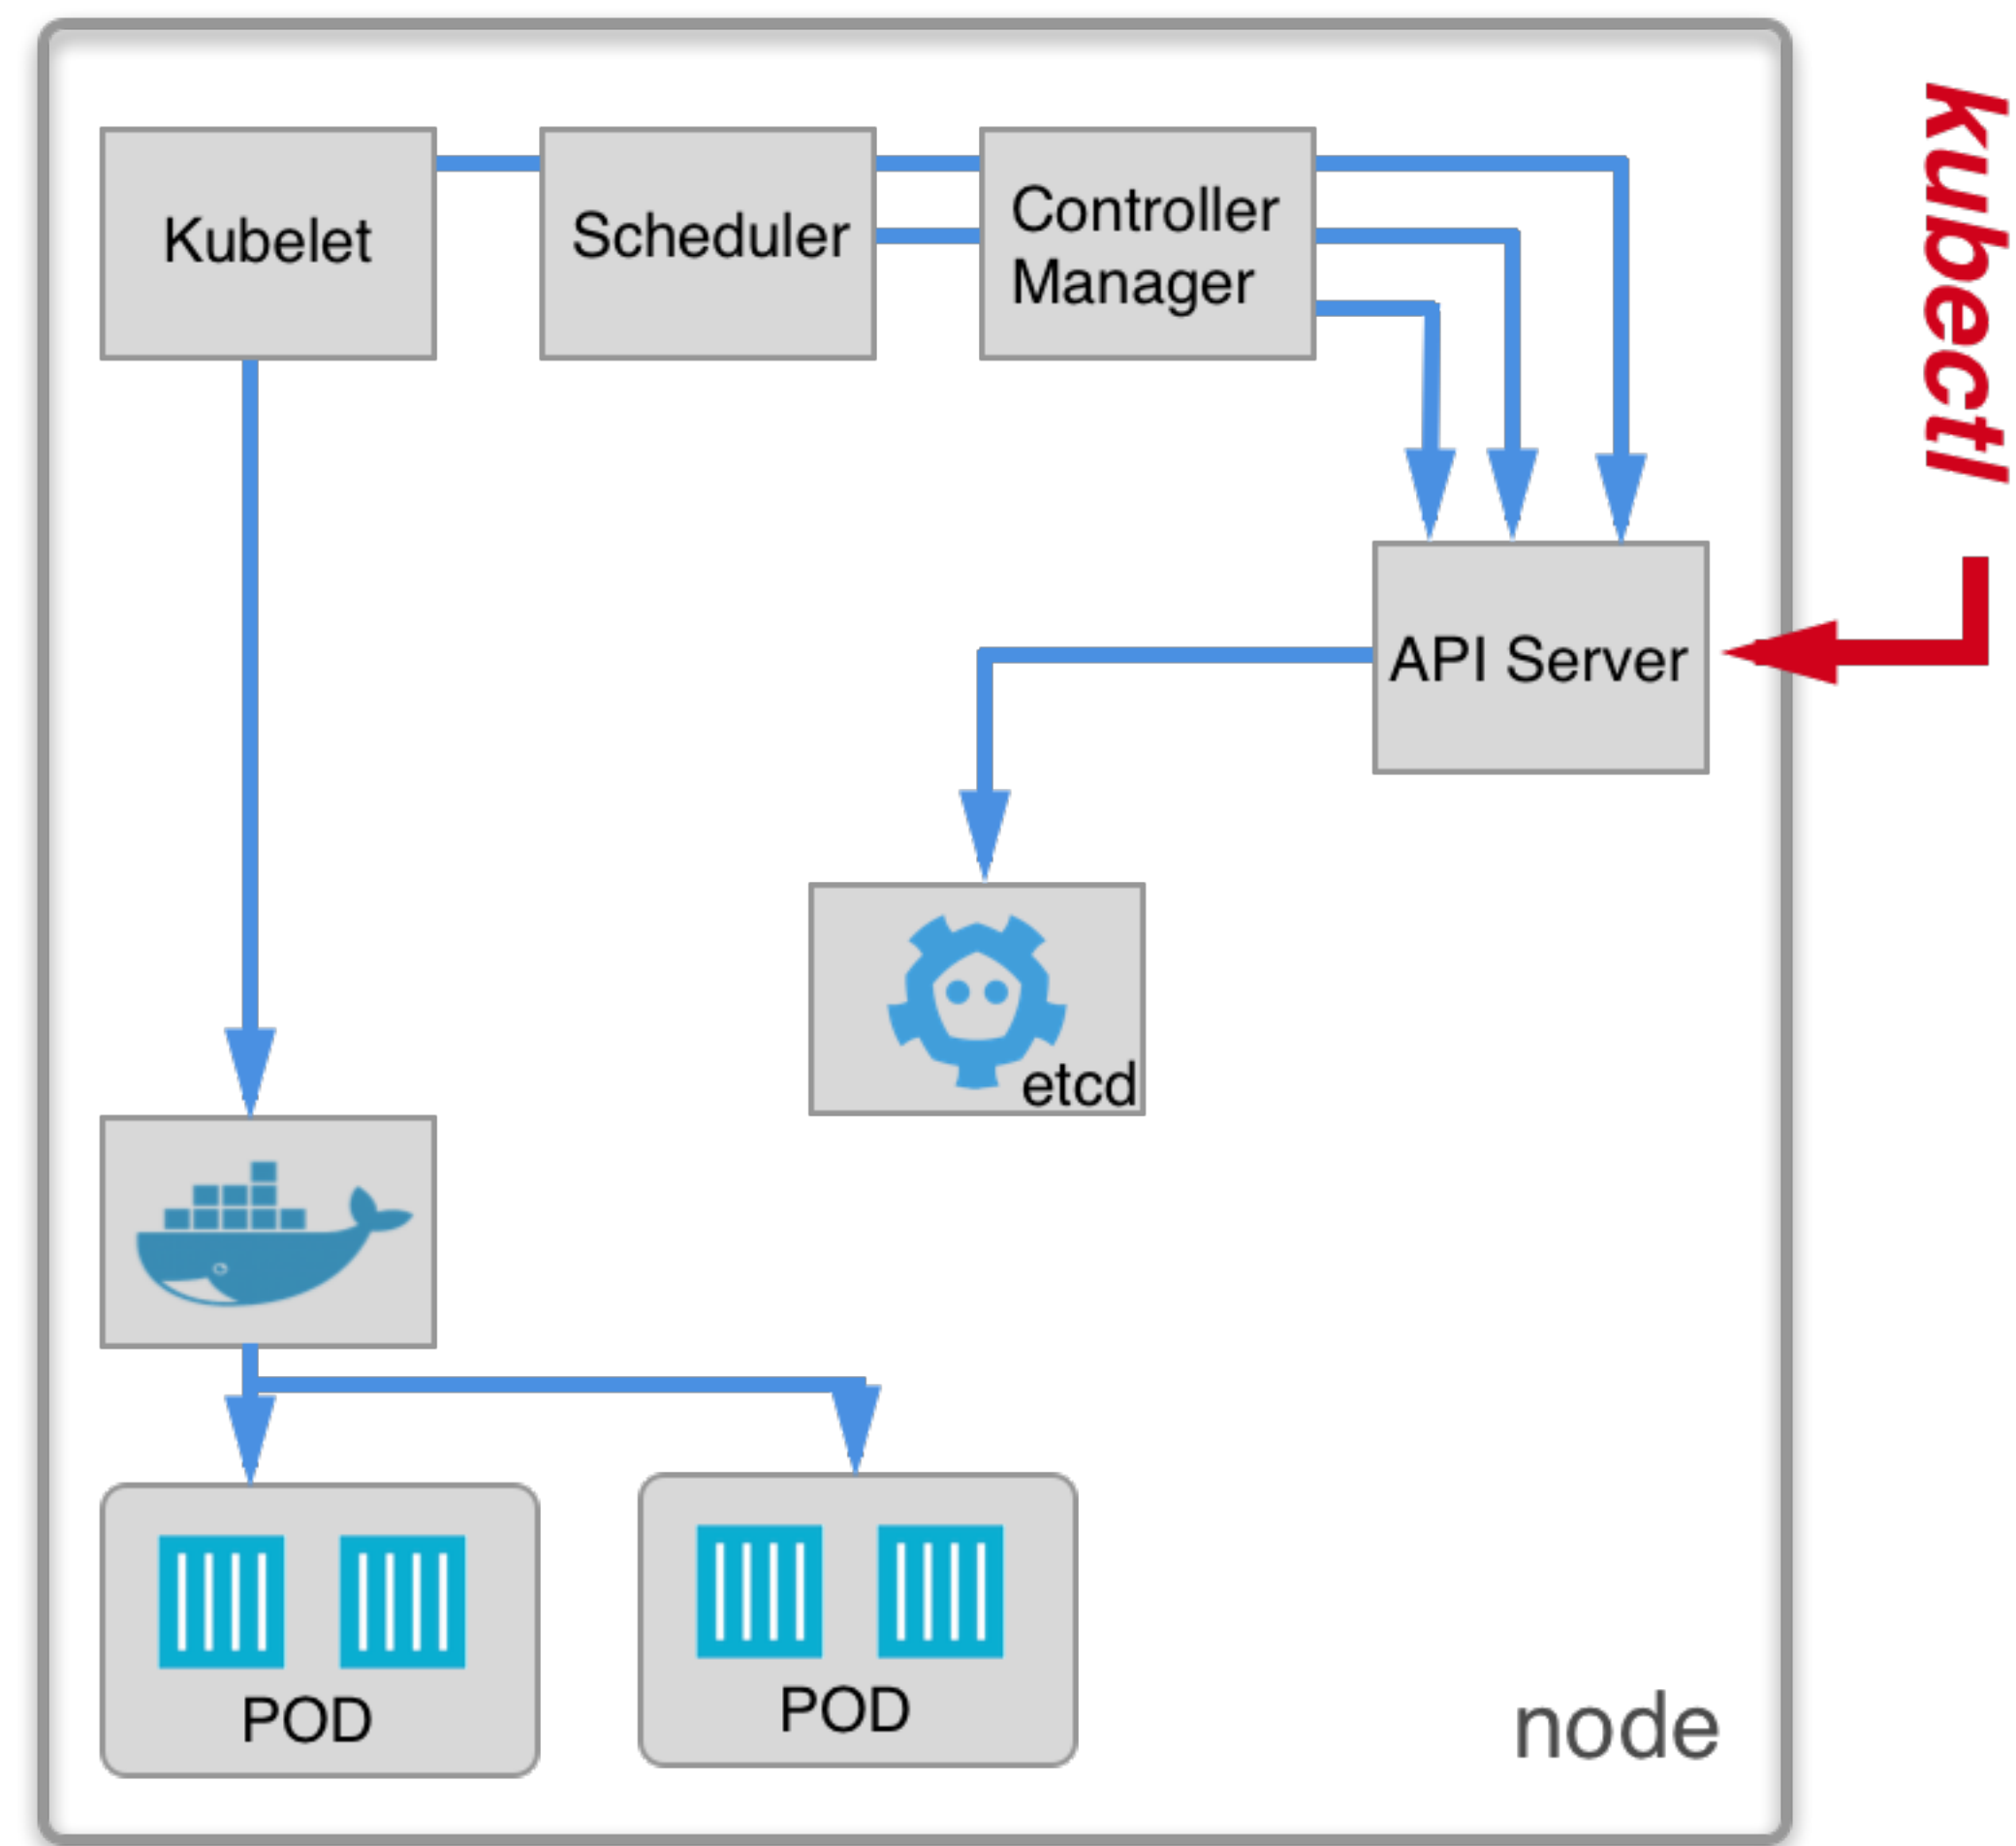


# SCHEDULER

- responsible for scheduling pods on nodes
- connects to API Server
- watches for pods that aren't bound to a node and assigns one

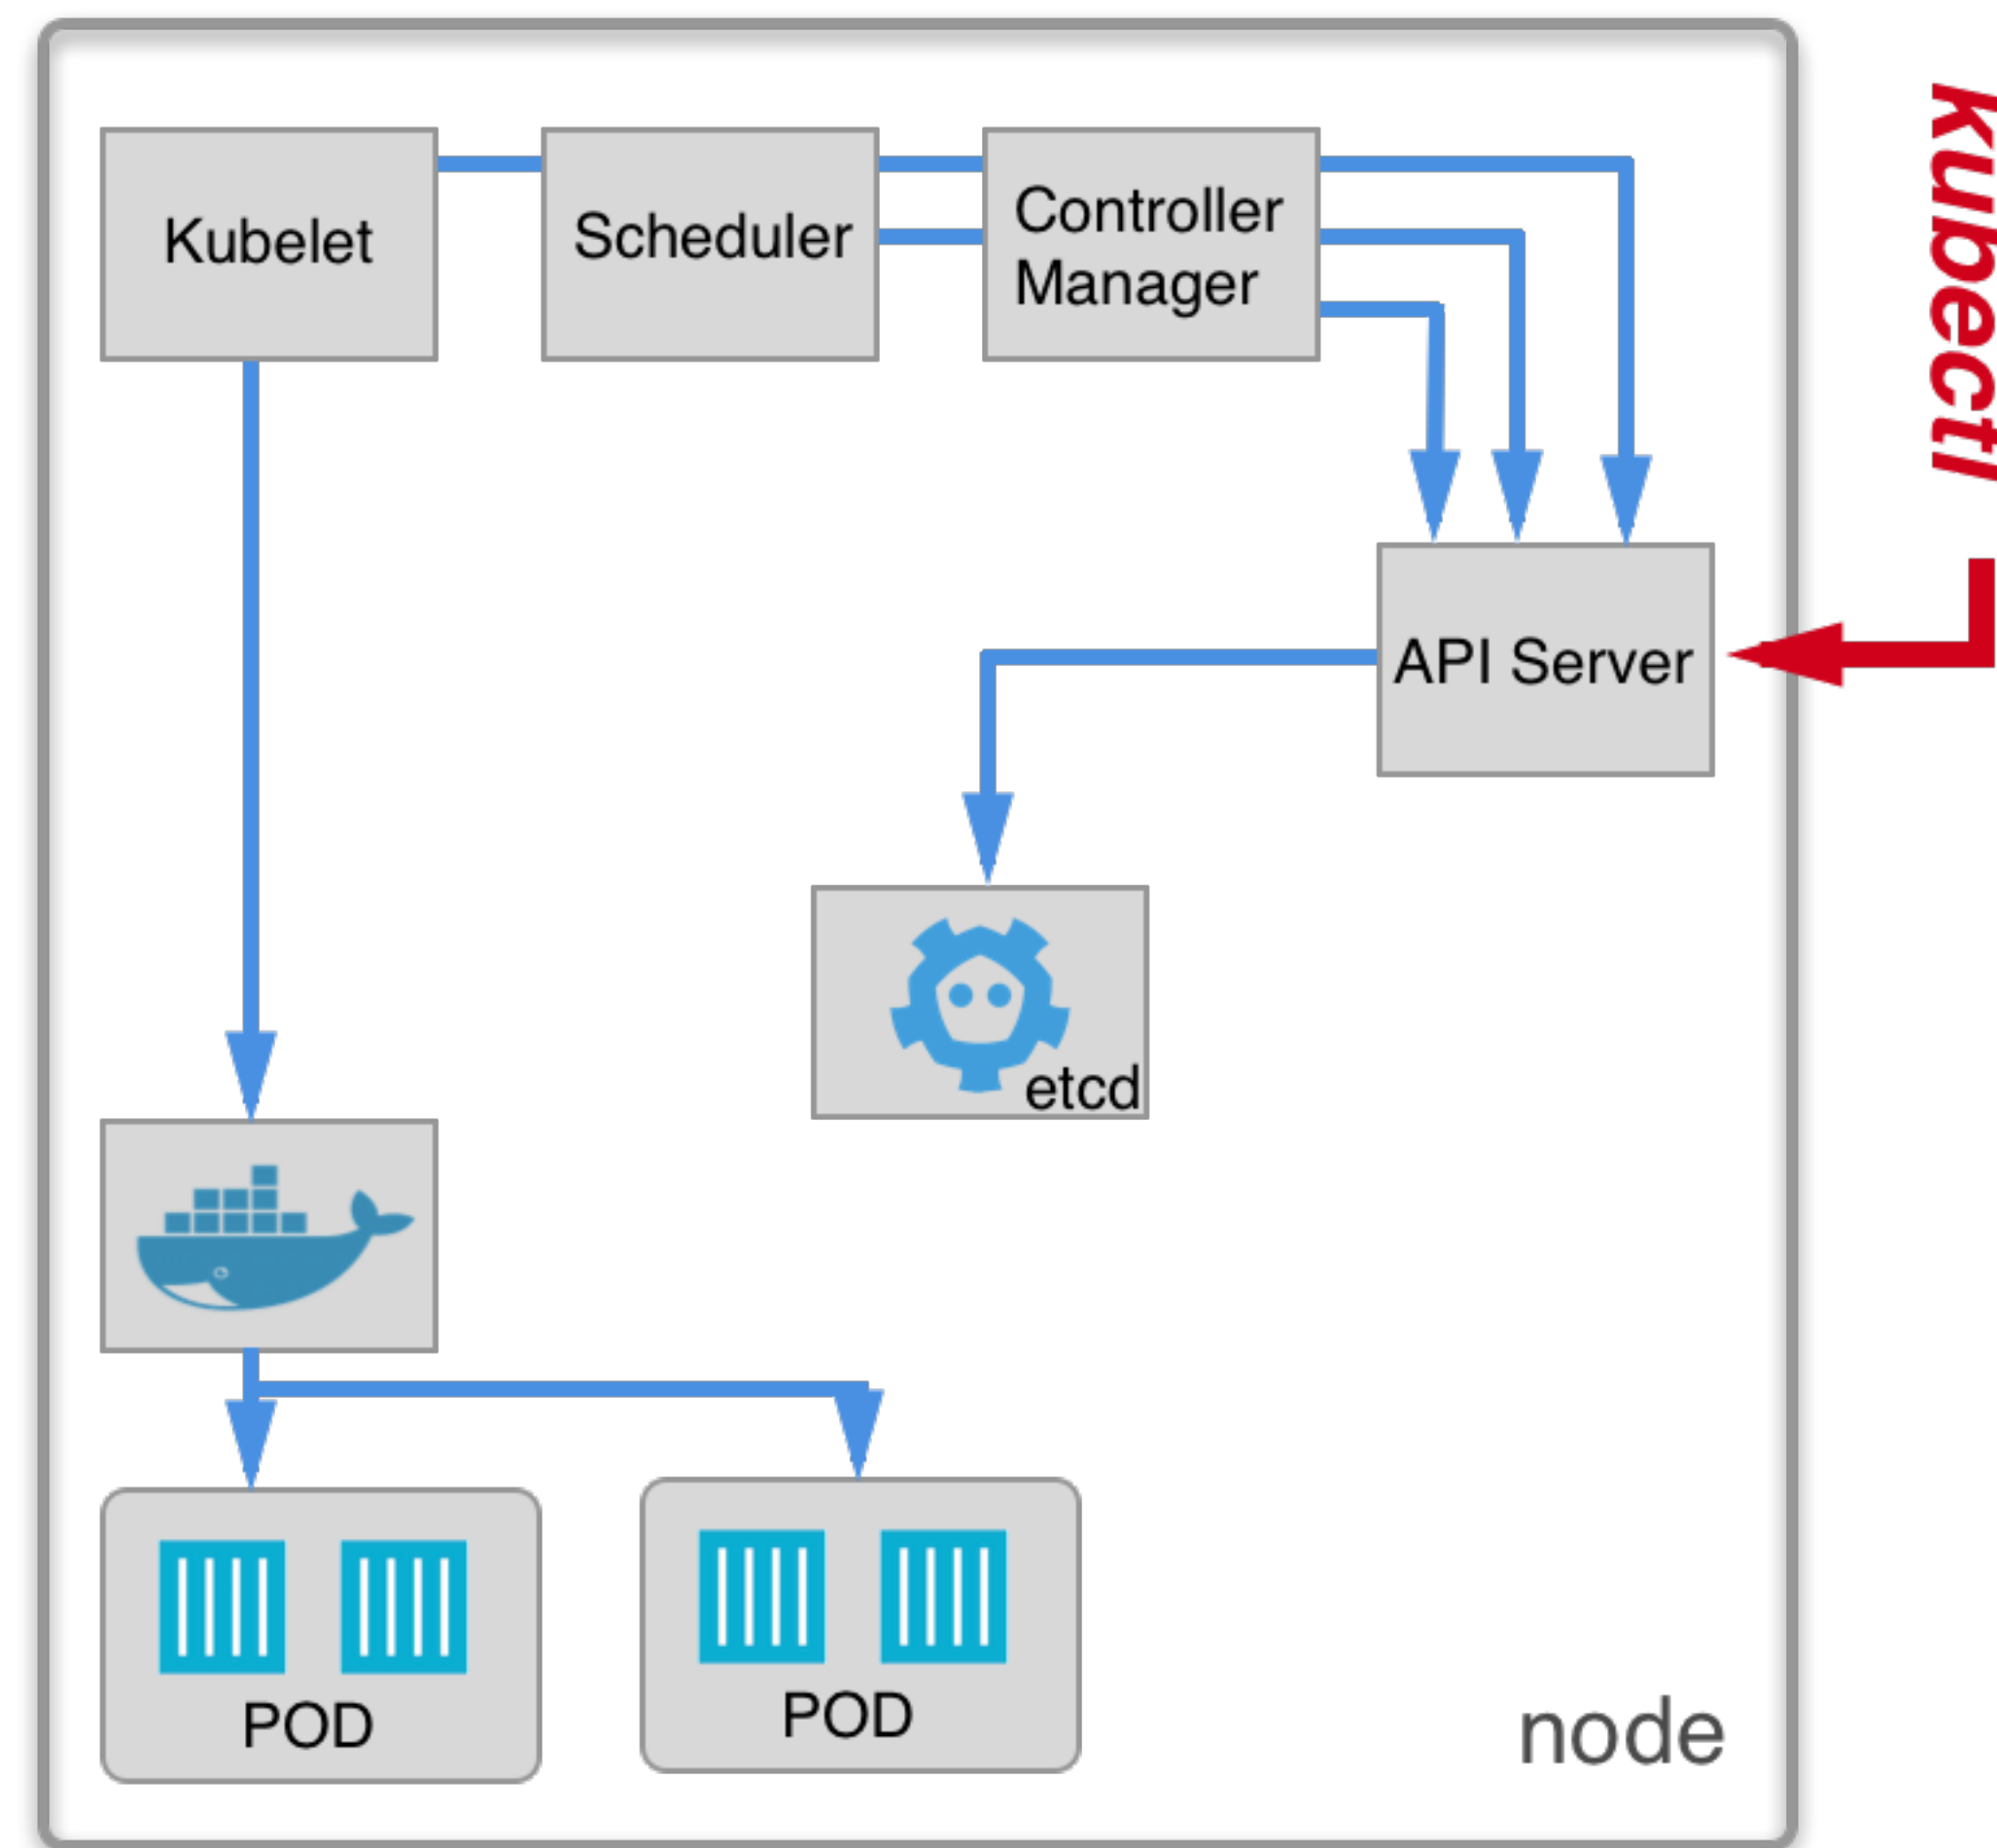


# CONTROLLER MANAGER



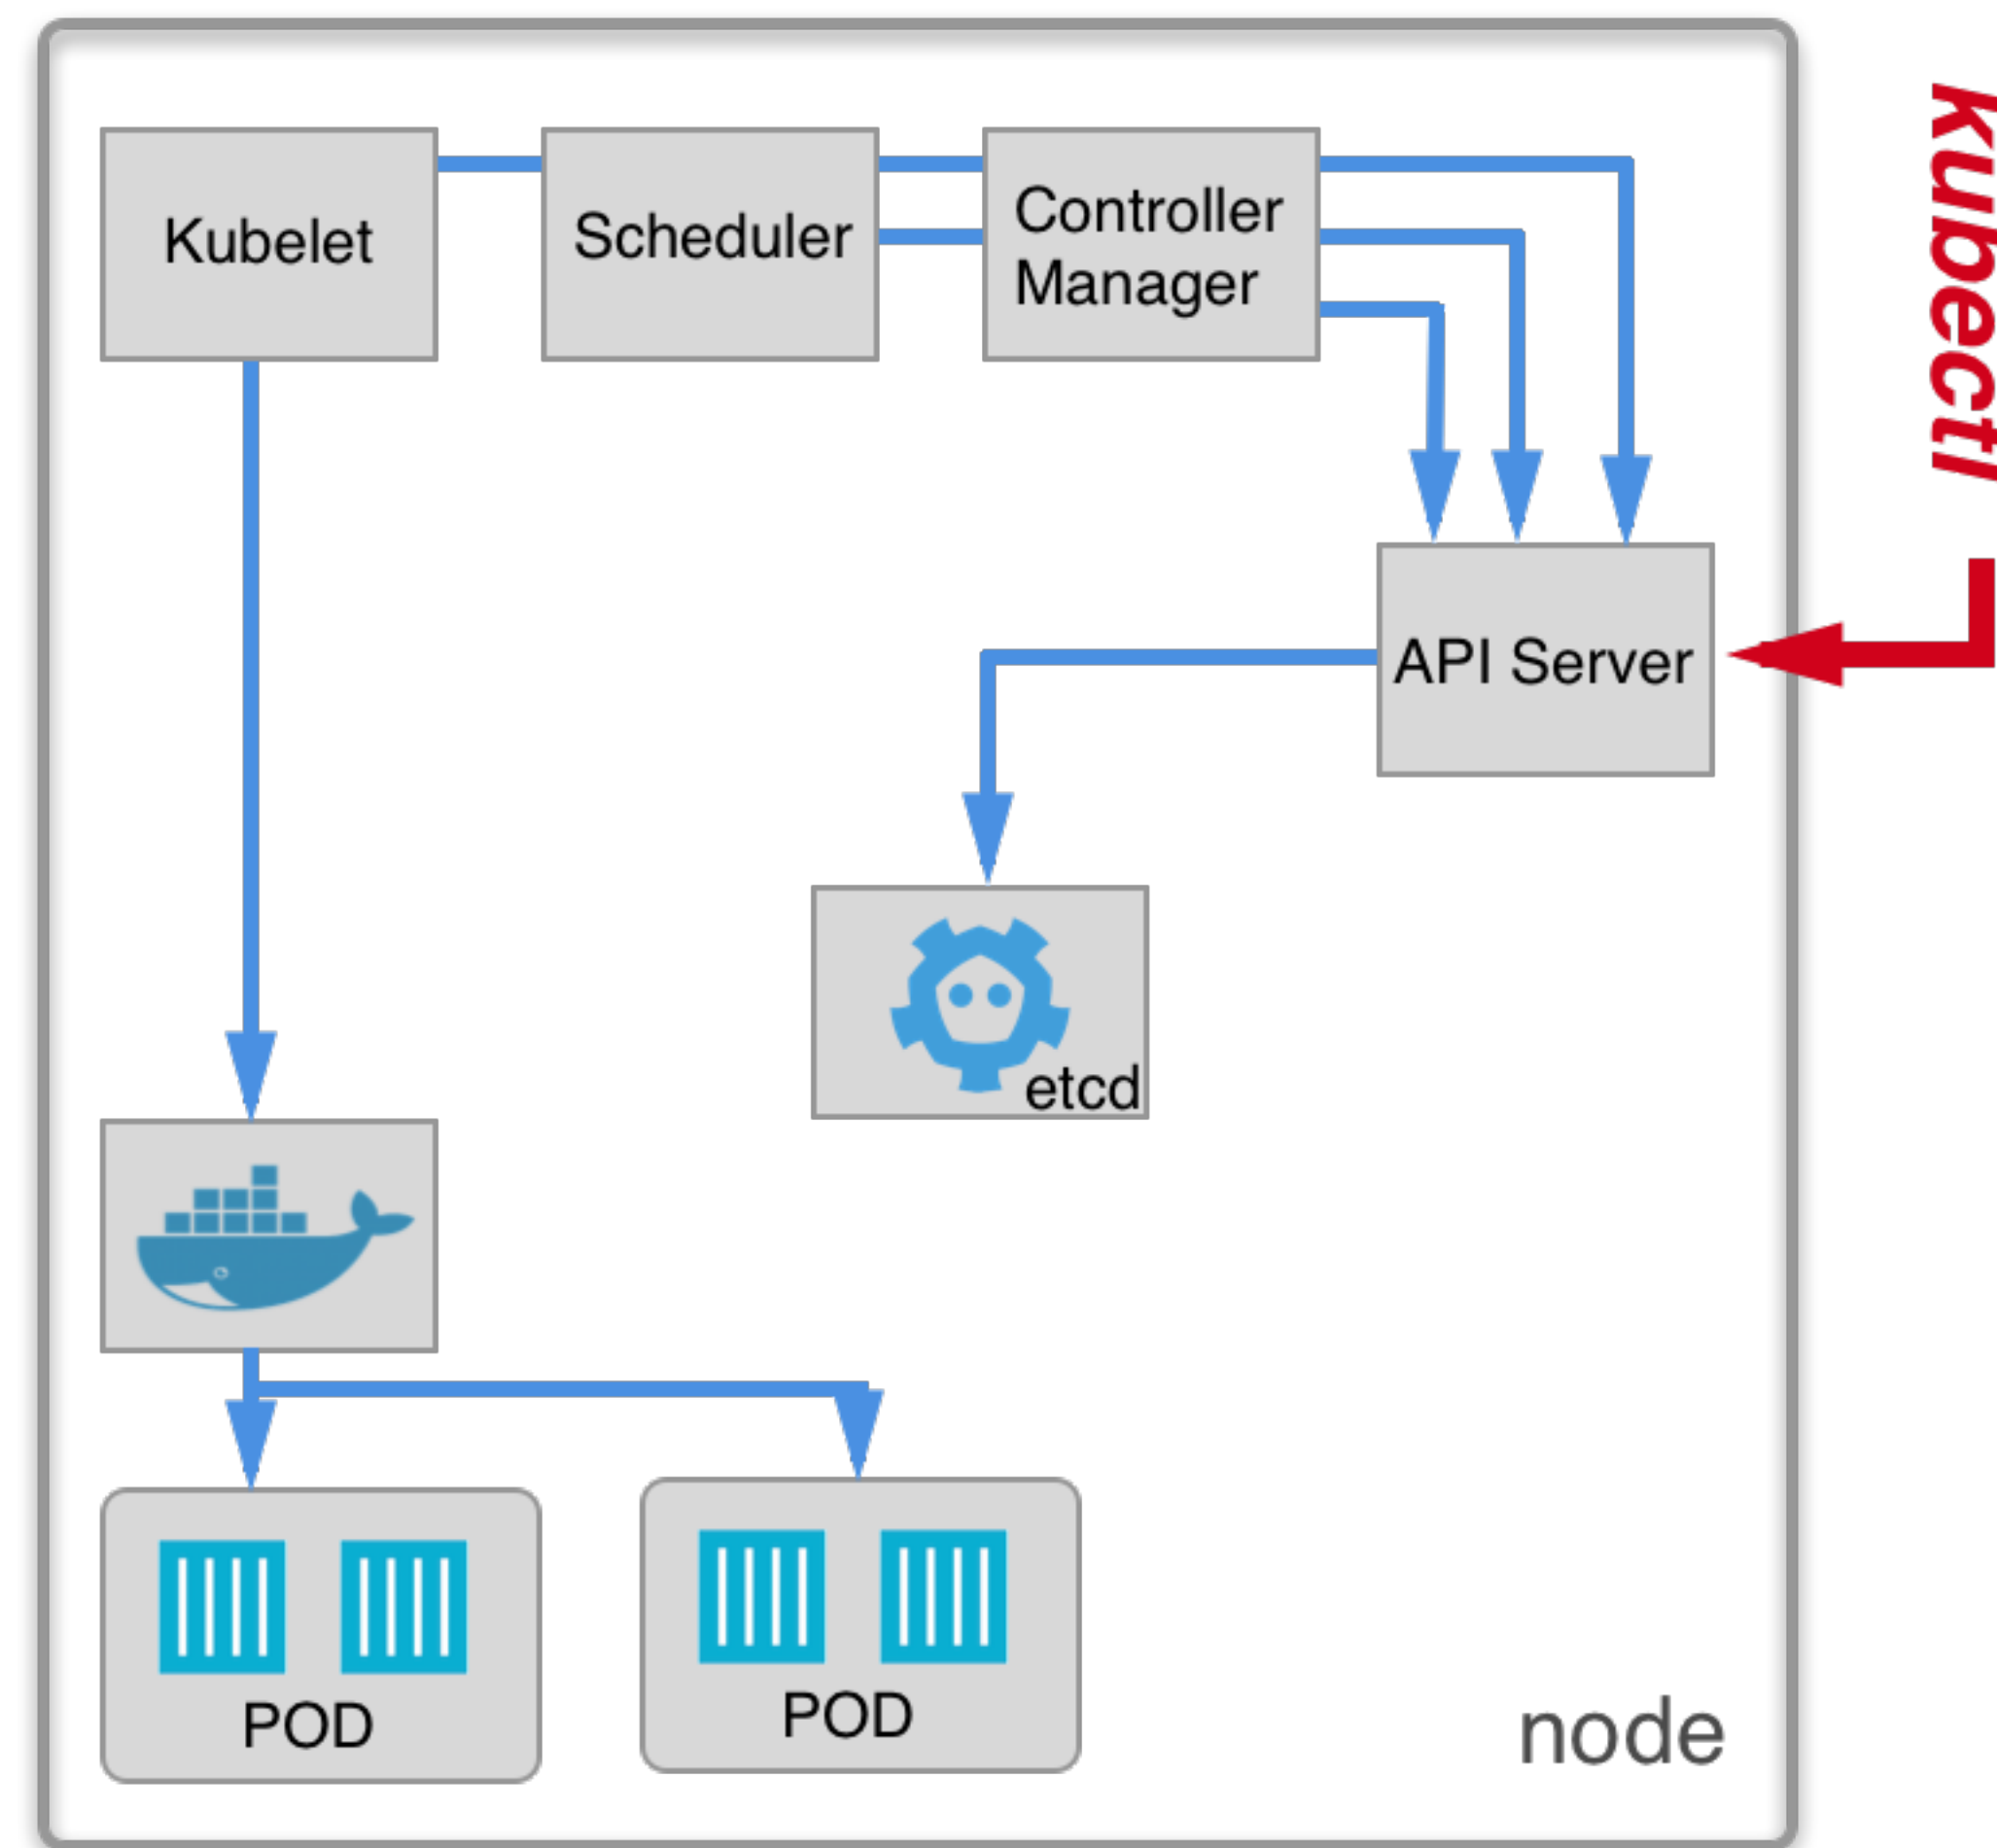
# CONTROLLER MANAGER

- responsible for managing Deployments and ReplicaSets



# CONTROLLER MANAGER

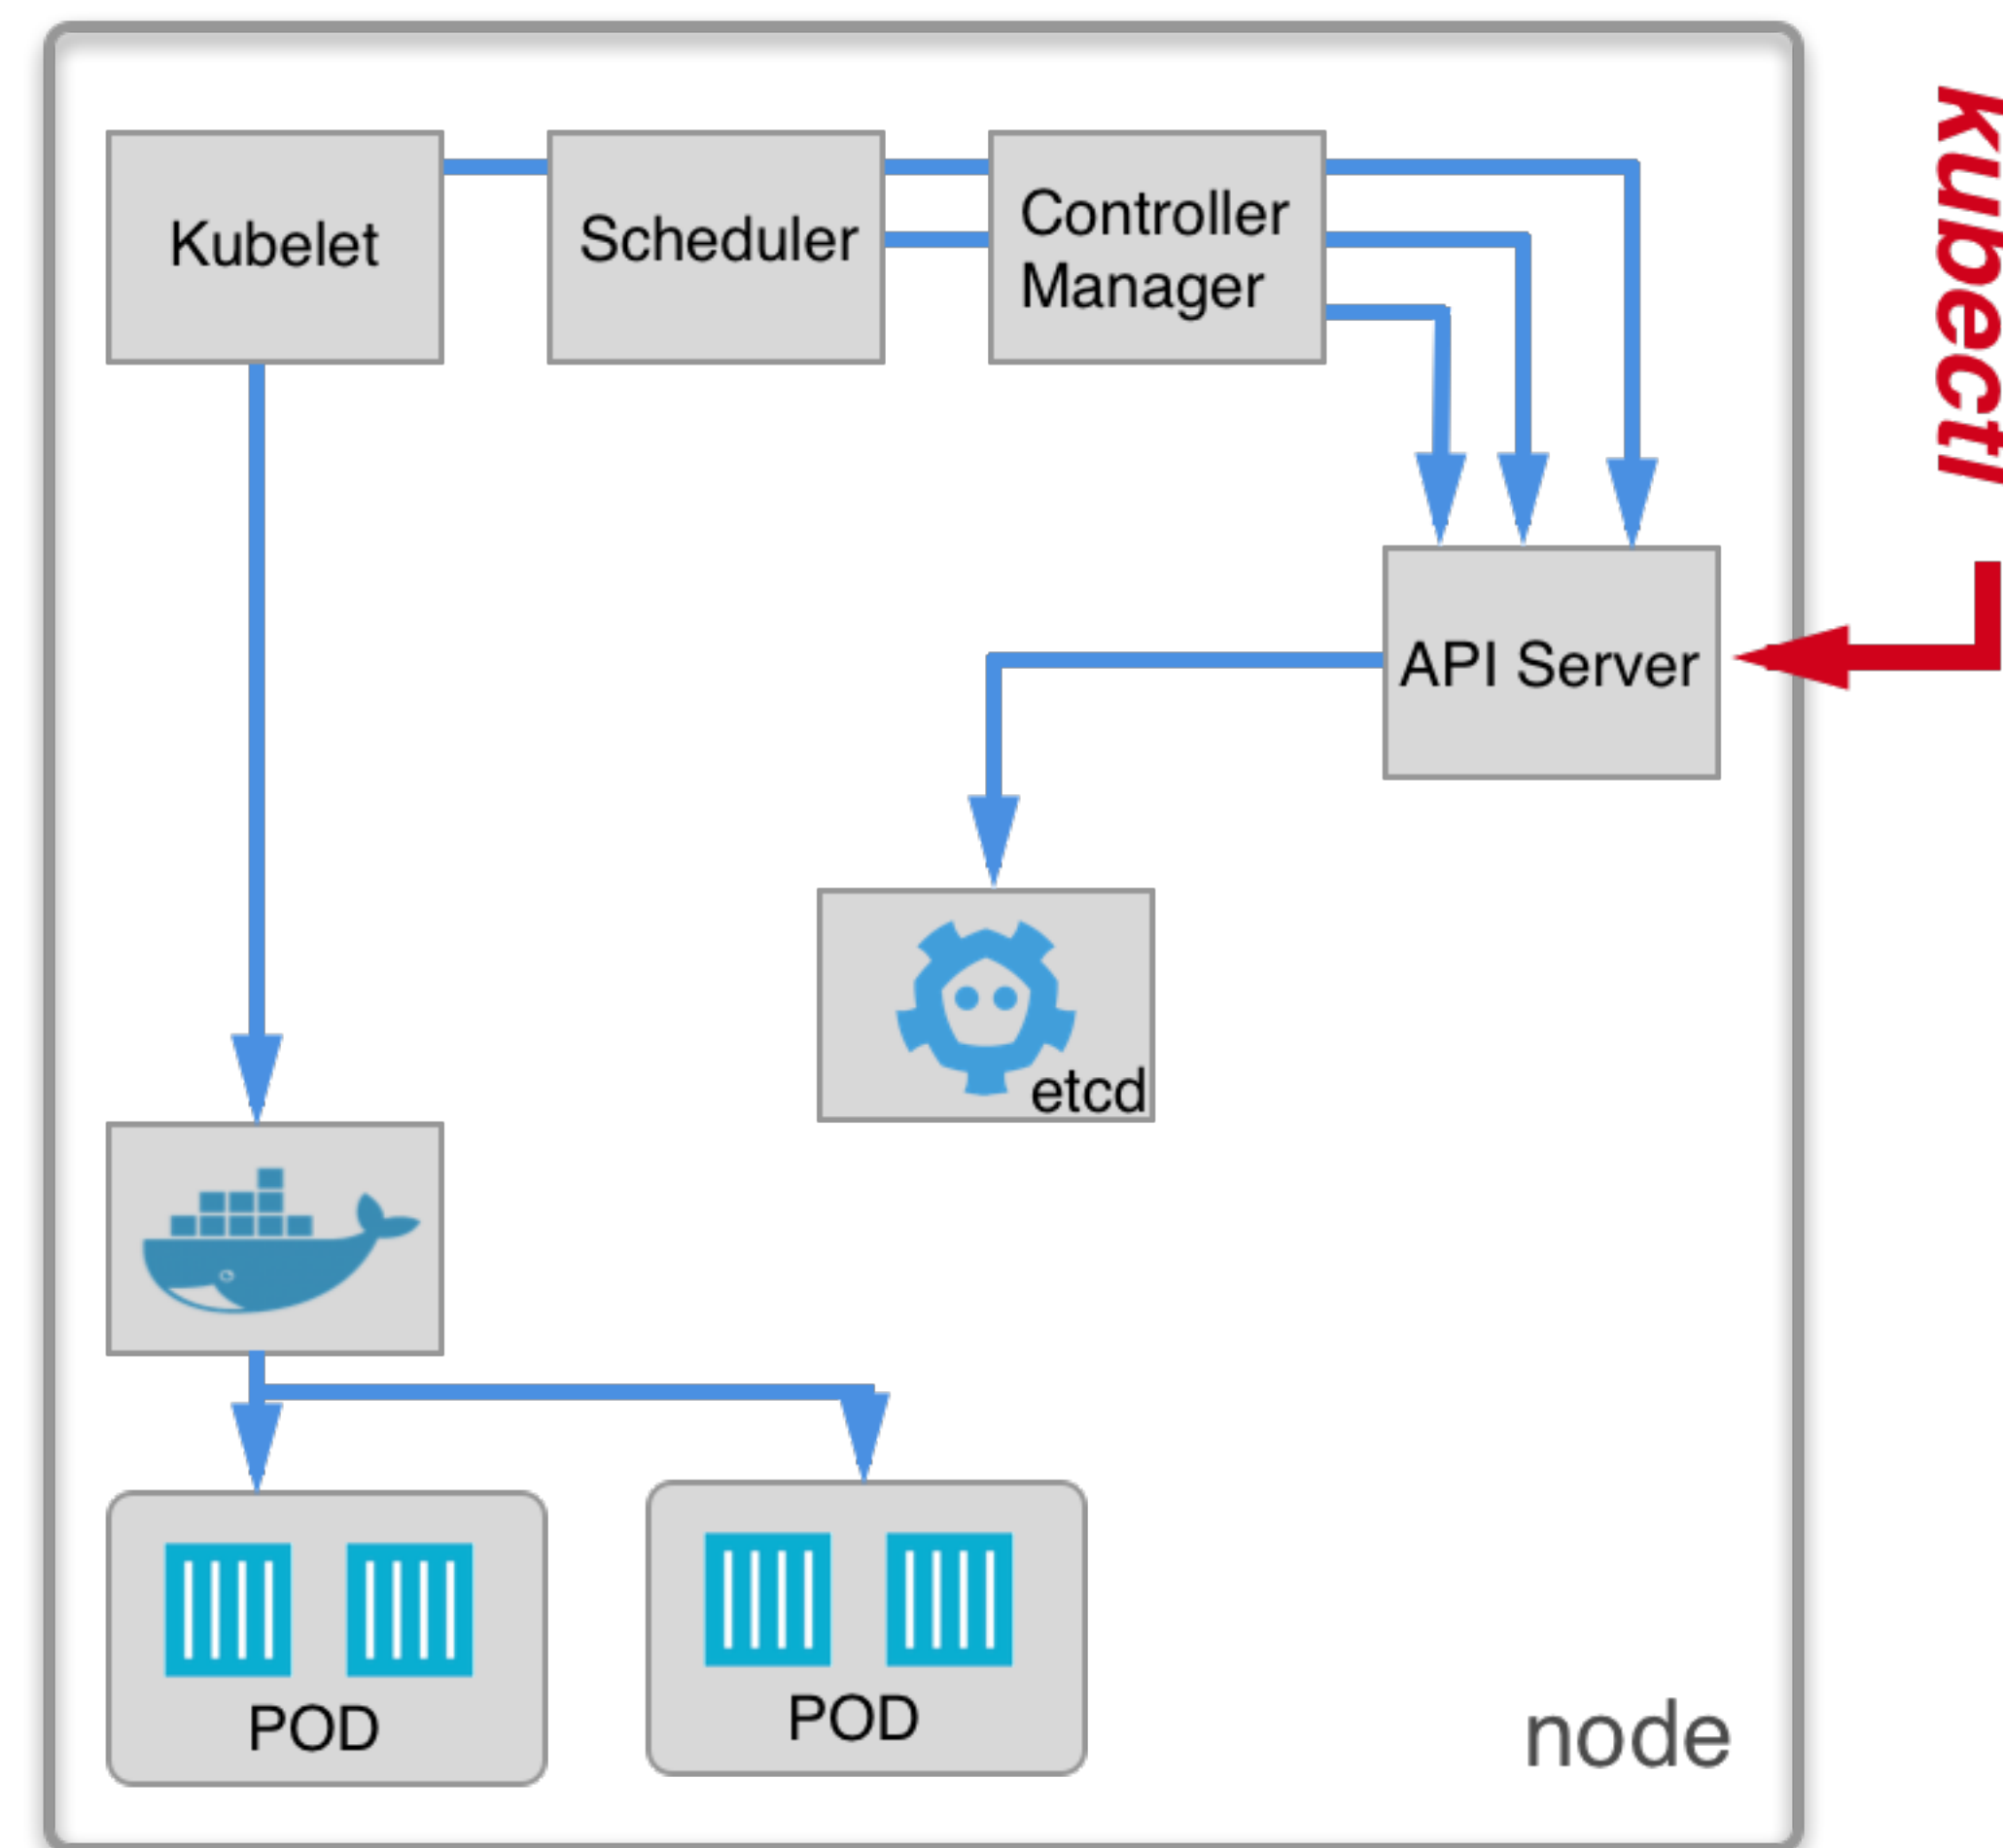
- responsible for managing Deployments and ReplicaSets
- connects to API Server



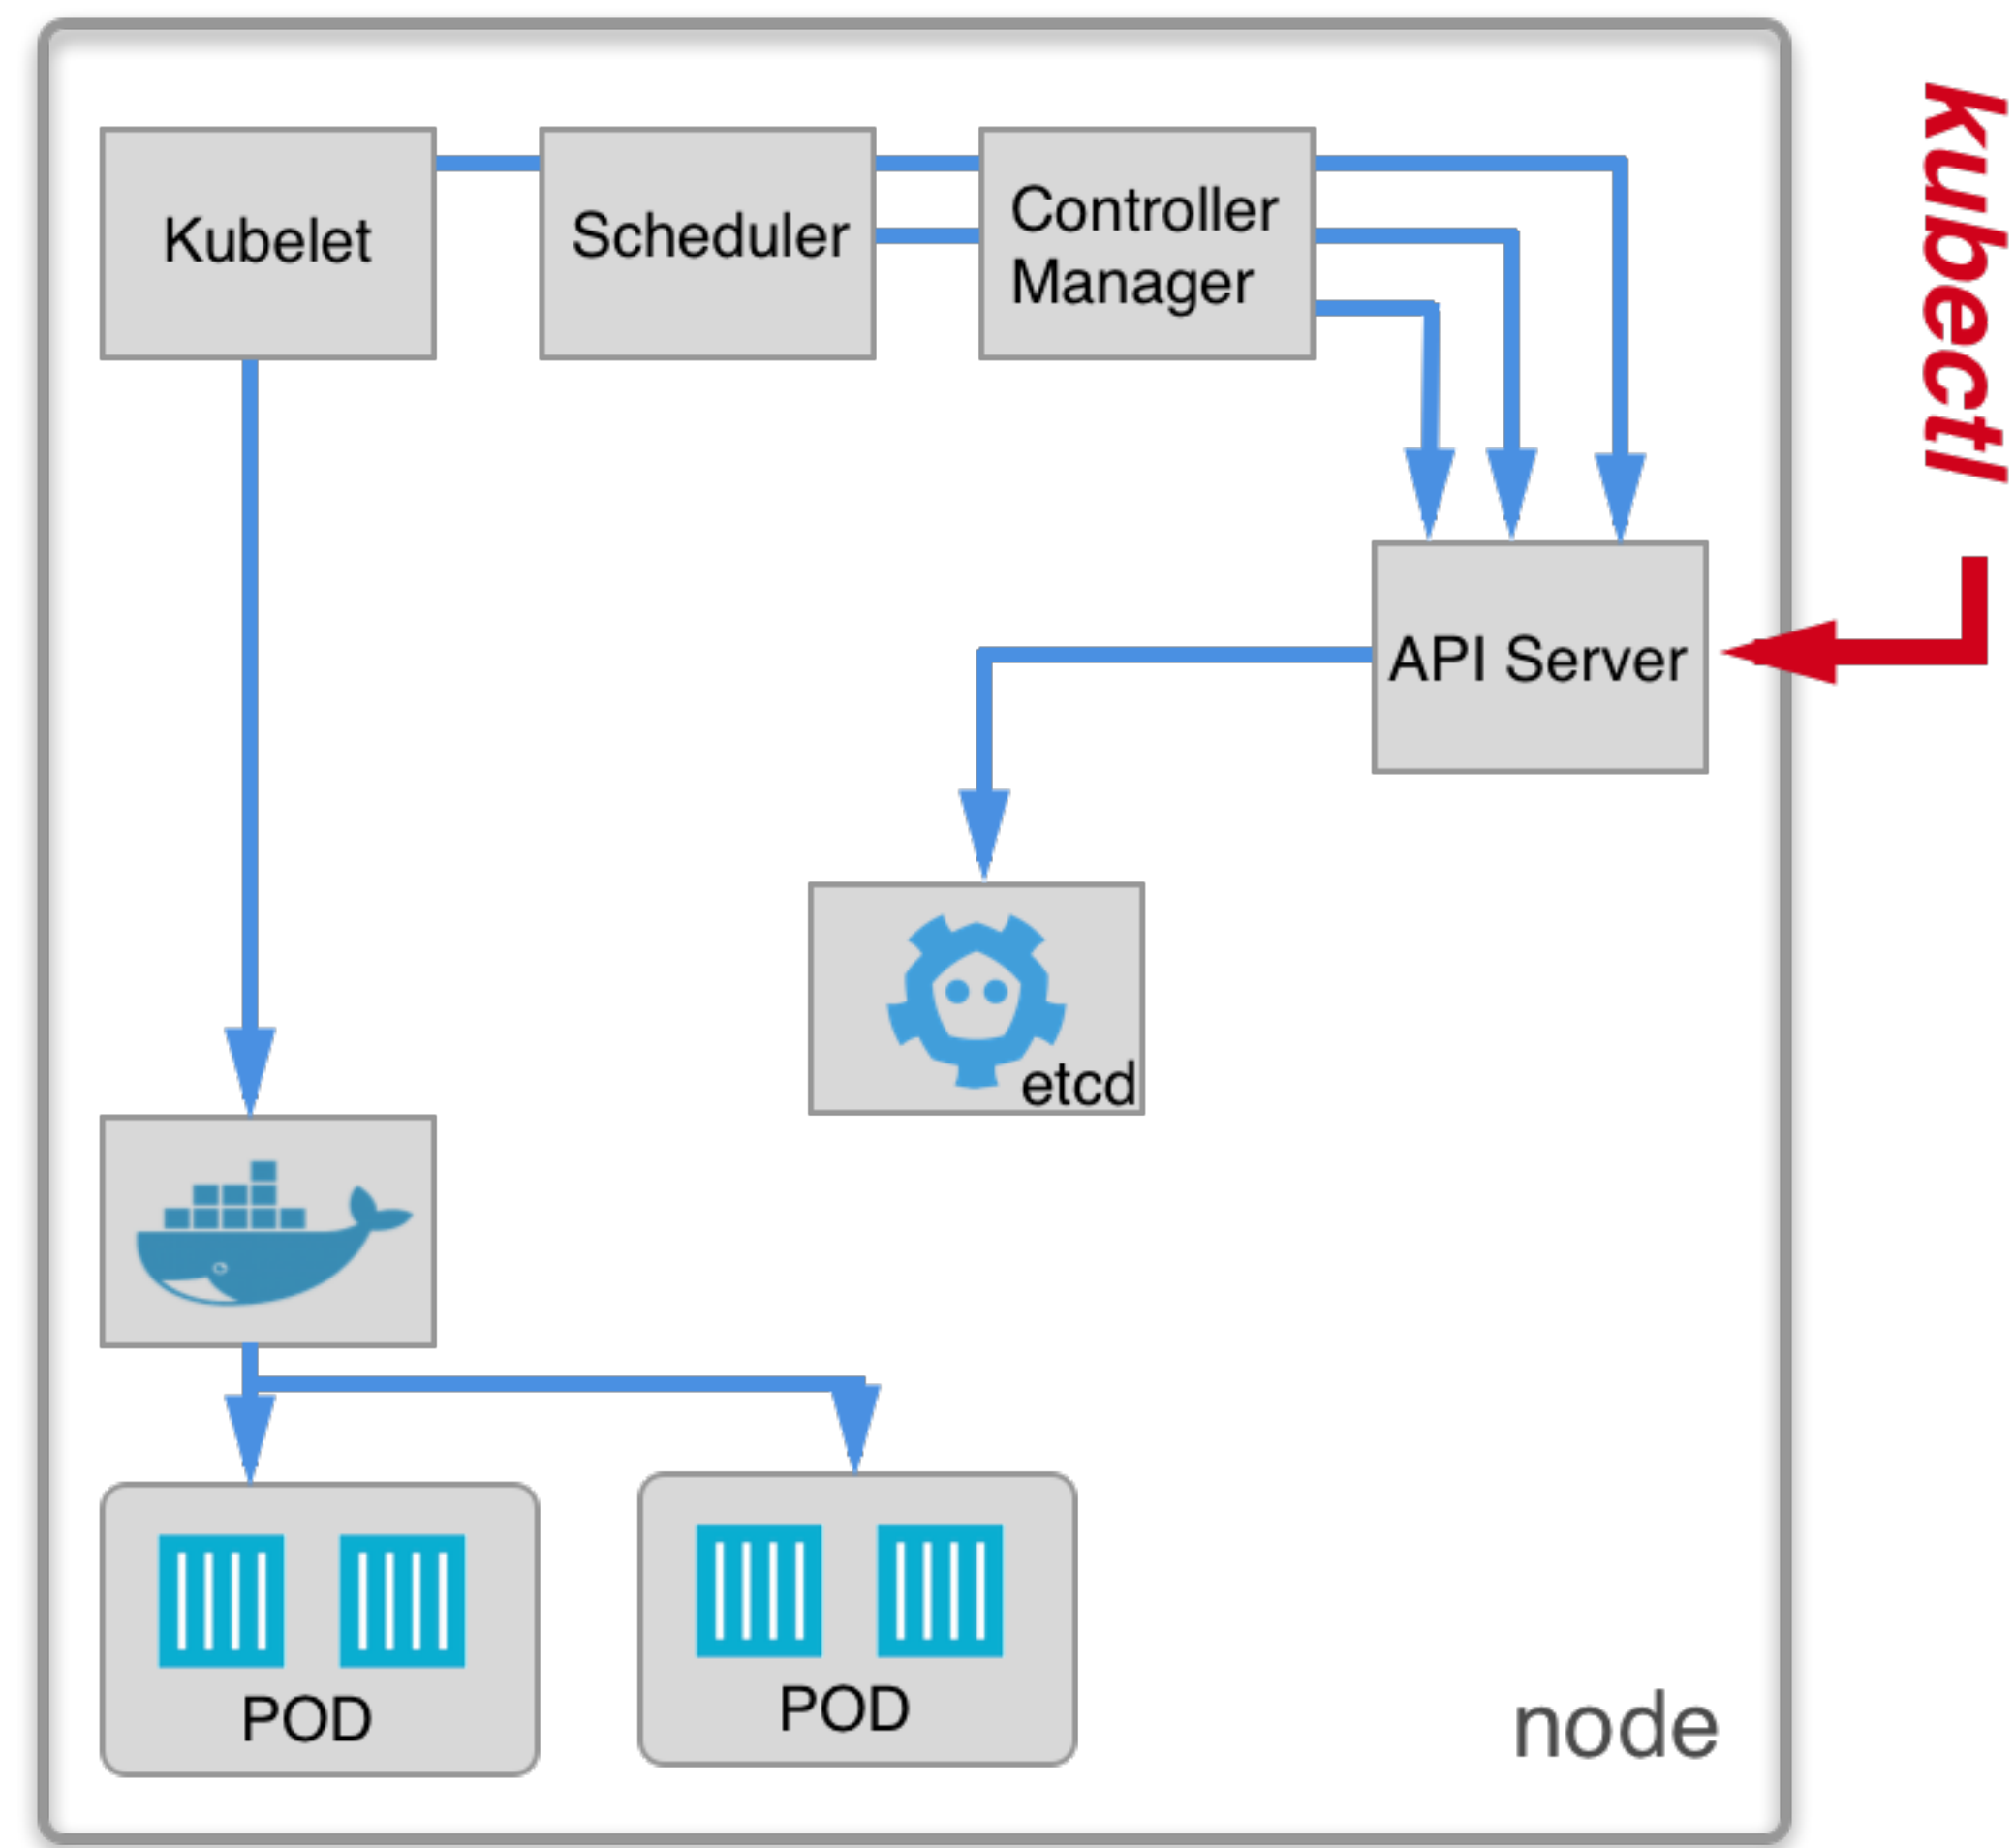


# CONTROLLER MANAGER

- responsible for managing Deployments and ReplicaSets
- connects to API Server
- takes care of keeping pods counts at given number

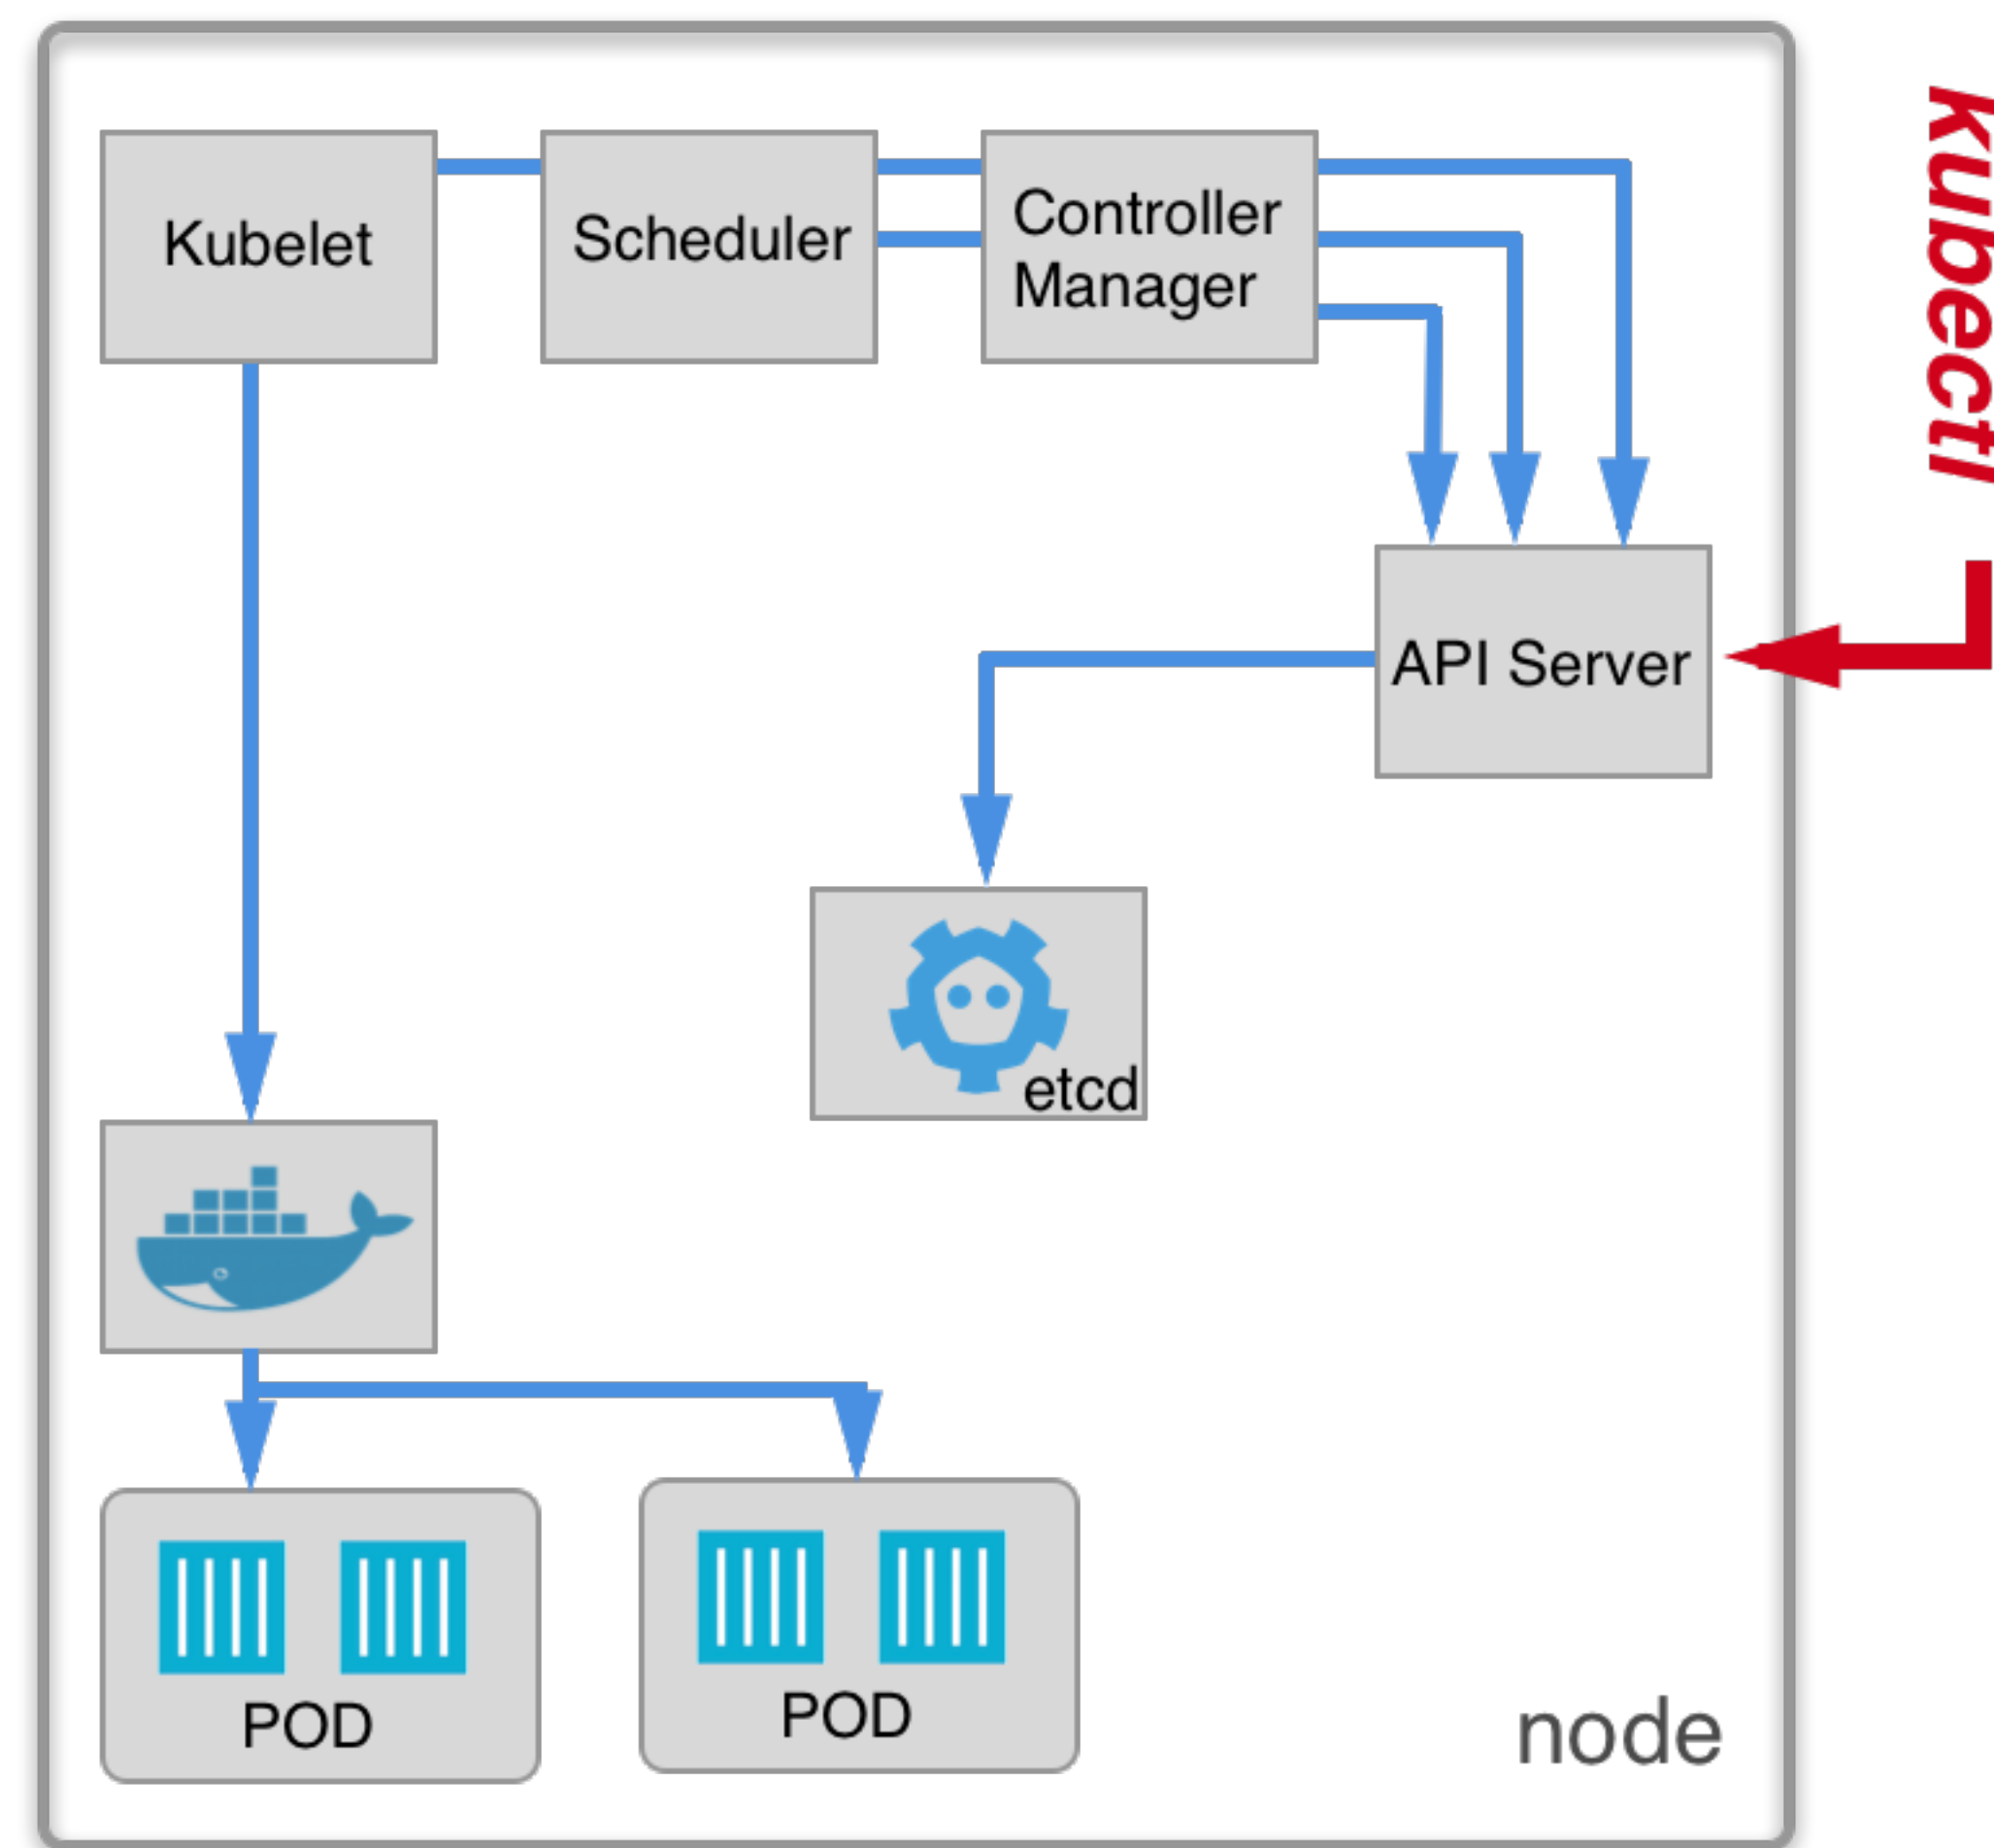


# BASIC FLOW



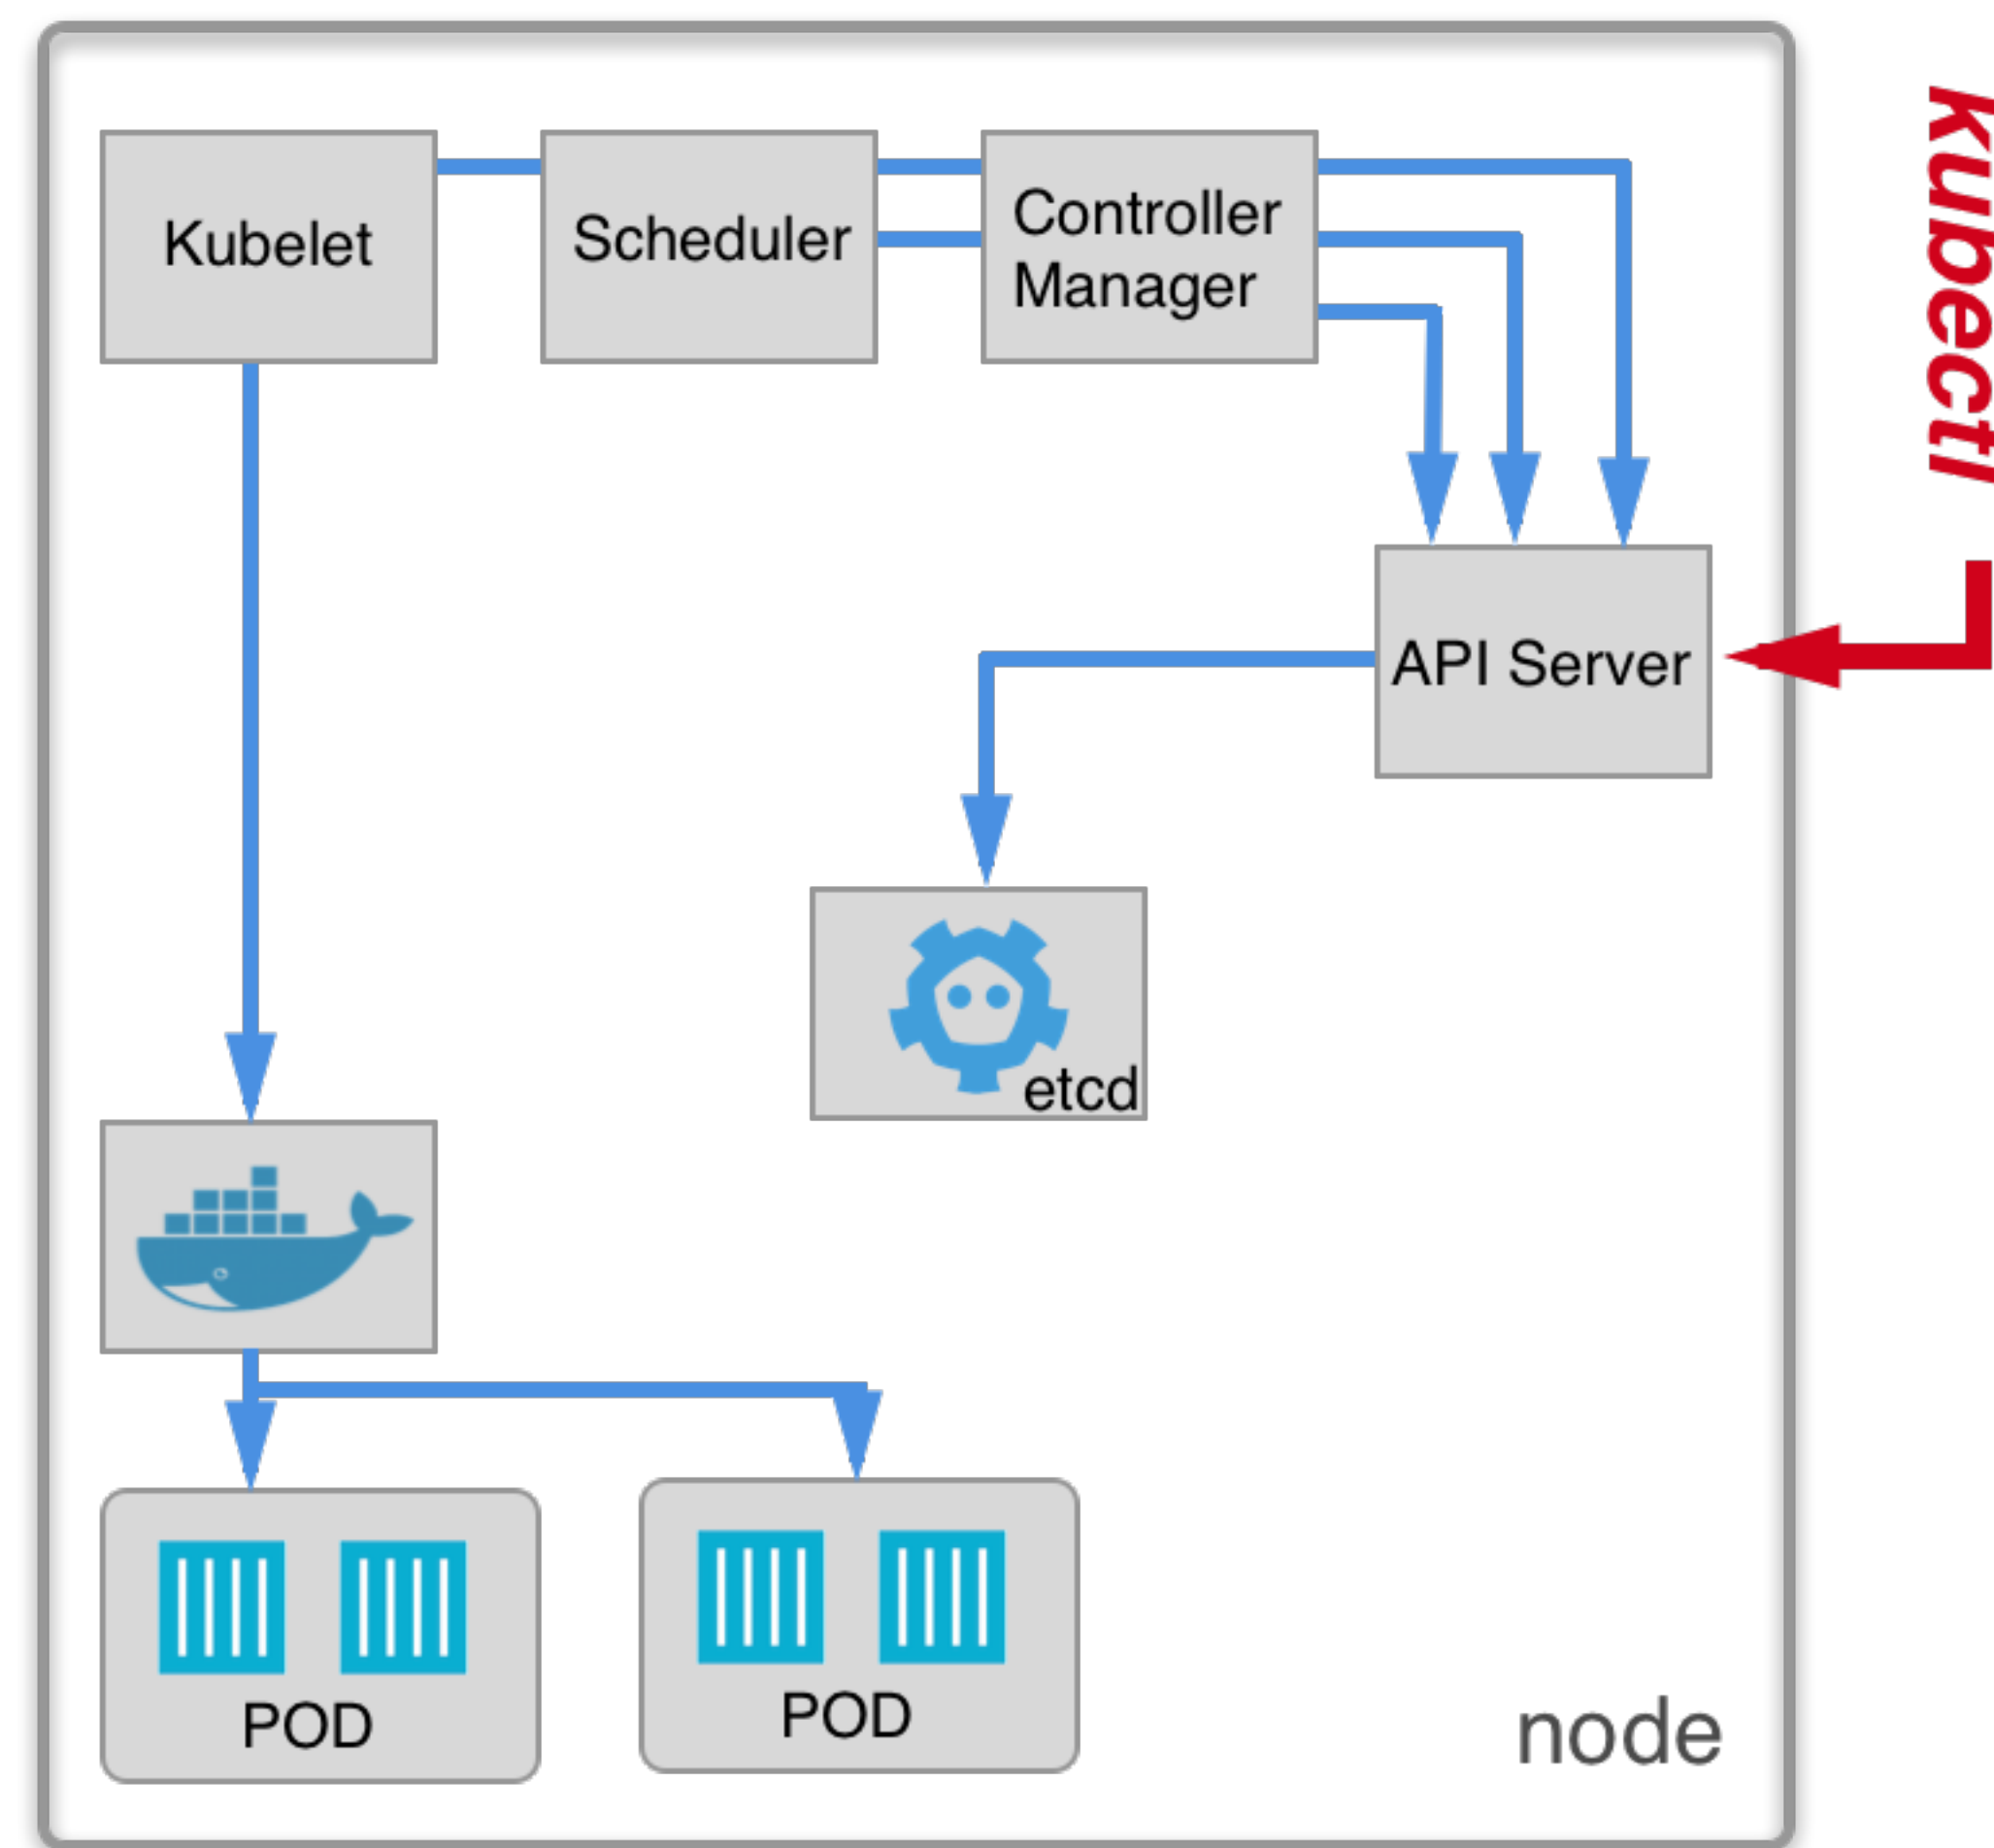
# BASIC FLOW

- kubectl apply -f deploy.yaml (5 replicas)



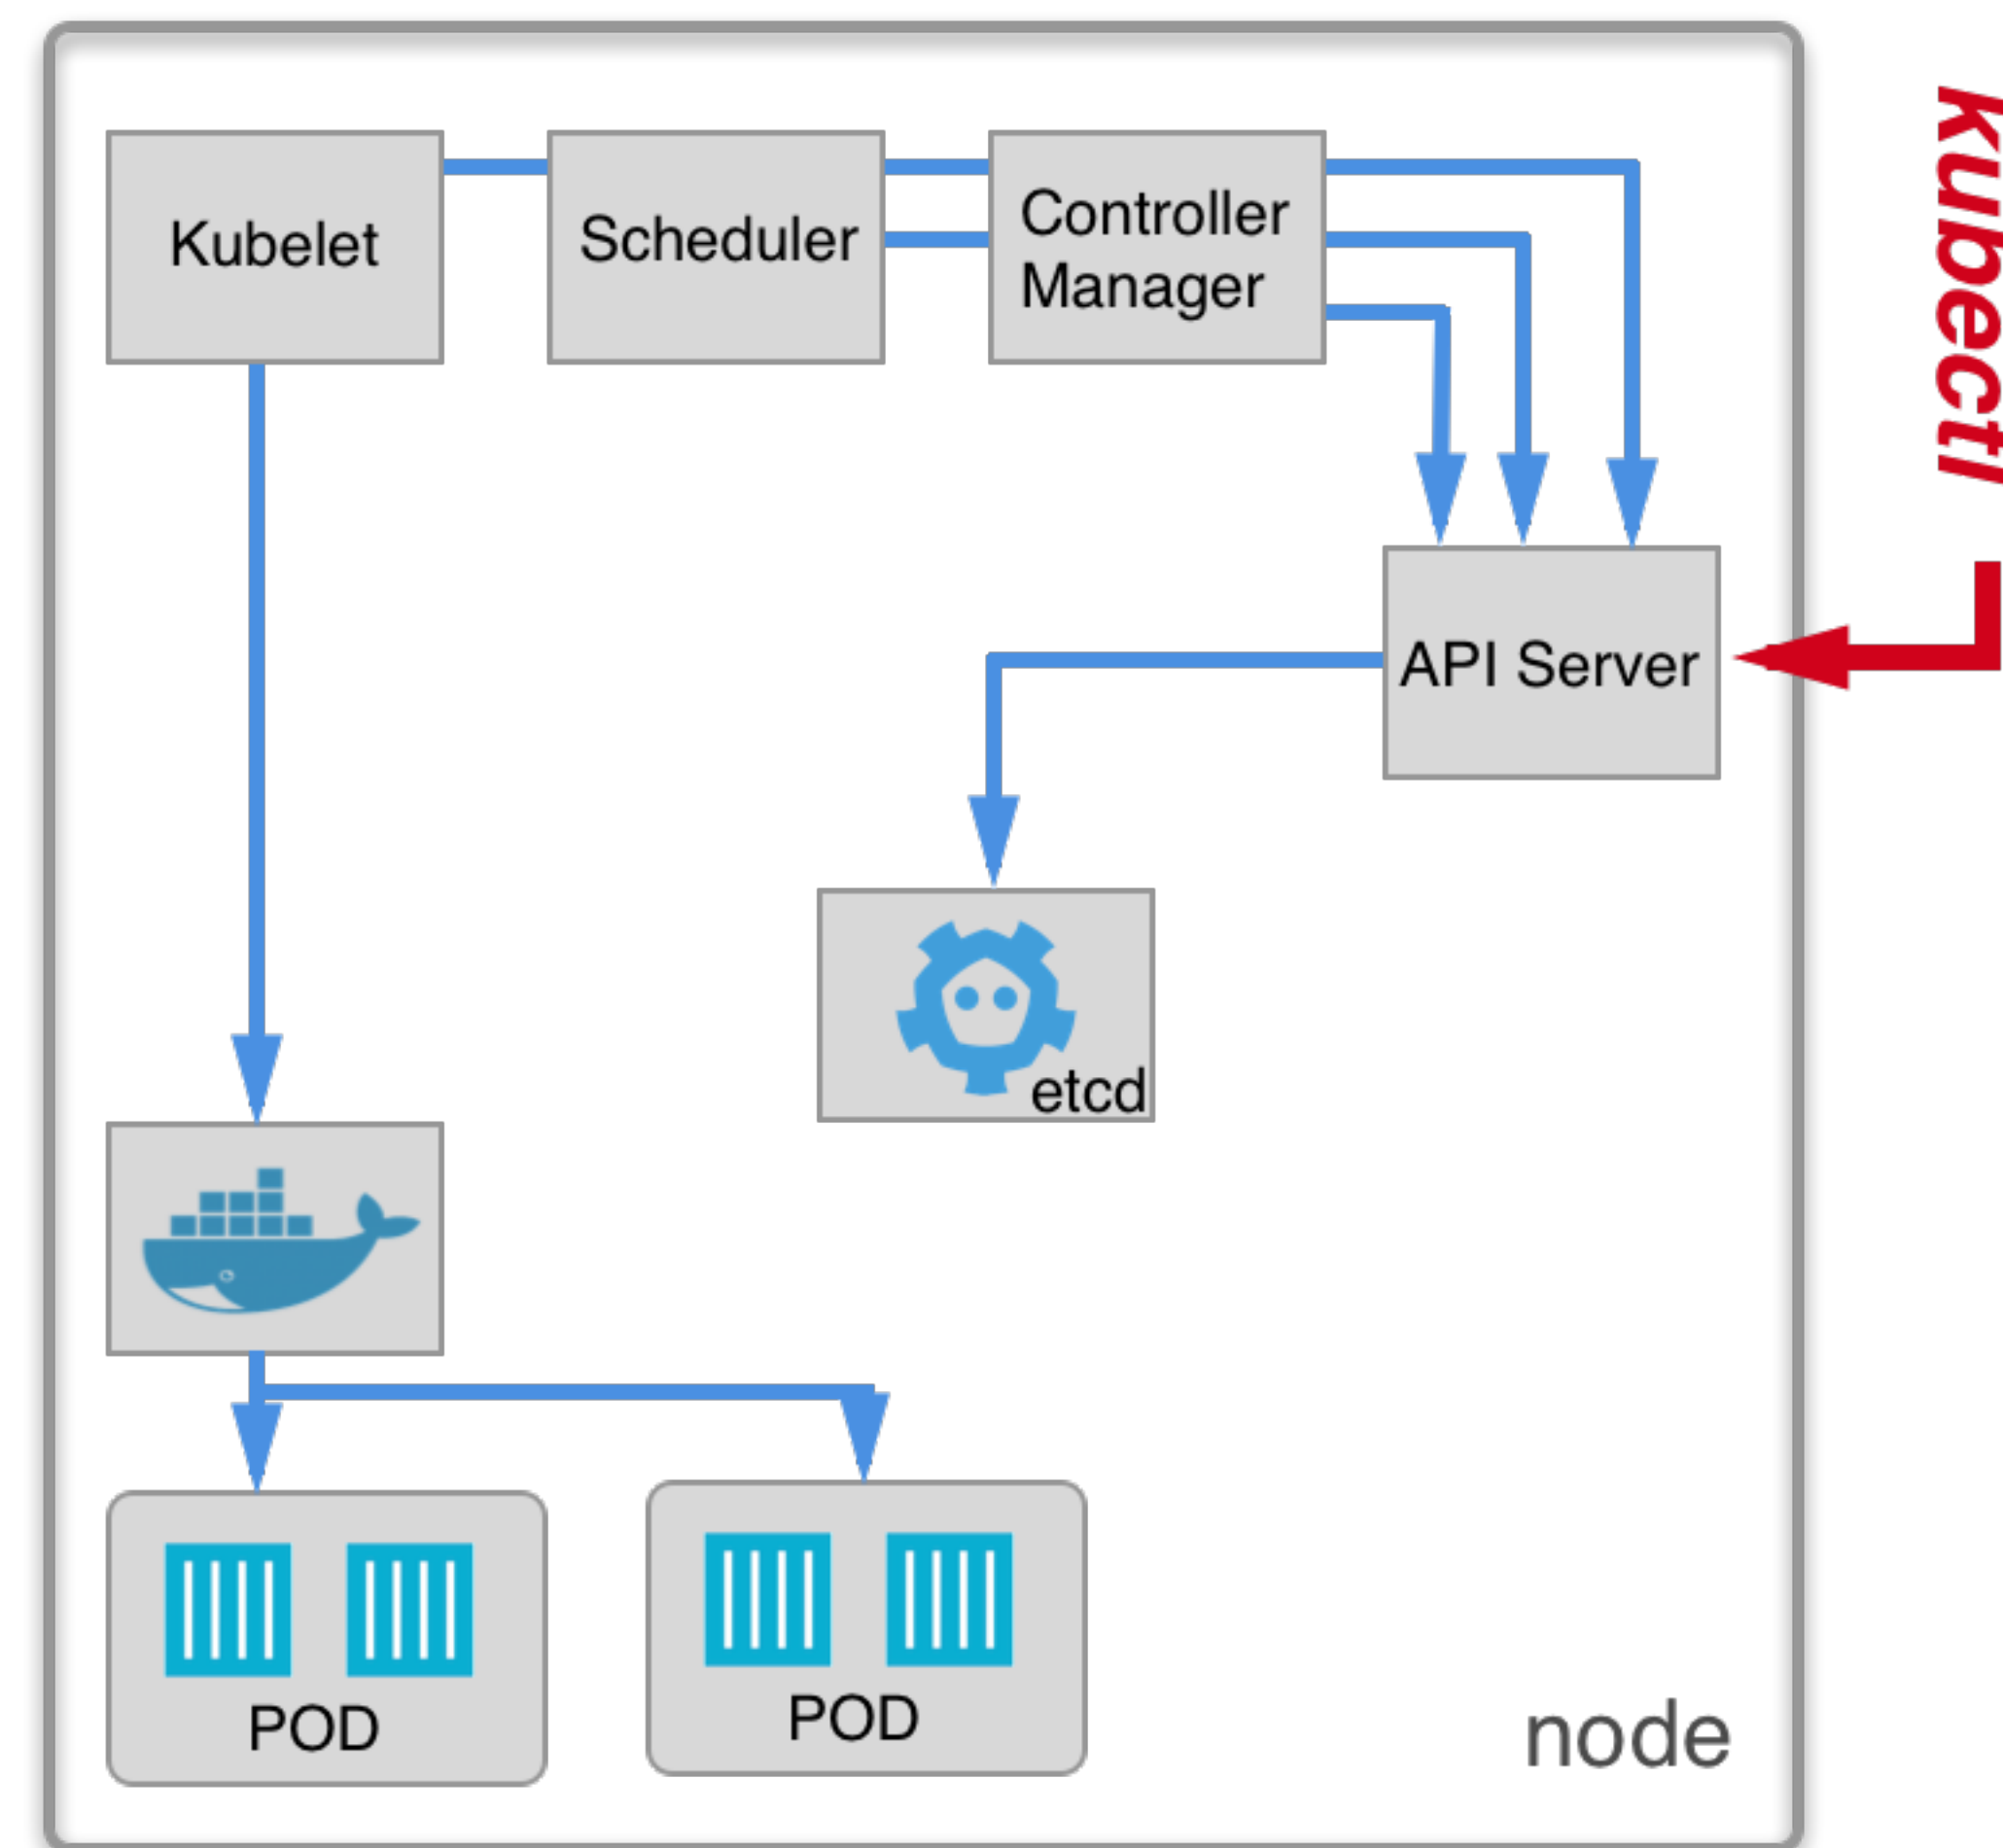
# BASIC FLOW

- `kubectl apply -f deploy.yaml` (5 replicas)
- CM creates replicaset object



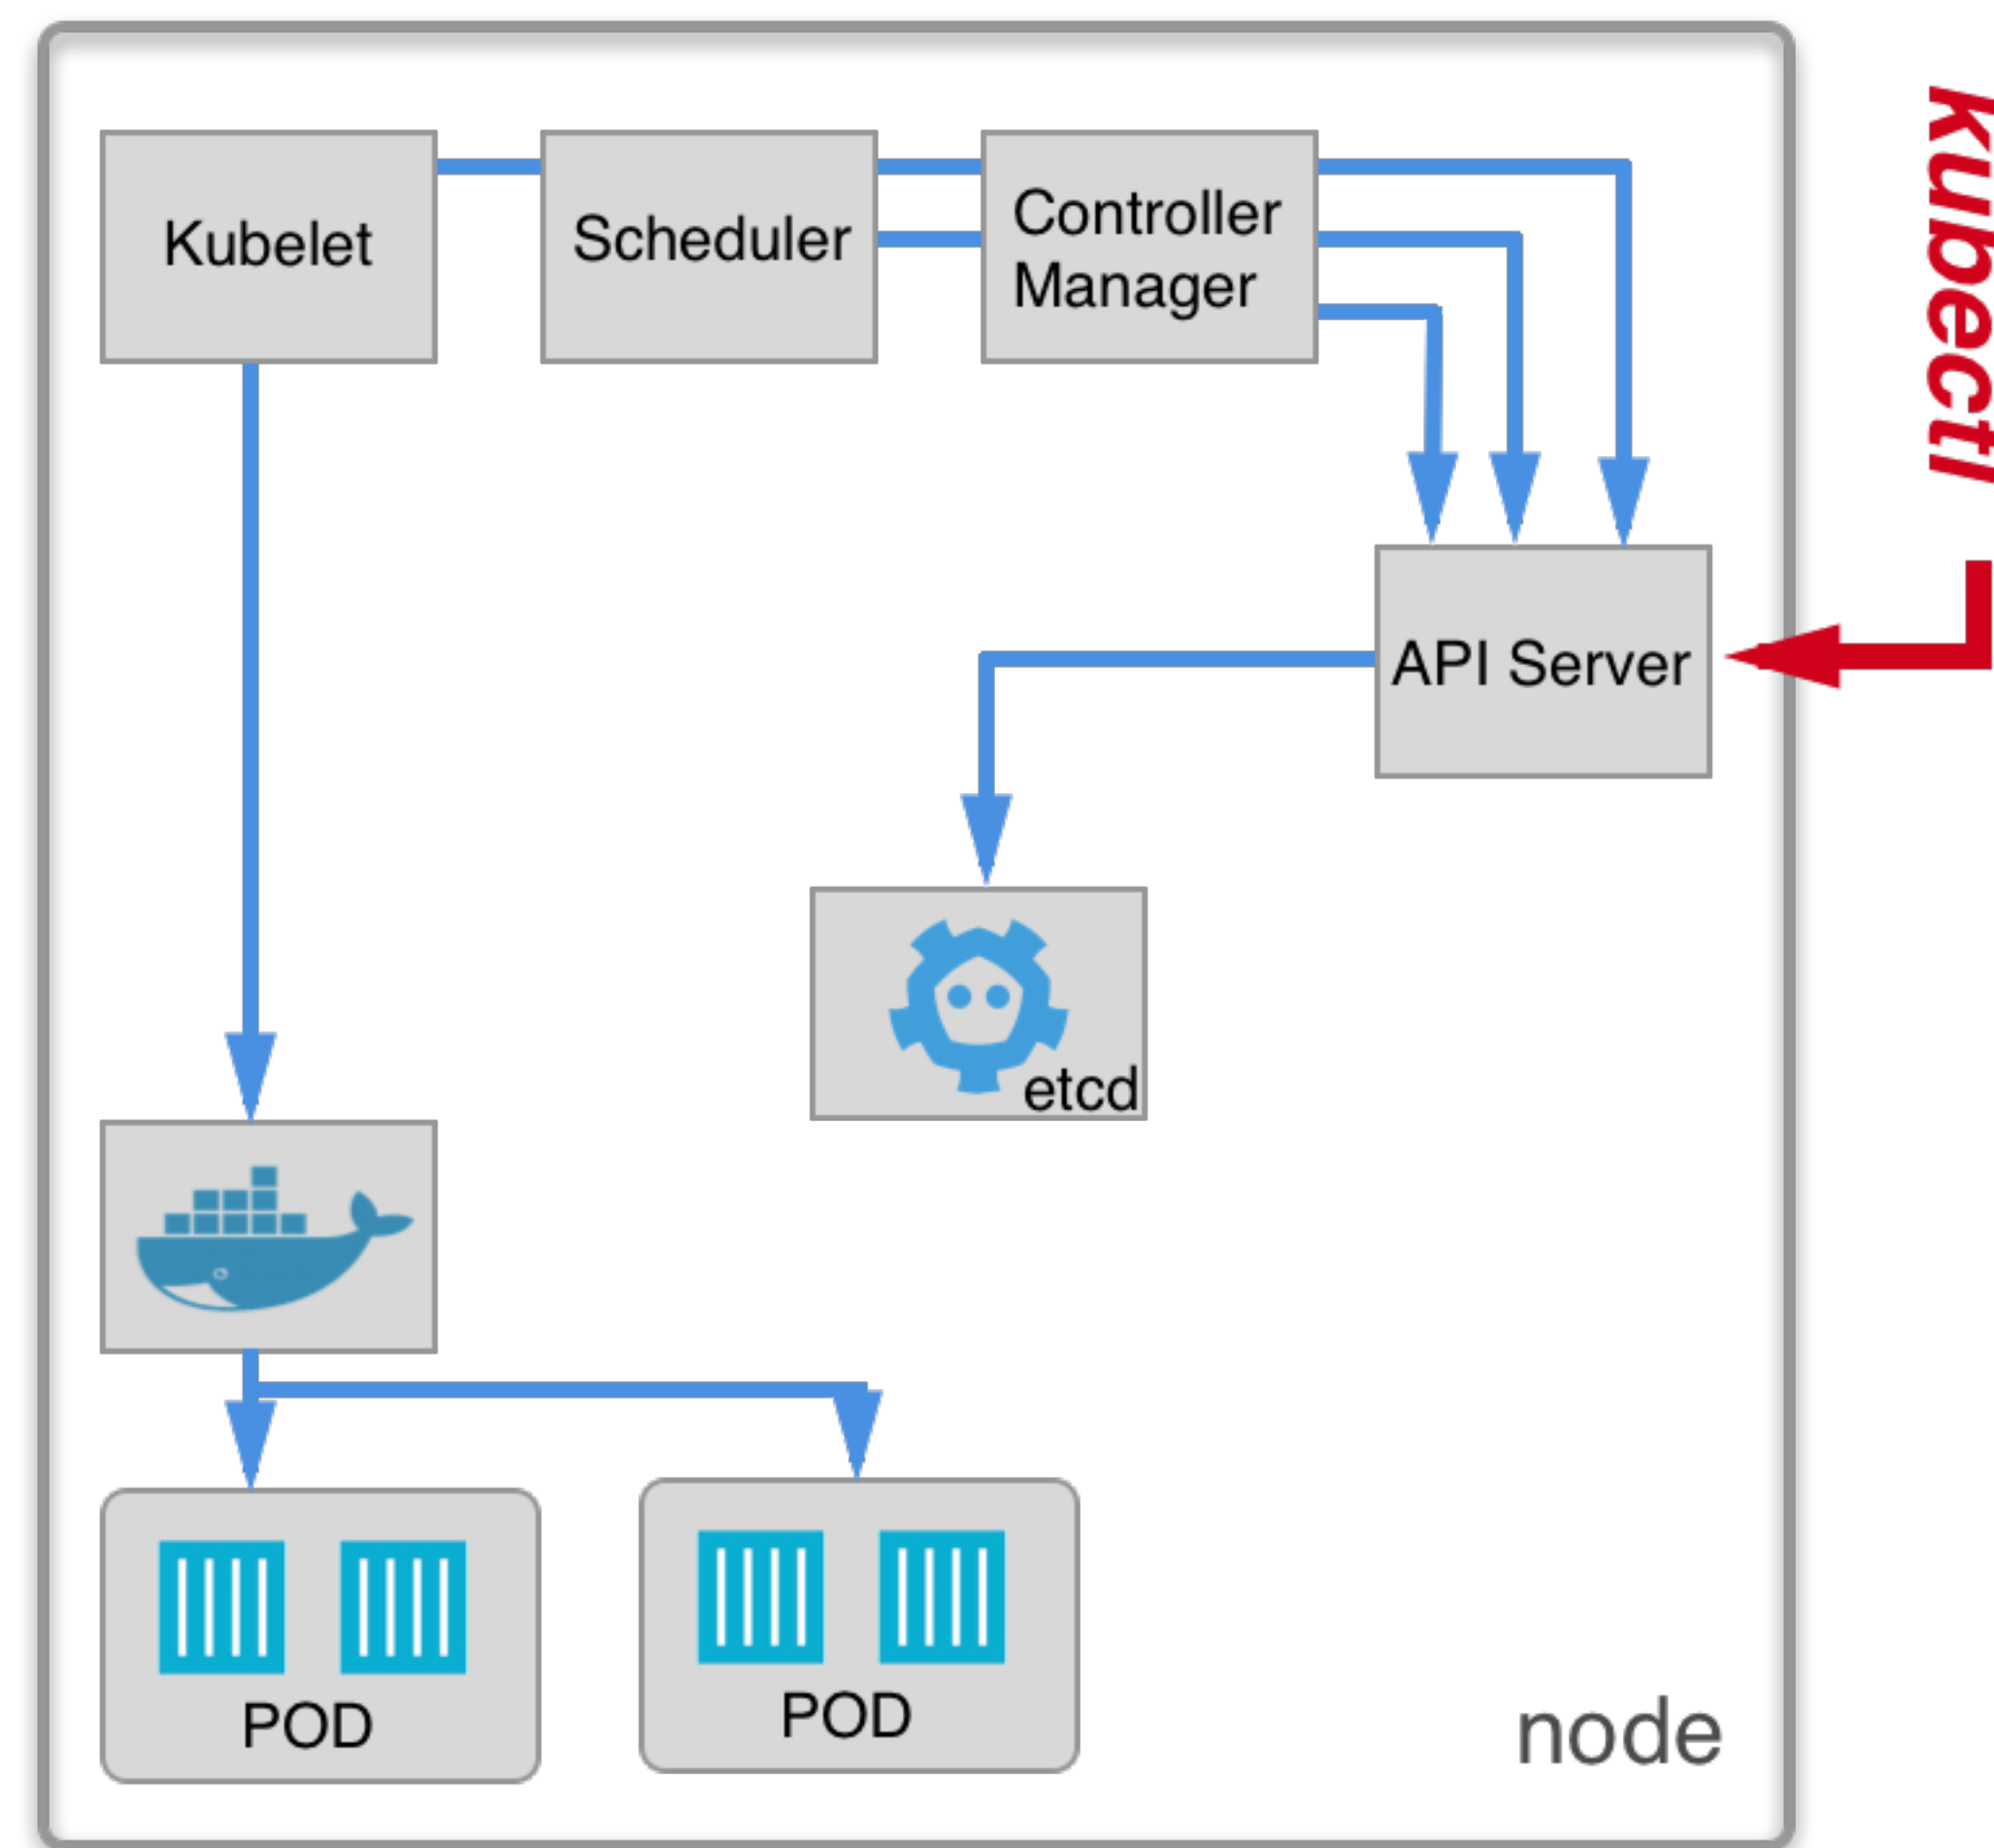
# BASIC FLOW

- kubectl apply -f deploy.yaml (5 replicas)
- CM creates replicaset object
- CM creates 5 pods' objects



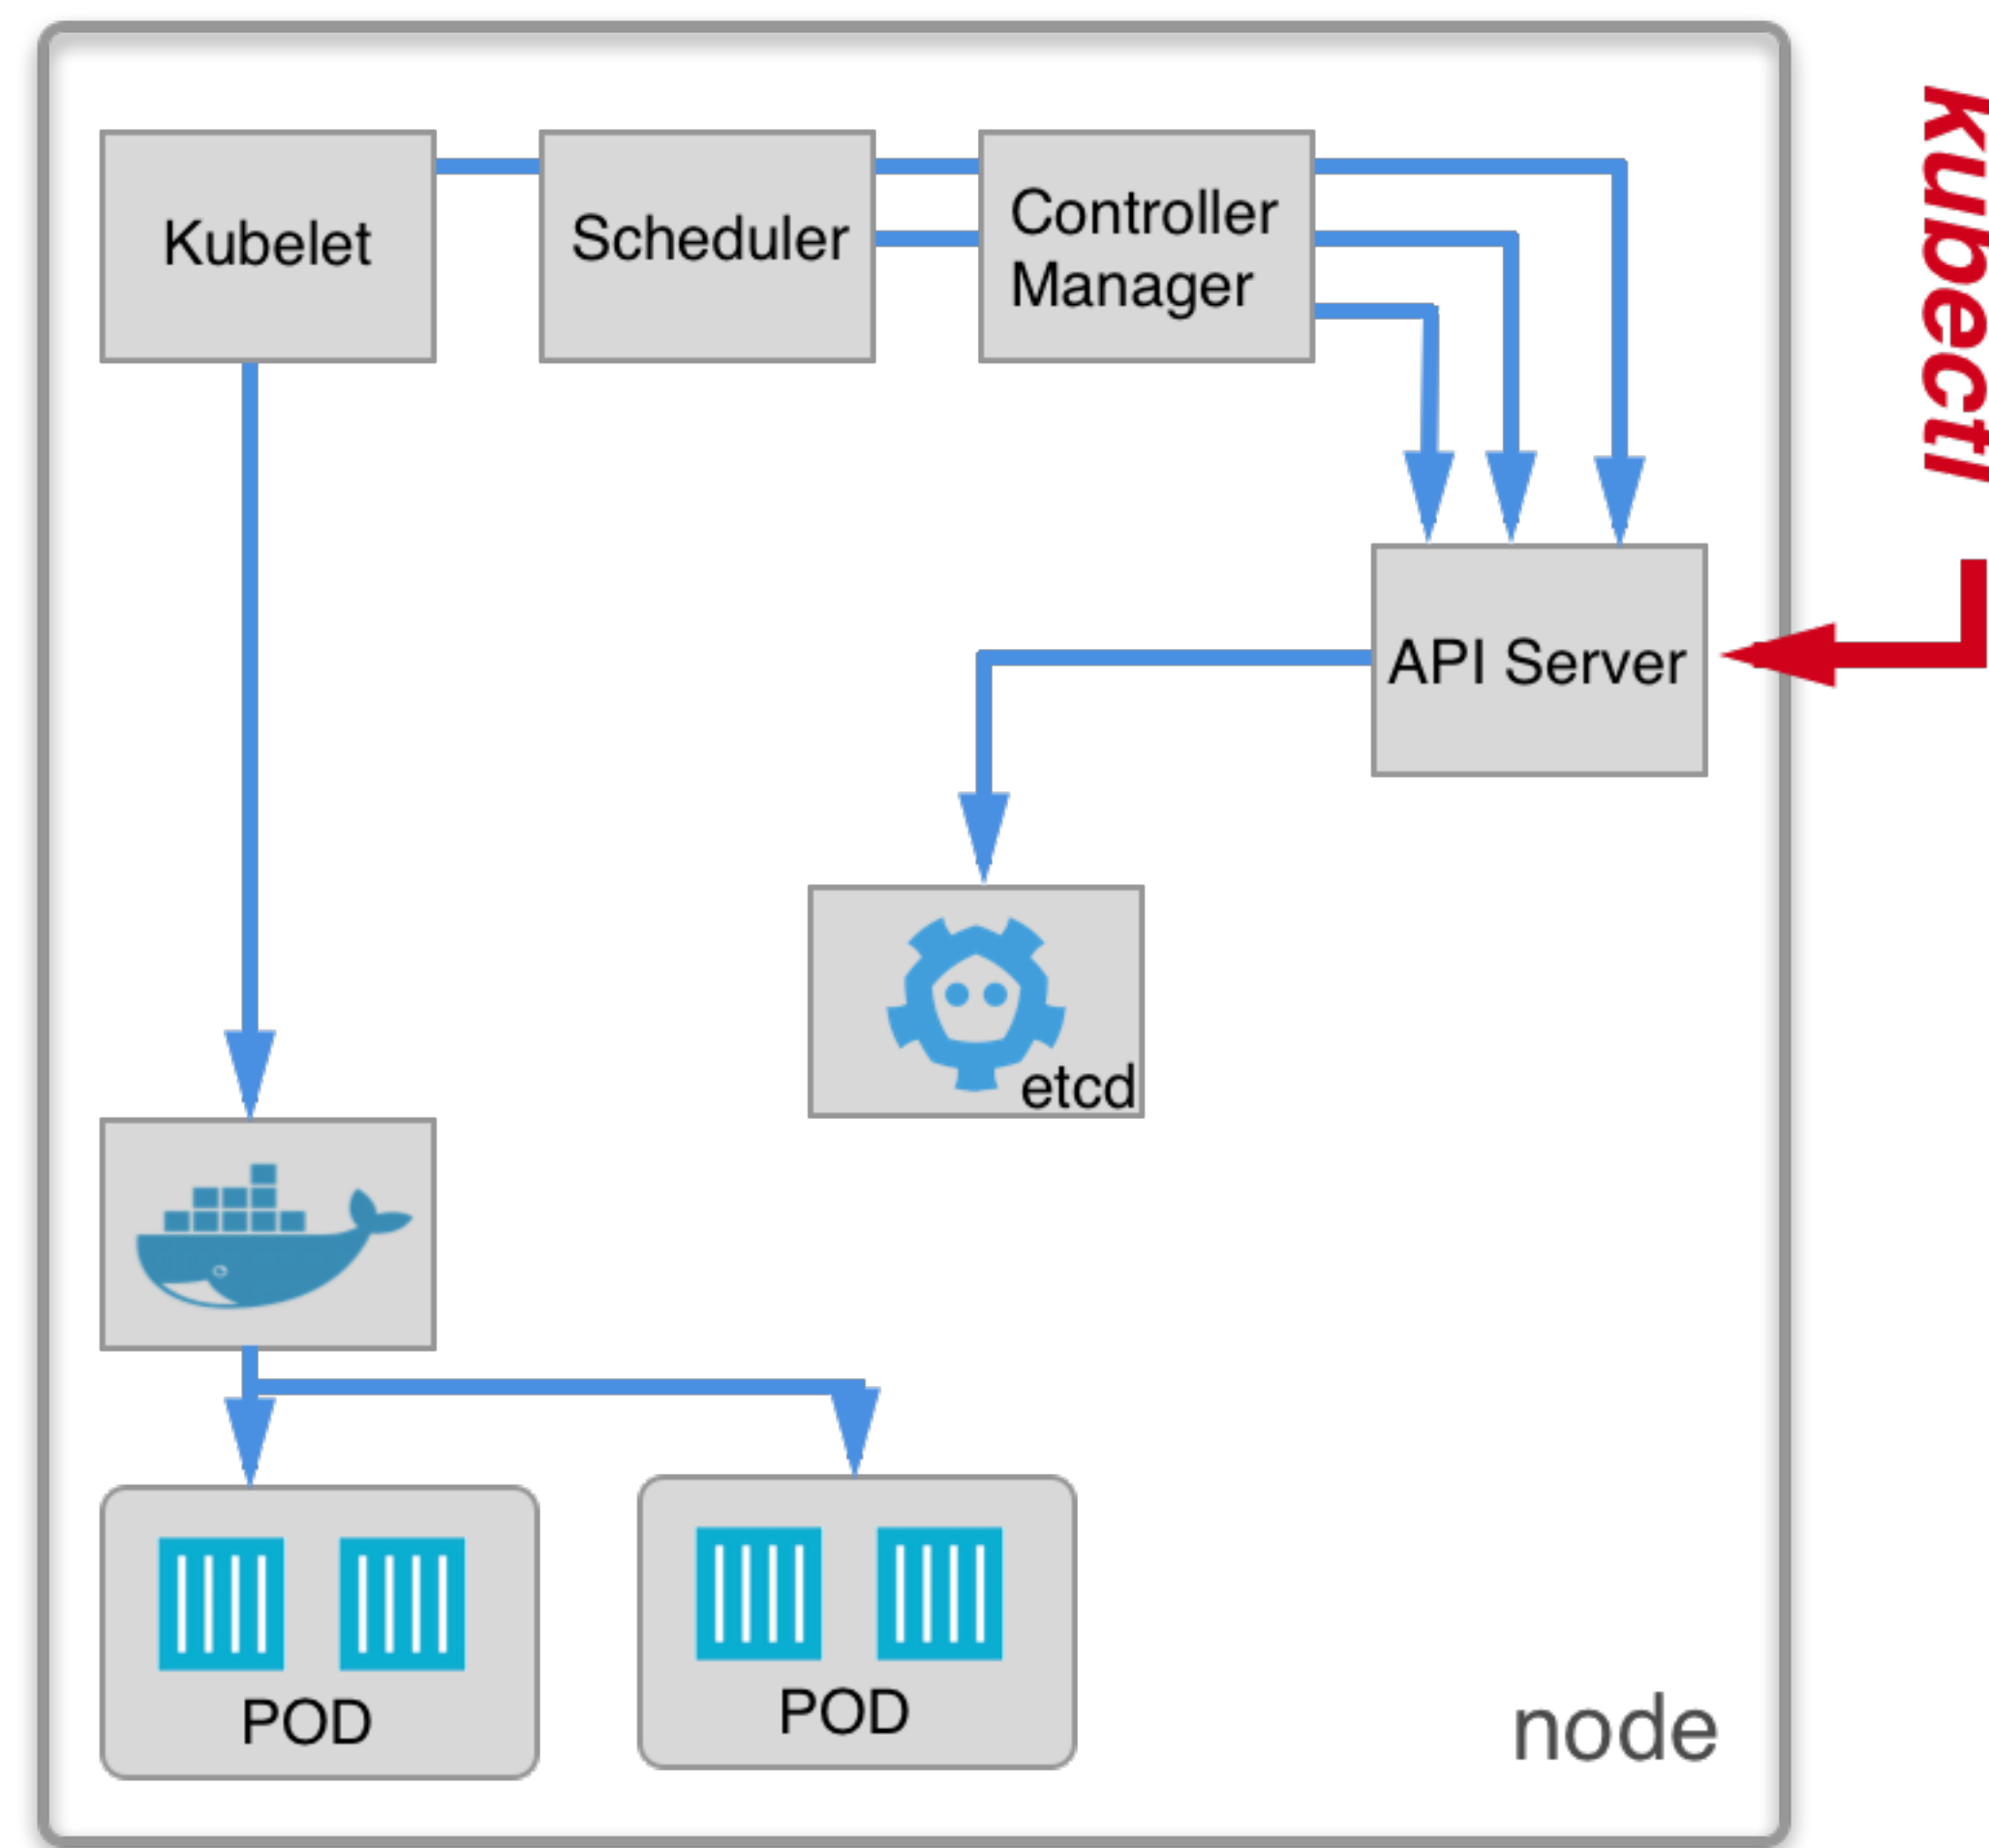
# BASIC FLOW

- kubectl apply -f deploy.yaml (5 replicas)
- CM creates replicaset object
- CM creates 5 pods' objects
- Scheduler assigns nodes

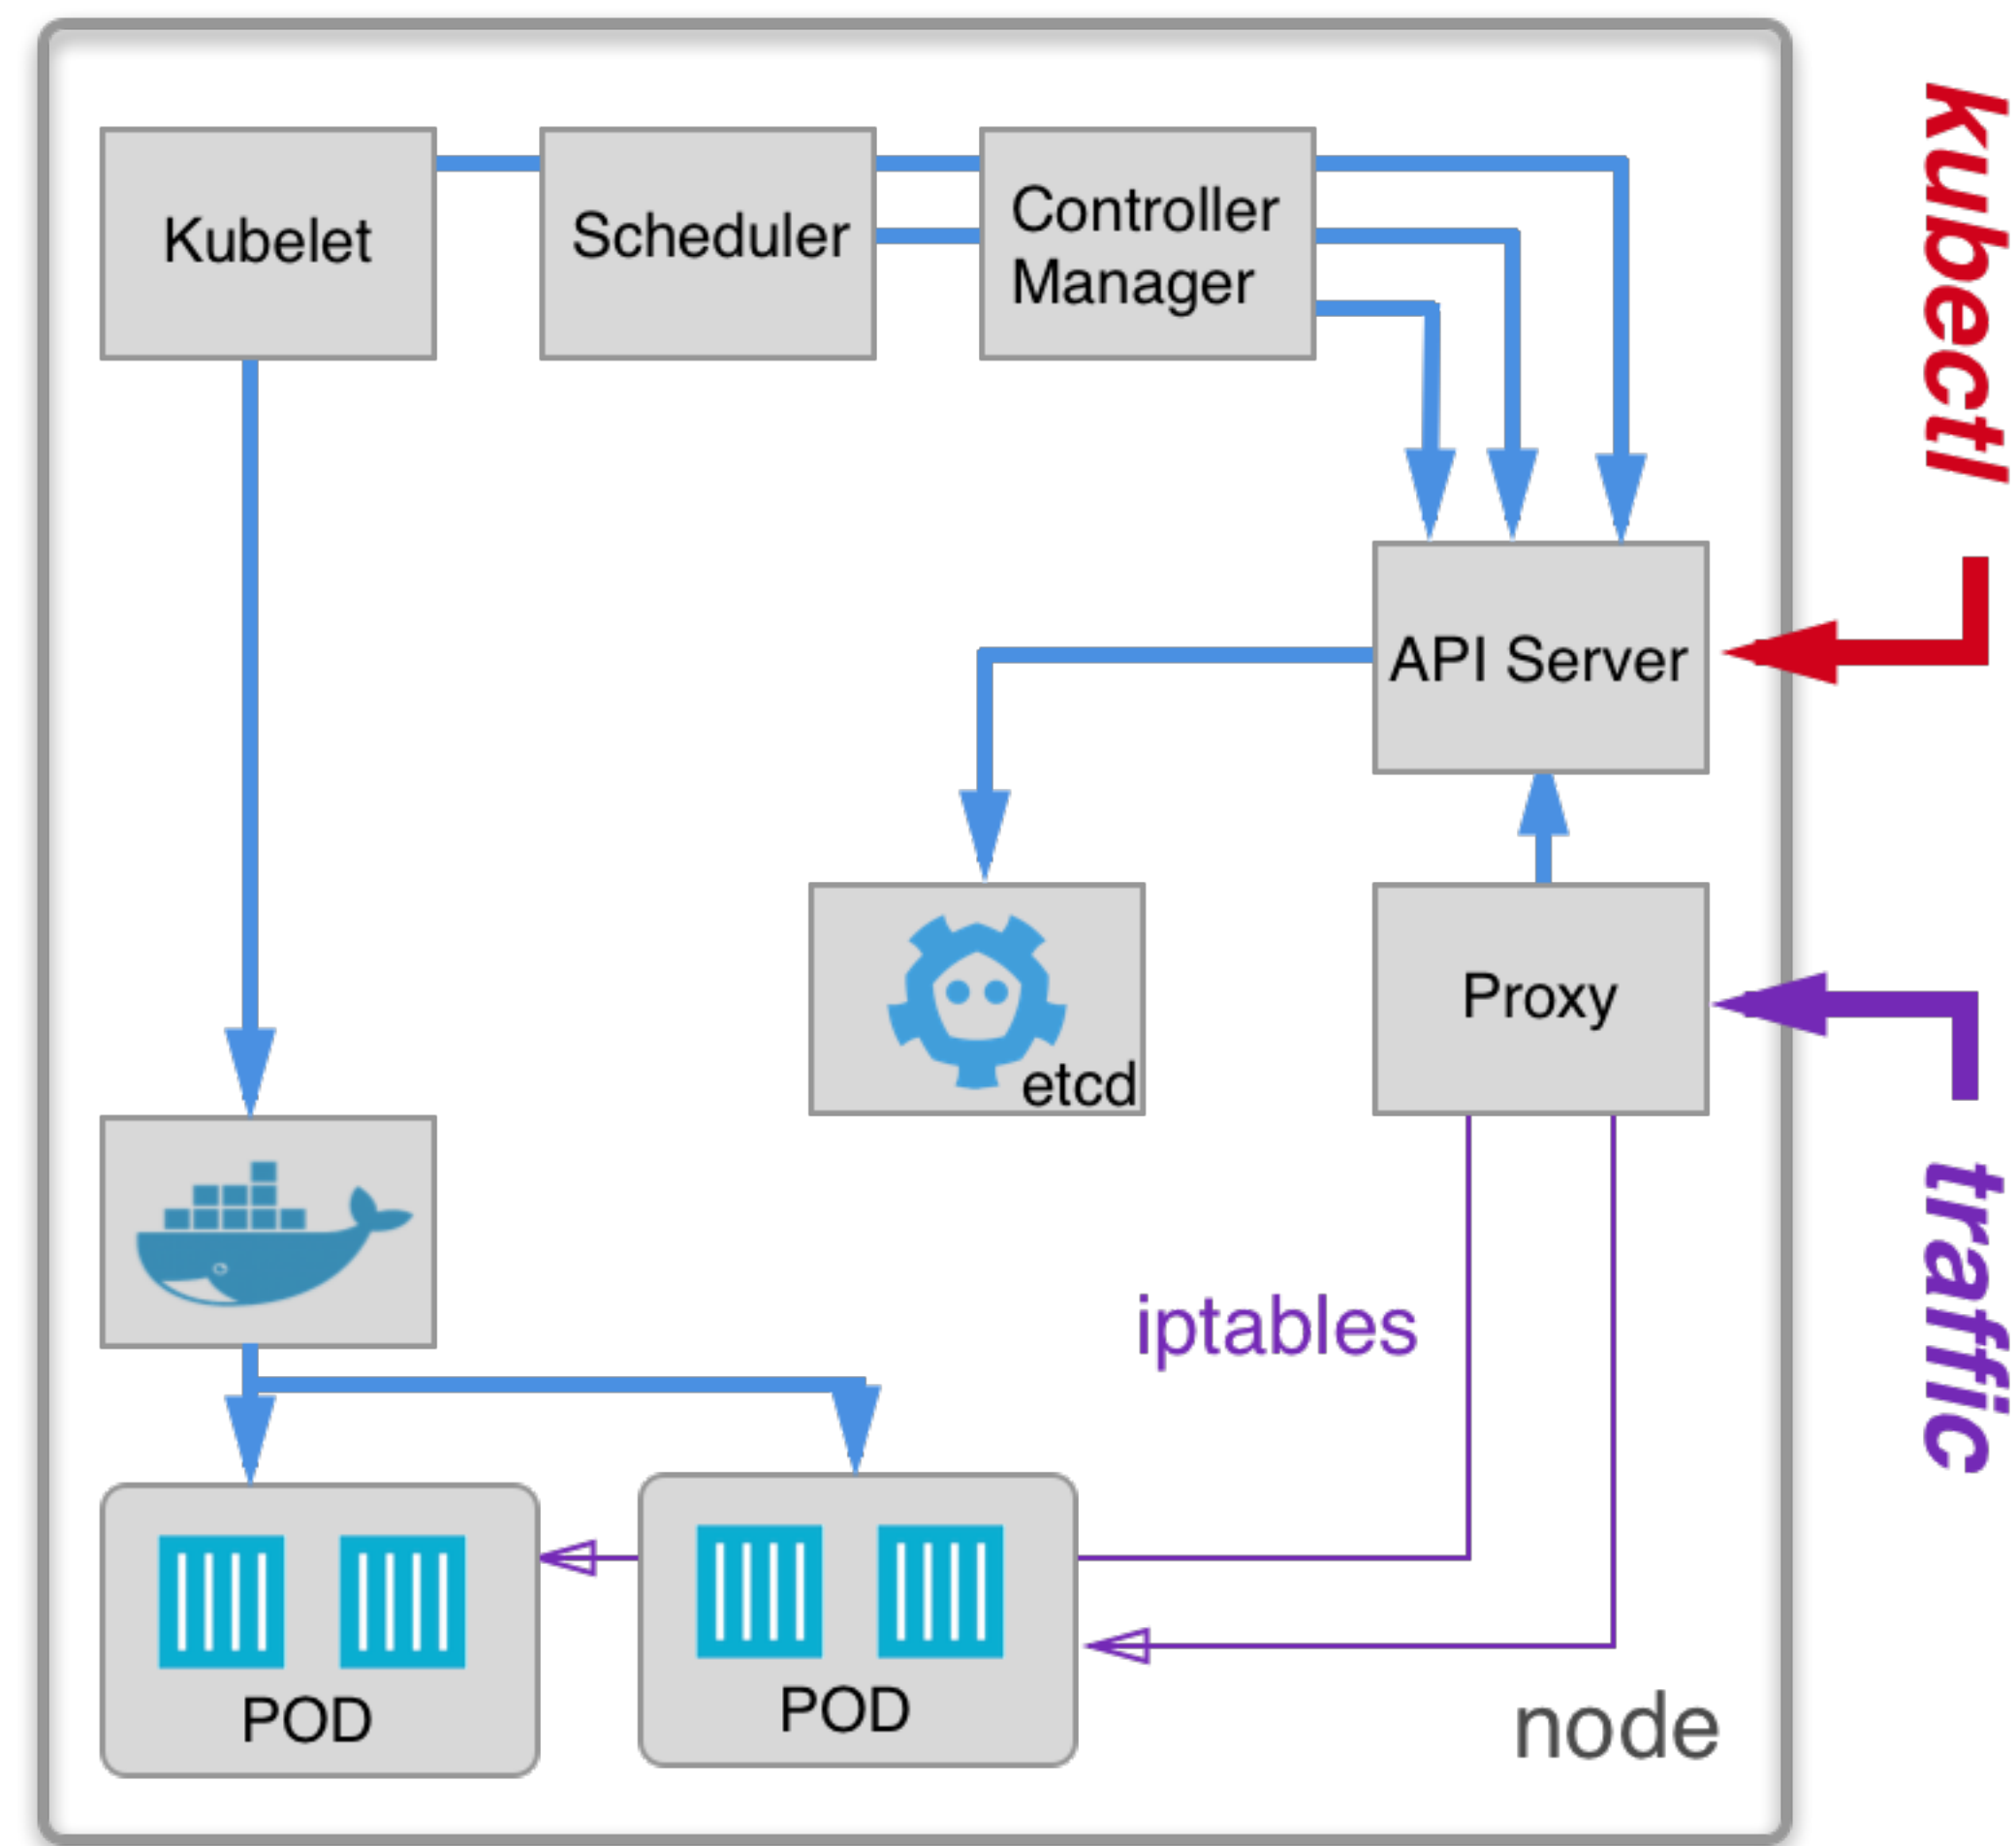


# BASIC FLOW

- `kubectl apply -f deploy.yaml` (5 replicas)
- CM creates replicaset object
- CM creates 5 pods' objects
- Scheduler assigns nodes
- Kubelets create containers



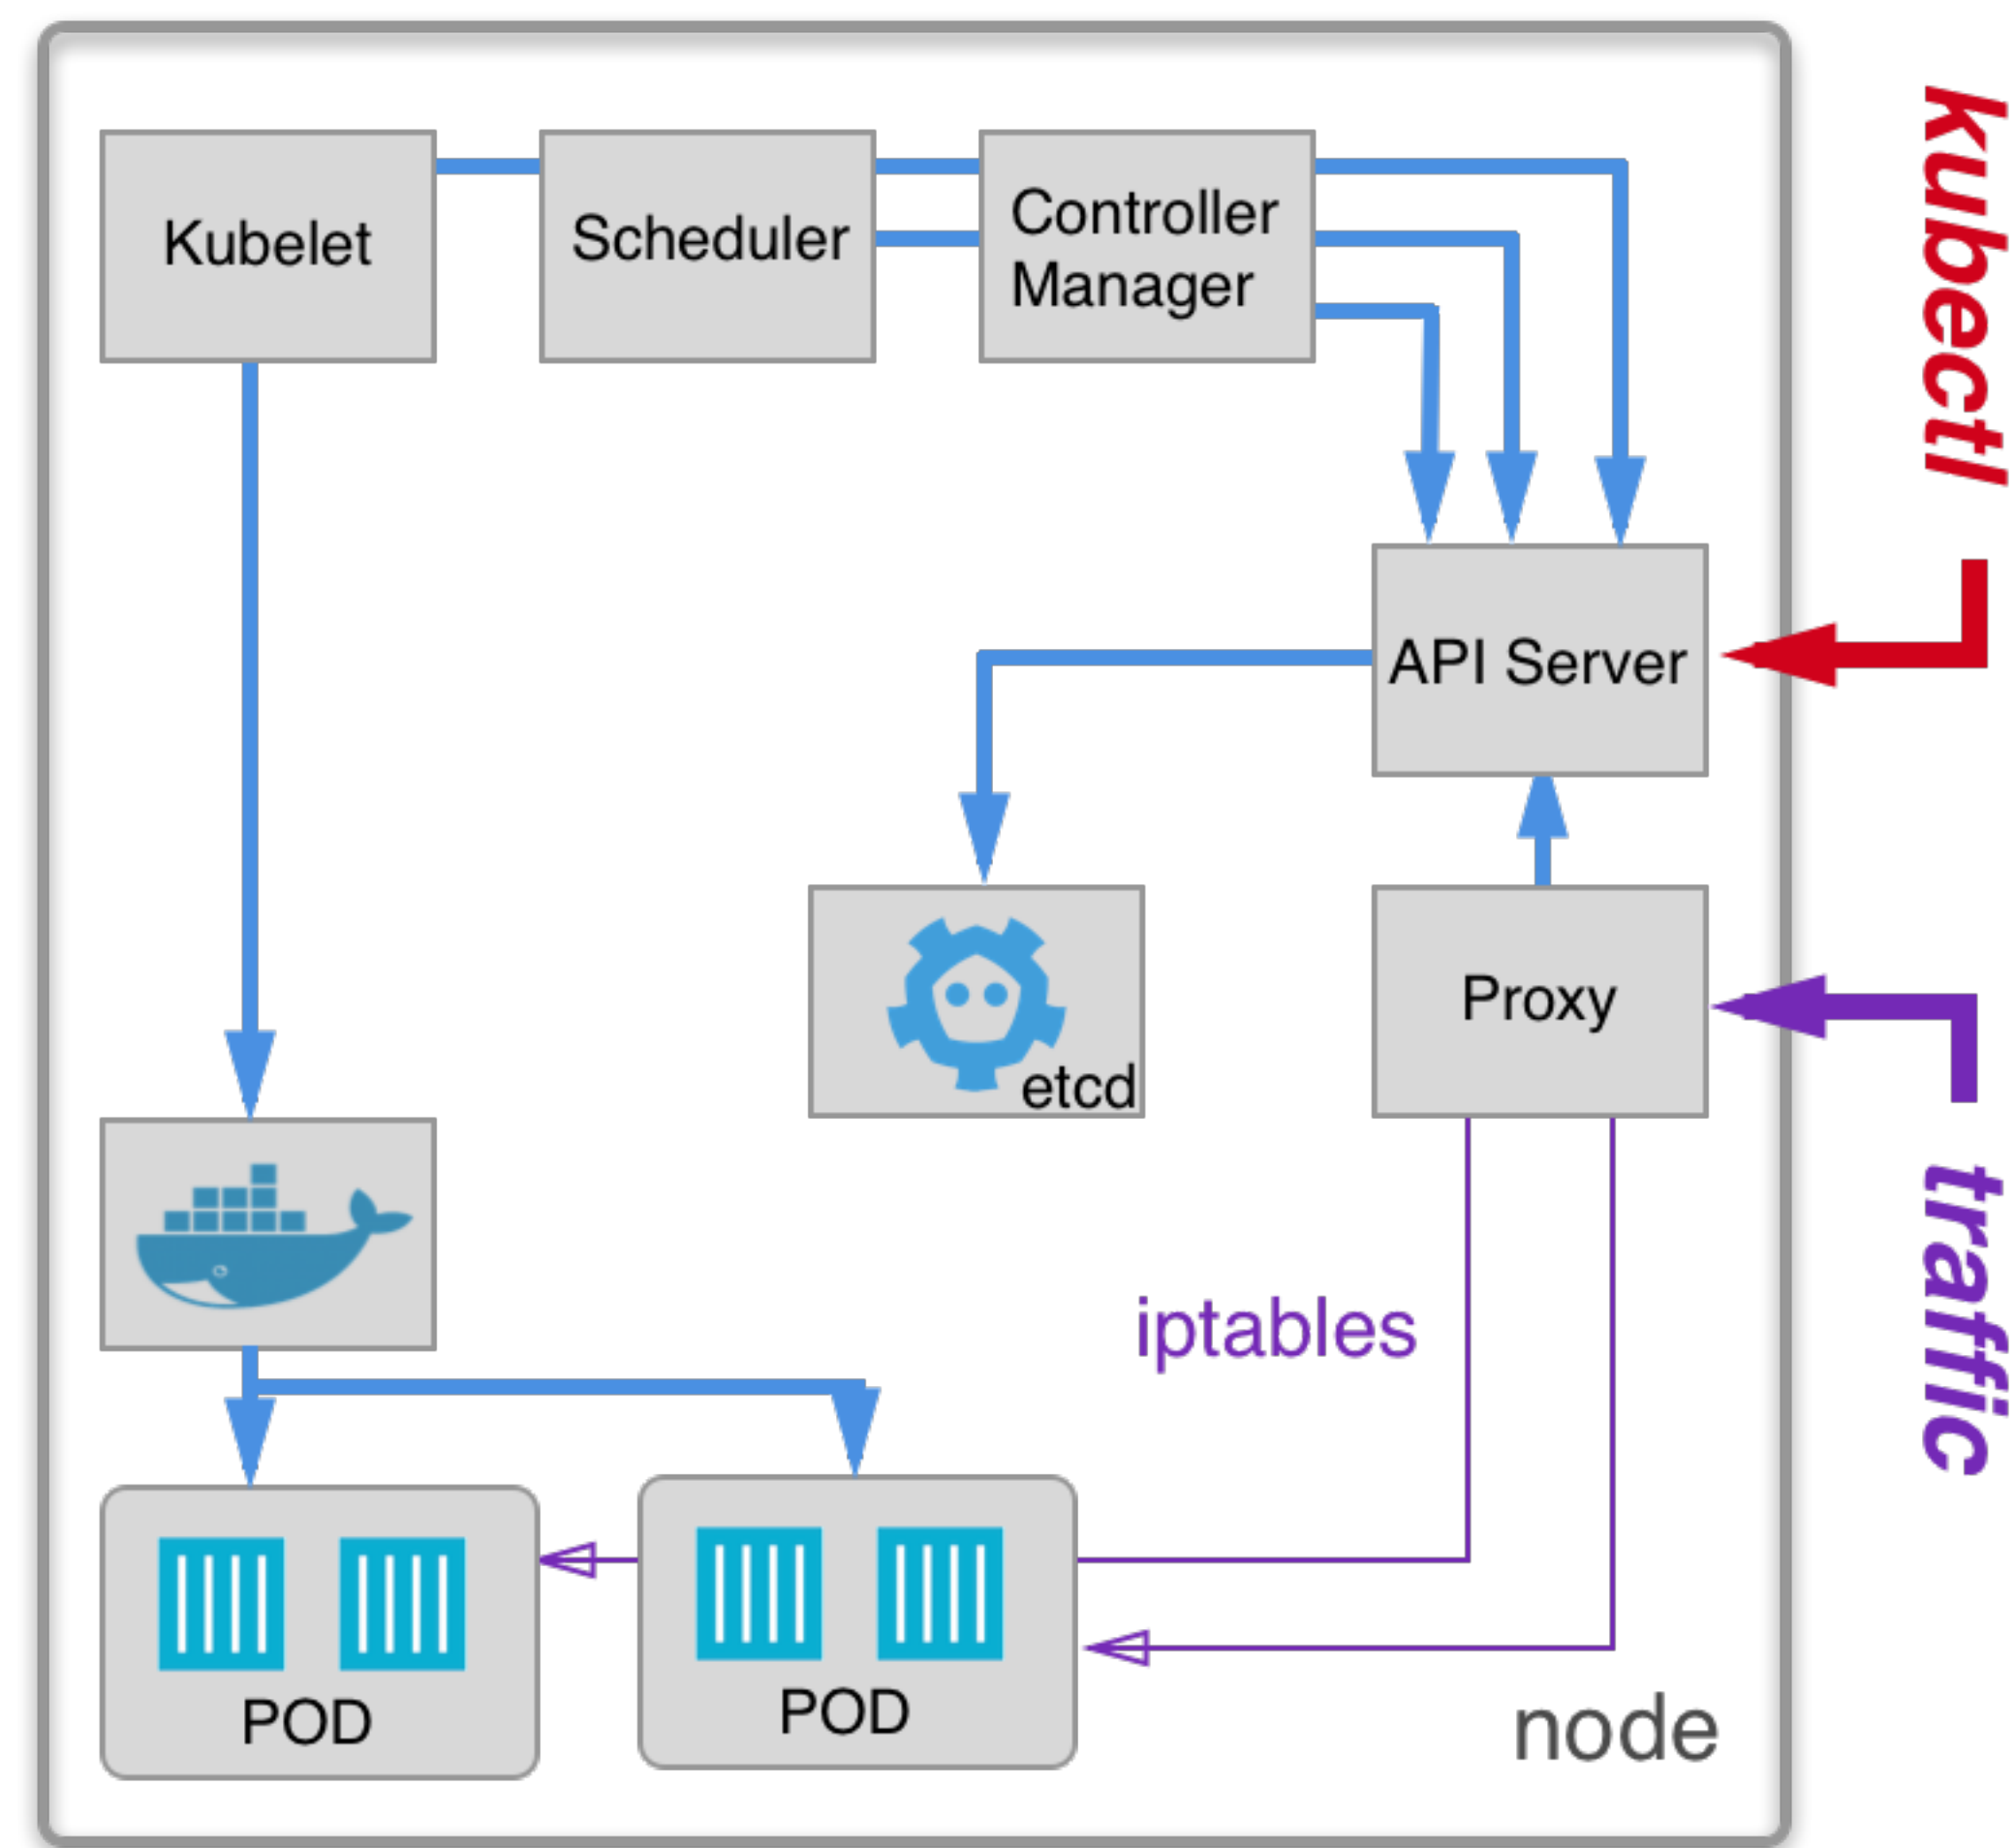
# PROXY





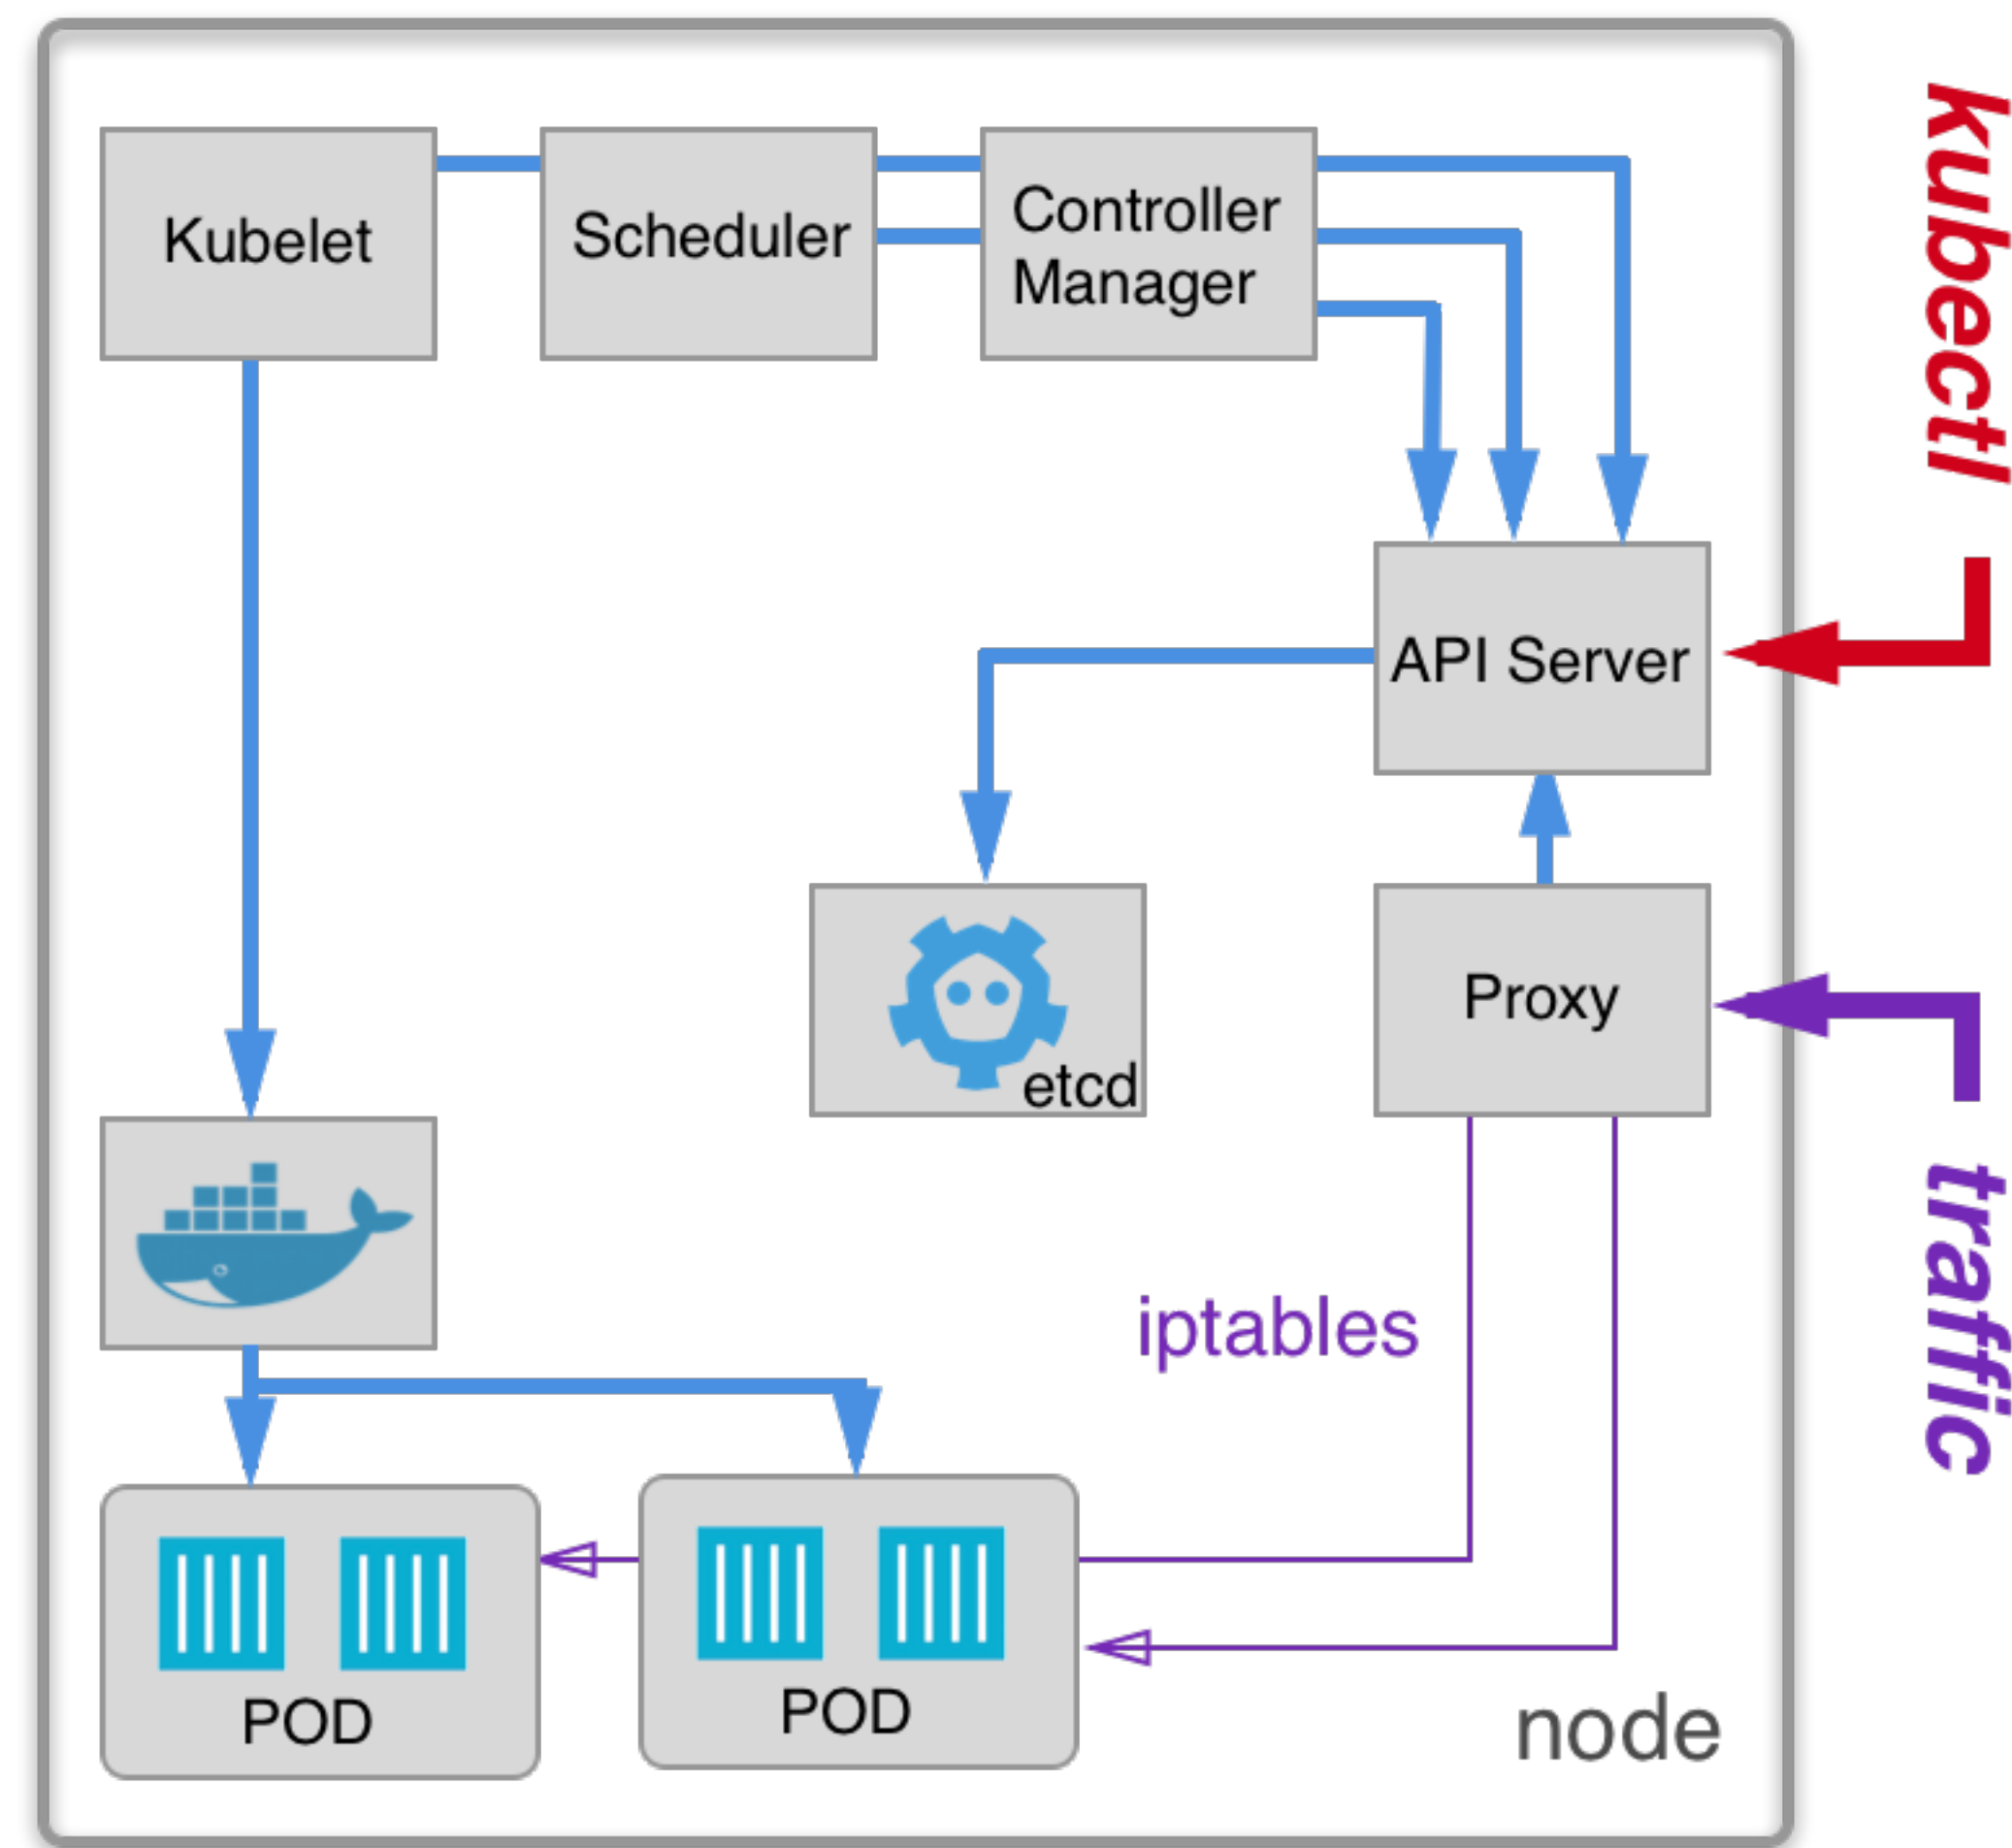
# PROXY

- responsible for managing Services



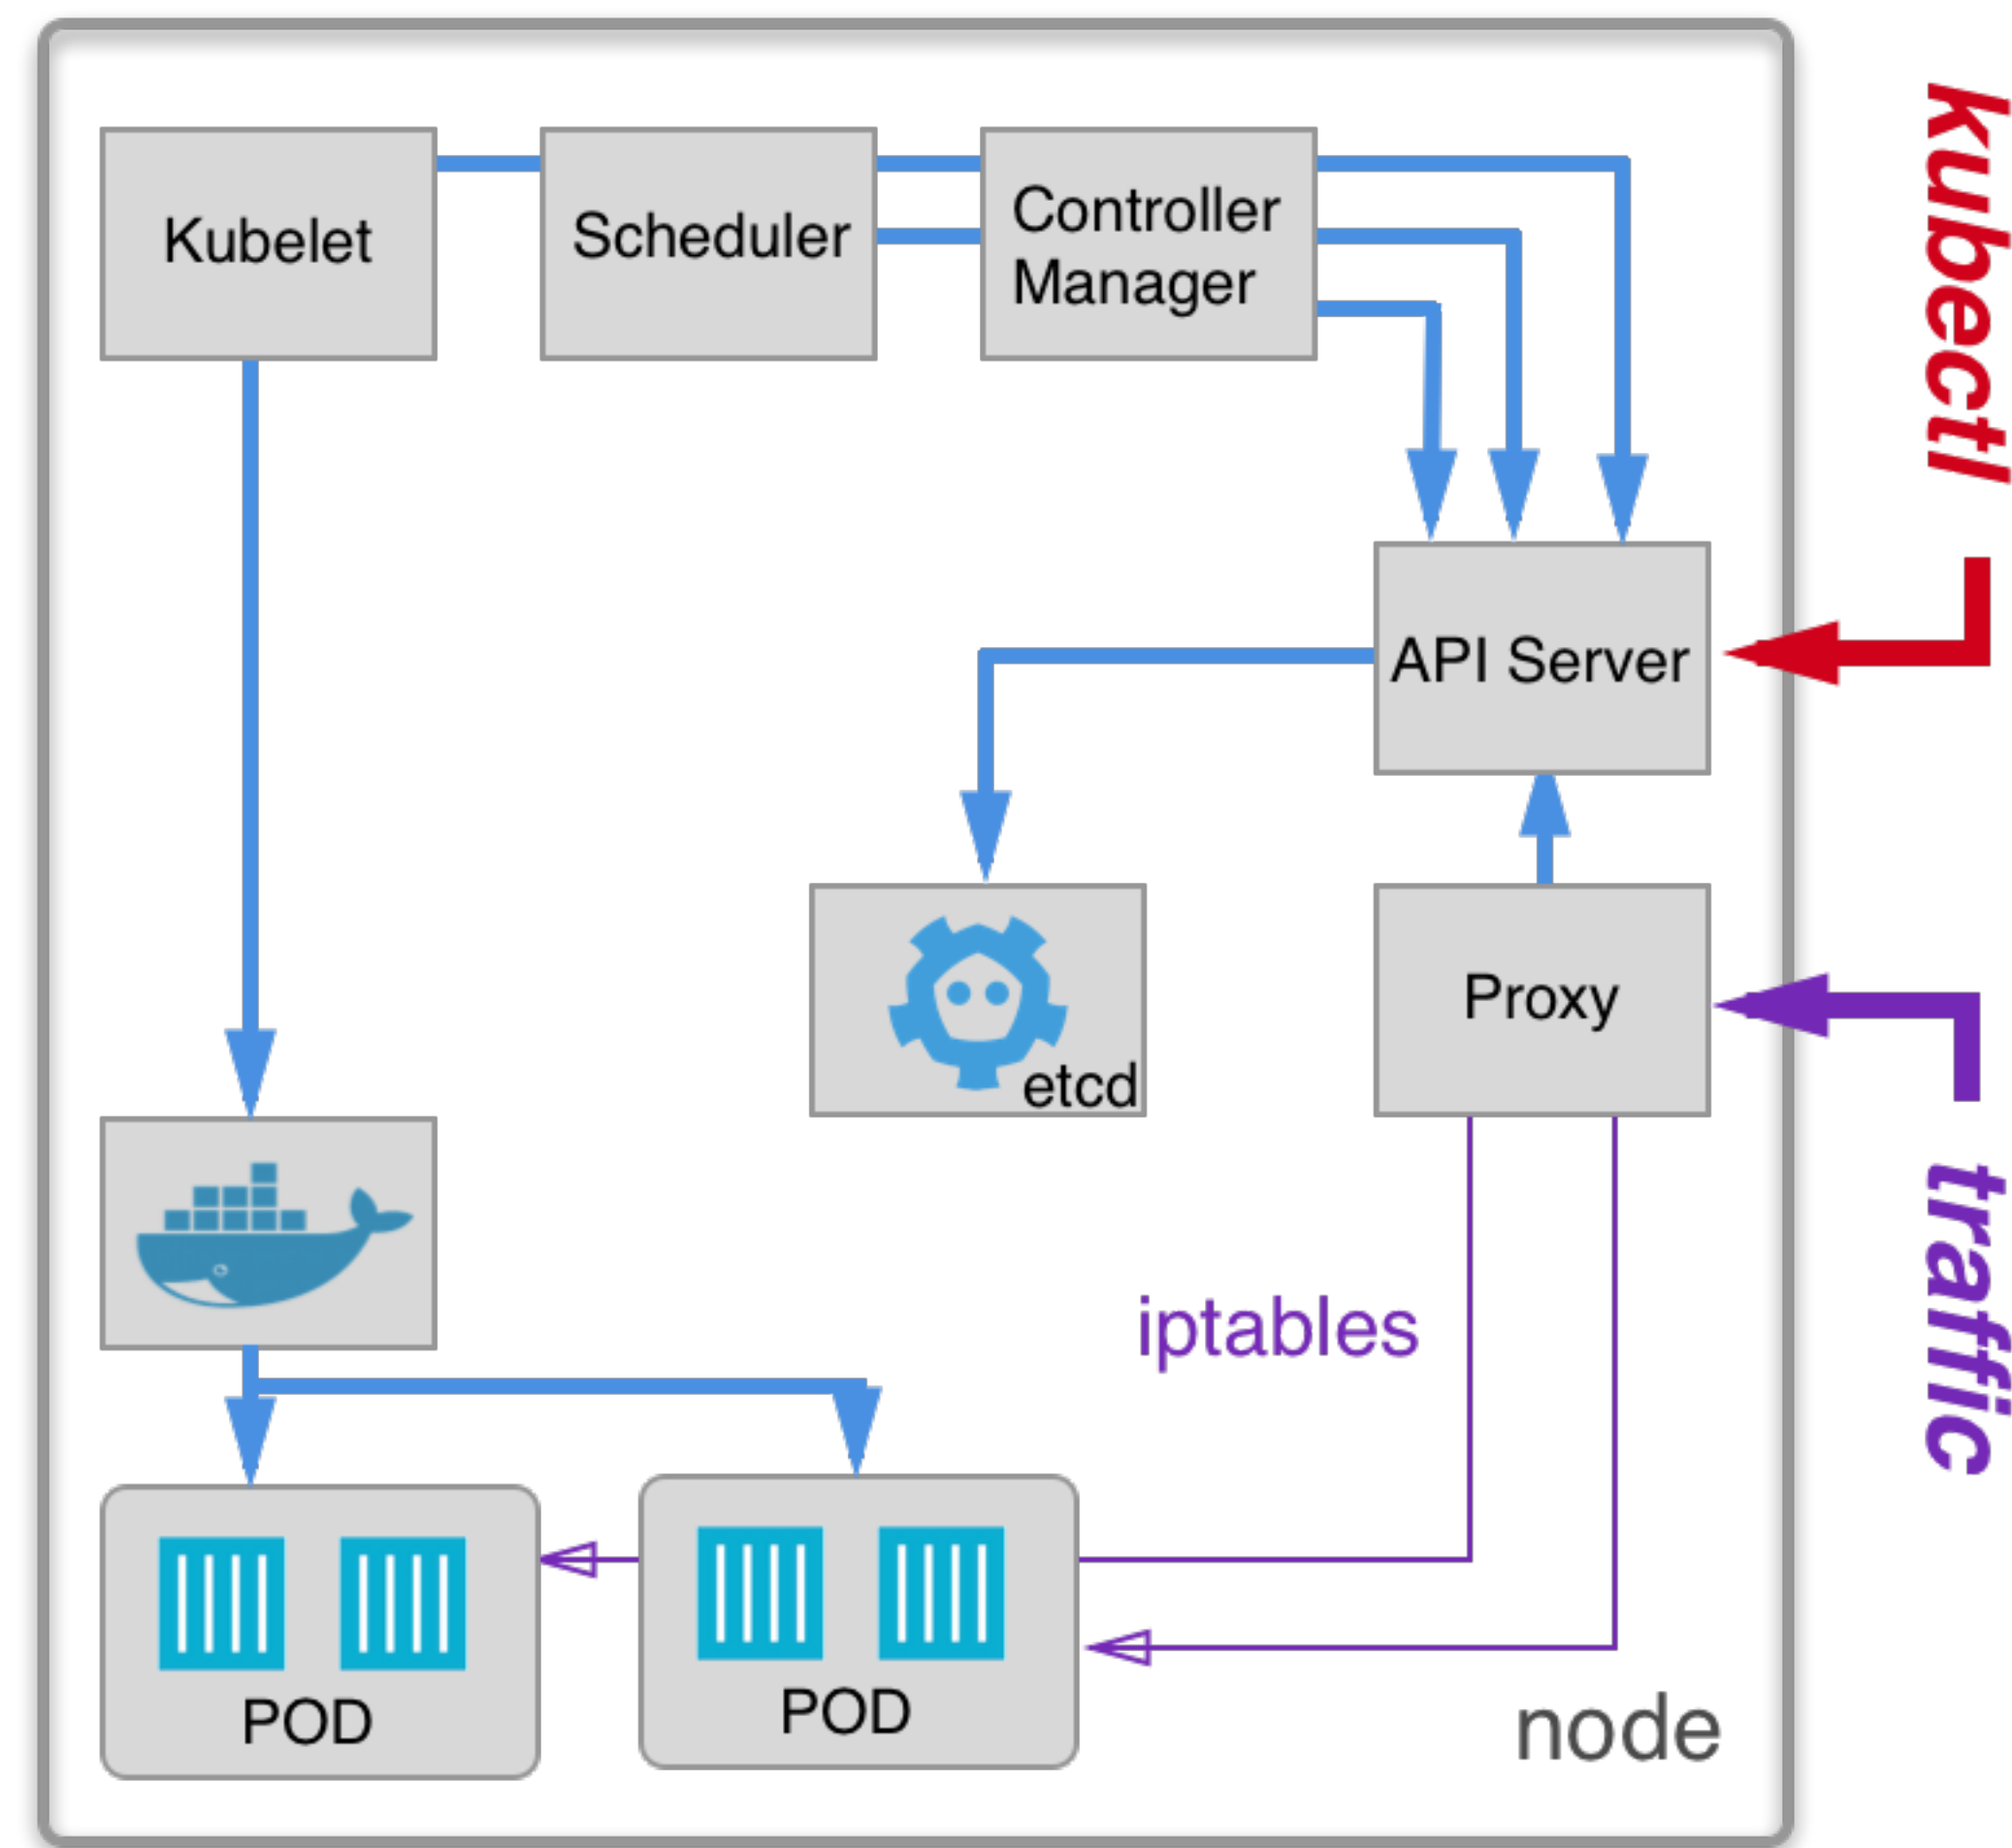
# PROXY

- responsible for managing Services
- accepts external traffic on NodePorts and distributes it to pods



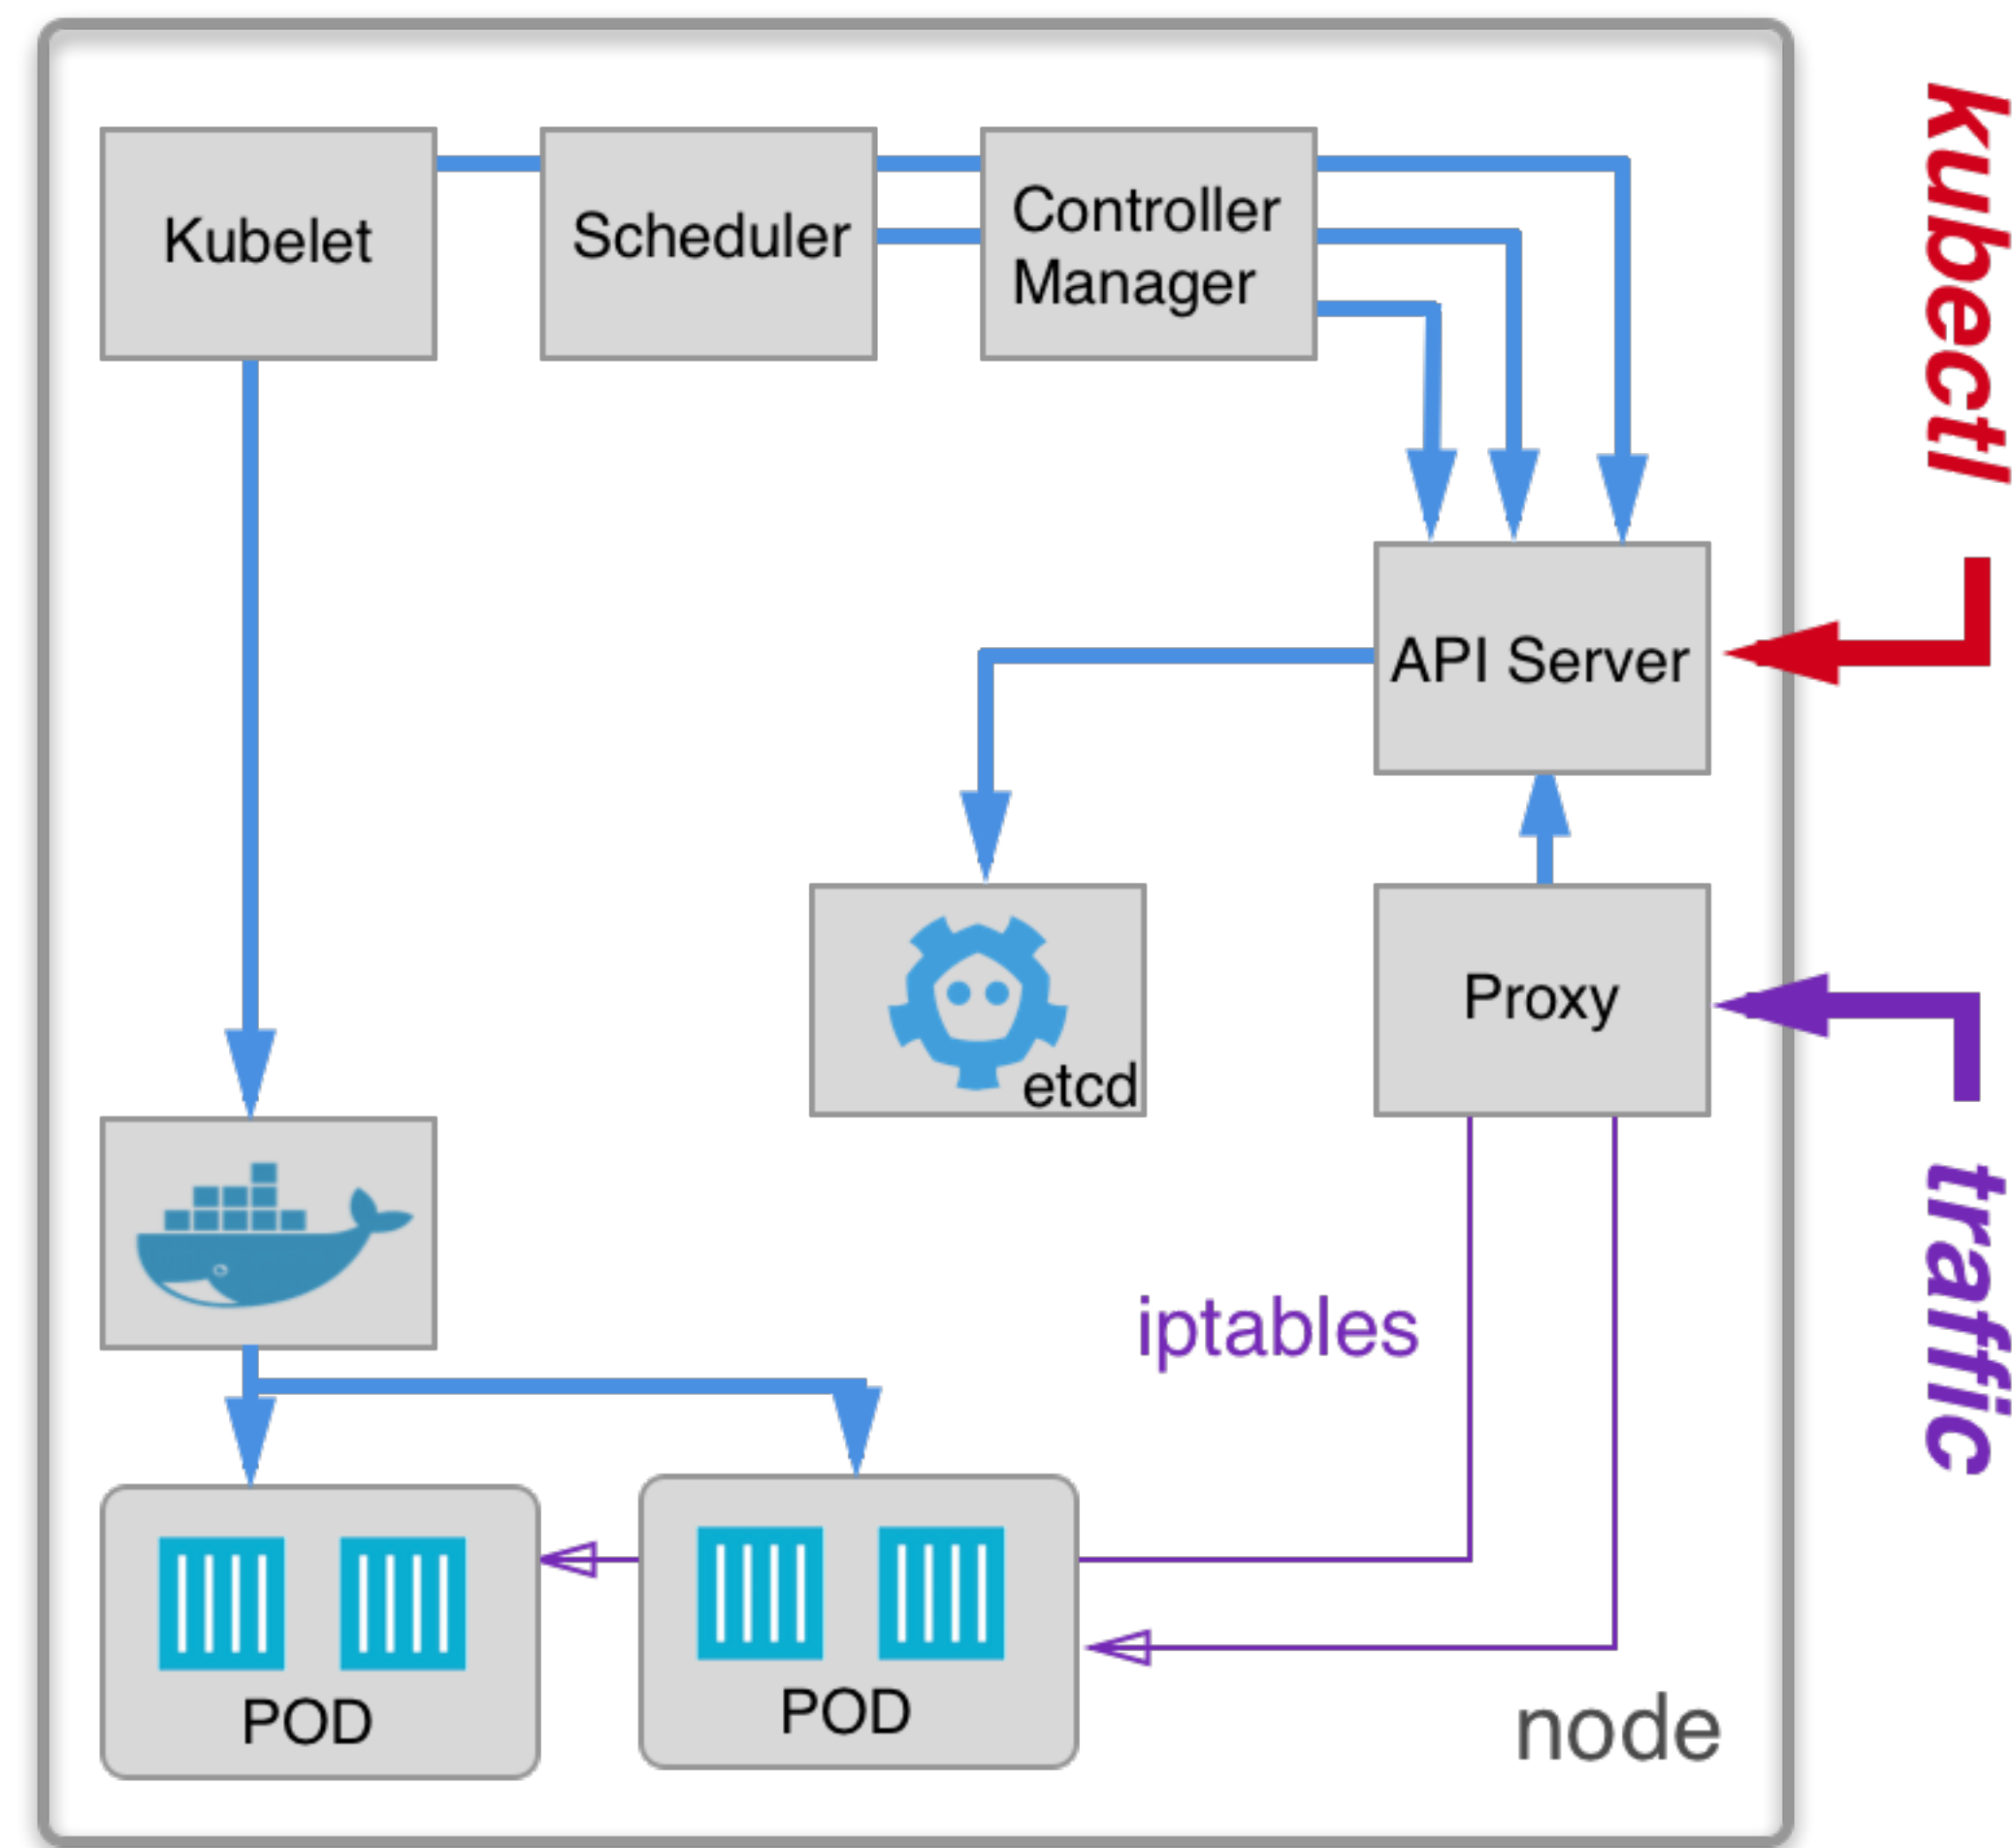
# PROXY

- responsible for managing Services
- accepts external traffic on NodePorts and distributes it to pods
- load balances traffic between Service's pods



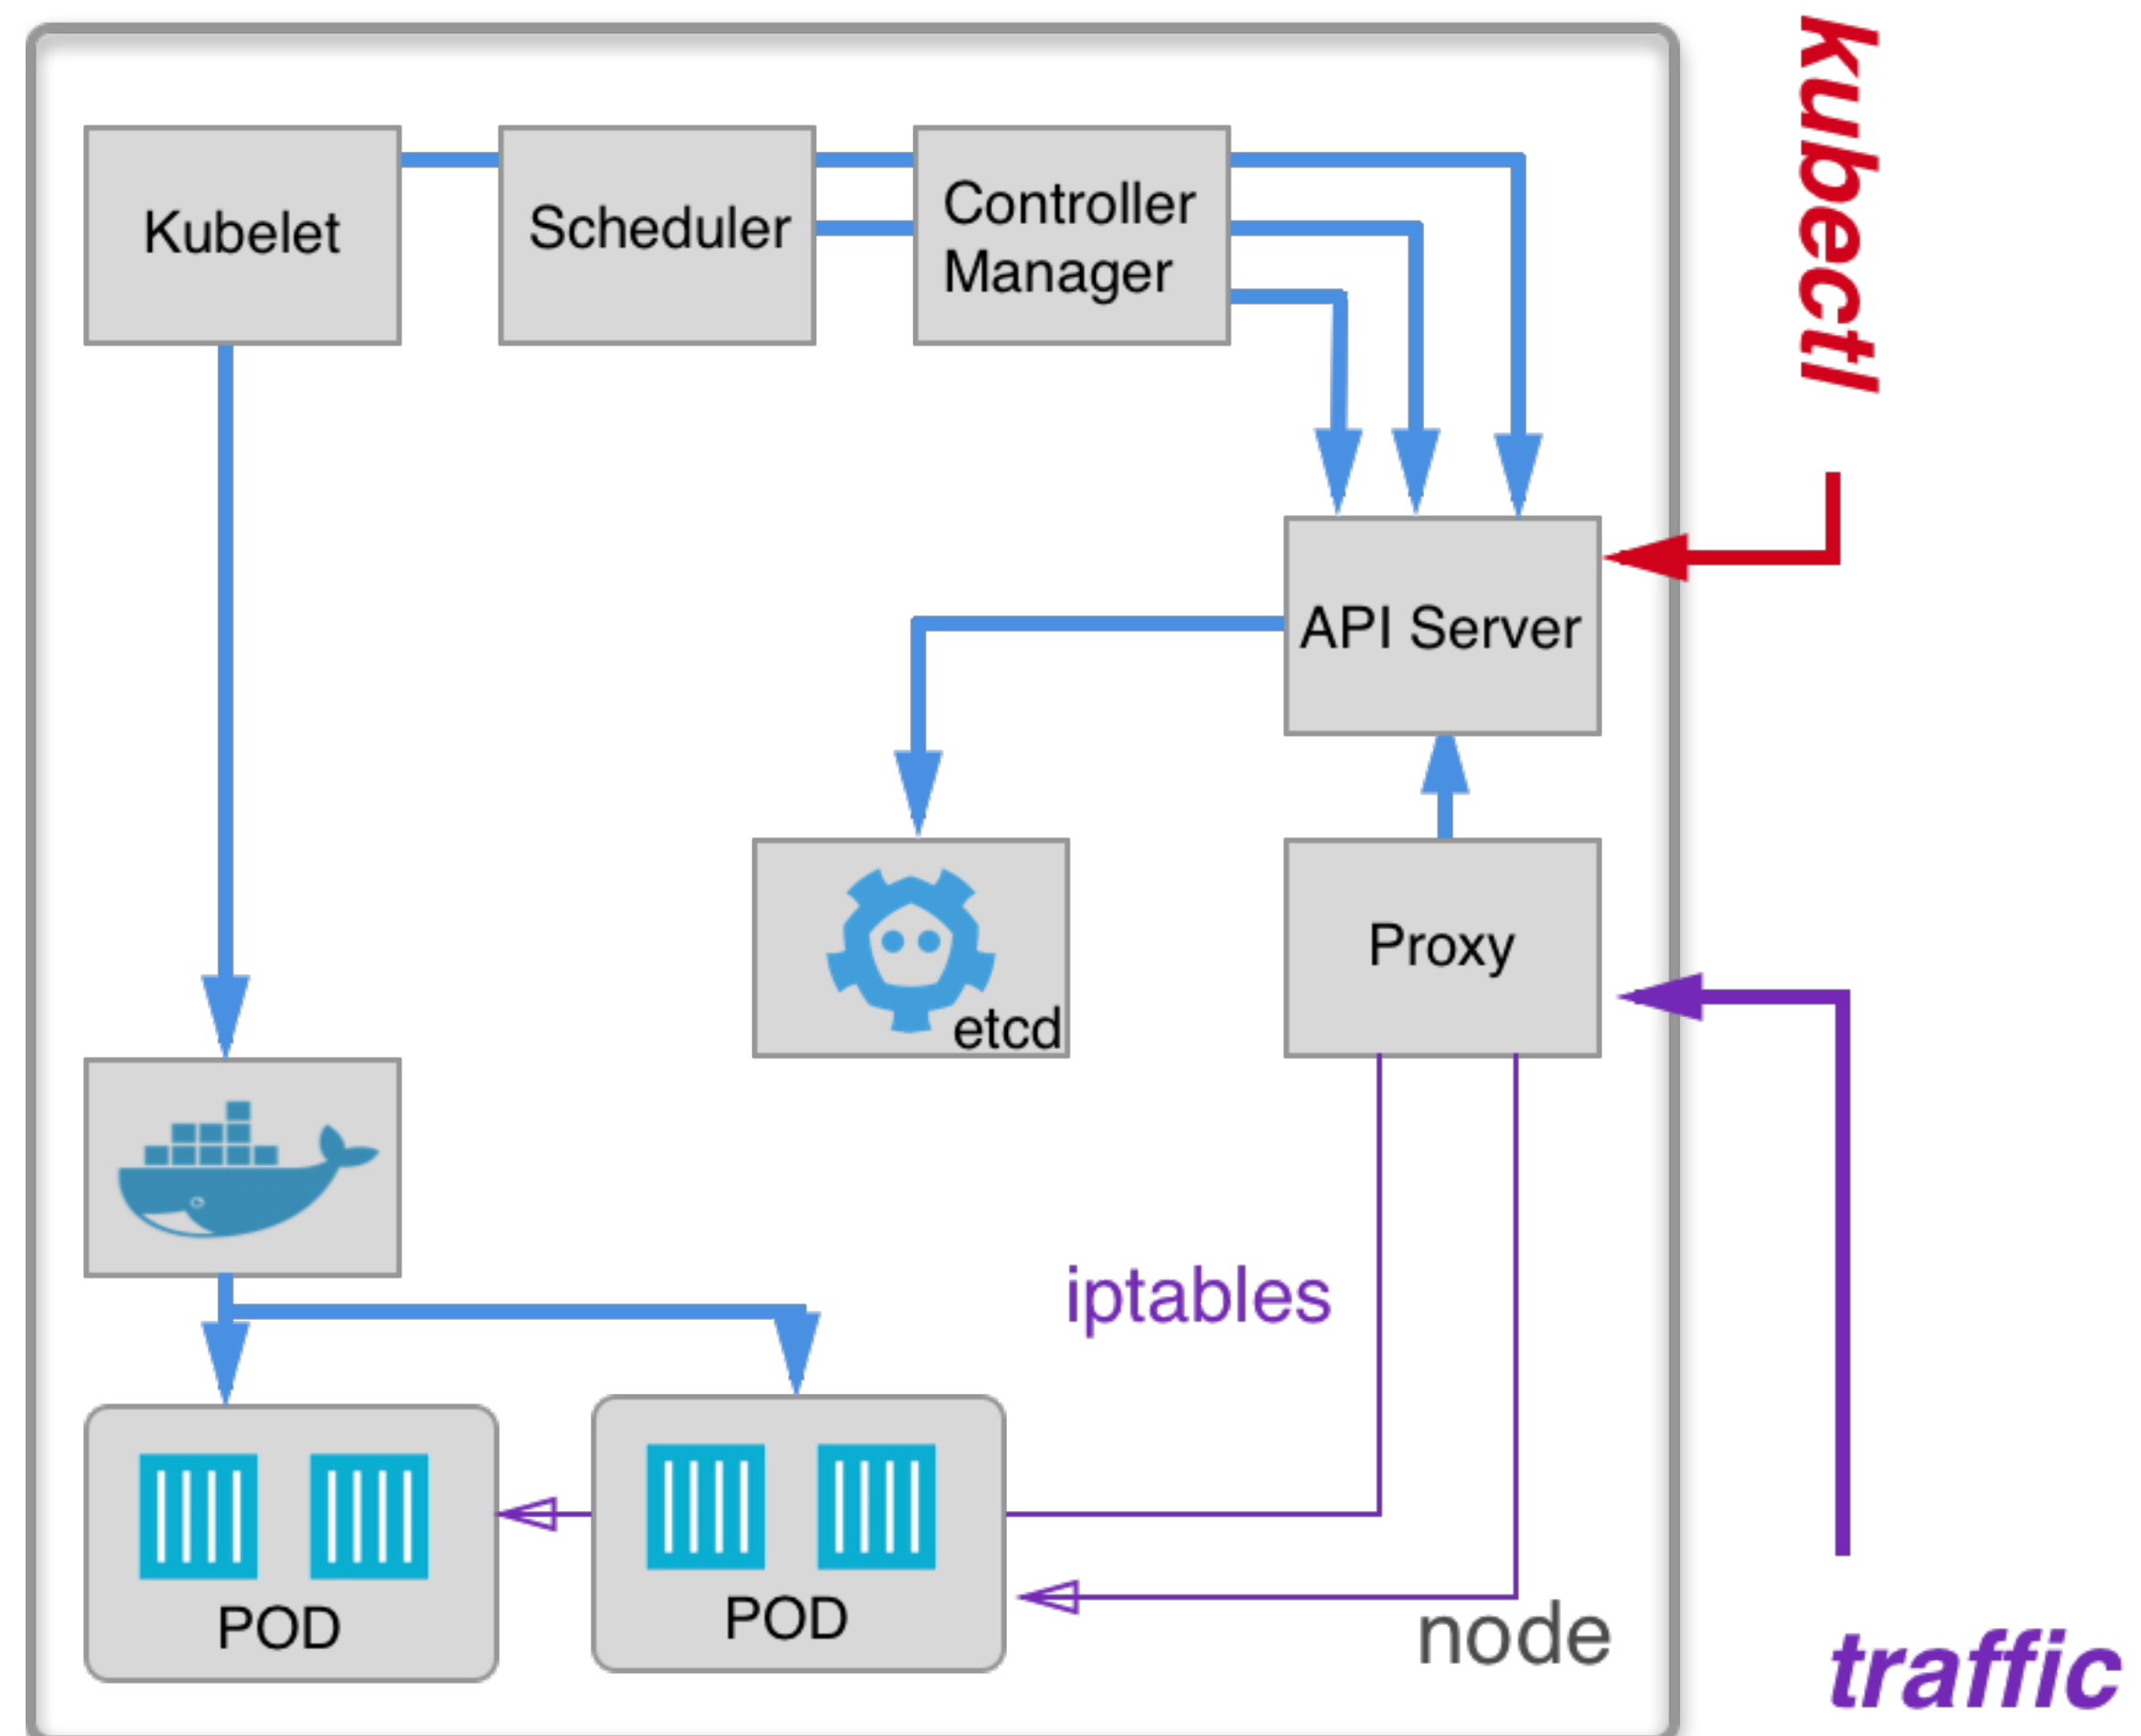
# PROXY

- responsible for managing Services
- accepts external traffic on NodePorts and distributes it to pods
- load balances traffic between Service's pods
- everything is done using *iptables*

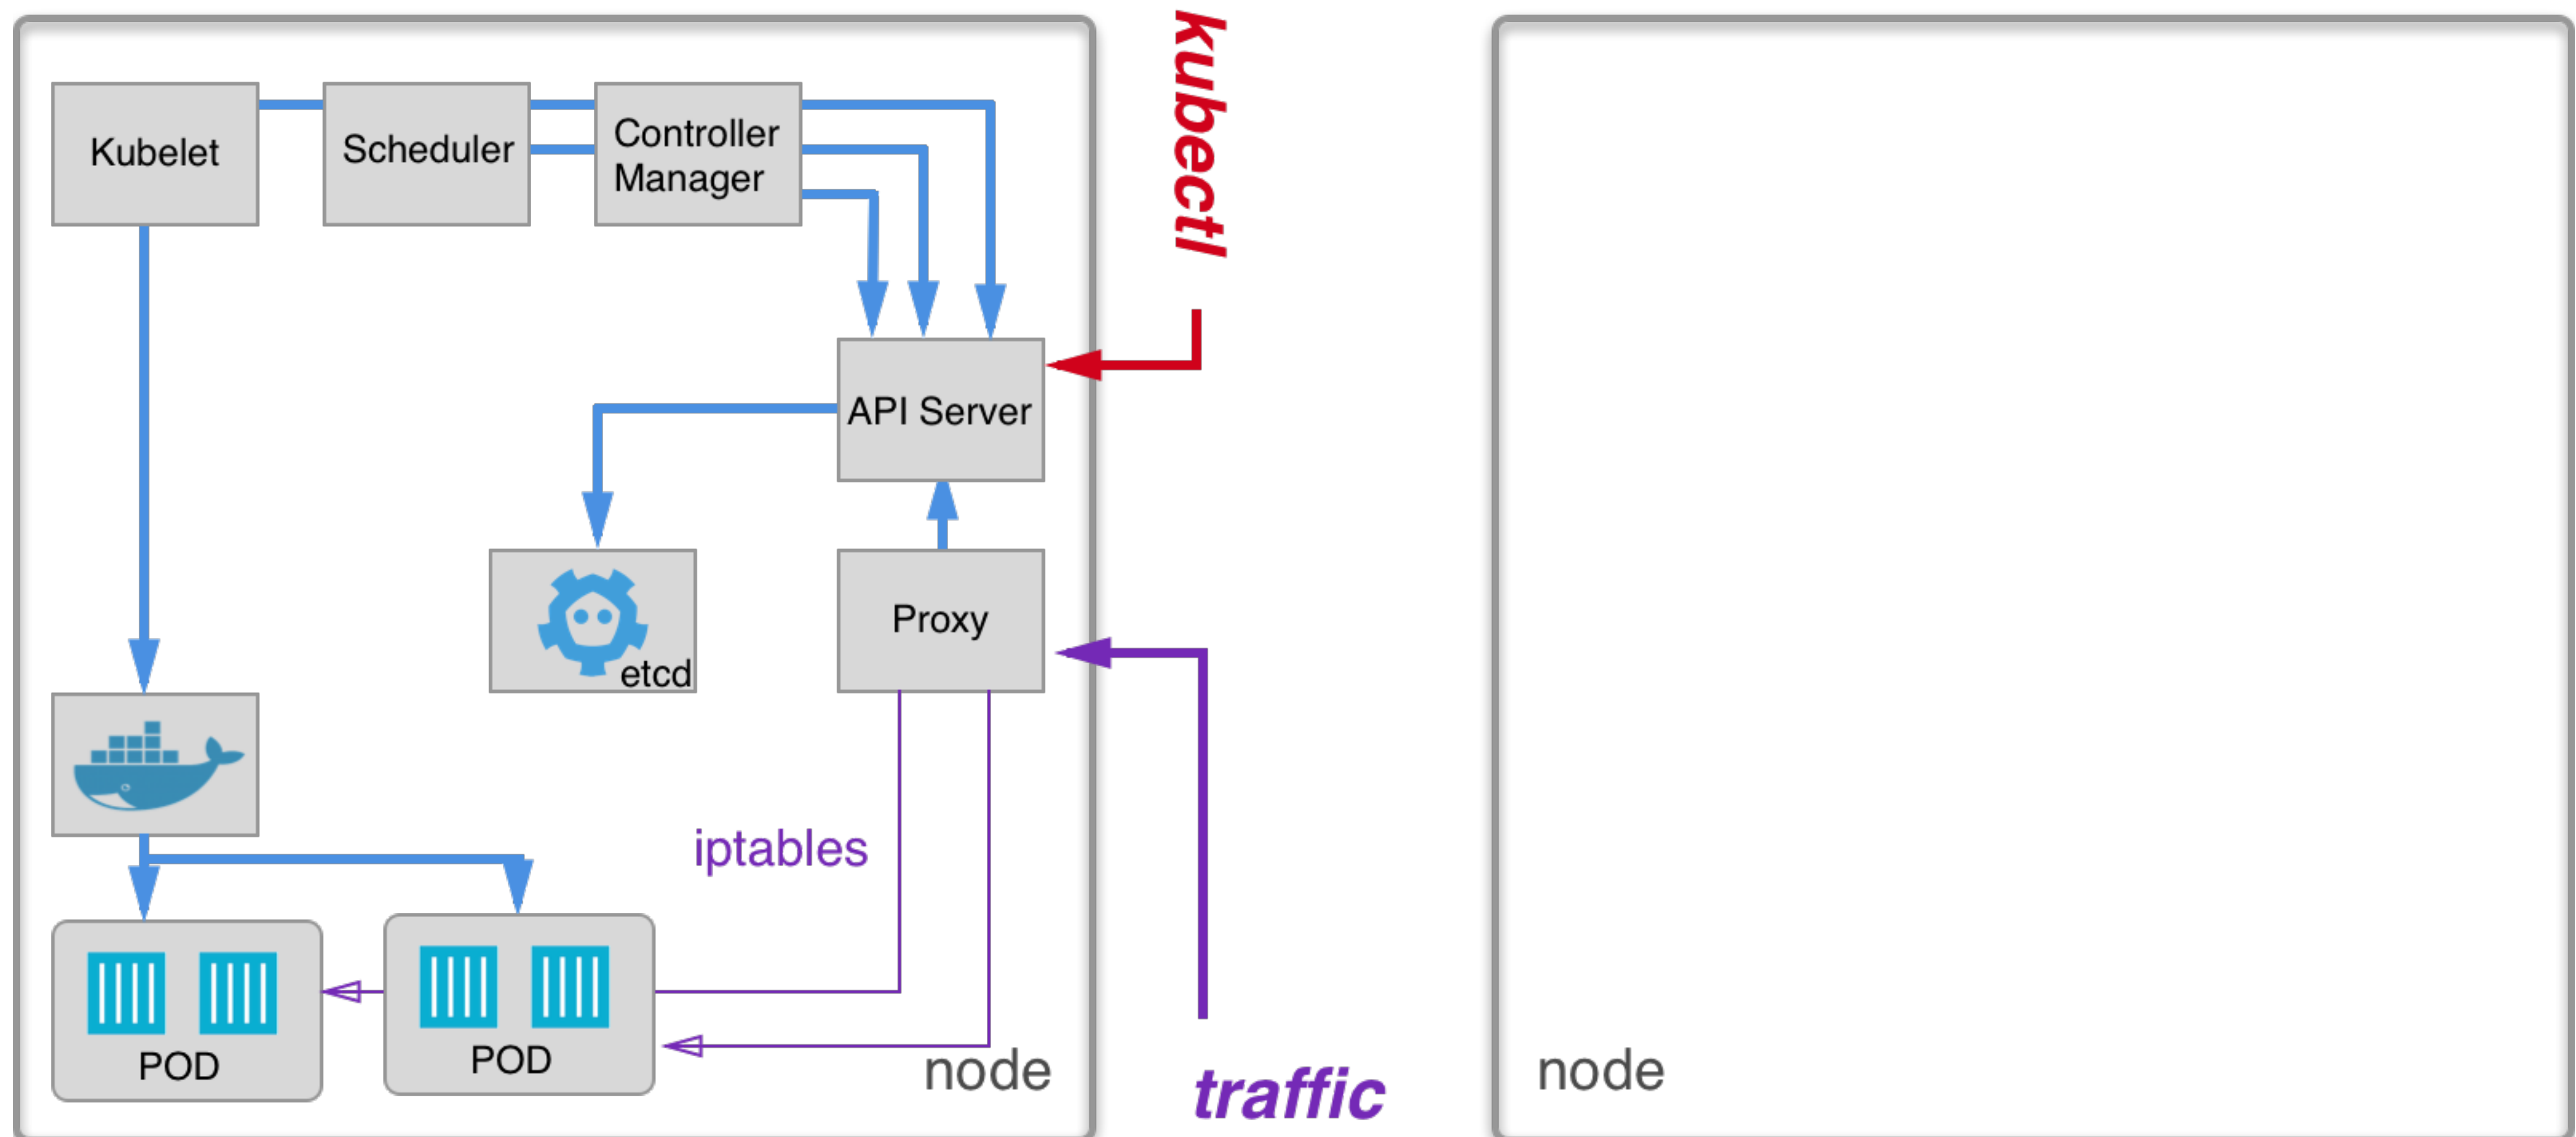


# DEMO

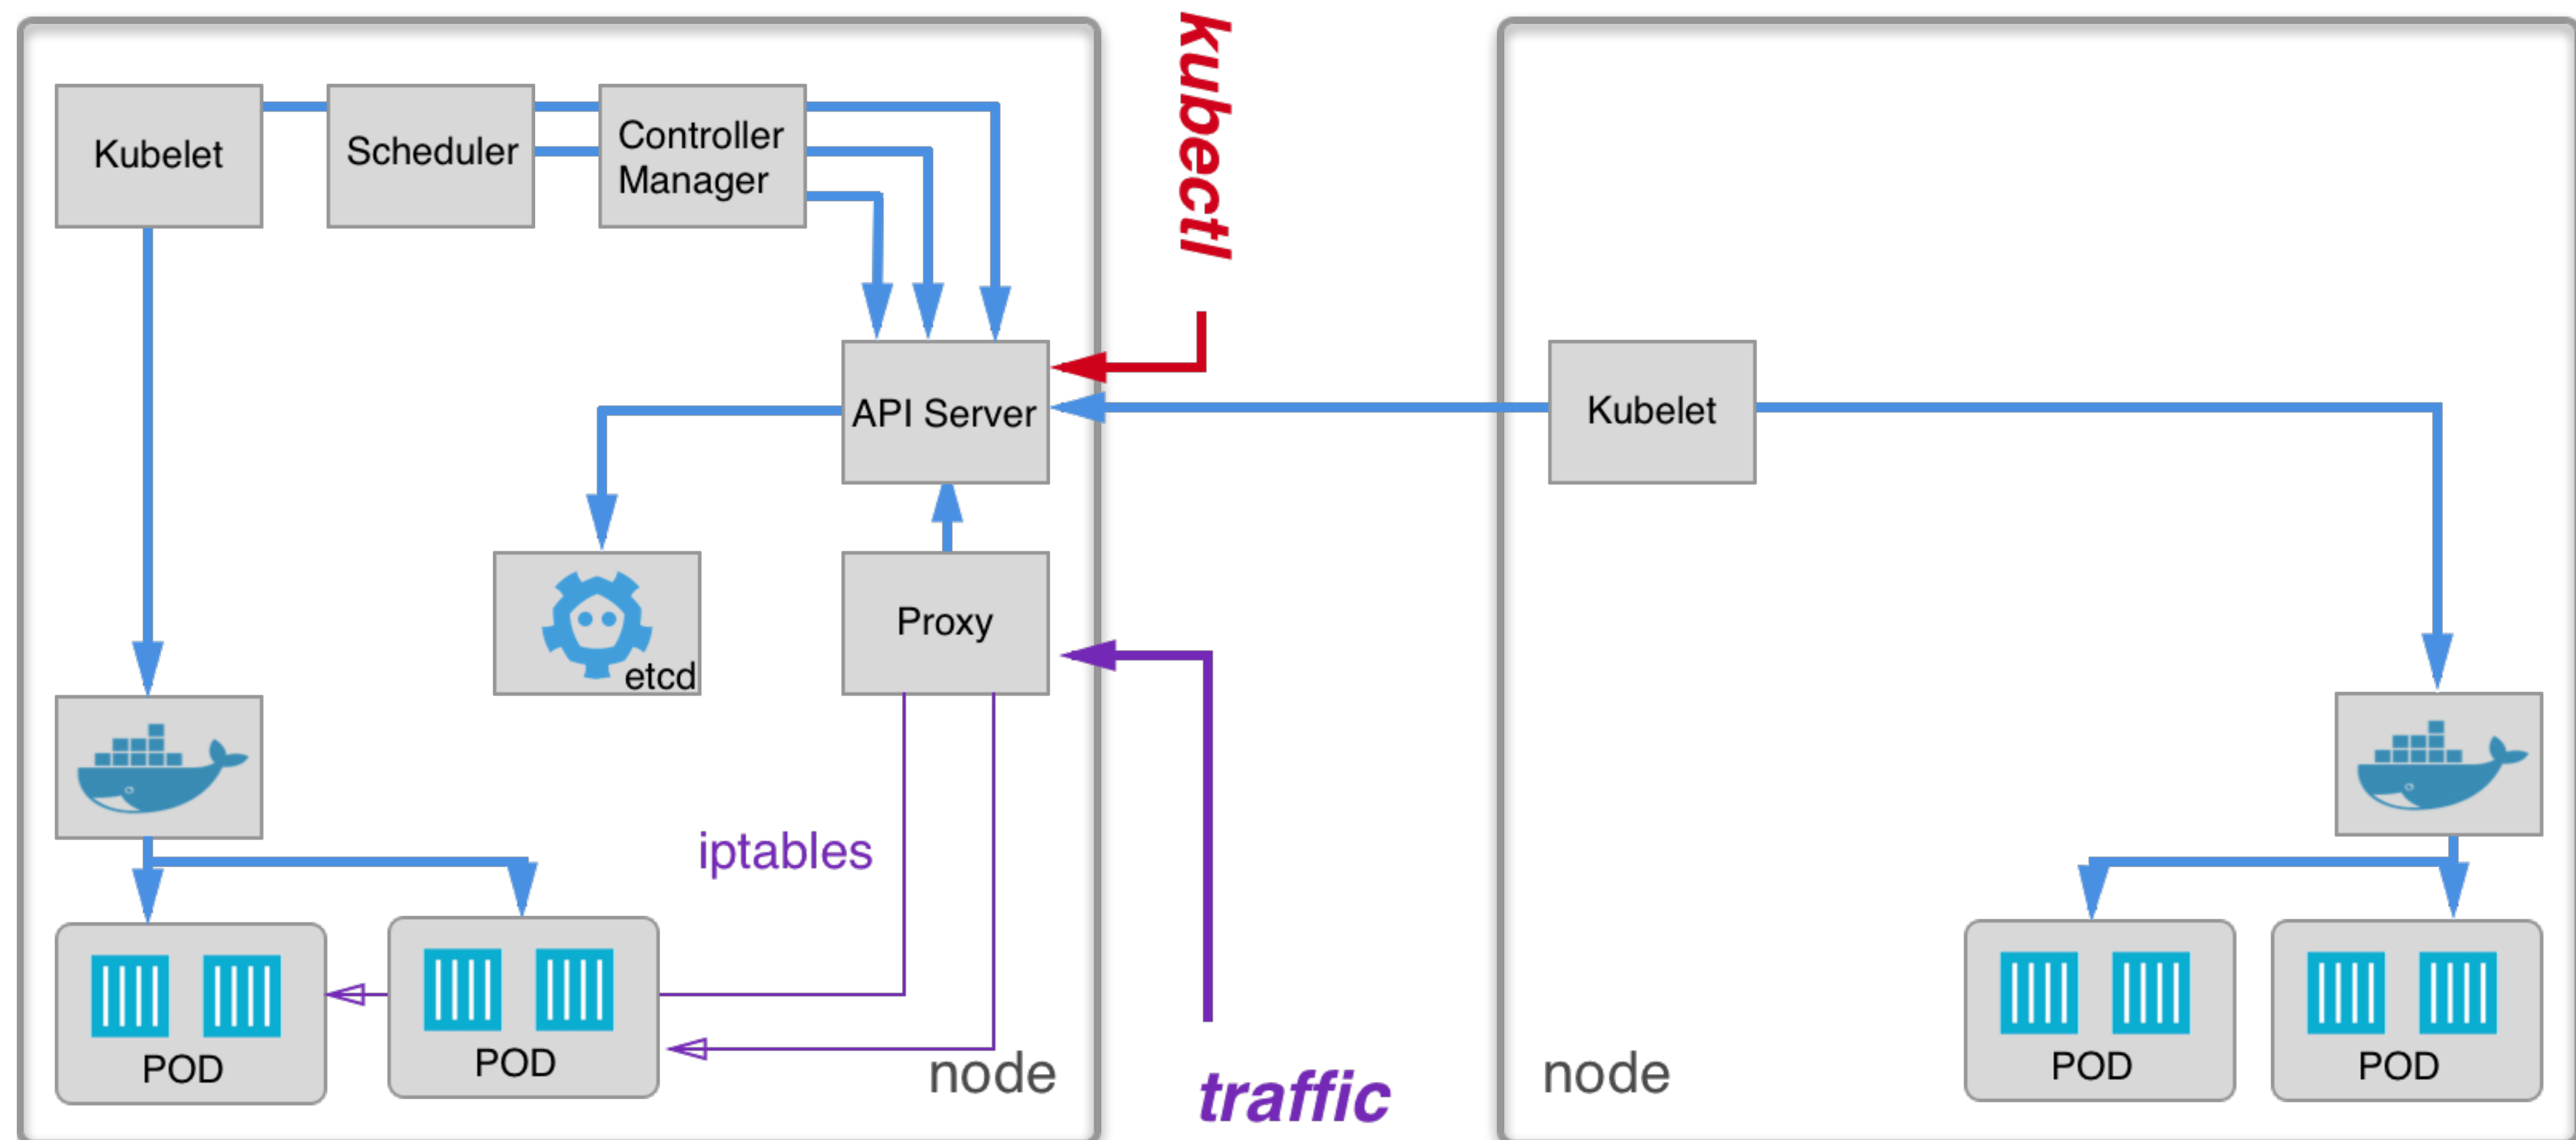
# NETWORKING



# NETWORKING



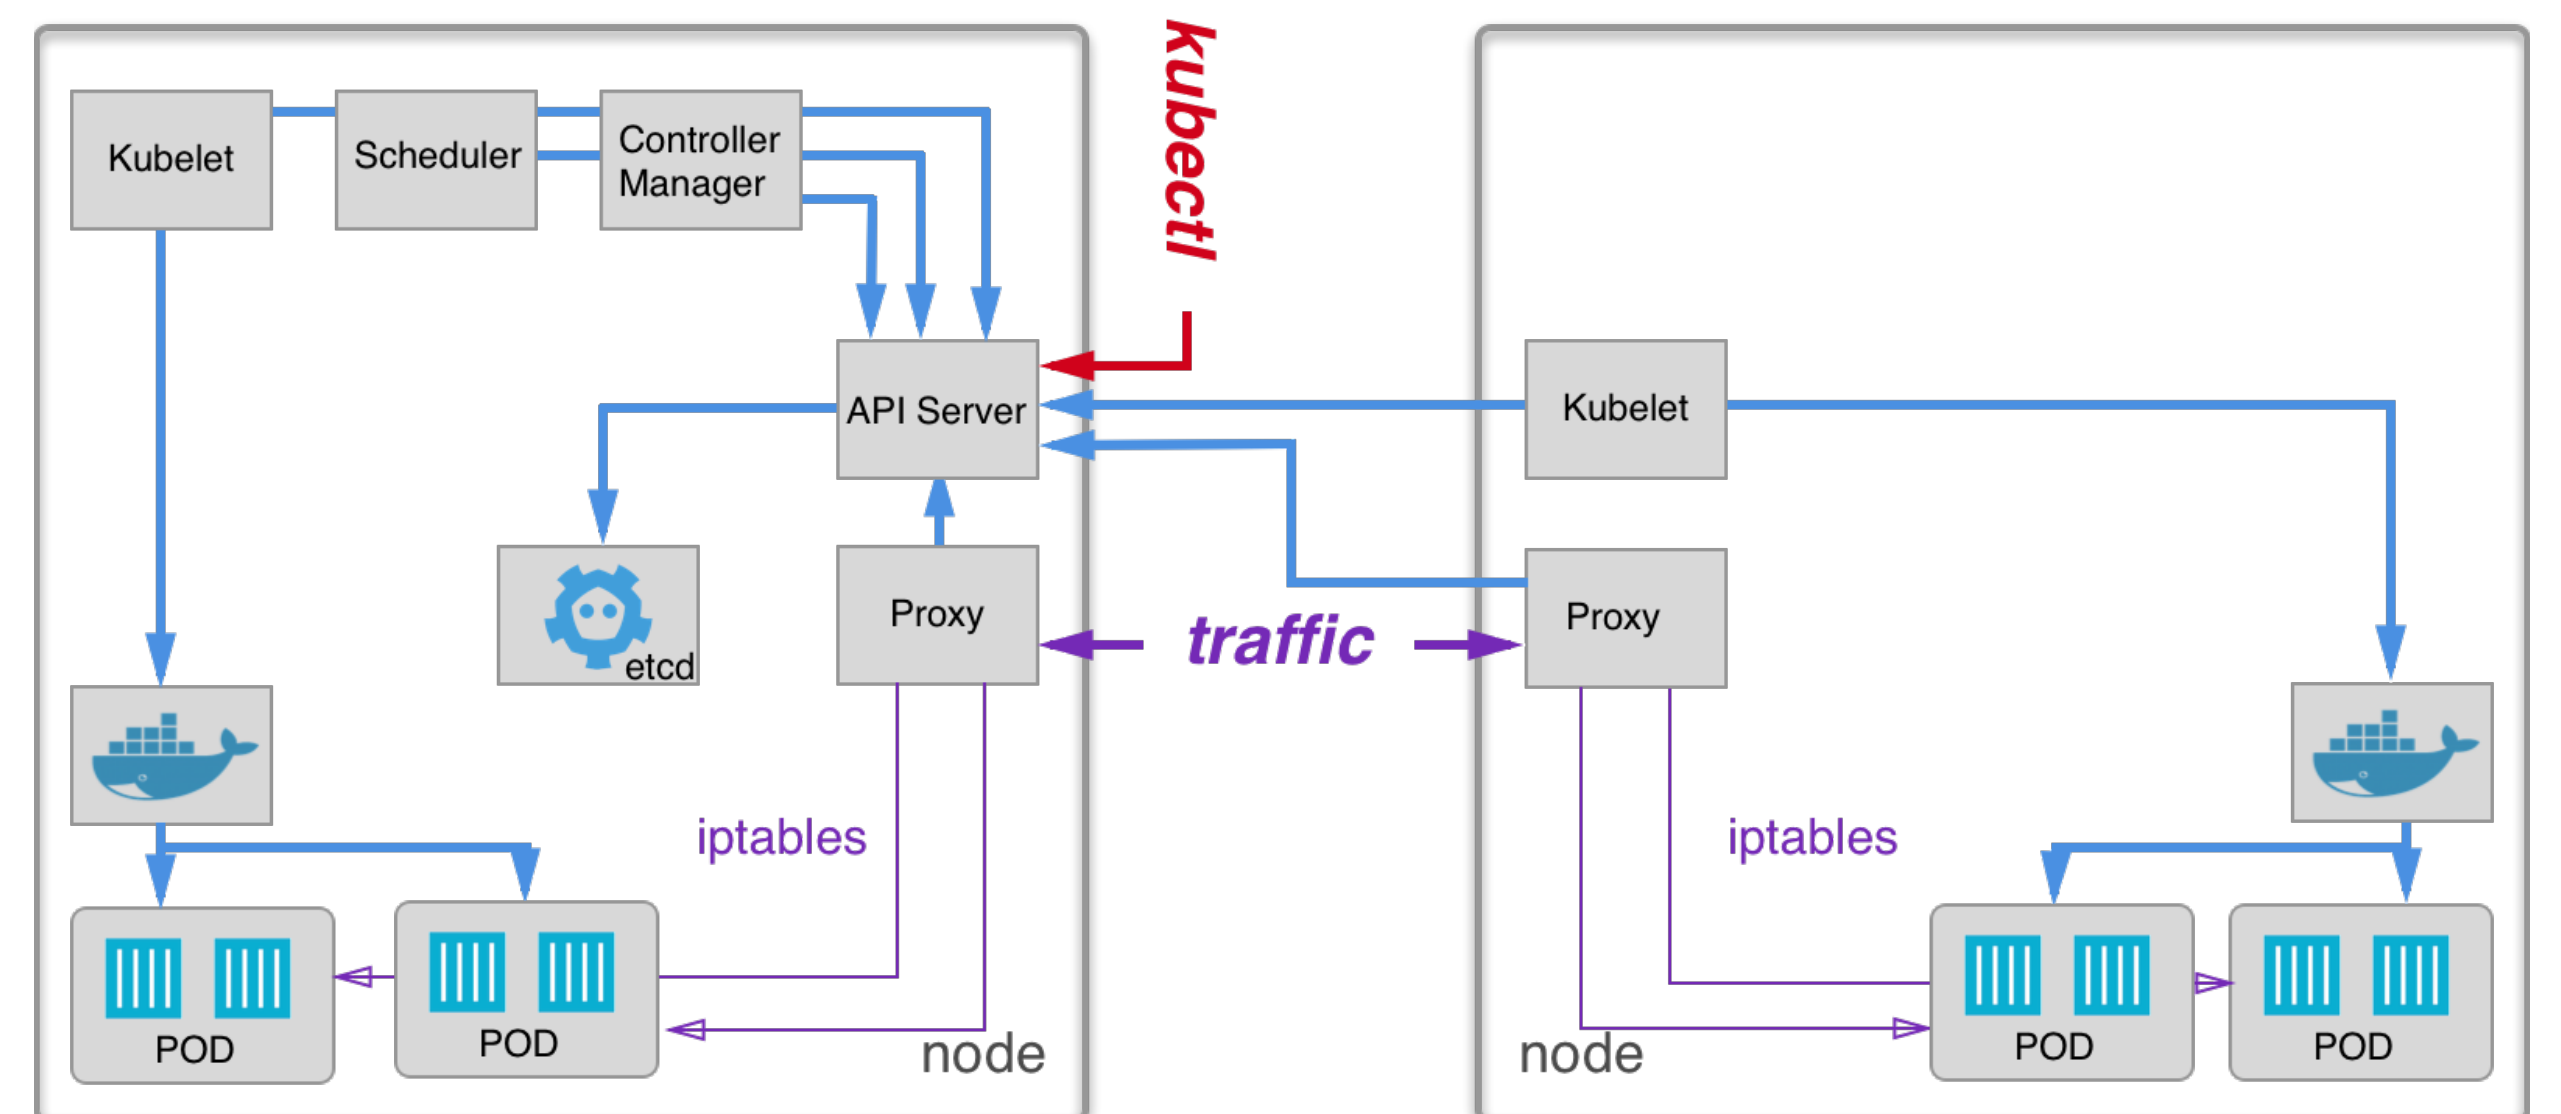
# NETWORKING





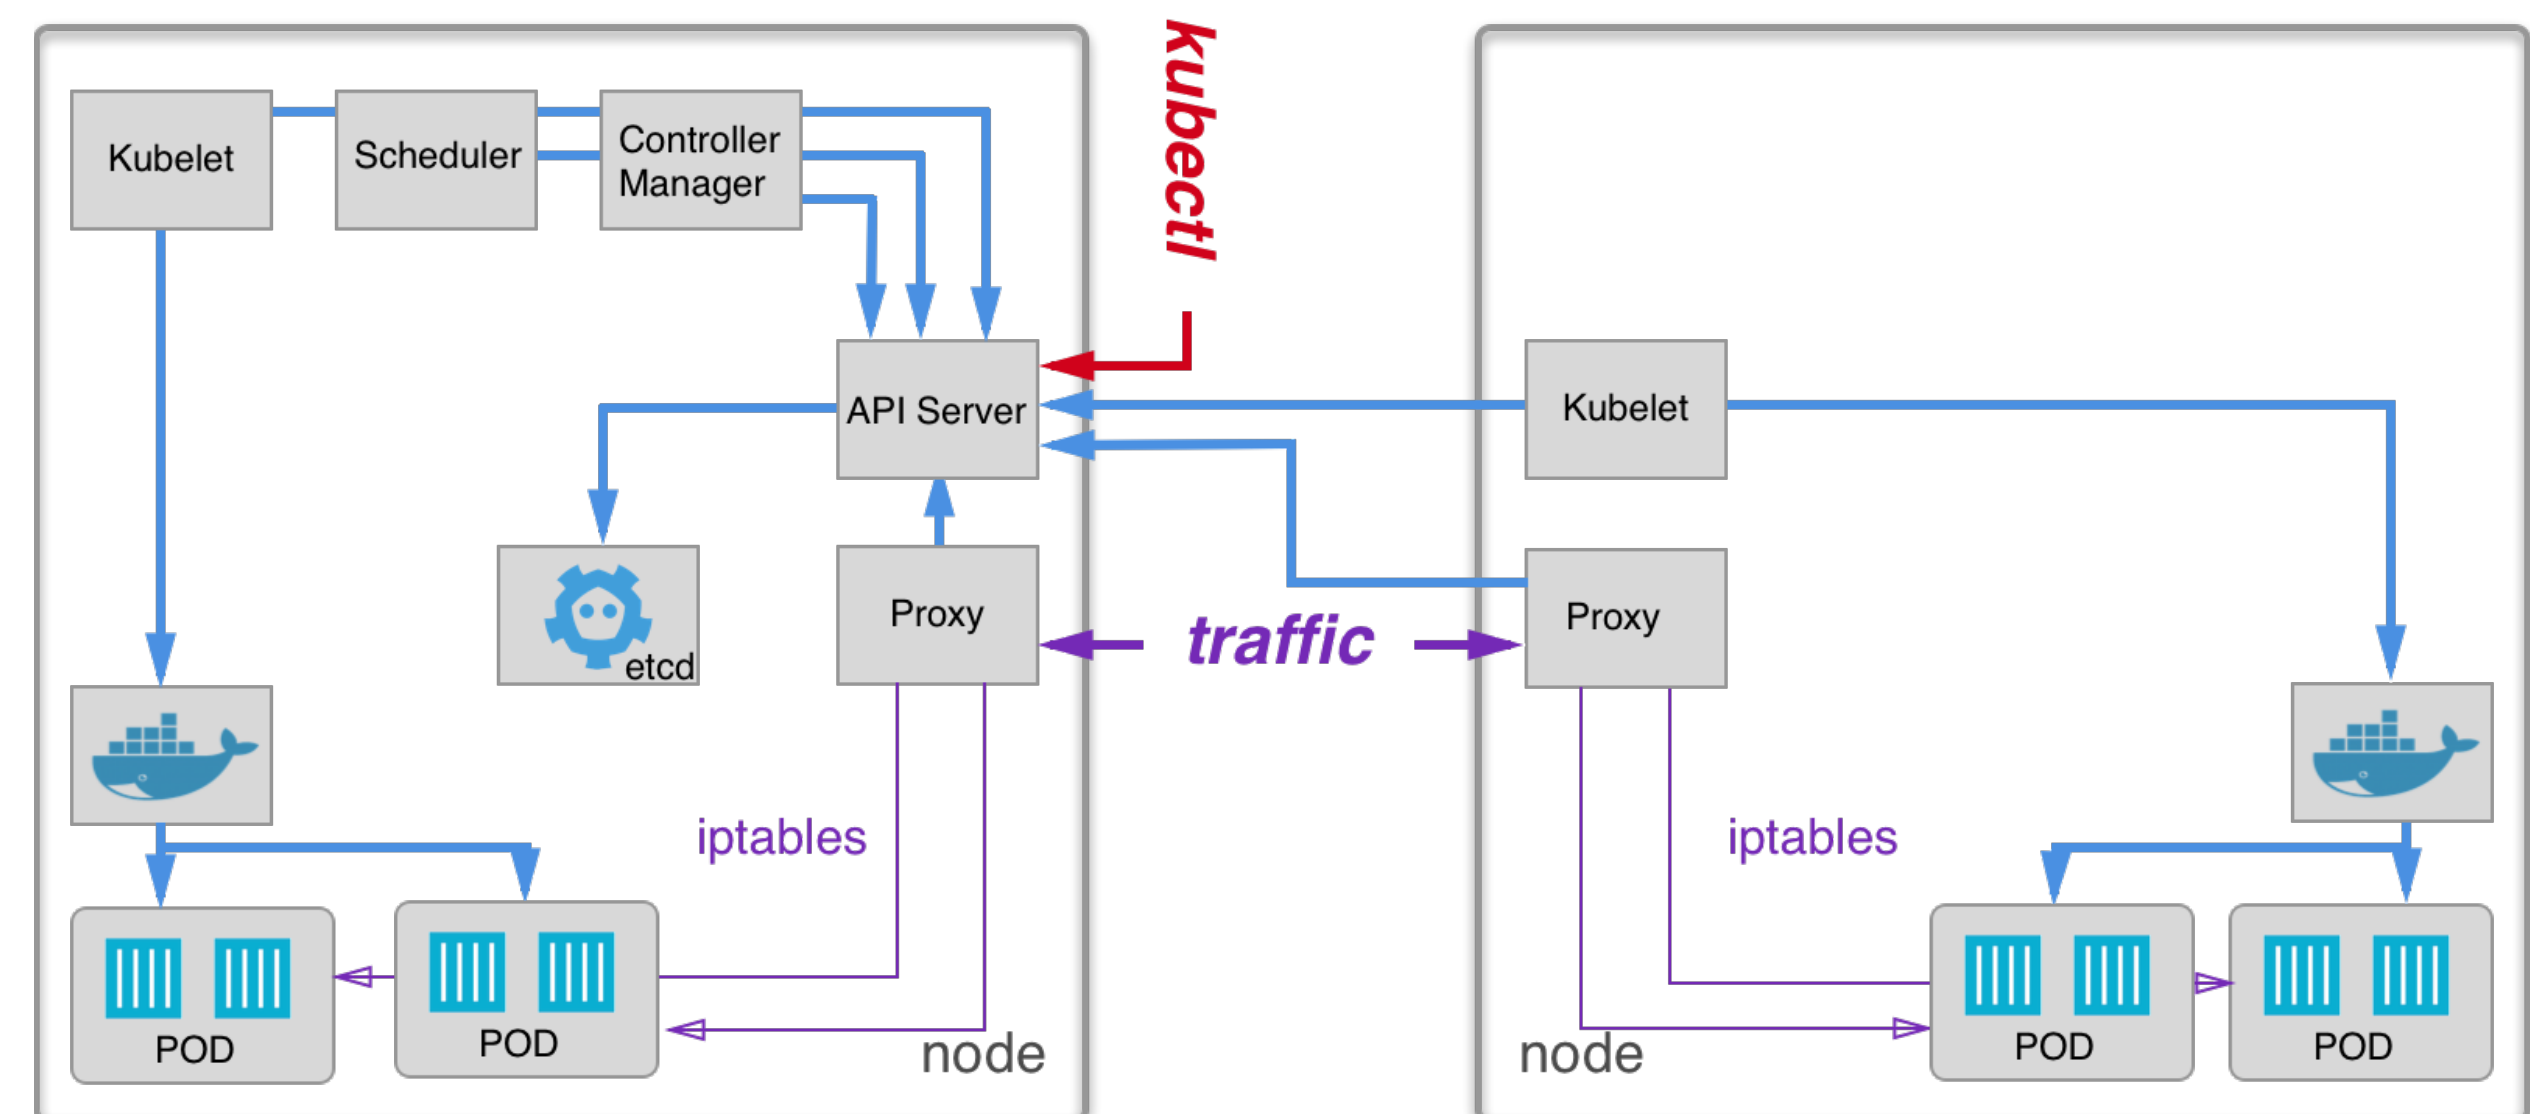


# NETWORKING



# NETWORKING

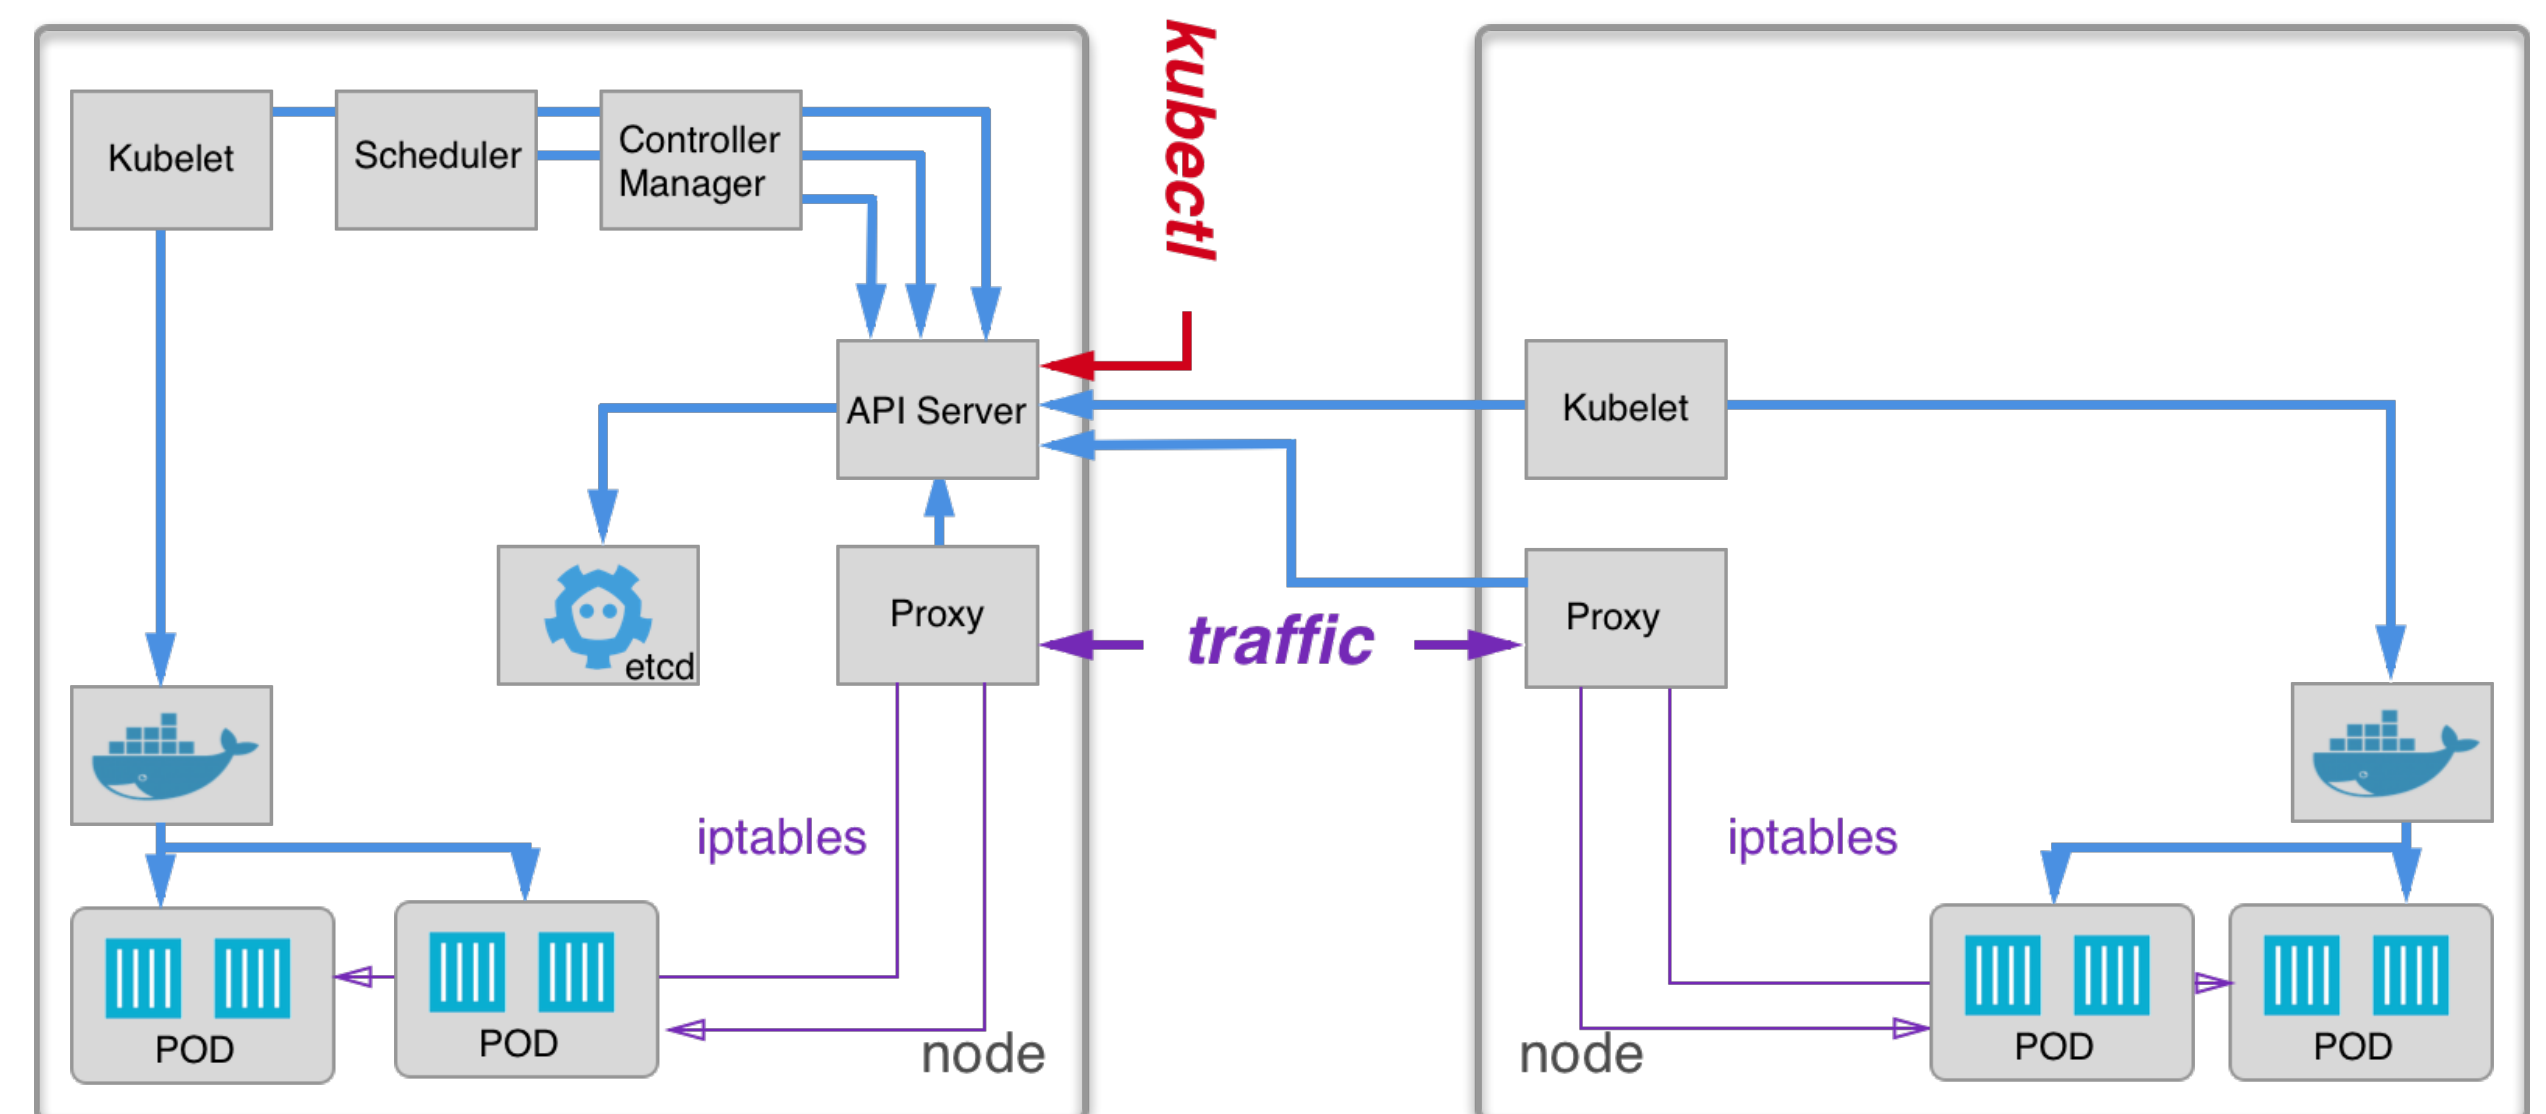
## PROBLEMS:



# NETWORKING

## PROBLEMS:

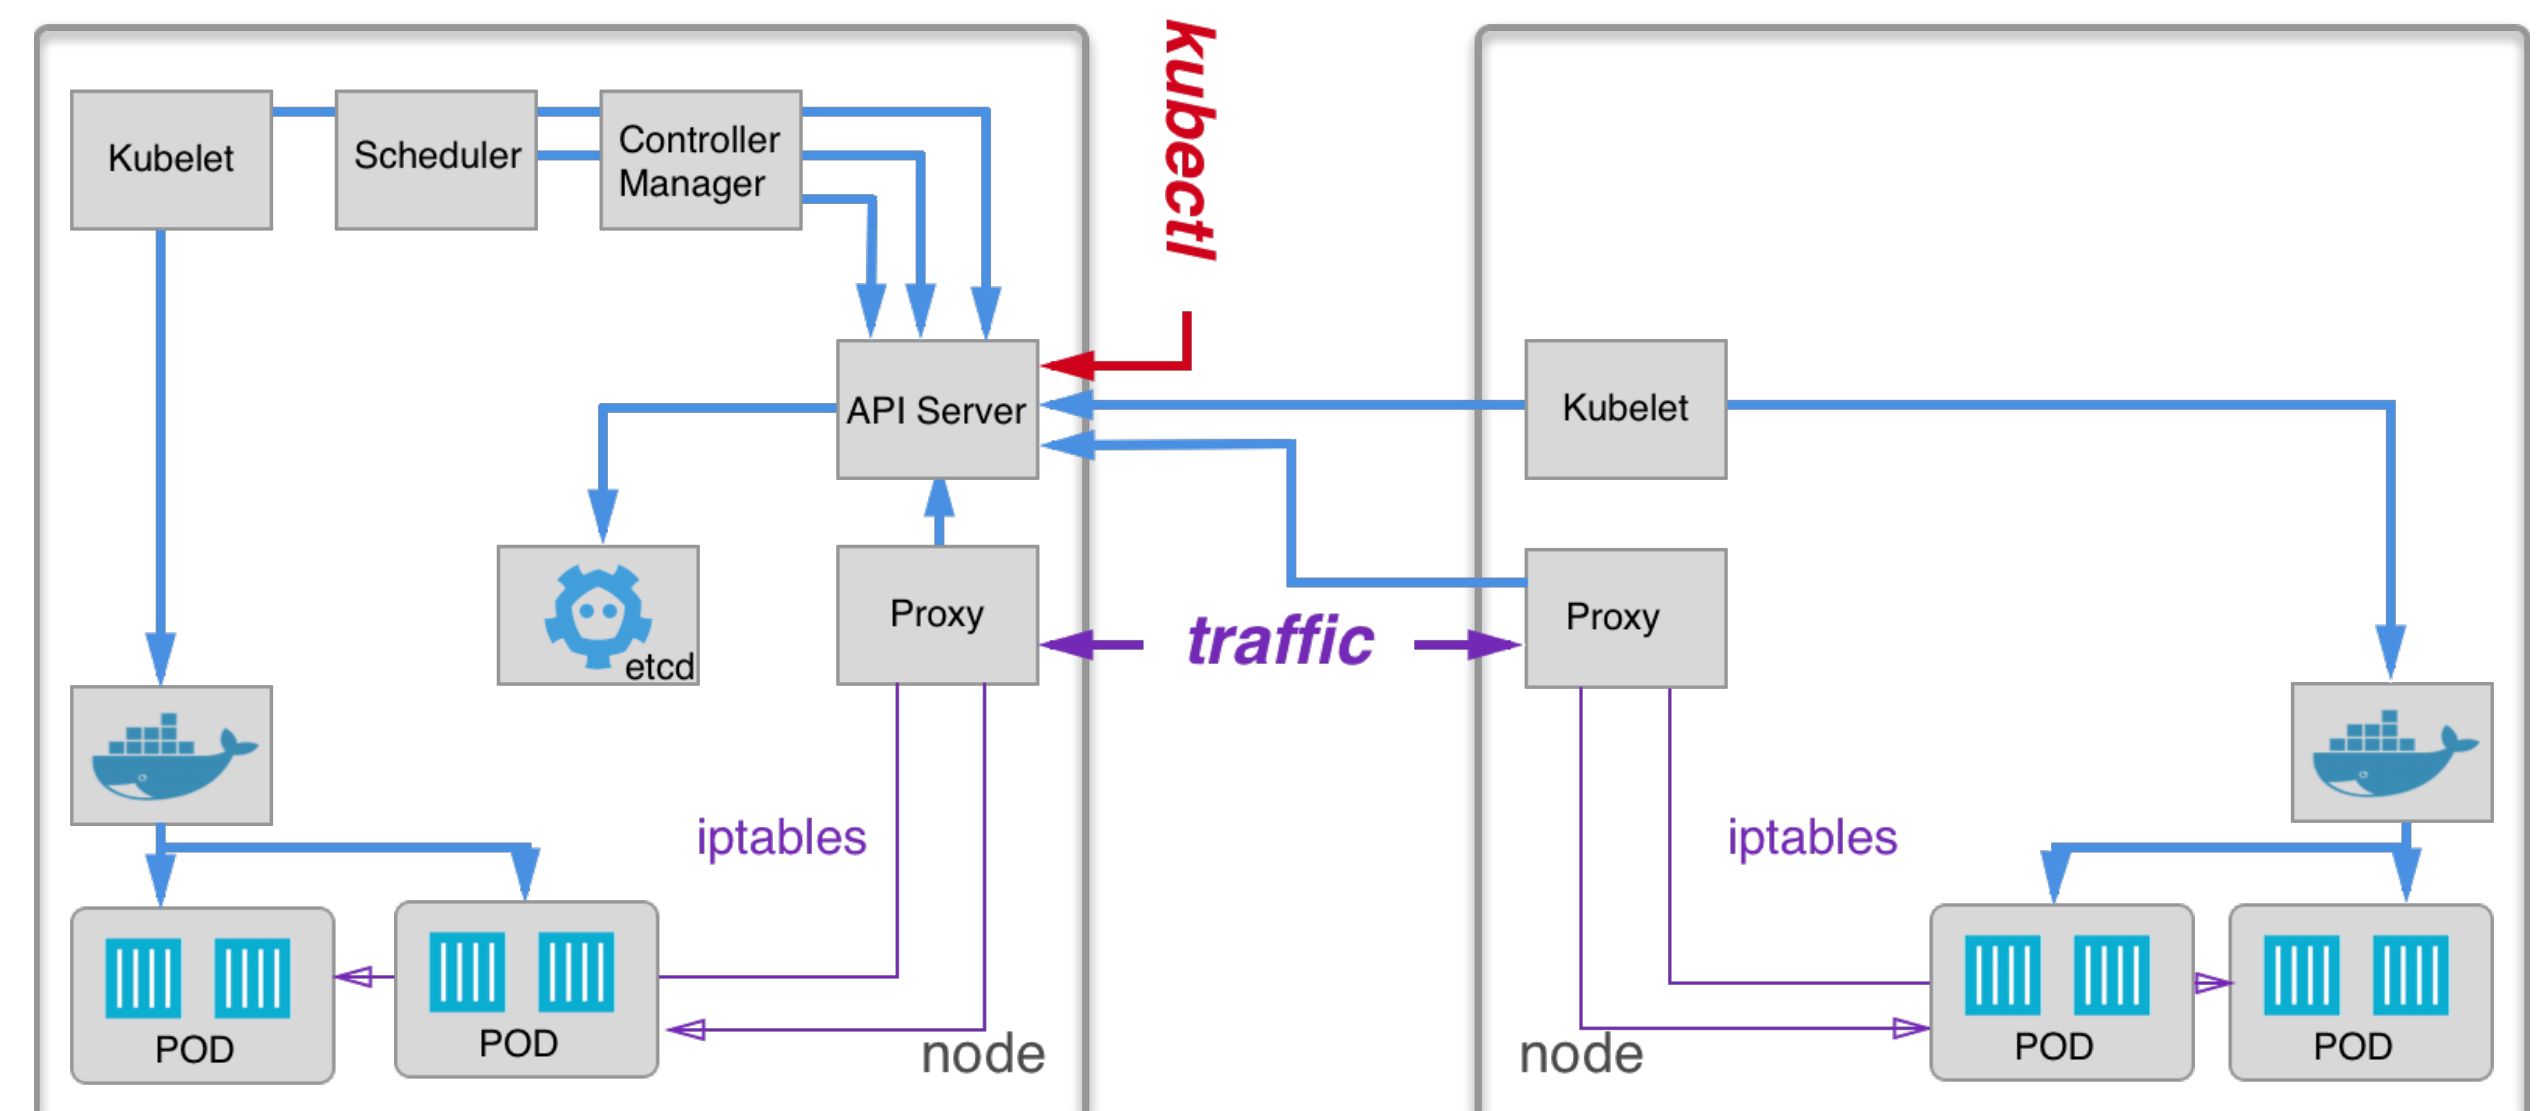
- no POD to POD communication



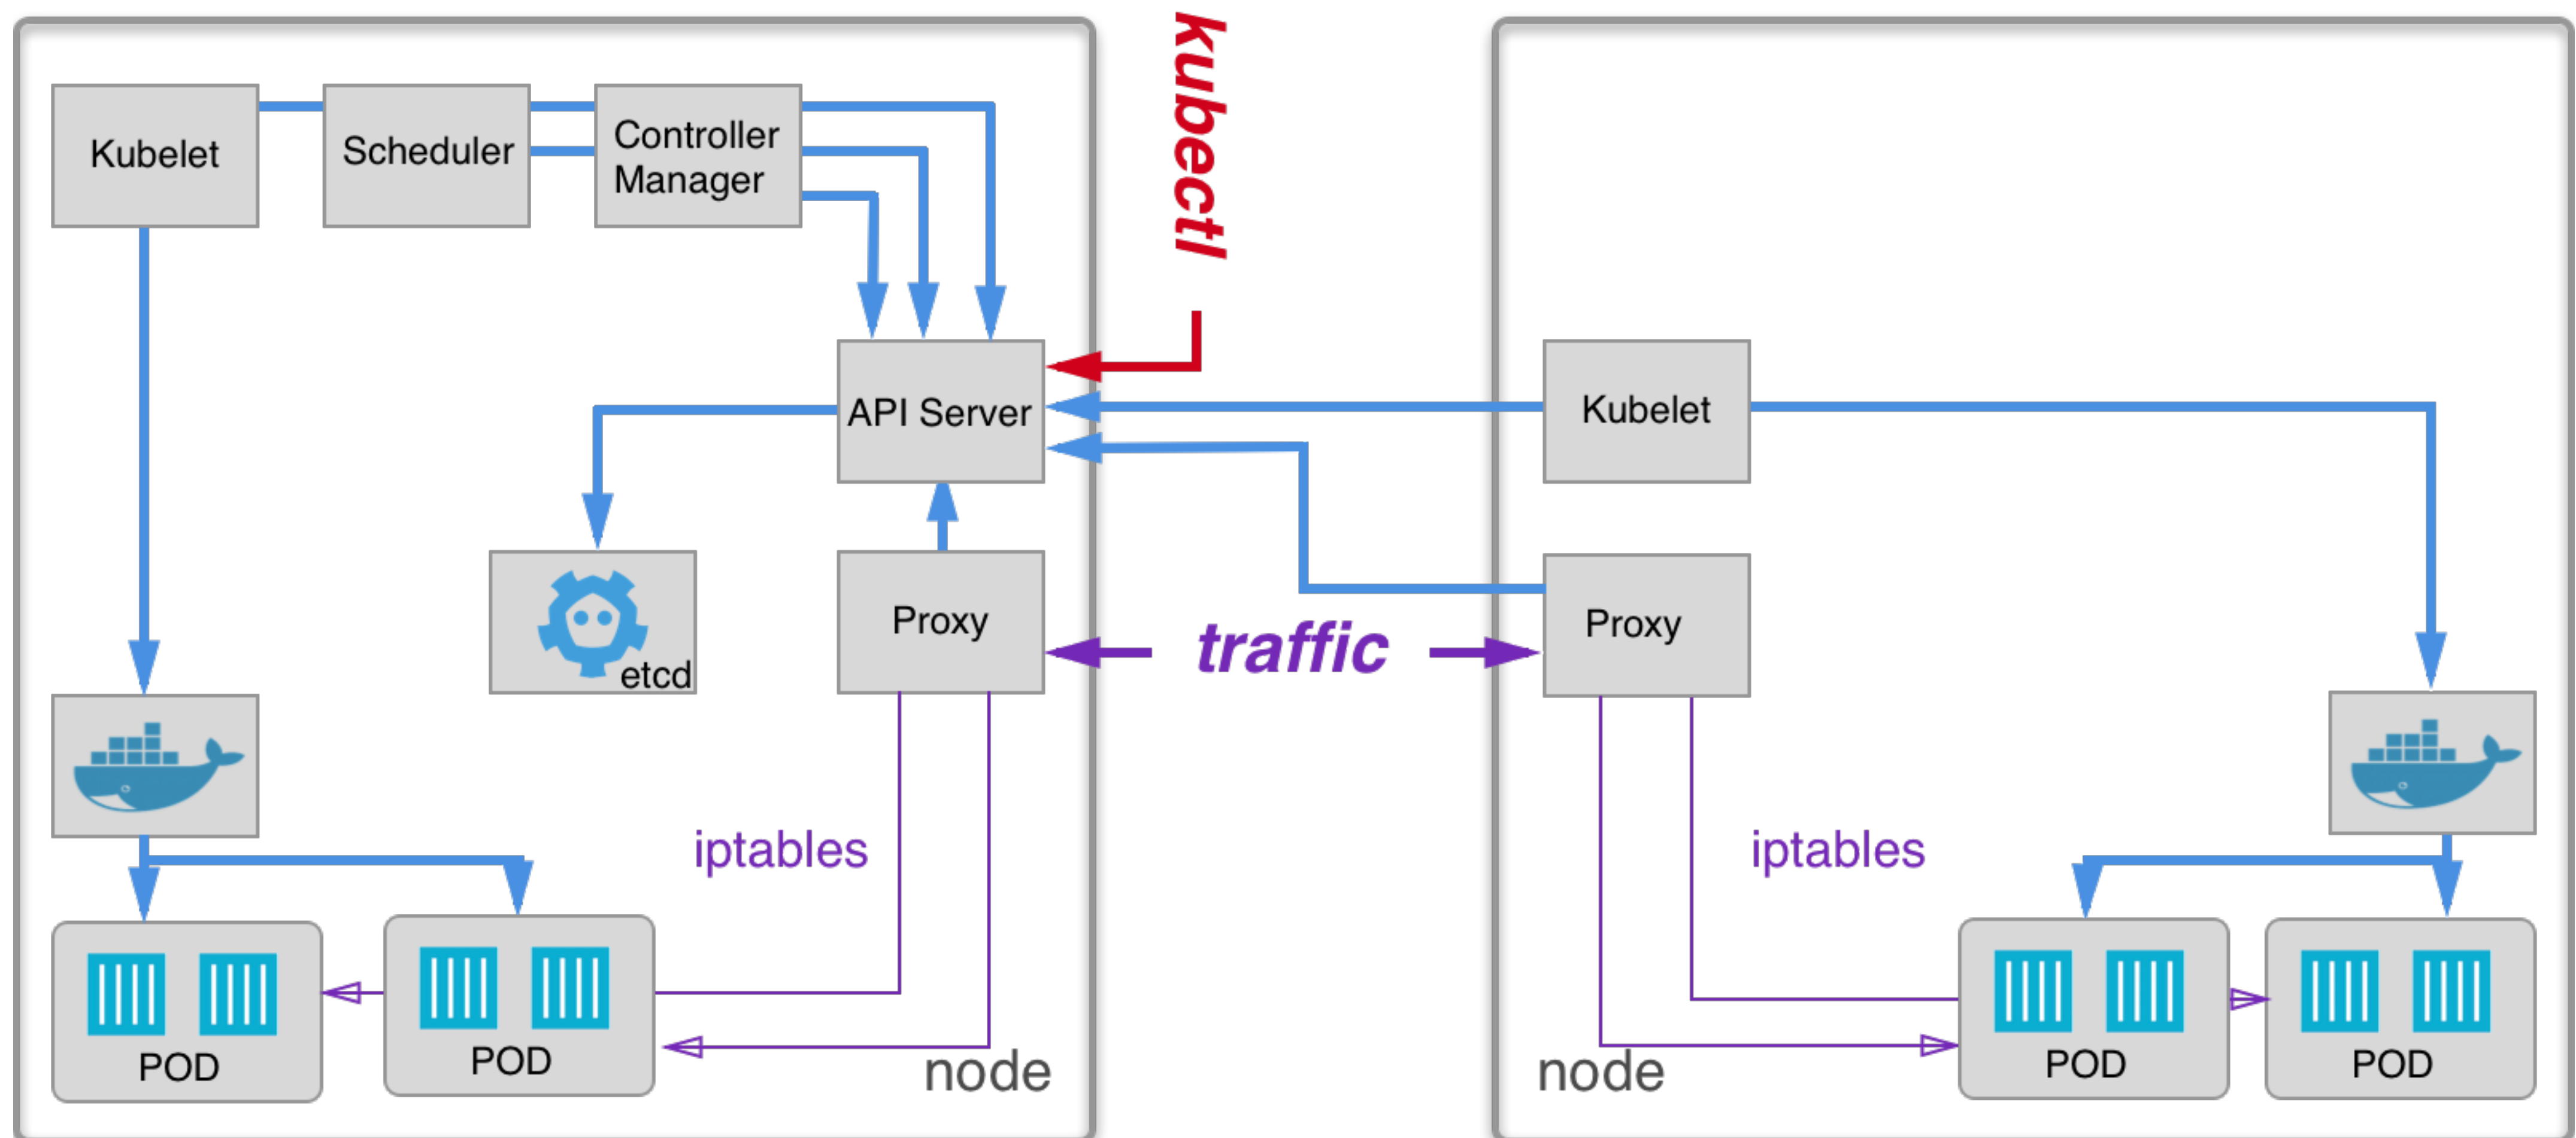
# NETWORKING

## PROBLEMS:

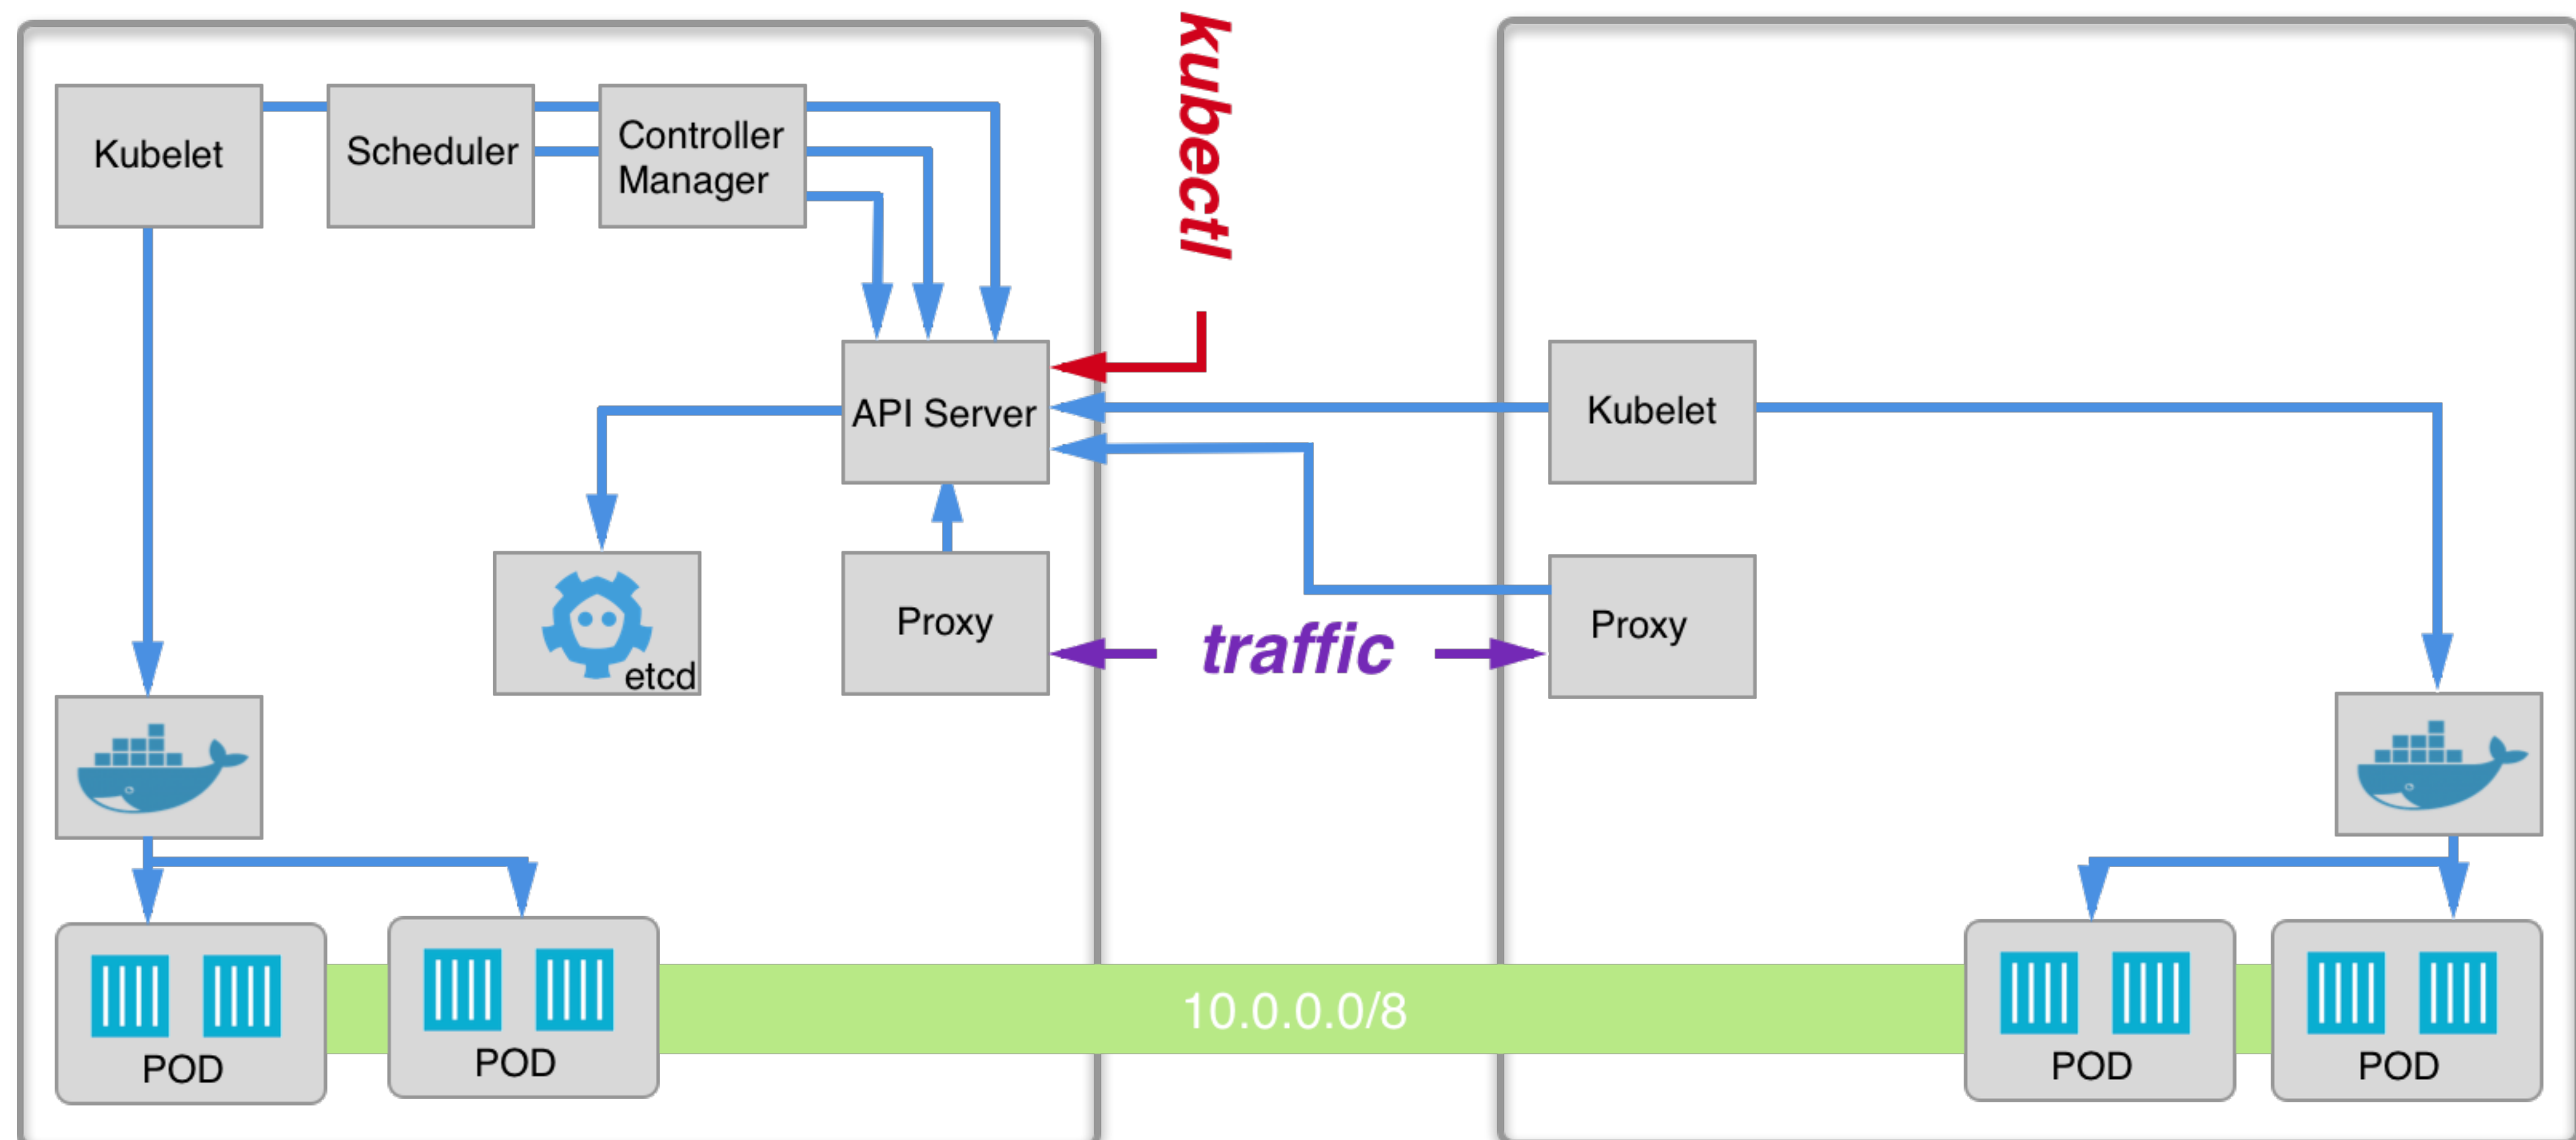
- no POD to POD communication
- traffic cannot be routed between nodes



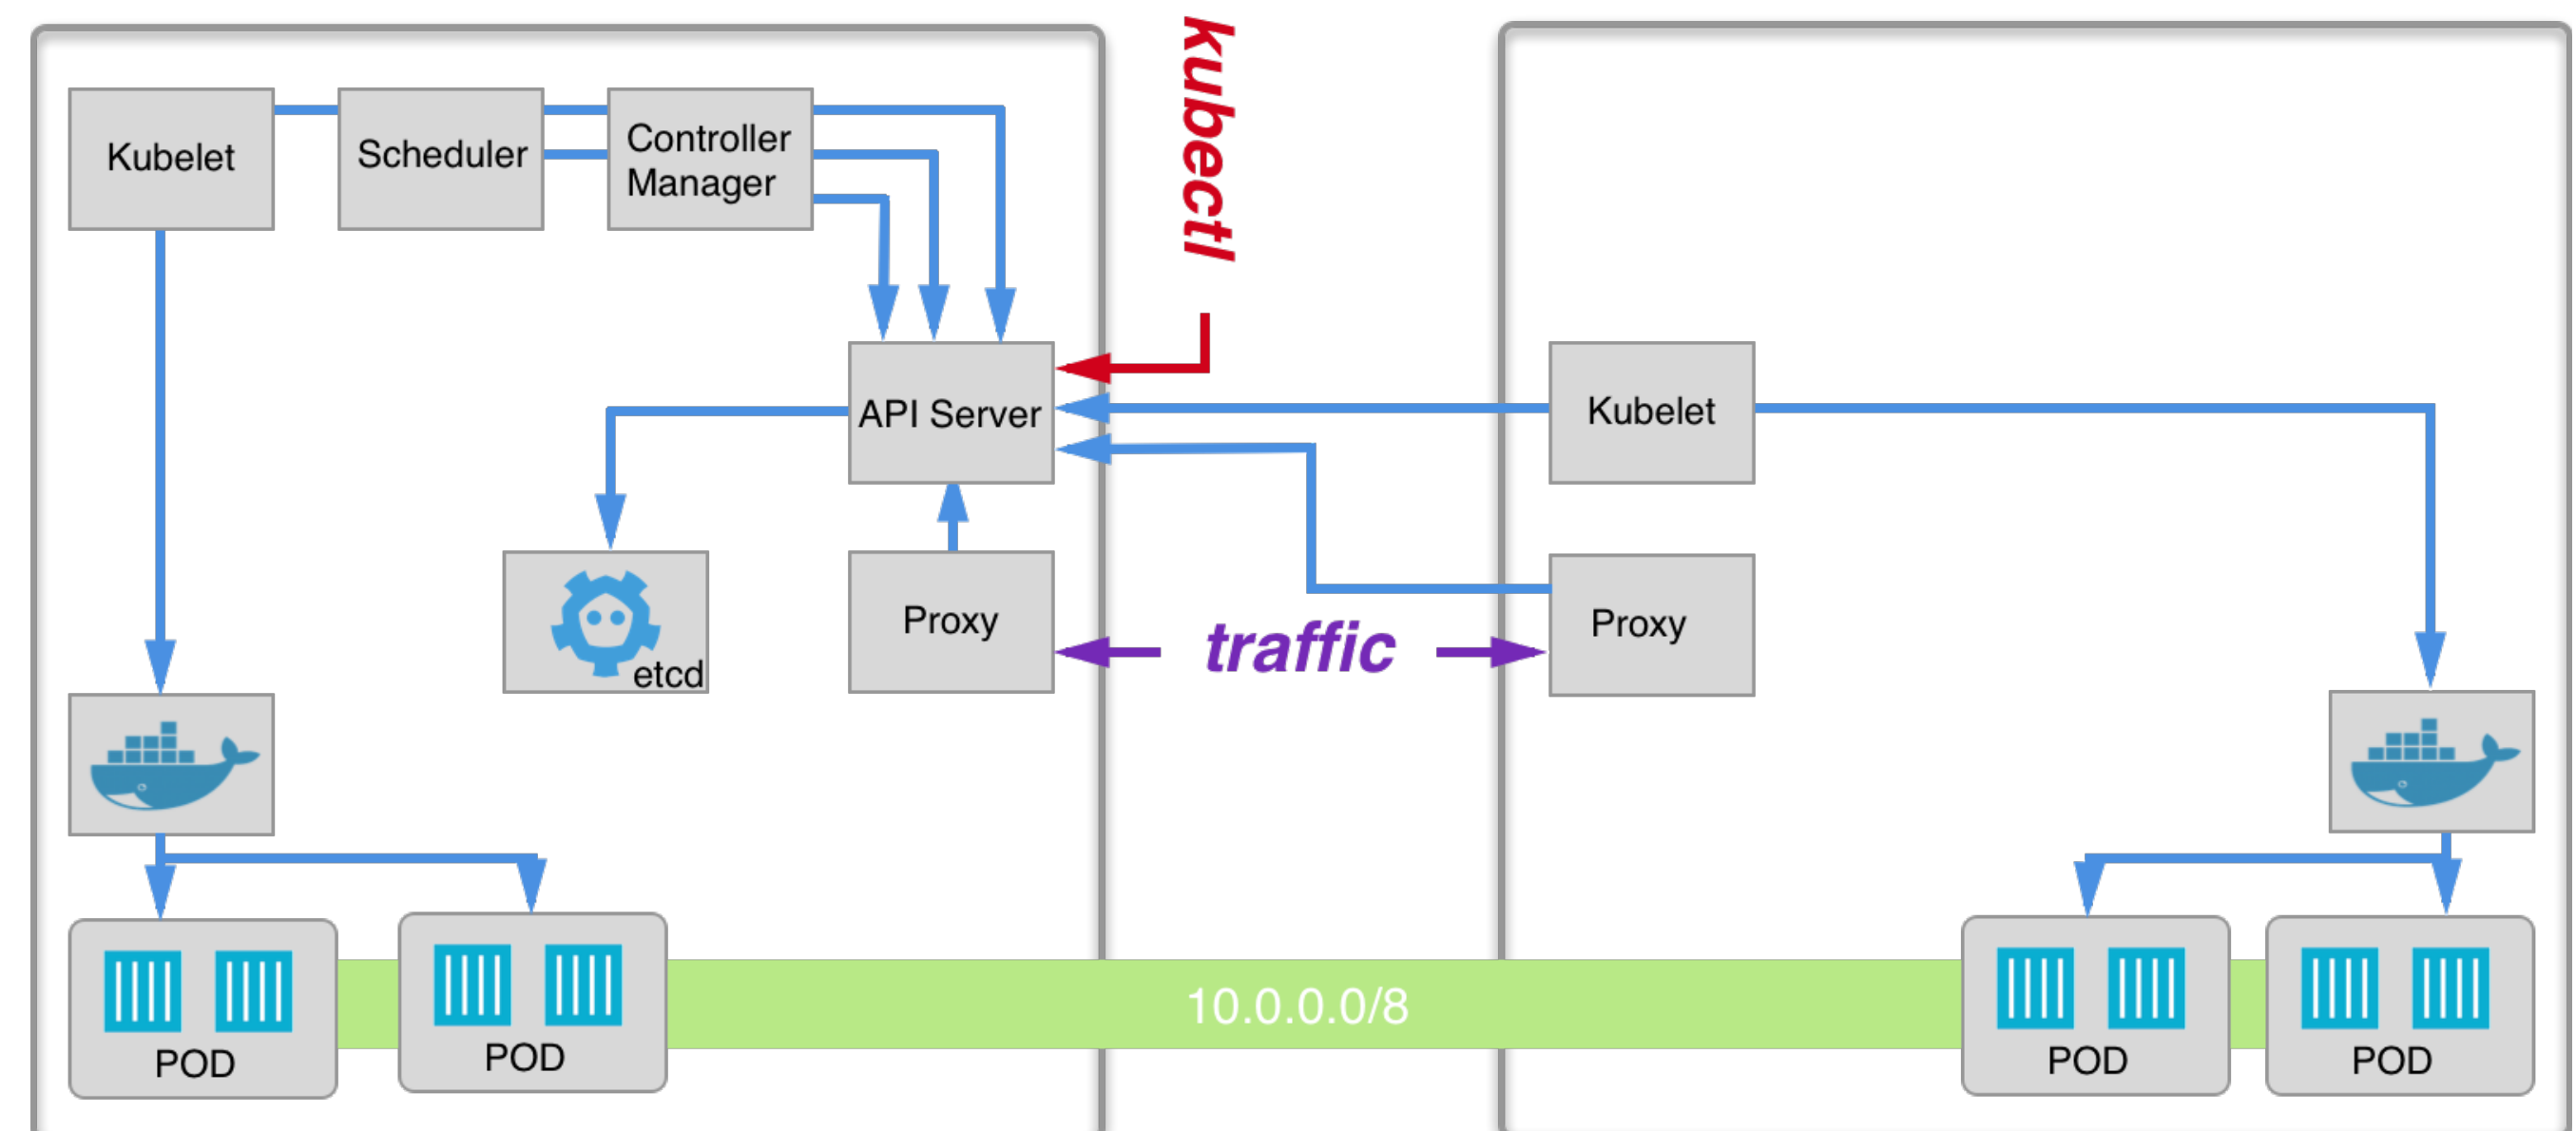
# NETWORKING



# NETWORKING



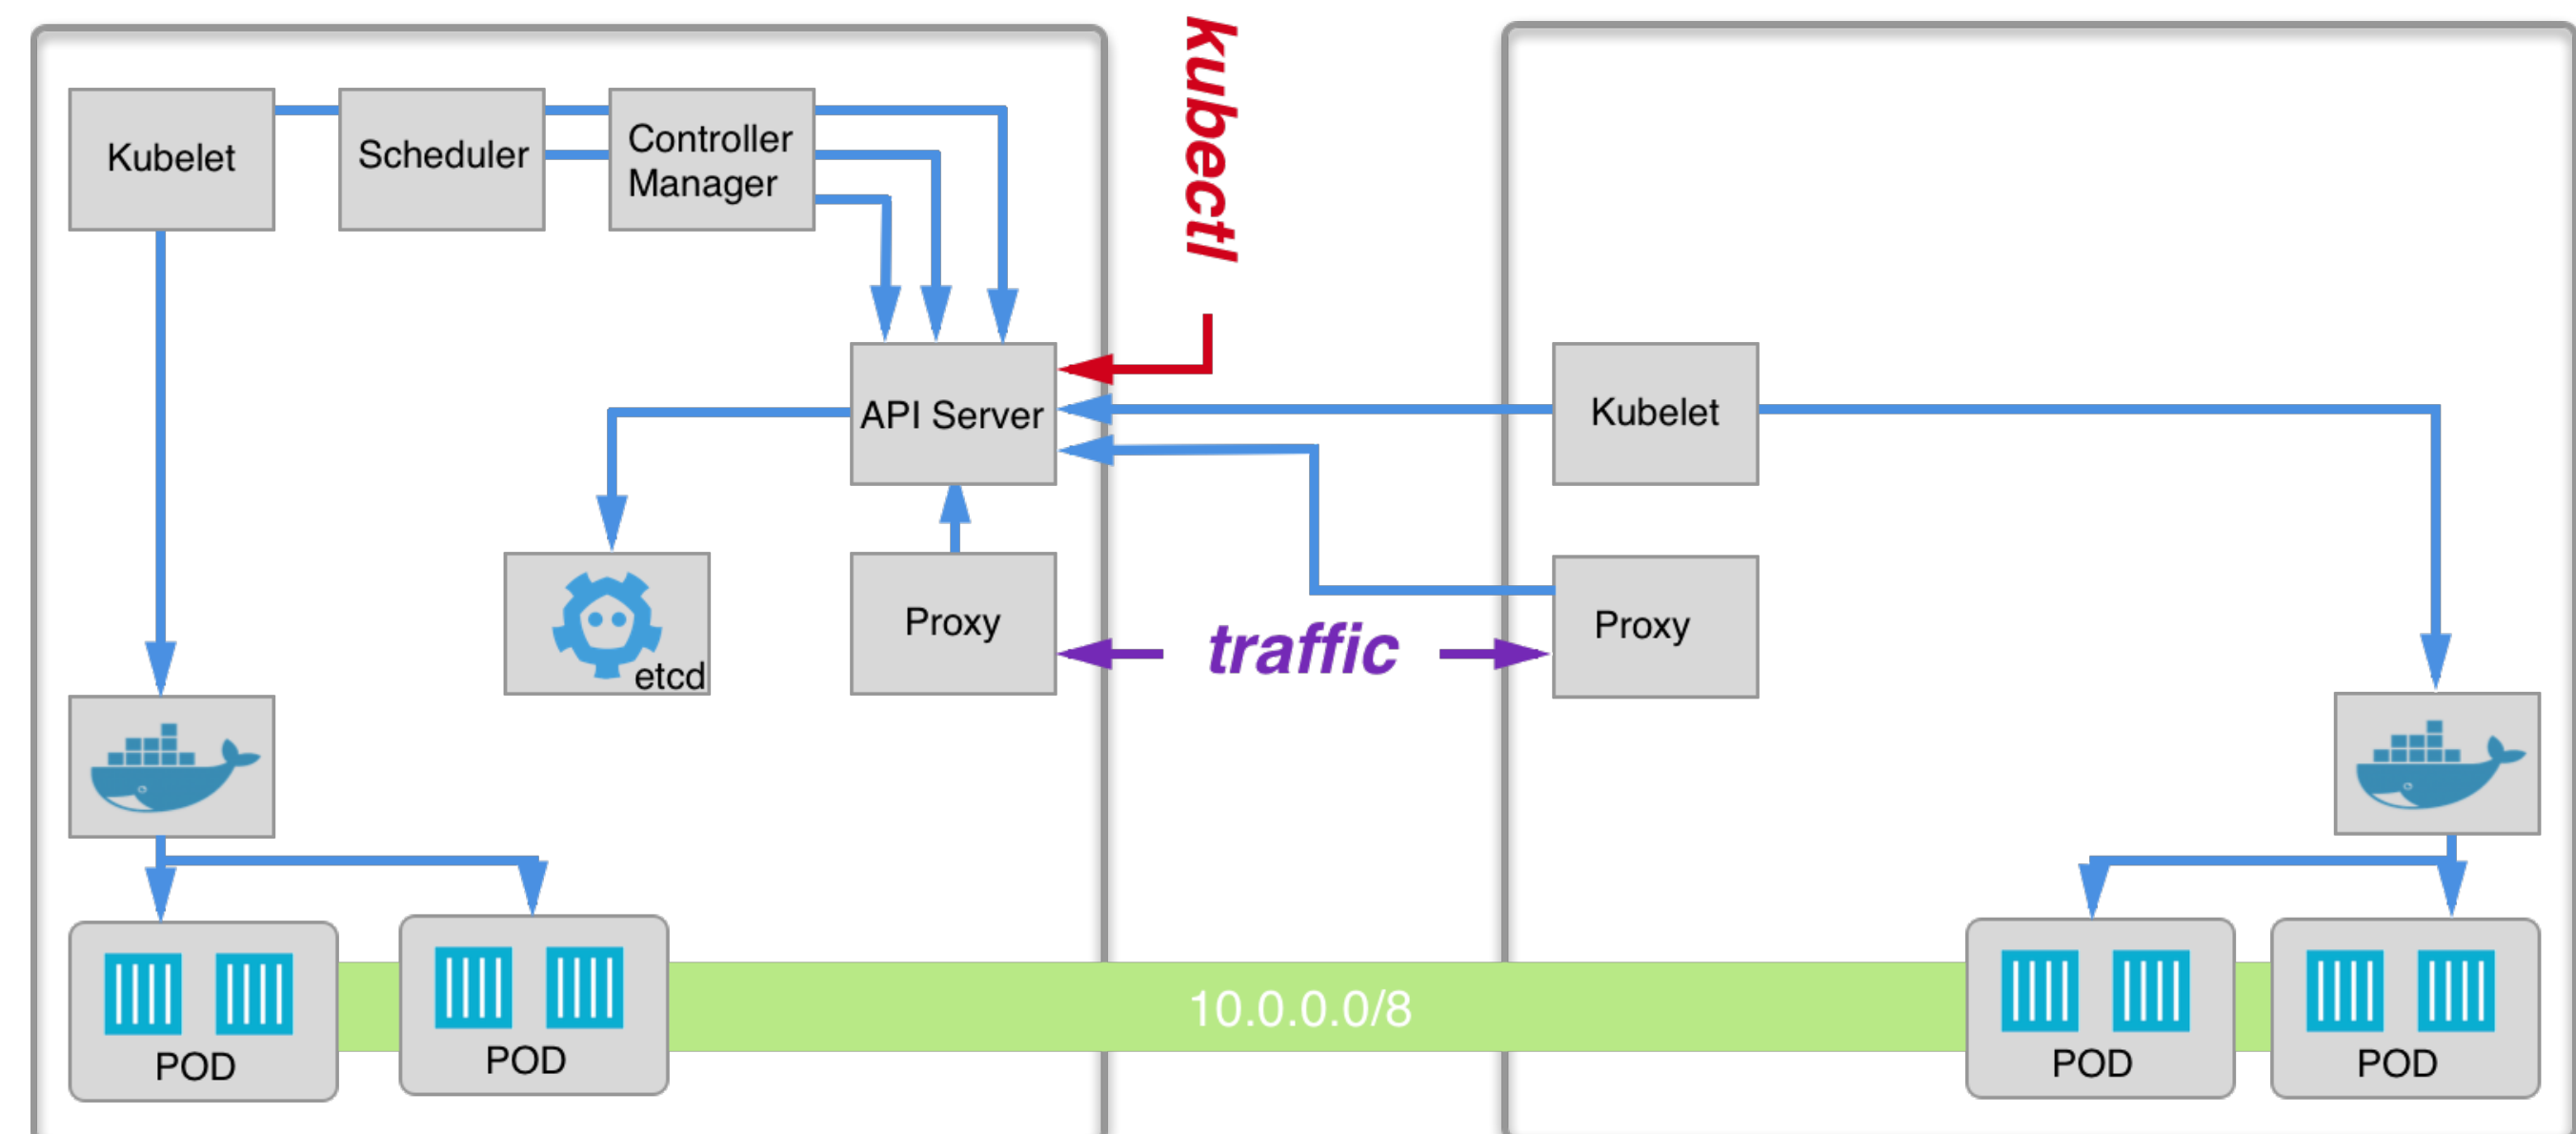
# NETWORKING





# NETWORKING

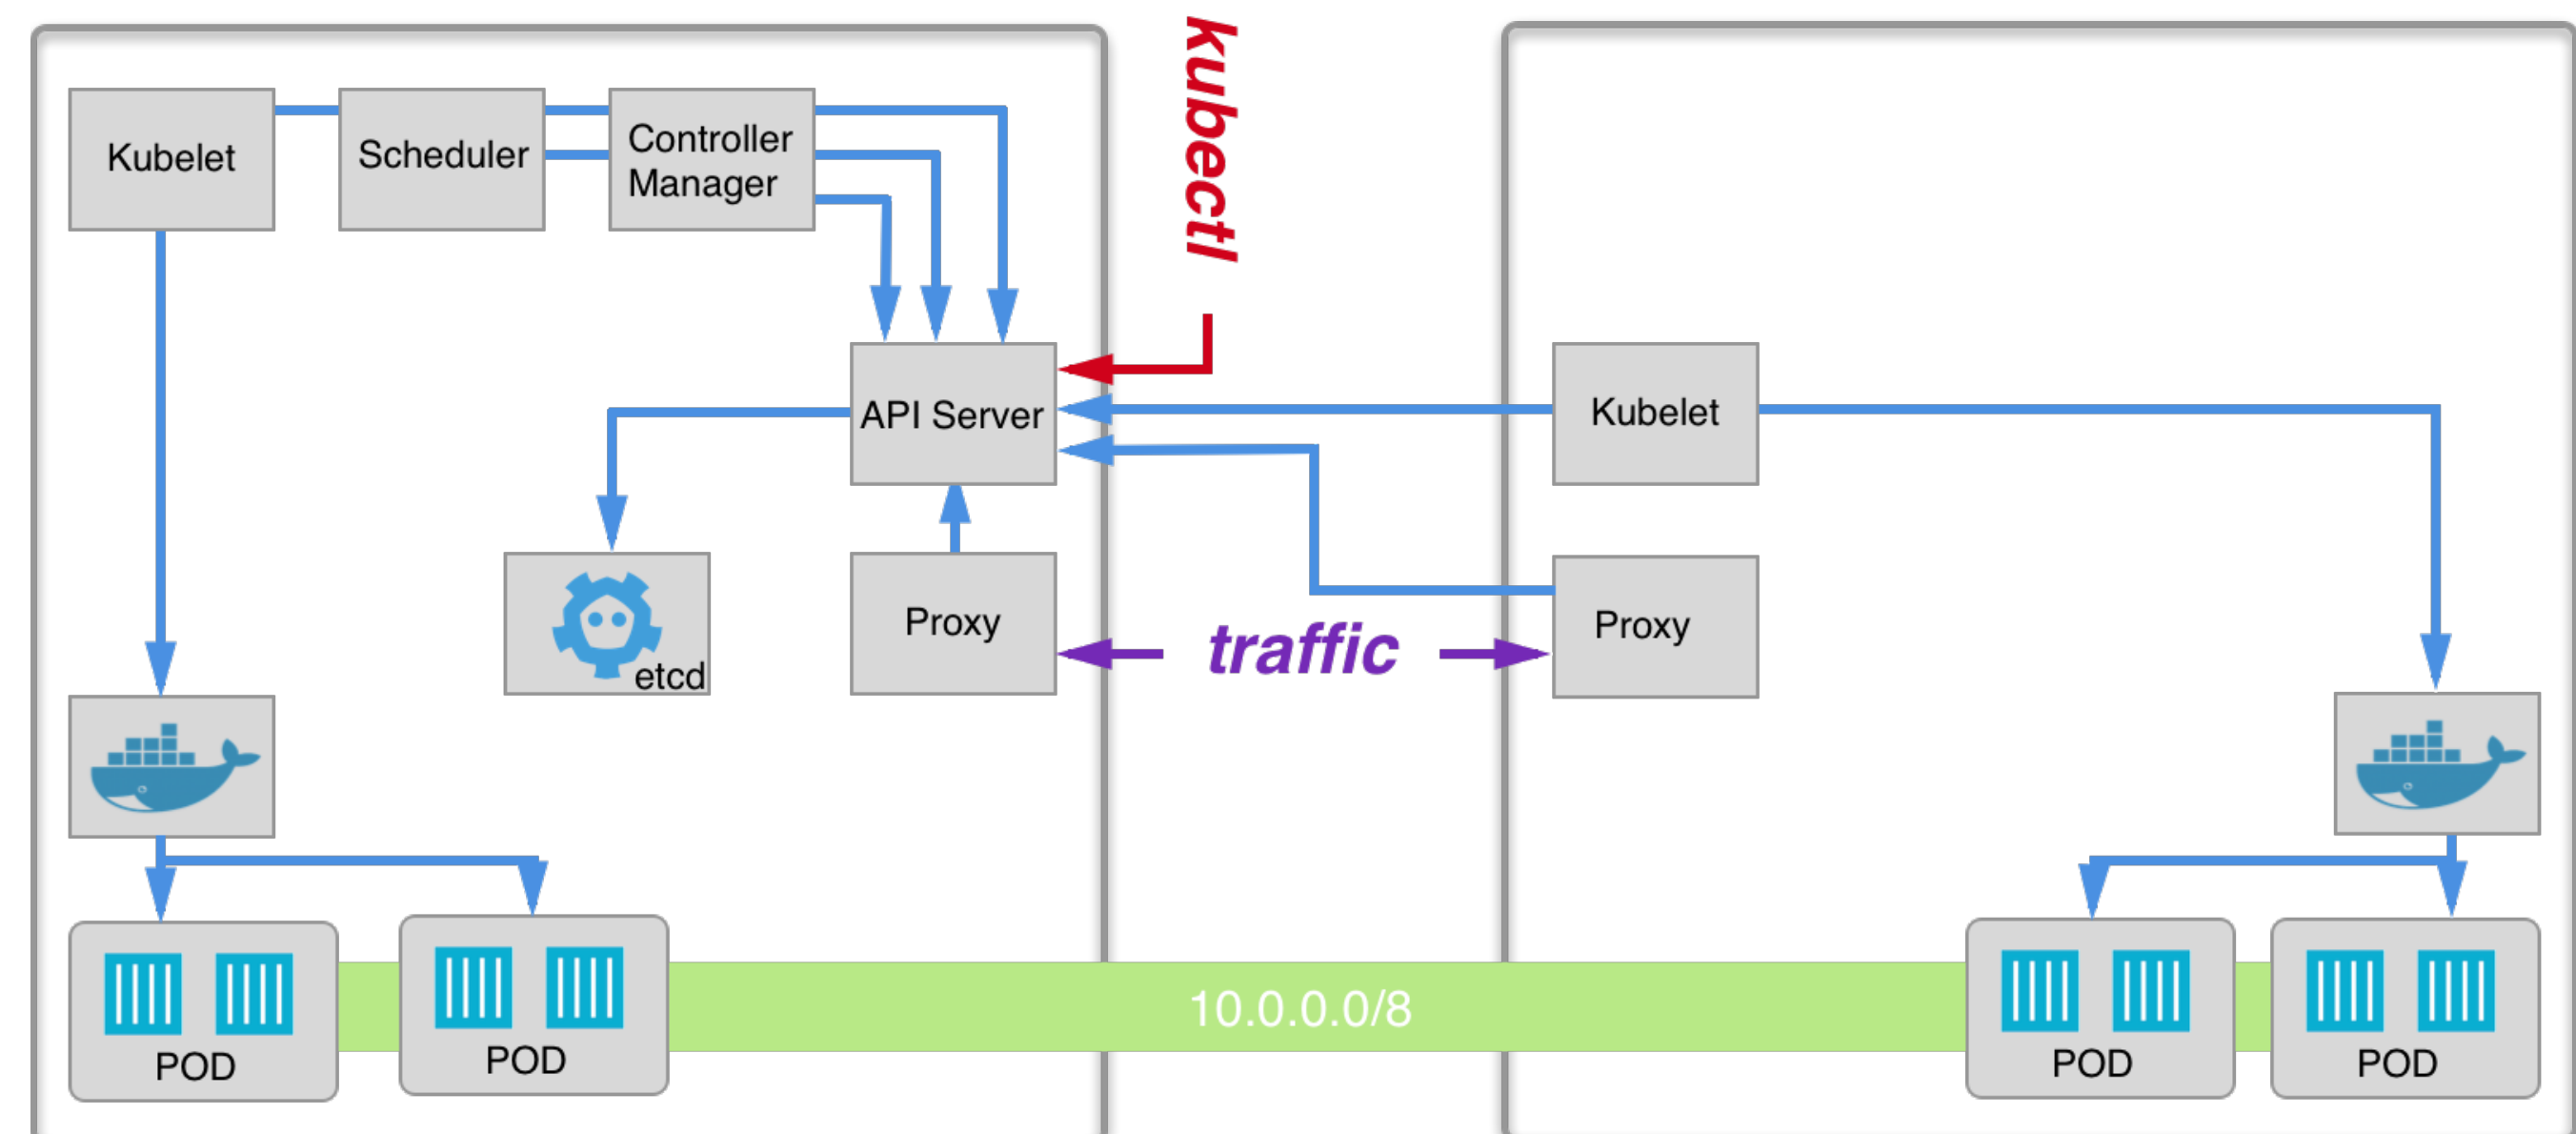
Kubernetes requirements:



# NETWORKING

Kubernetes requirements:

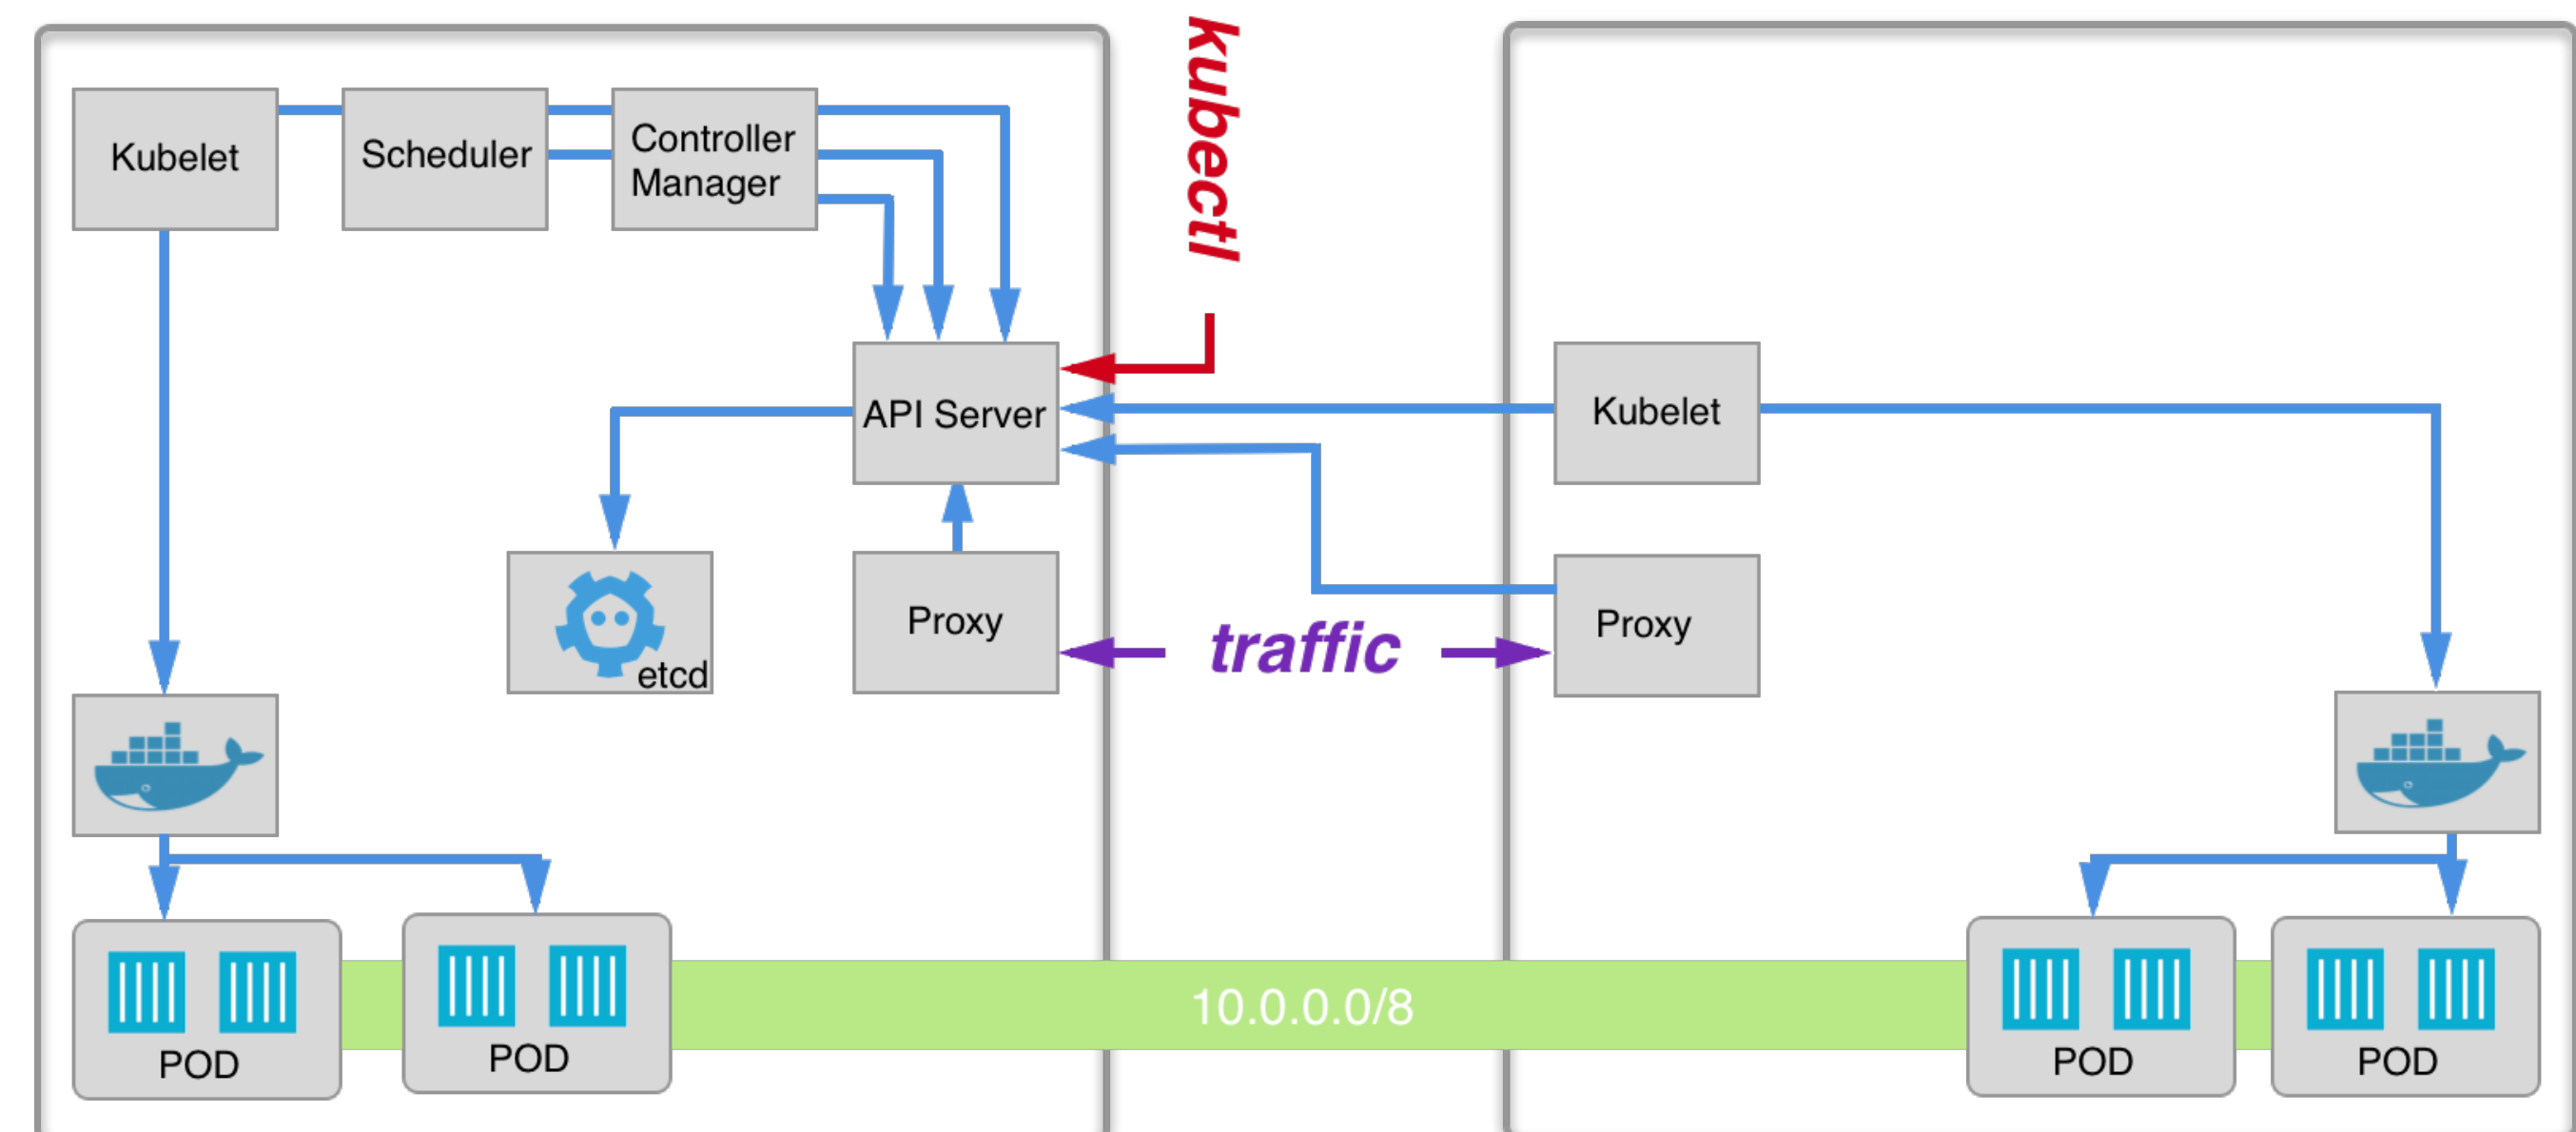
- each pod has it's own IP



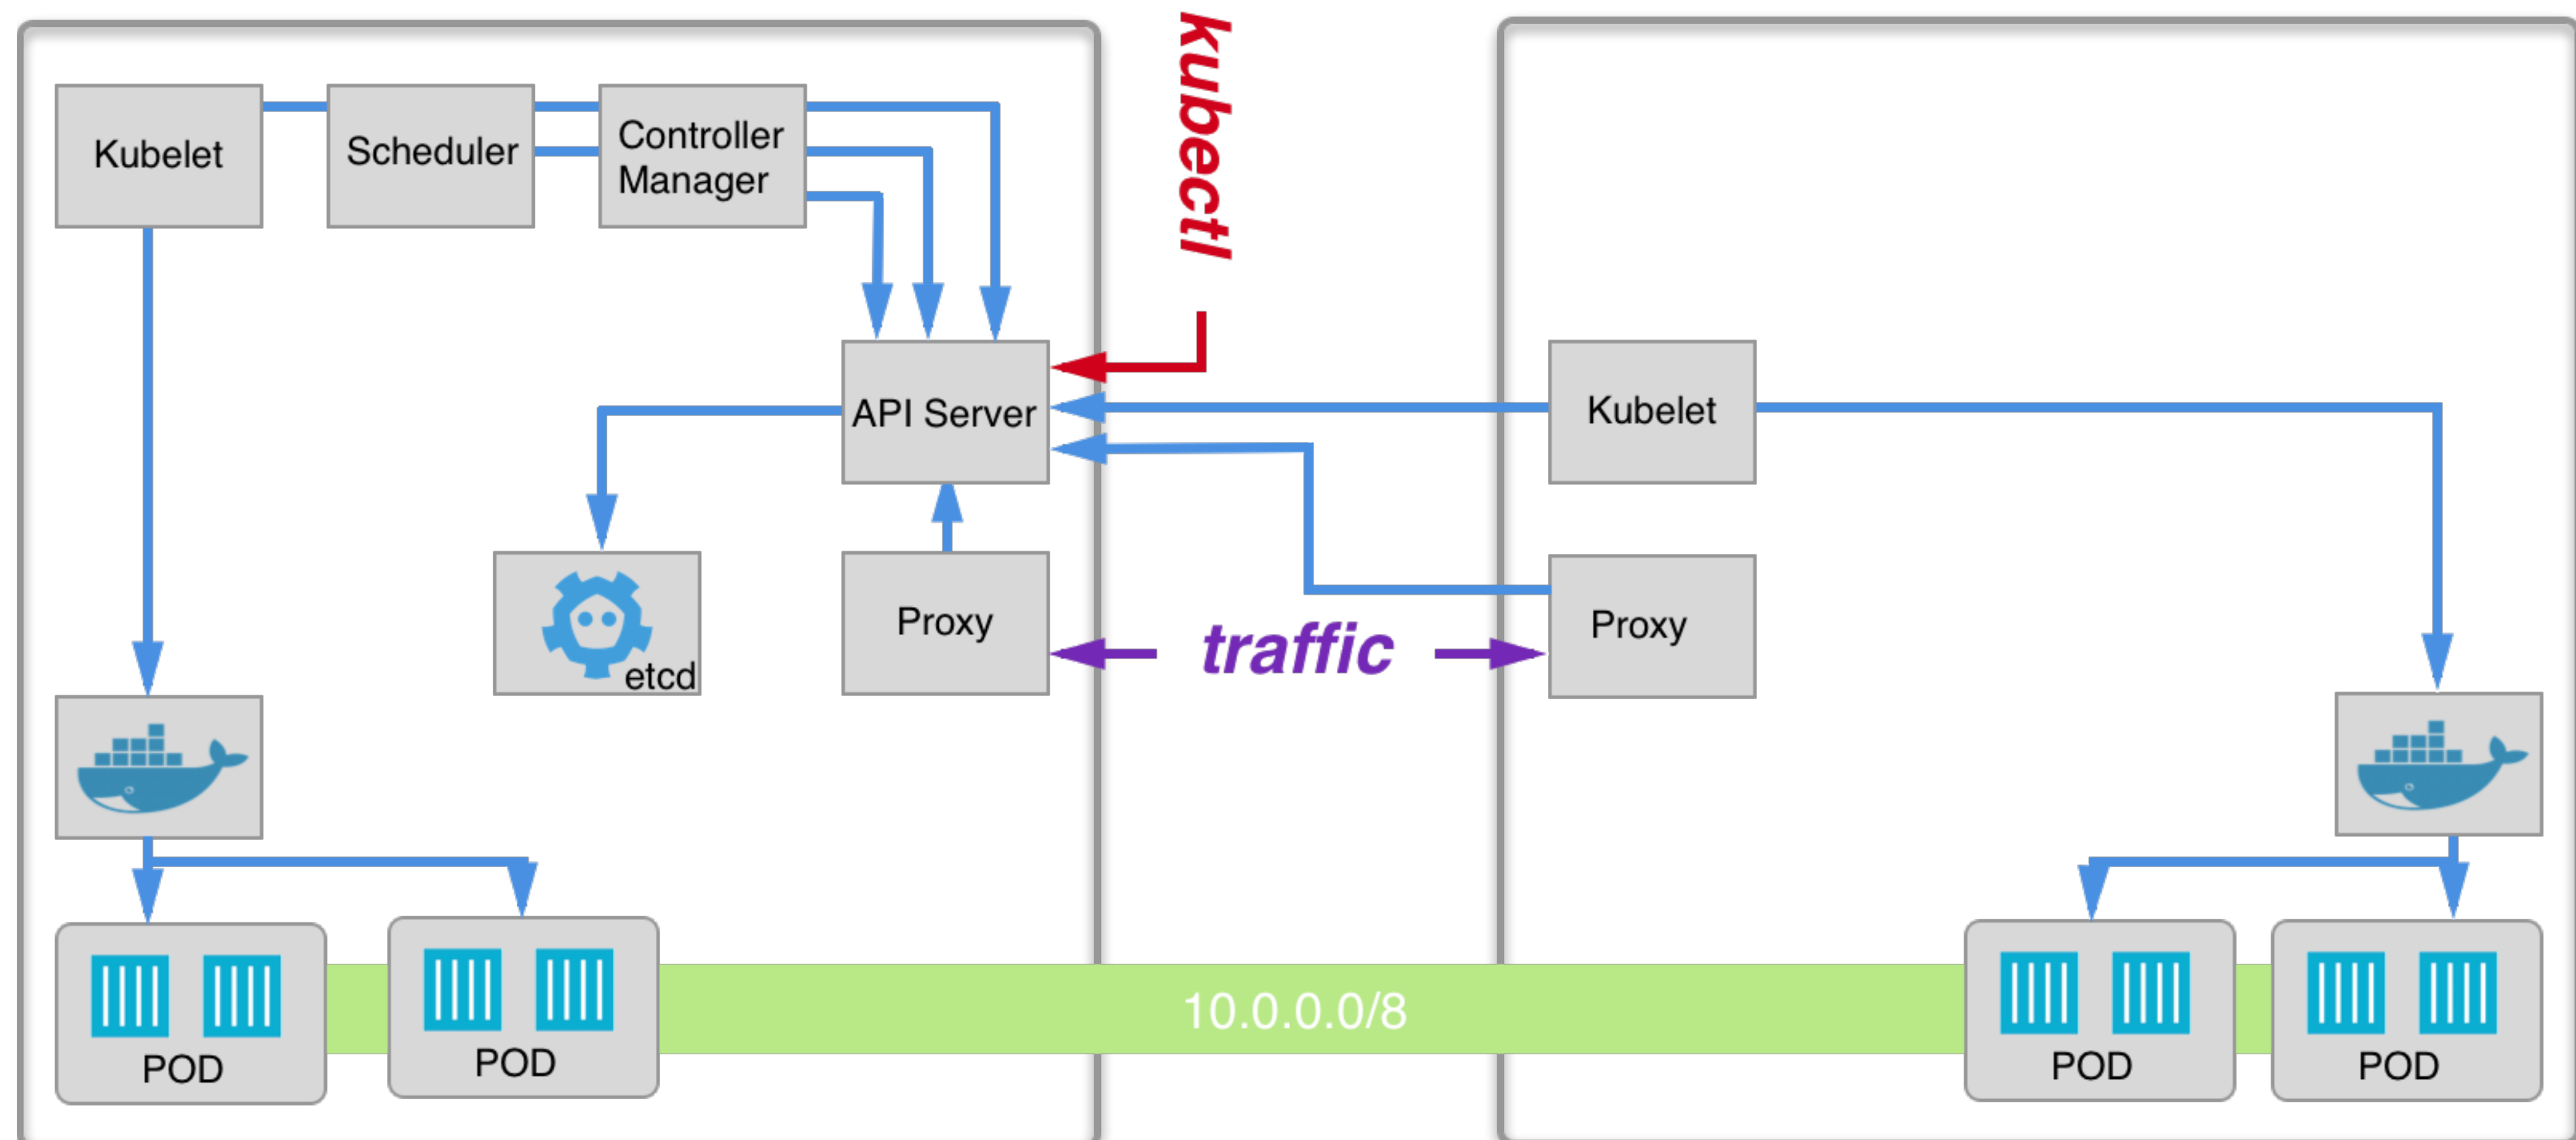
# NETWORKING

## Kubernetes requirements:

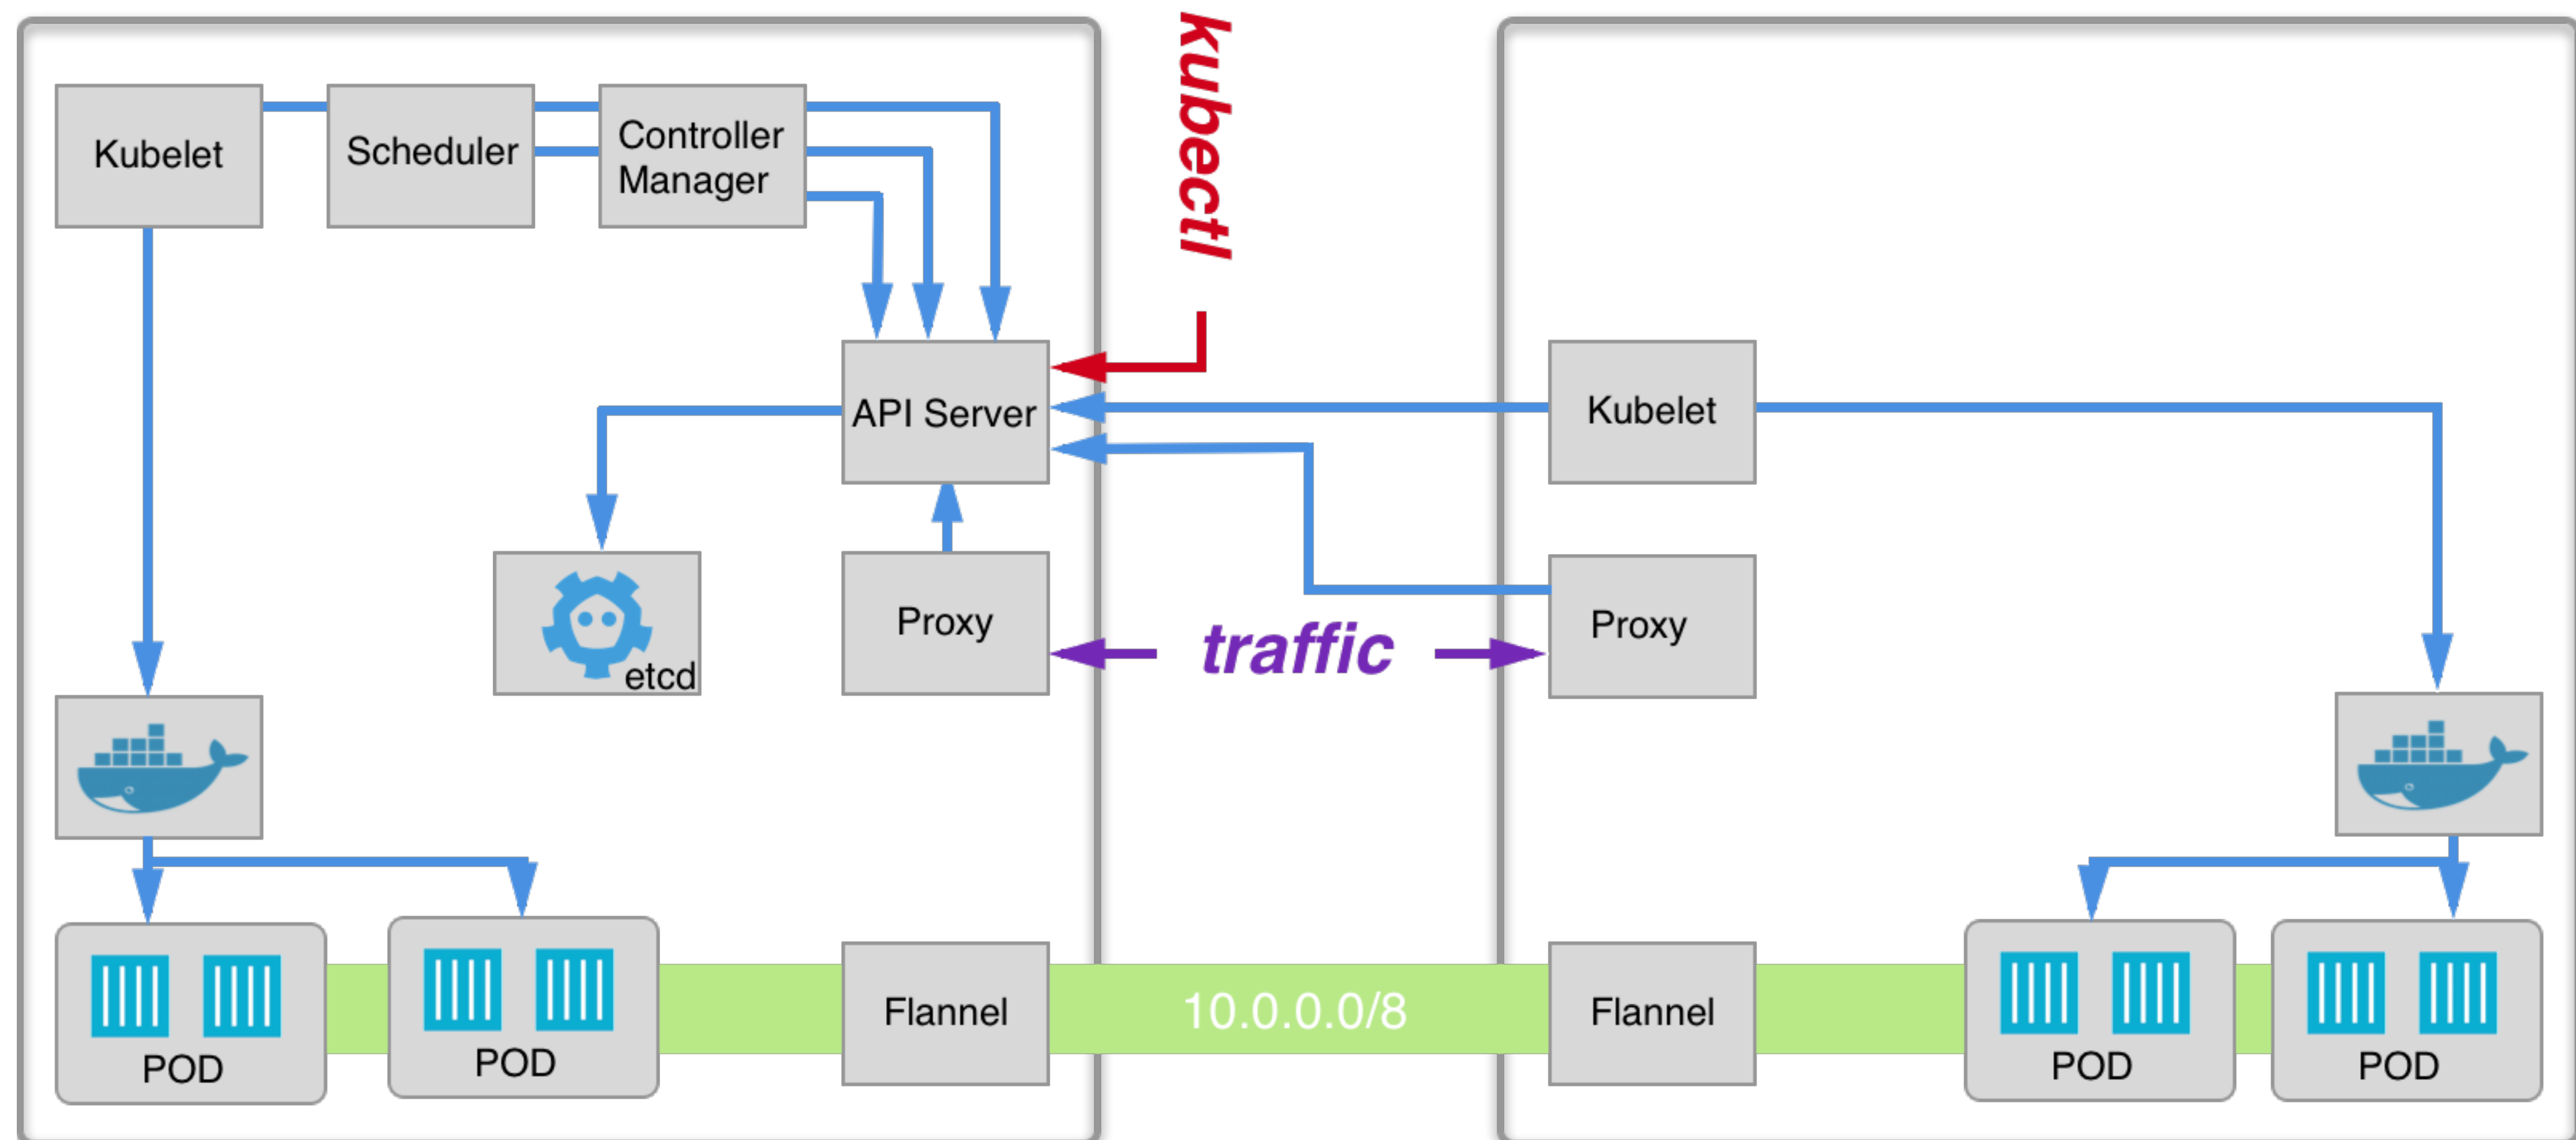
- each pod has it's own IP
- all containers can communicate with all other containers using their IPs



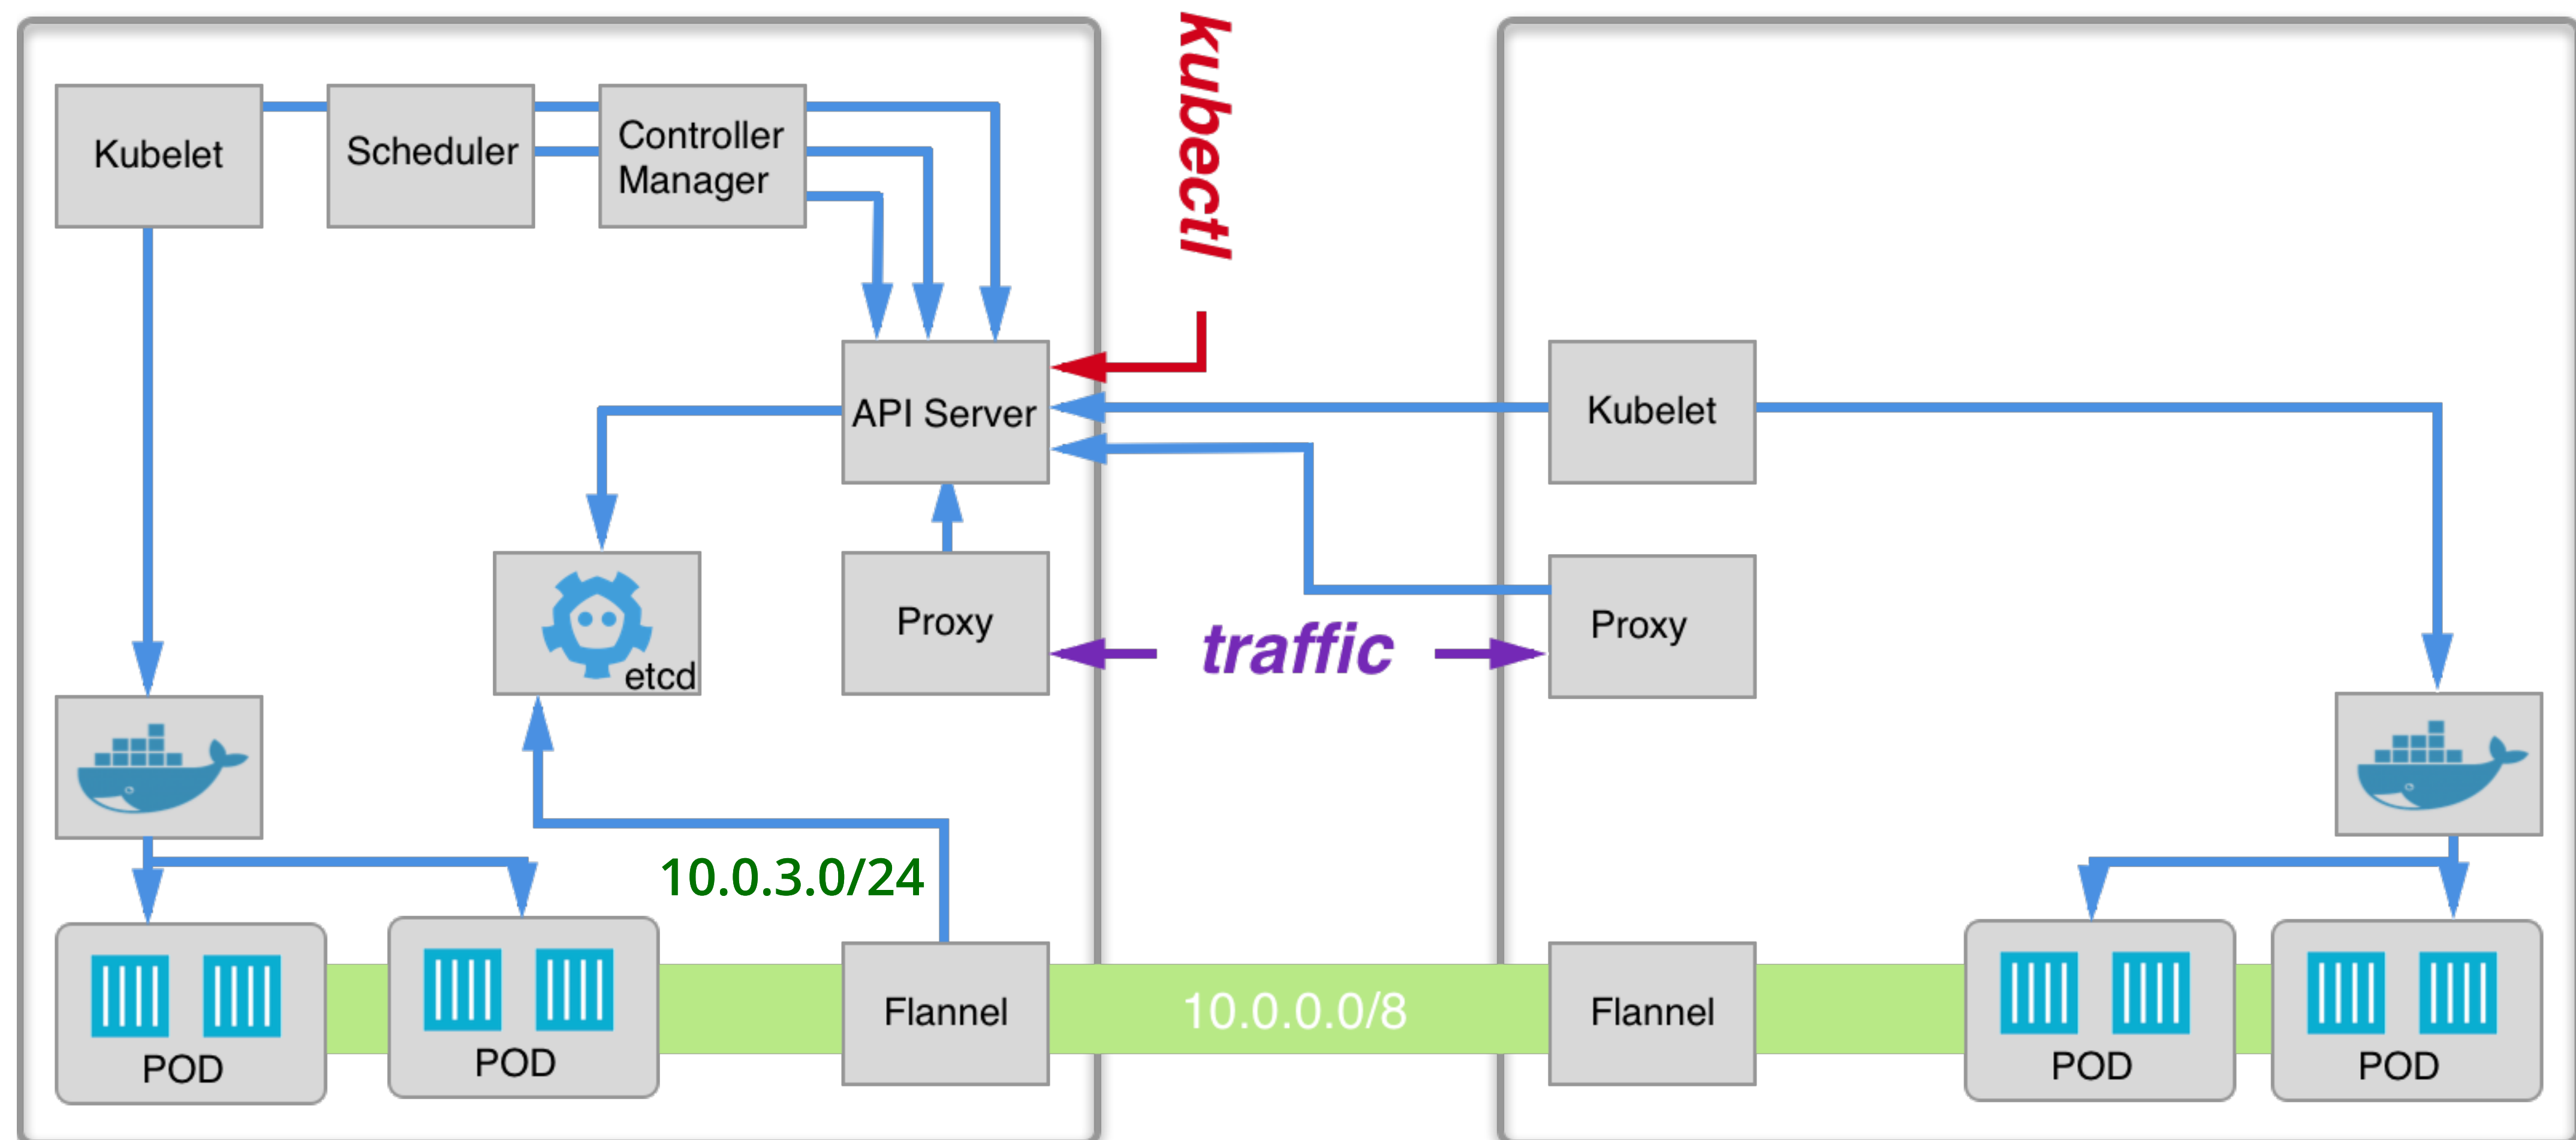
# NETWORKING



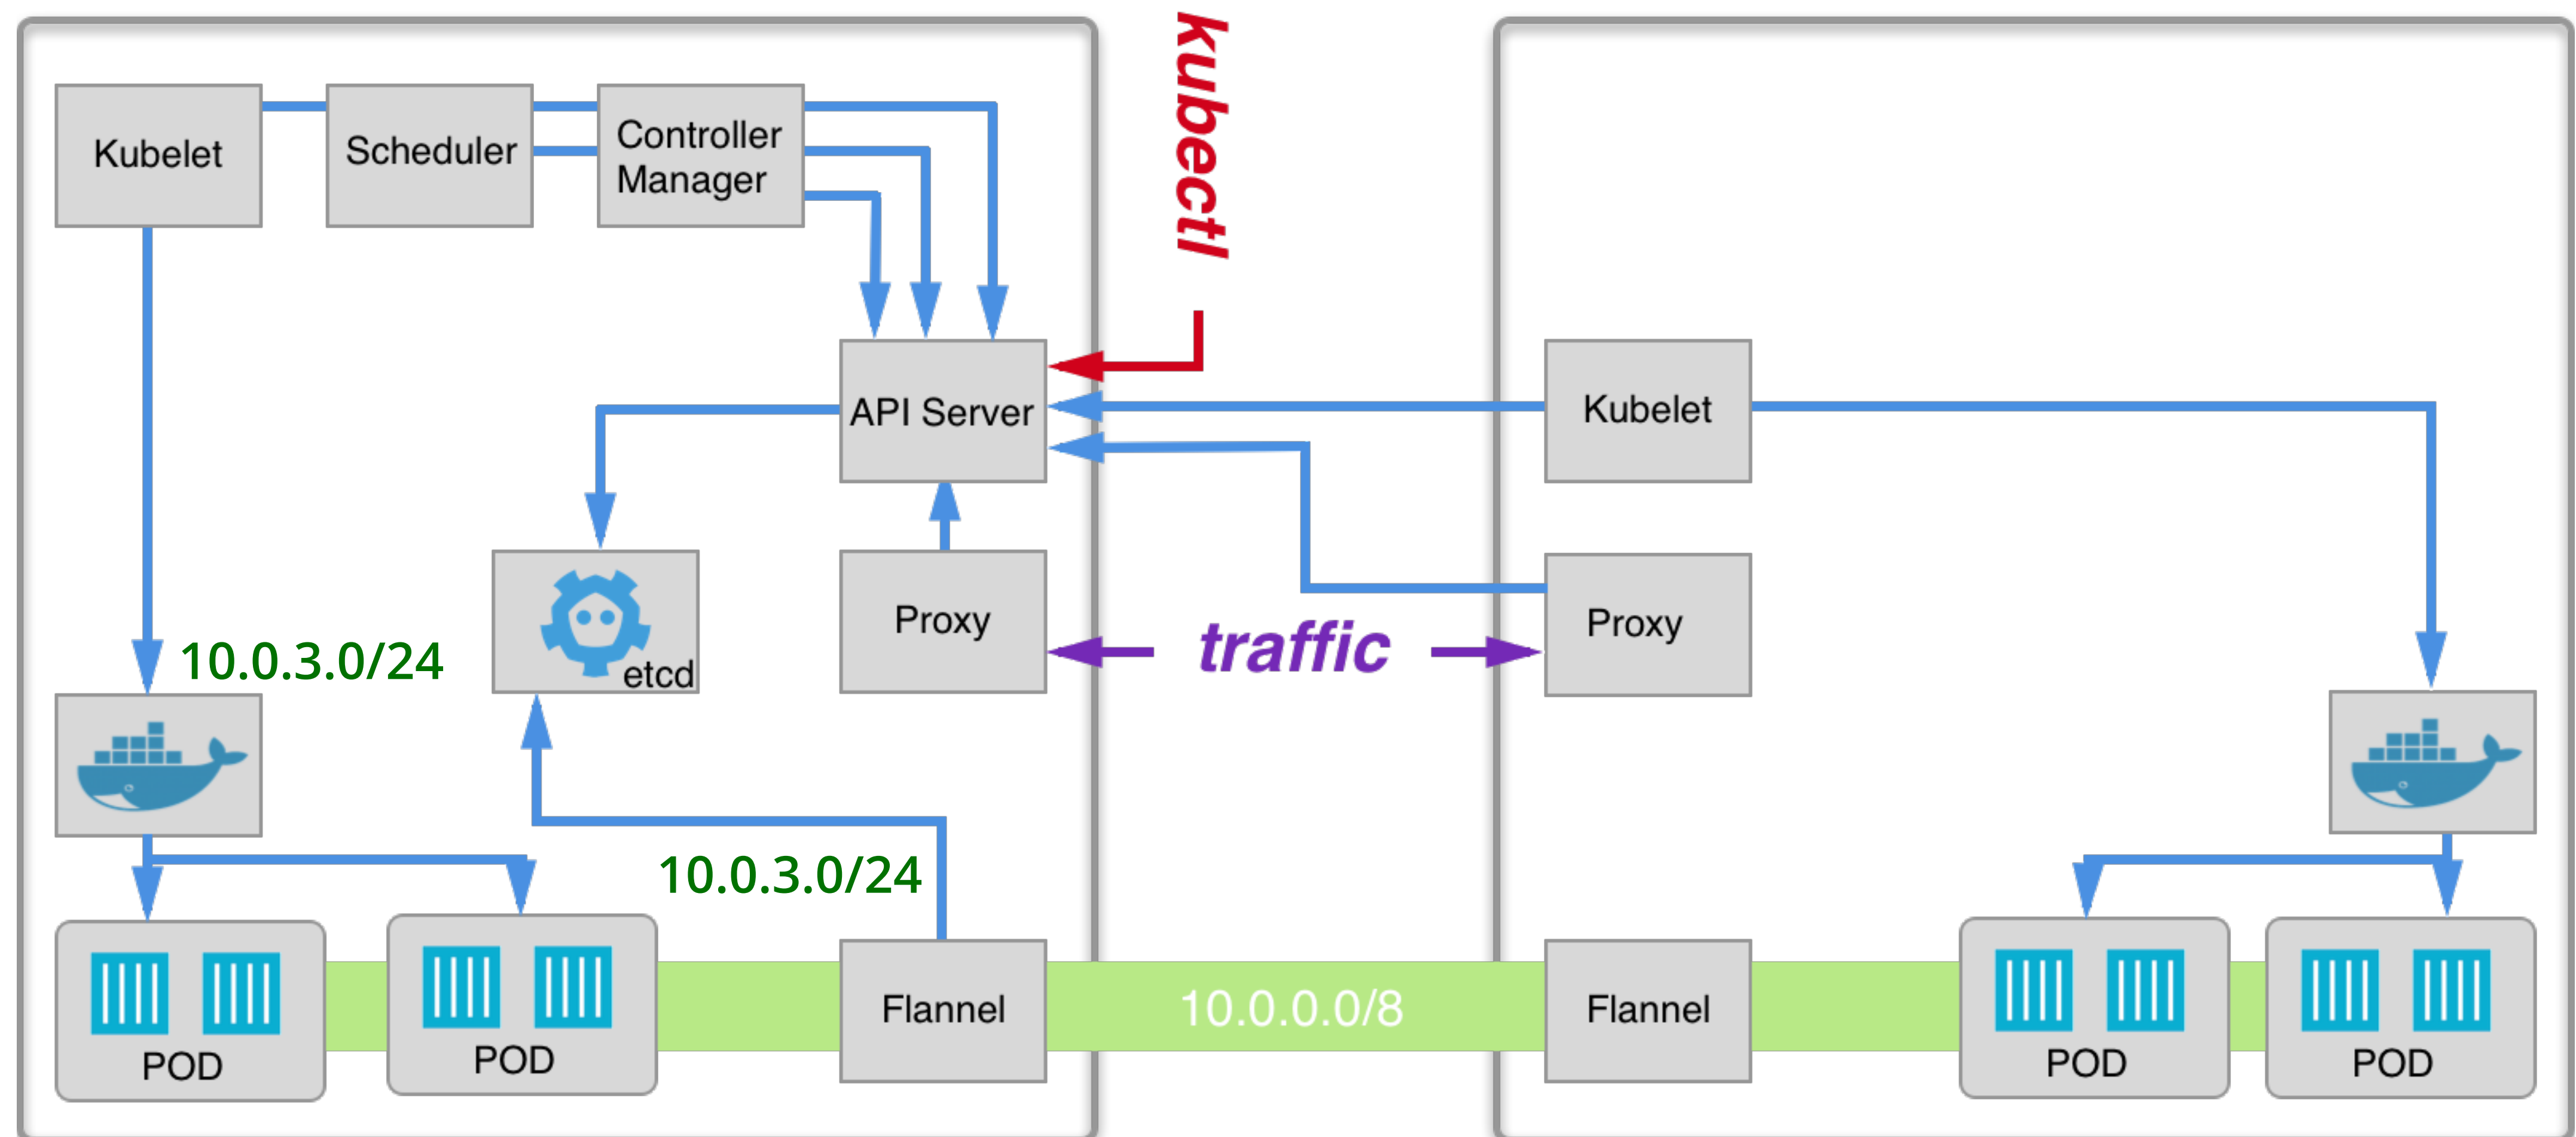
# NETWORKING



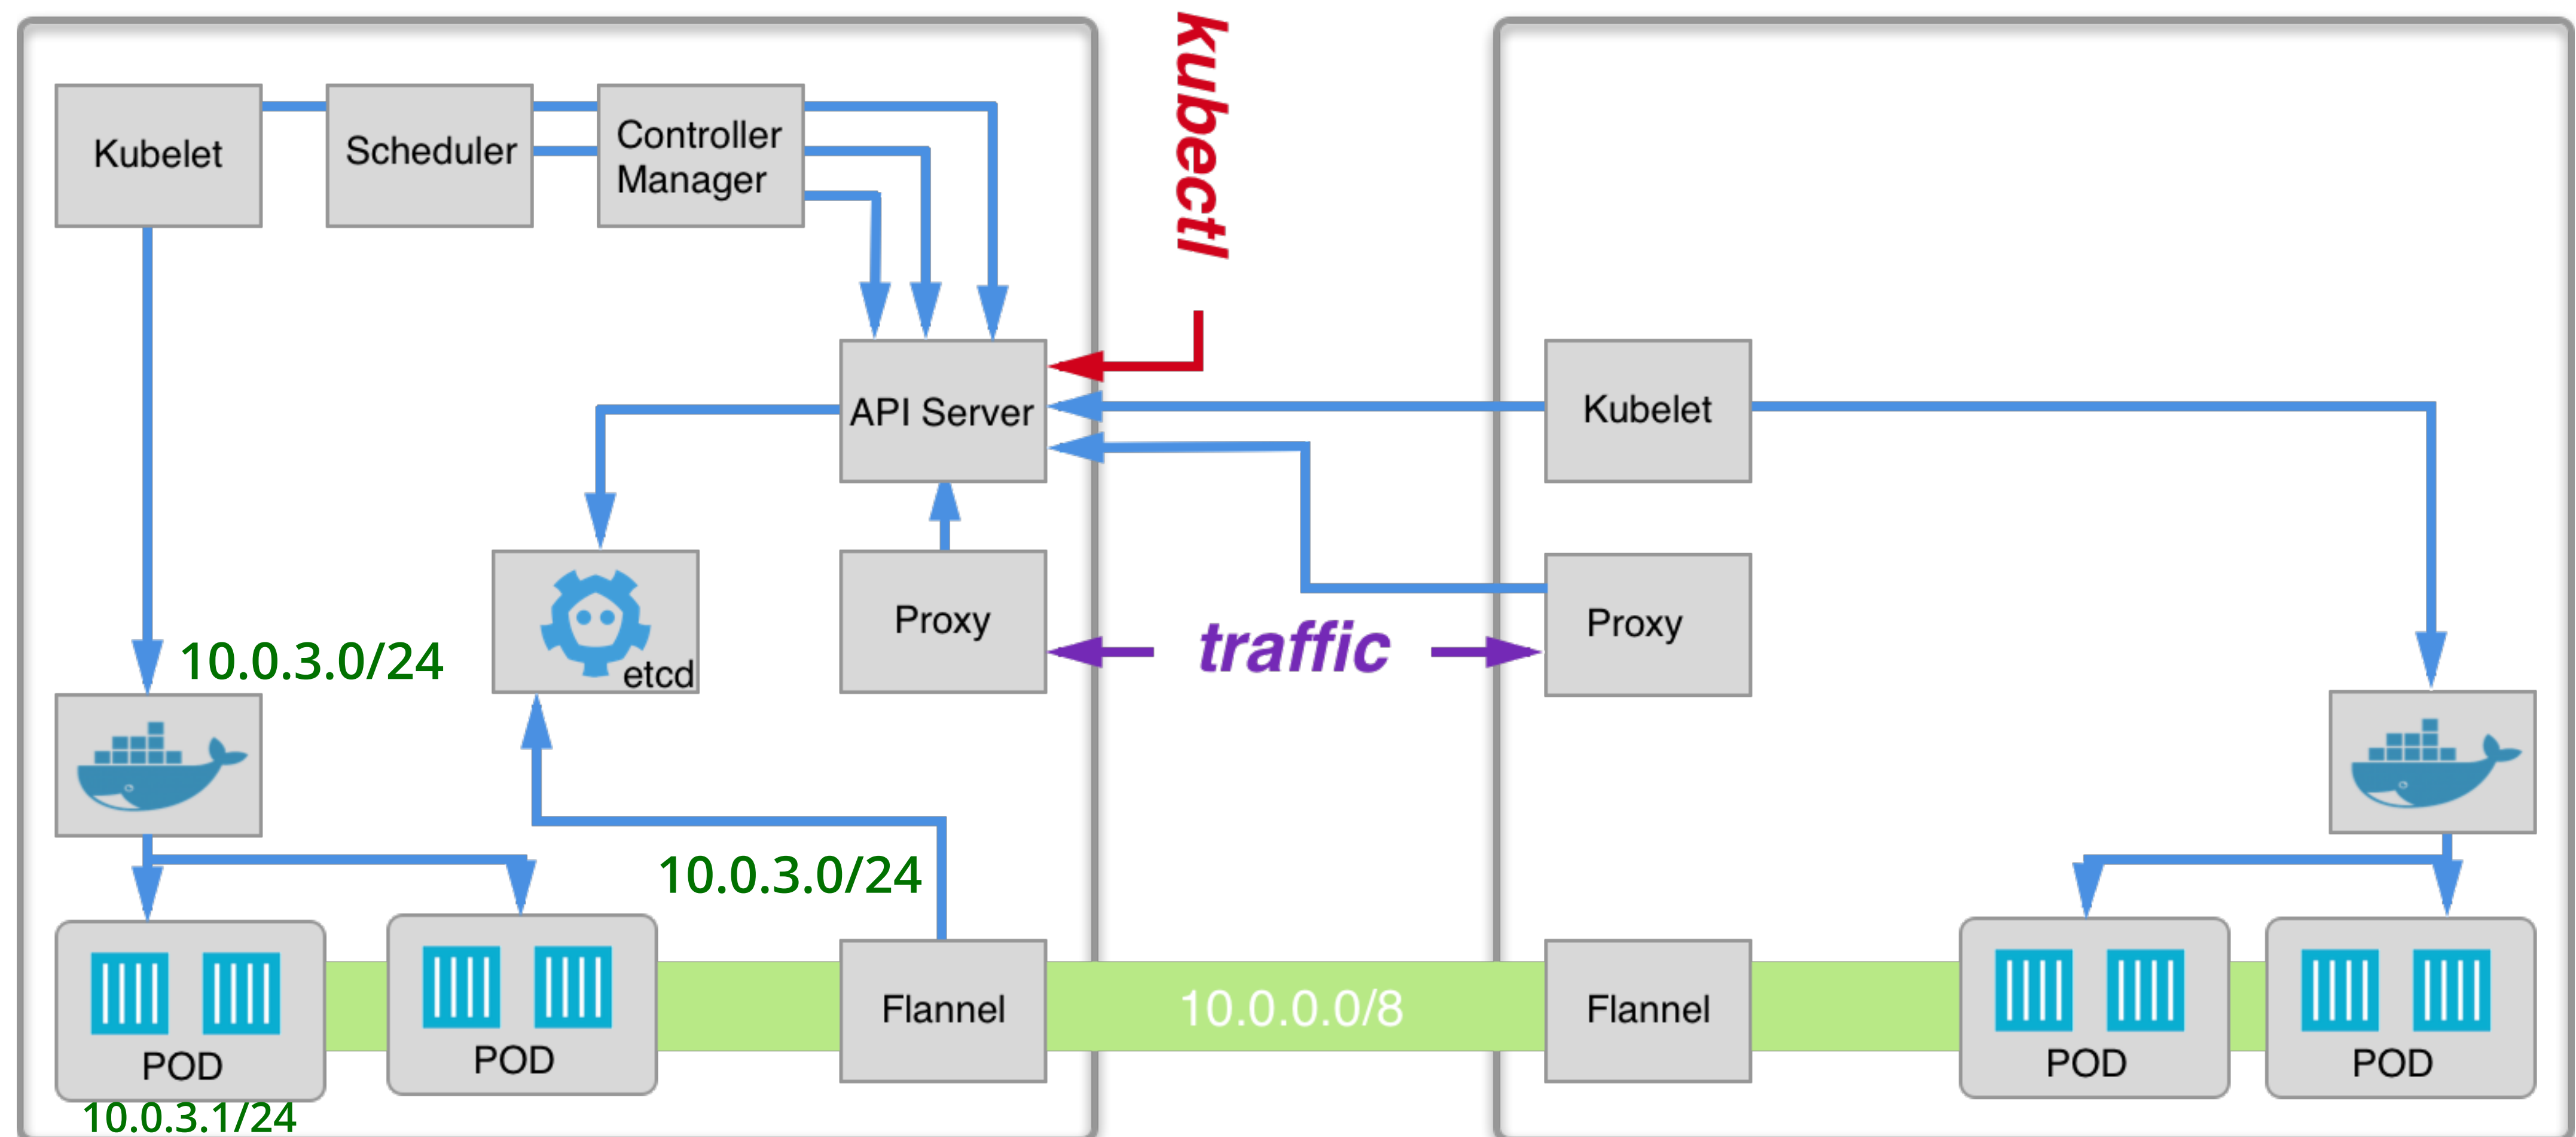
# NETWORKING



# NETWORKING

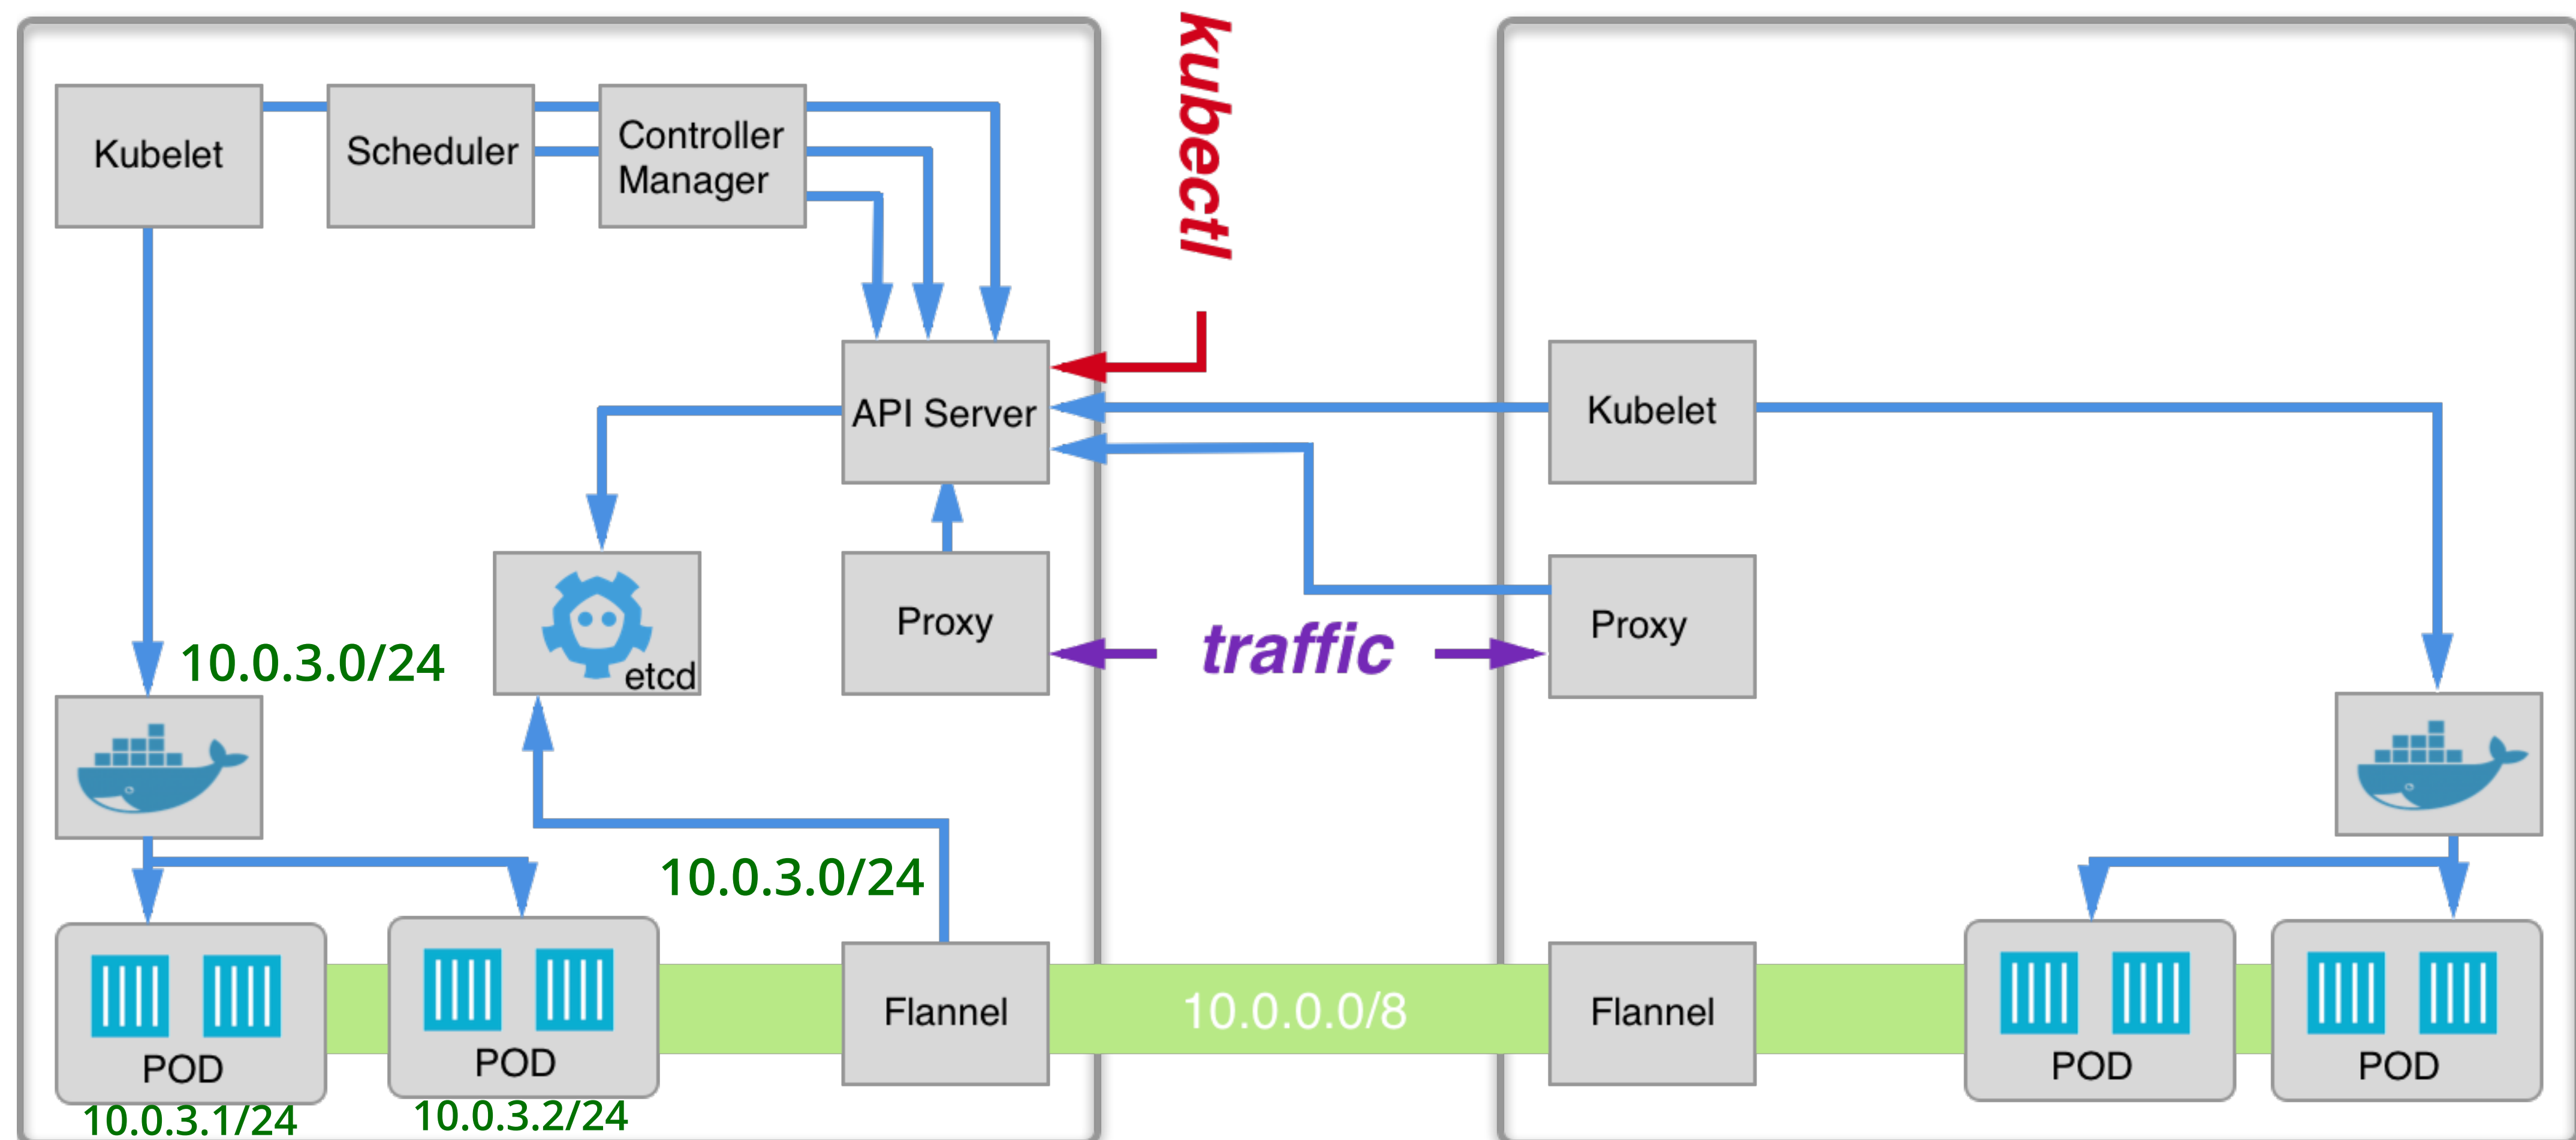


# NETWORKING

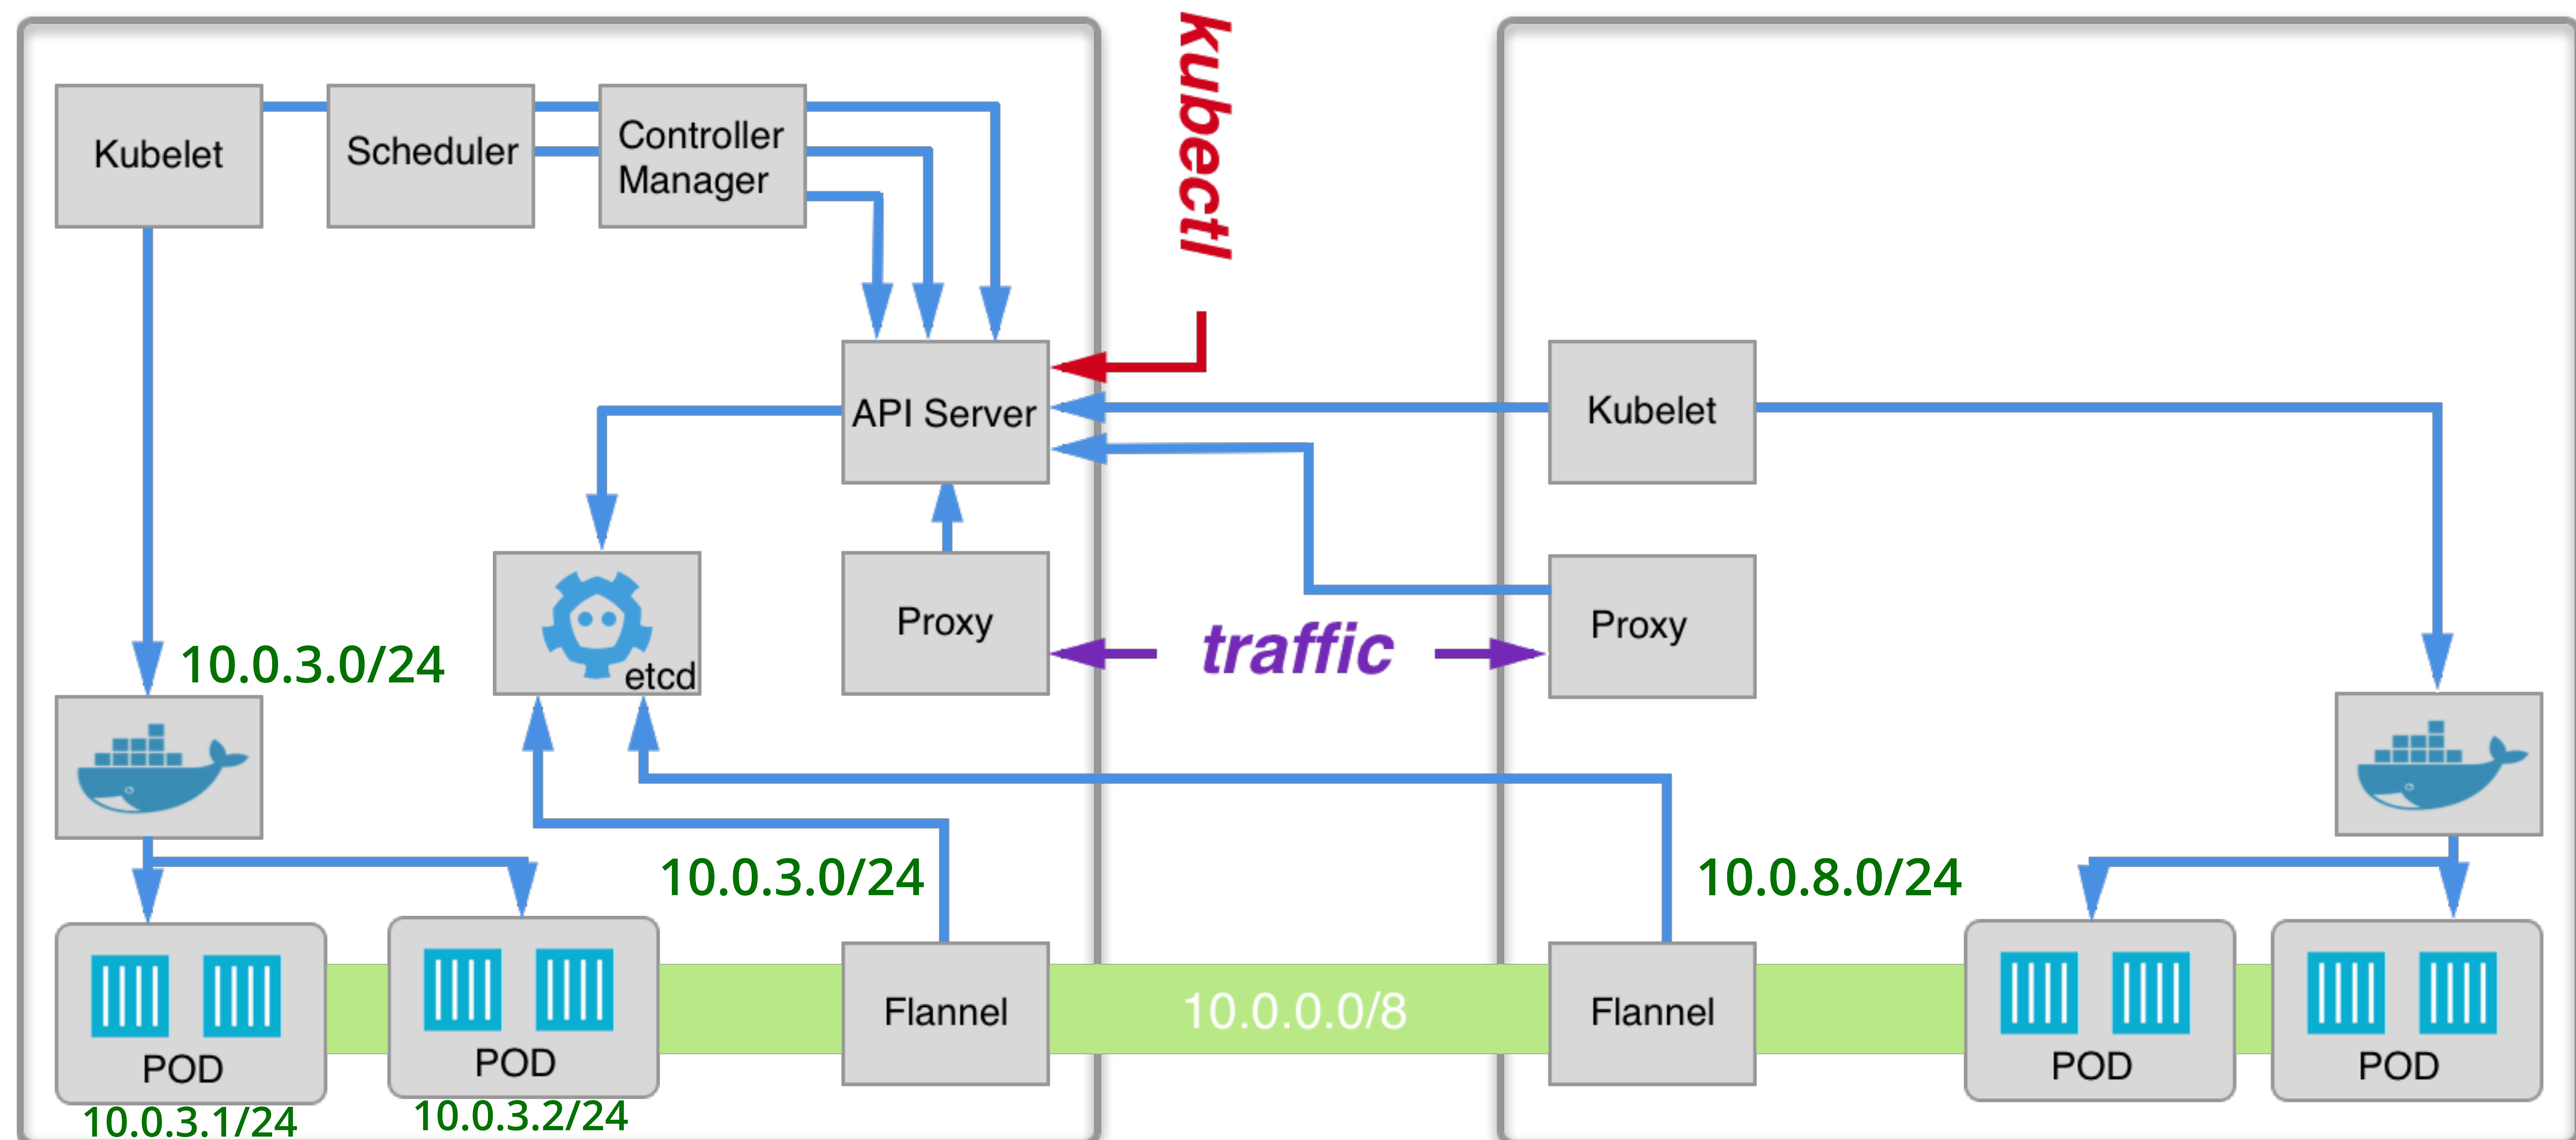




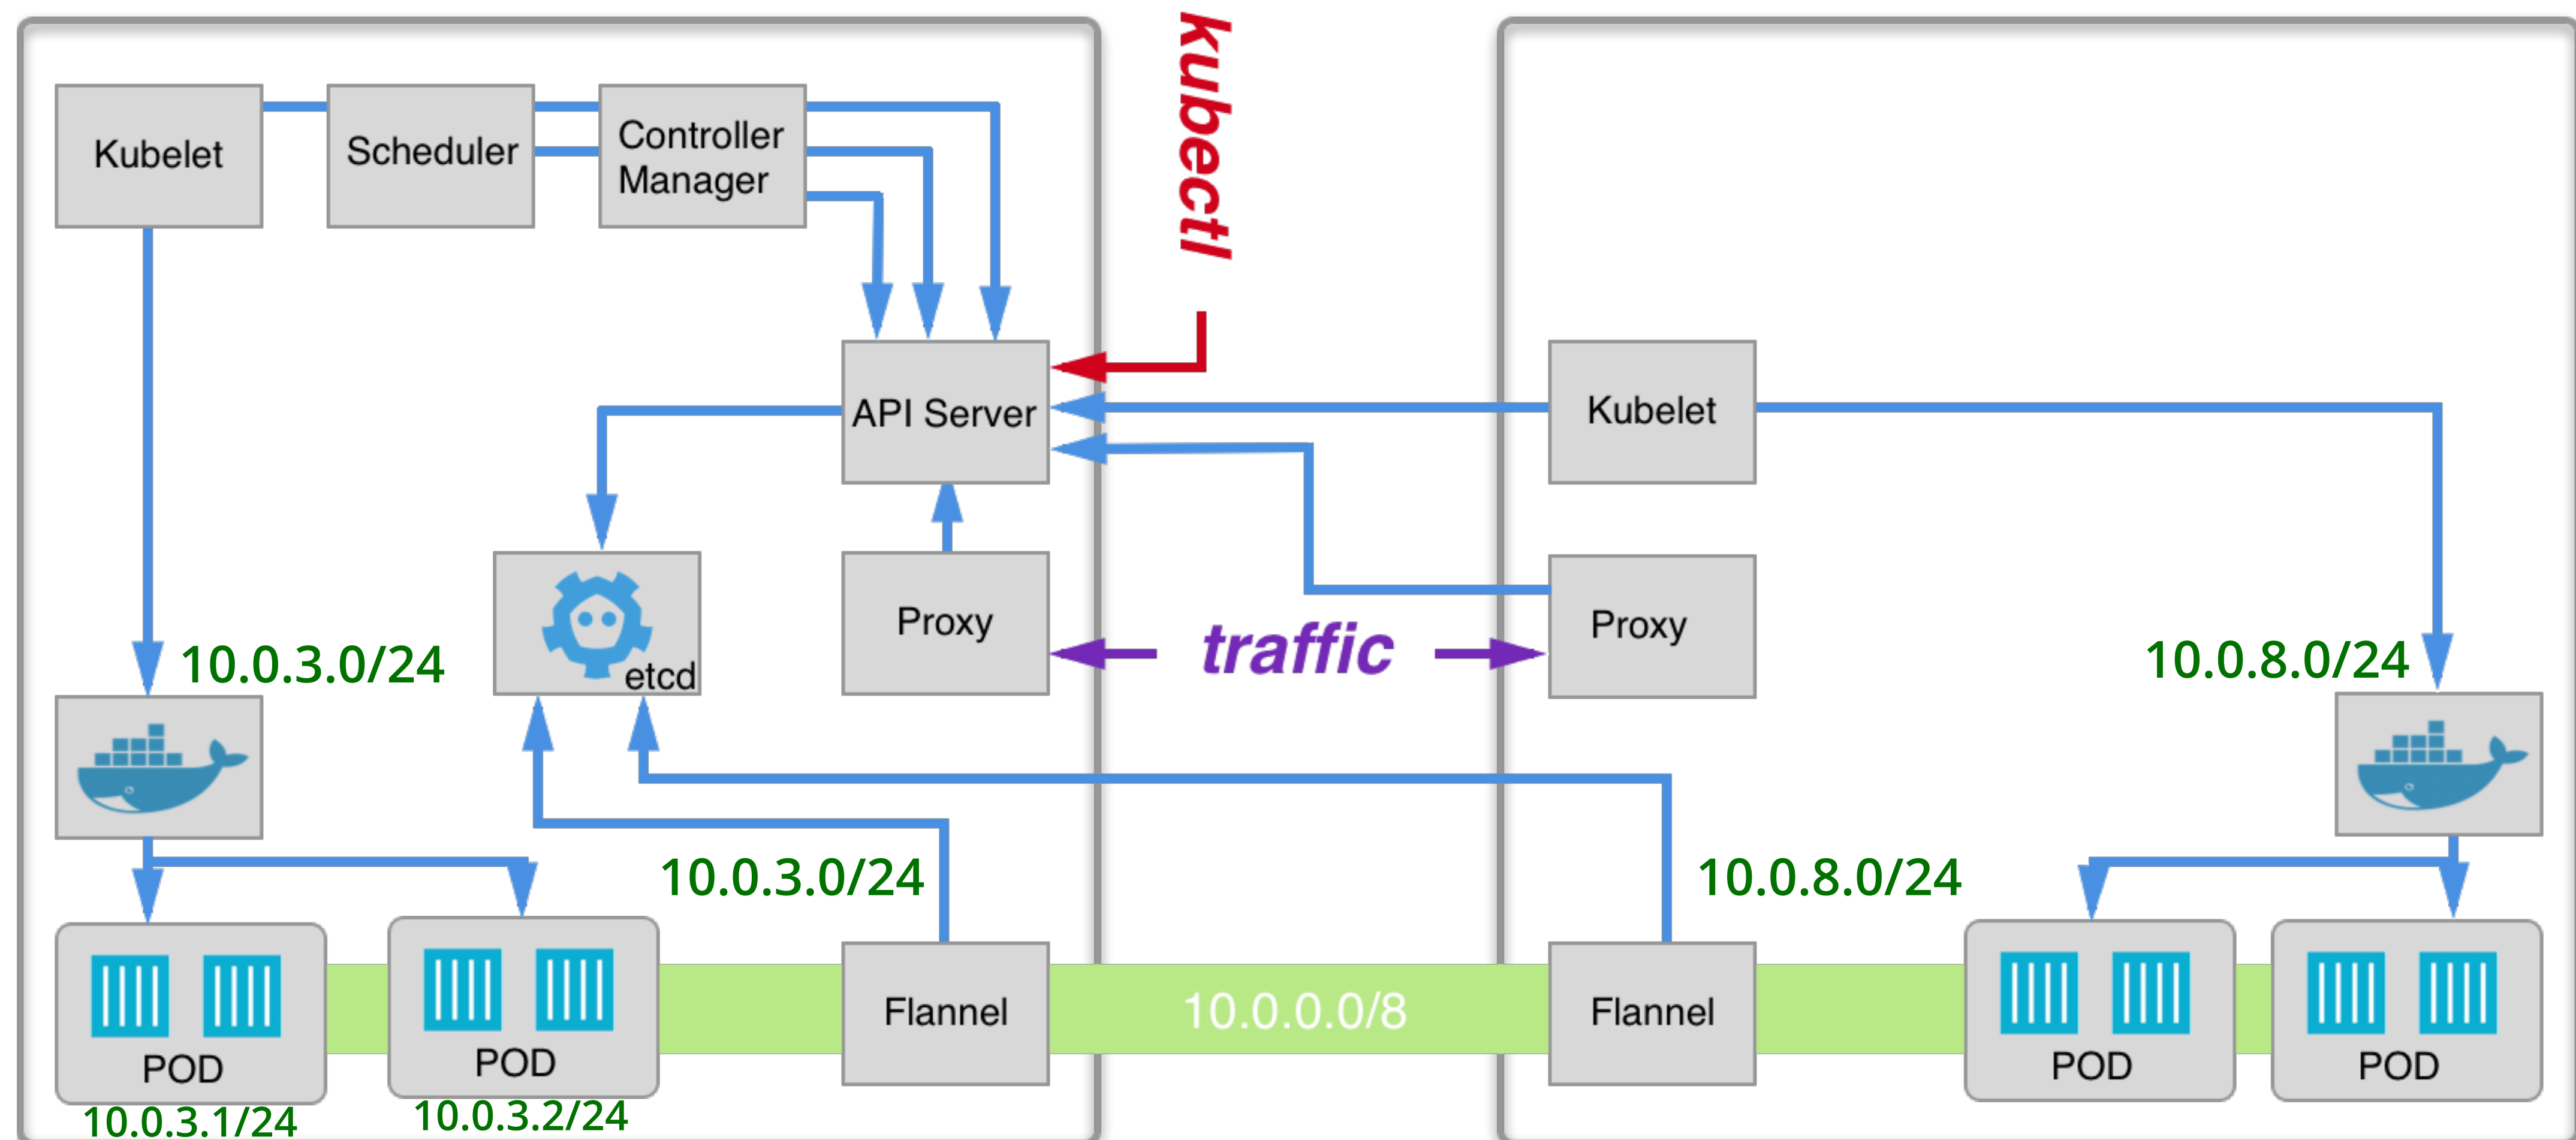
# NETWORKING



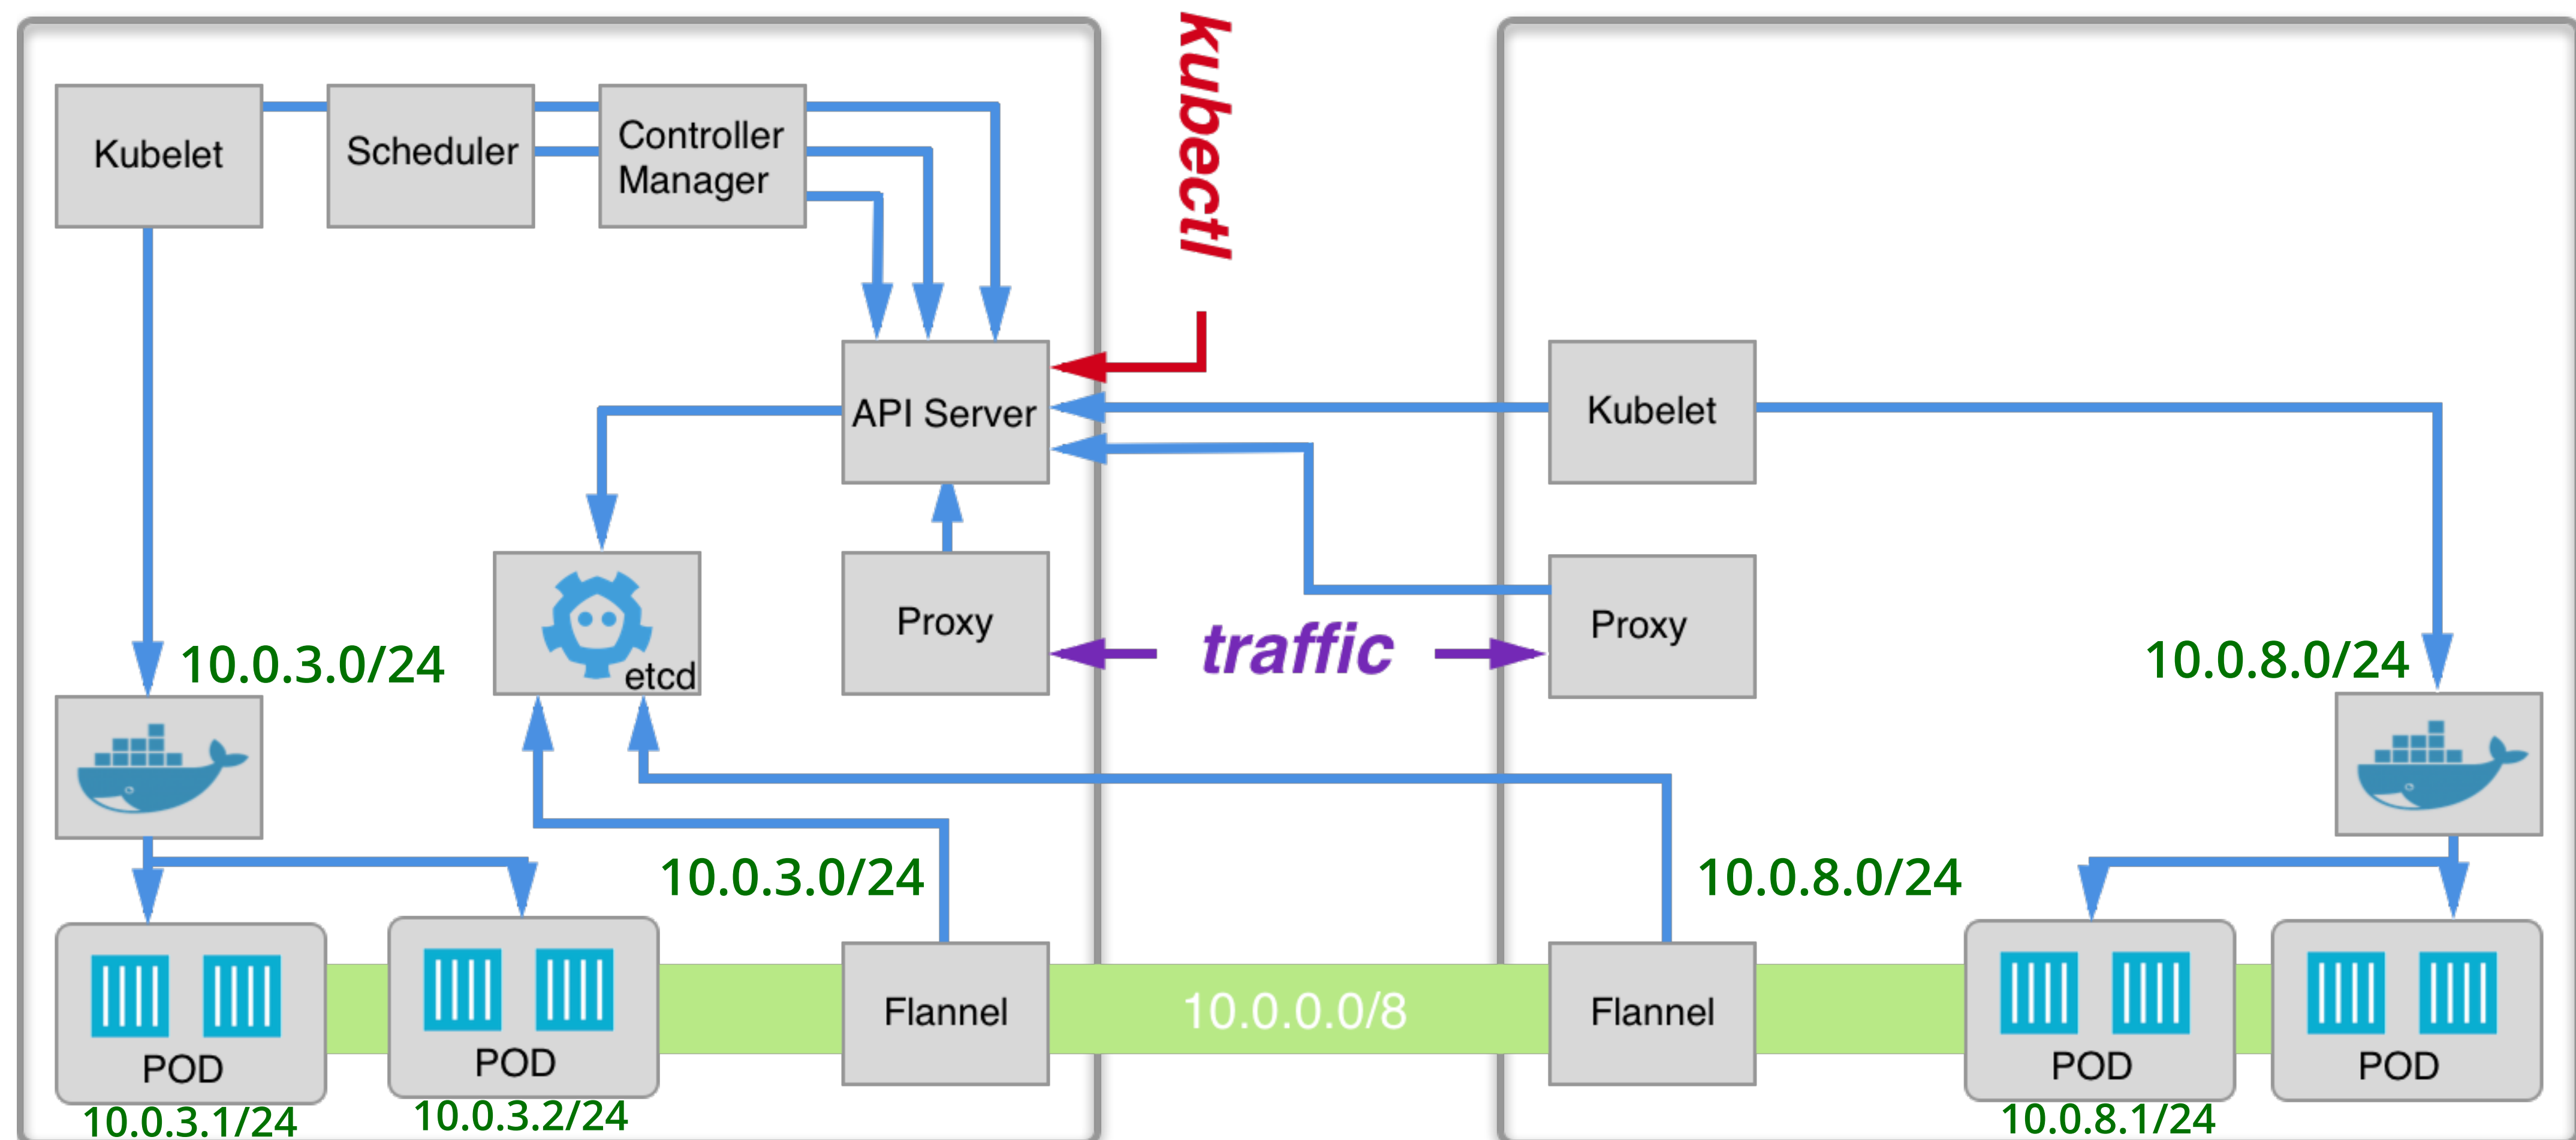
# NETWORKING



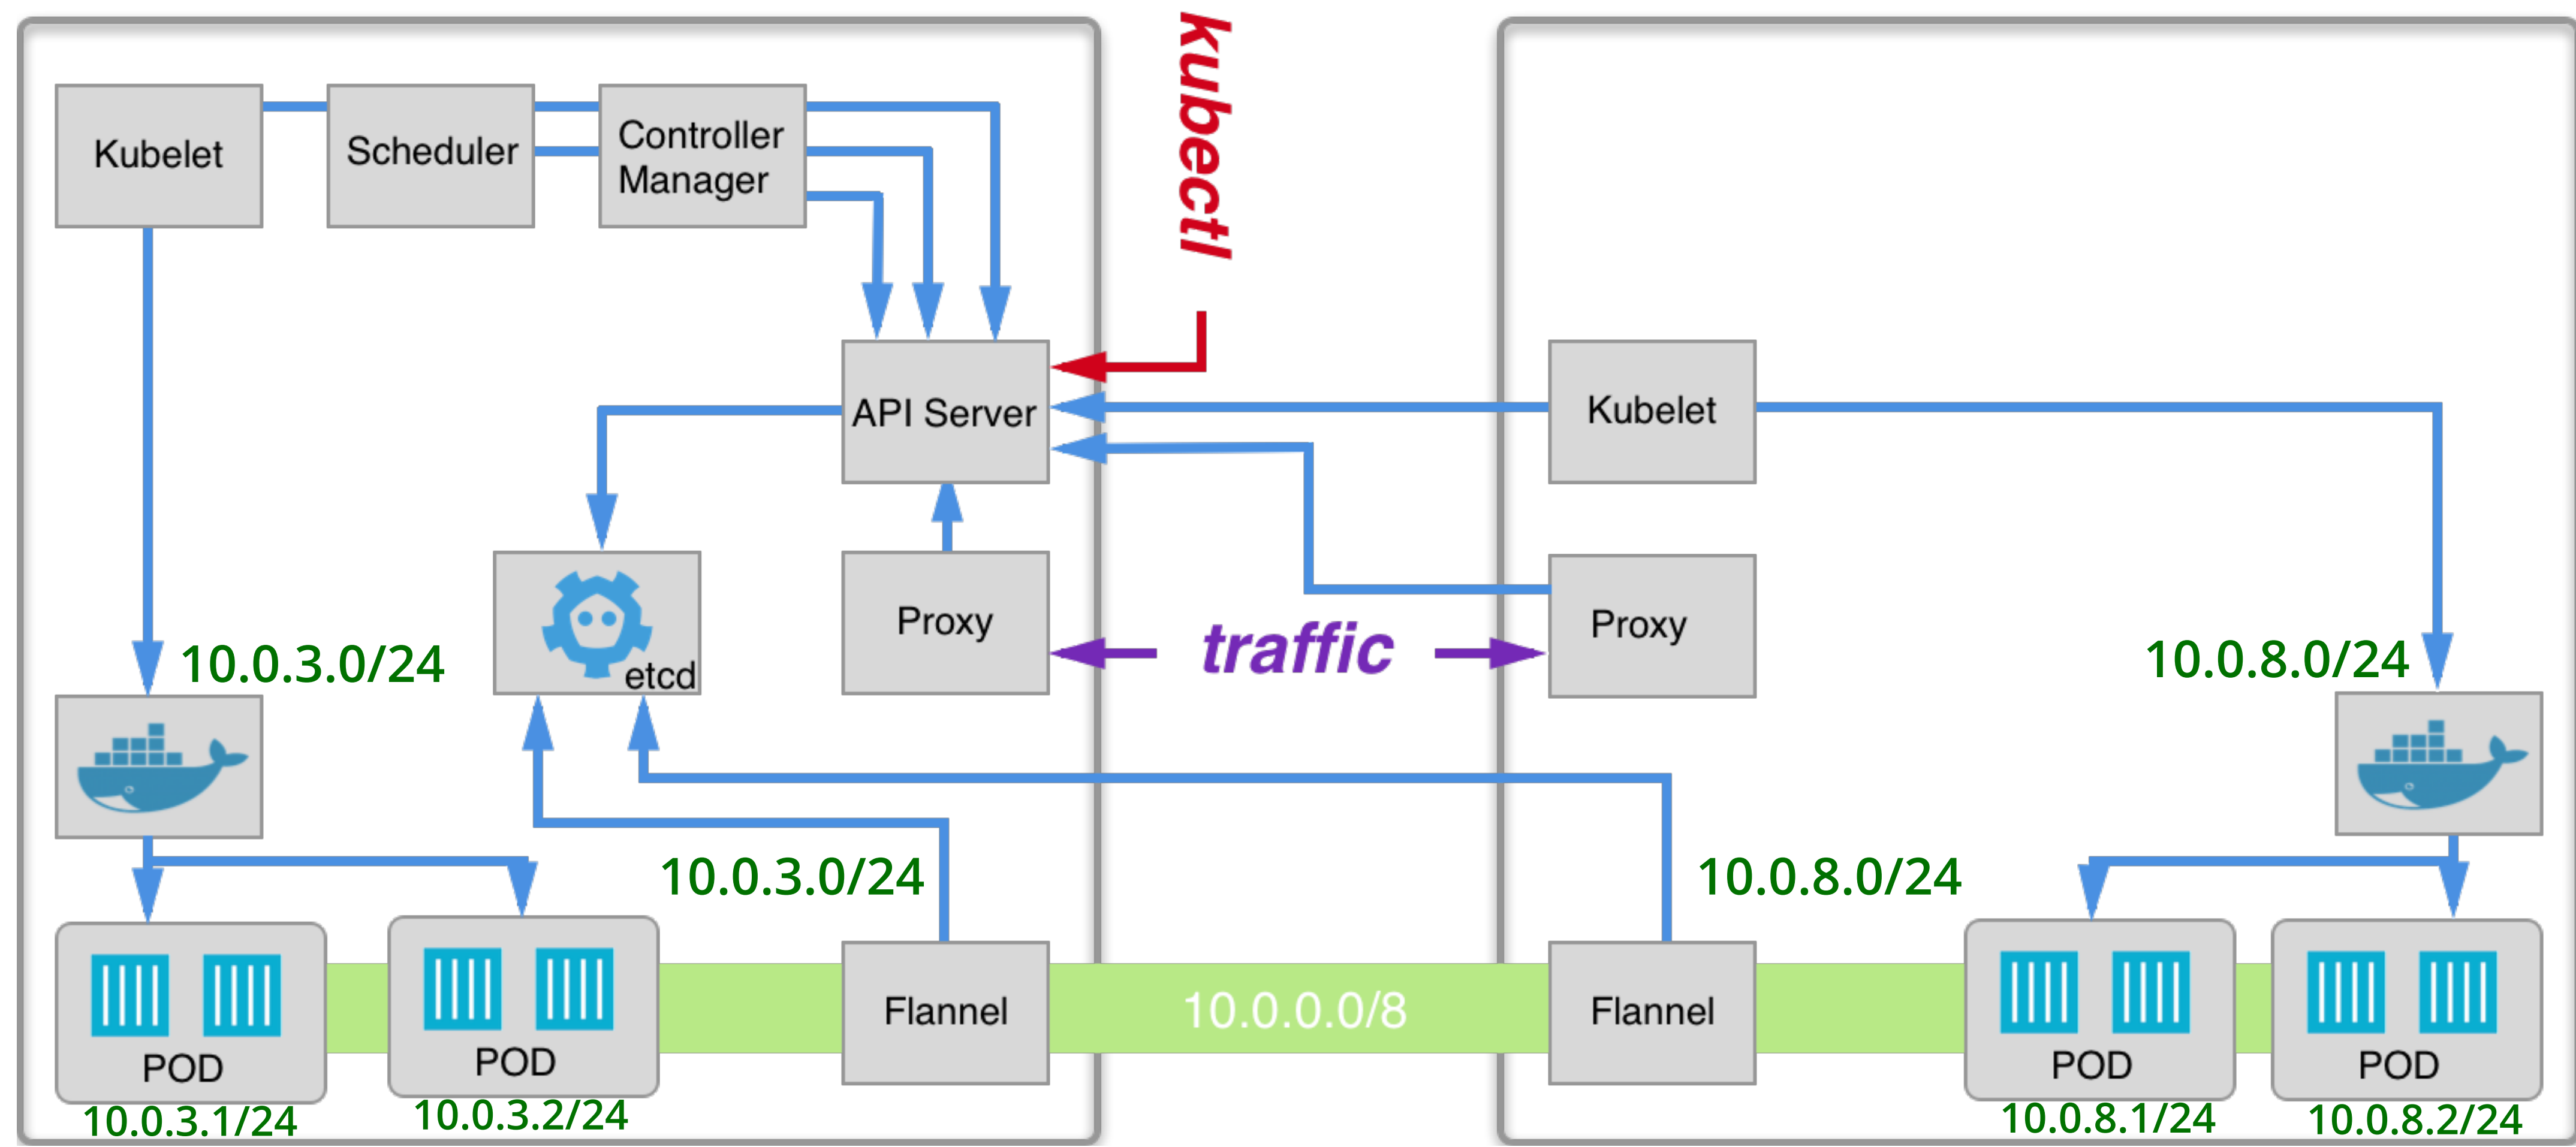
# NETWORKING



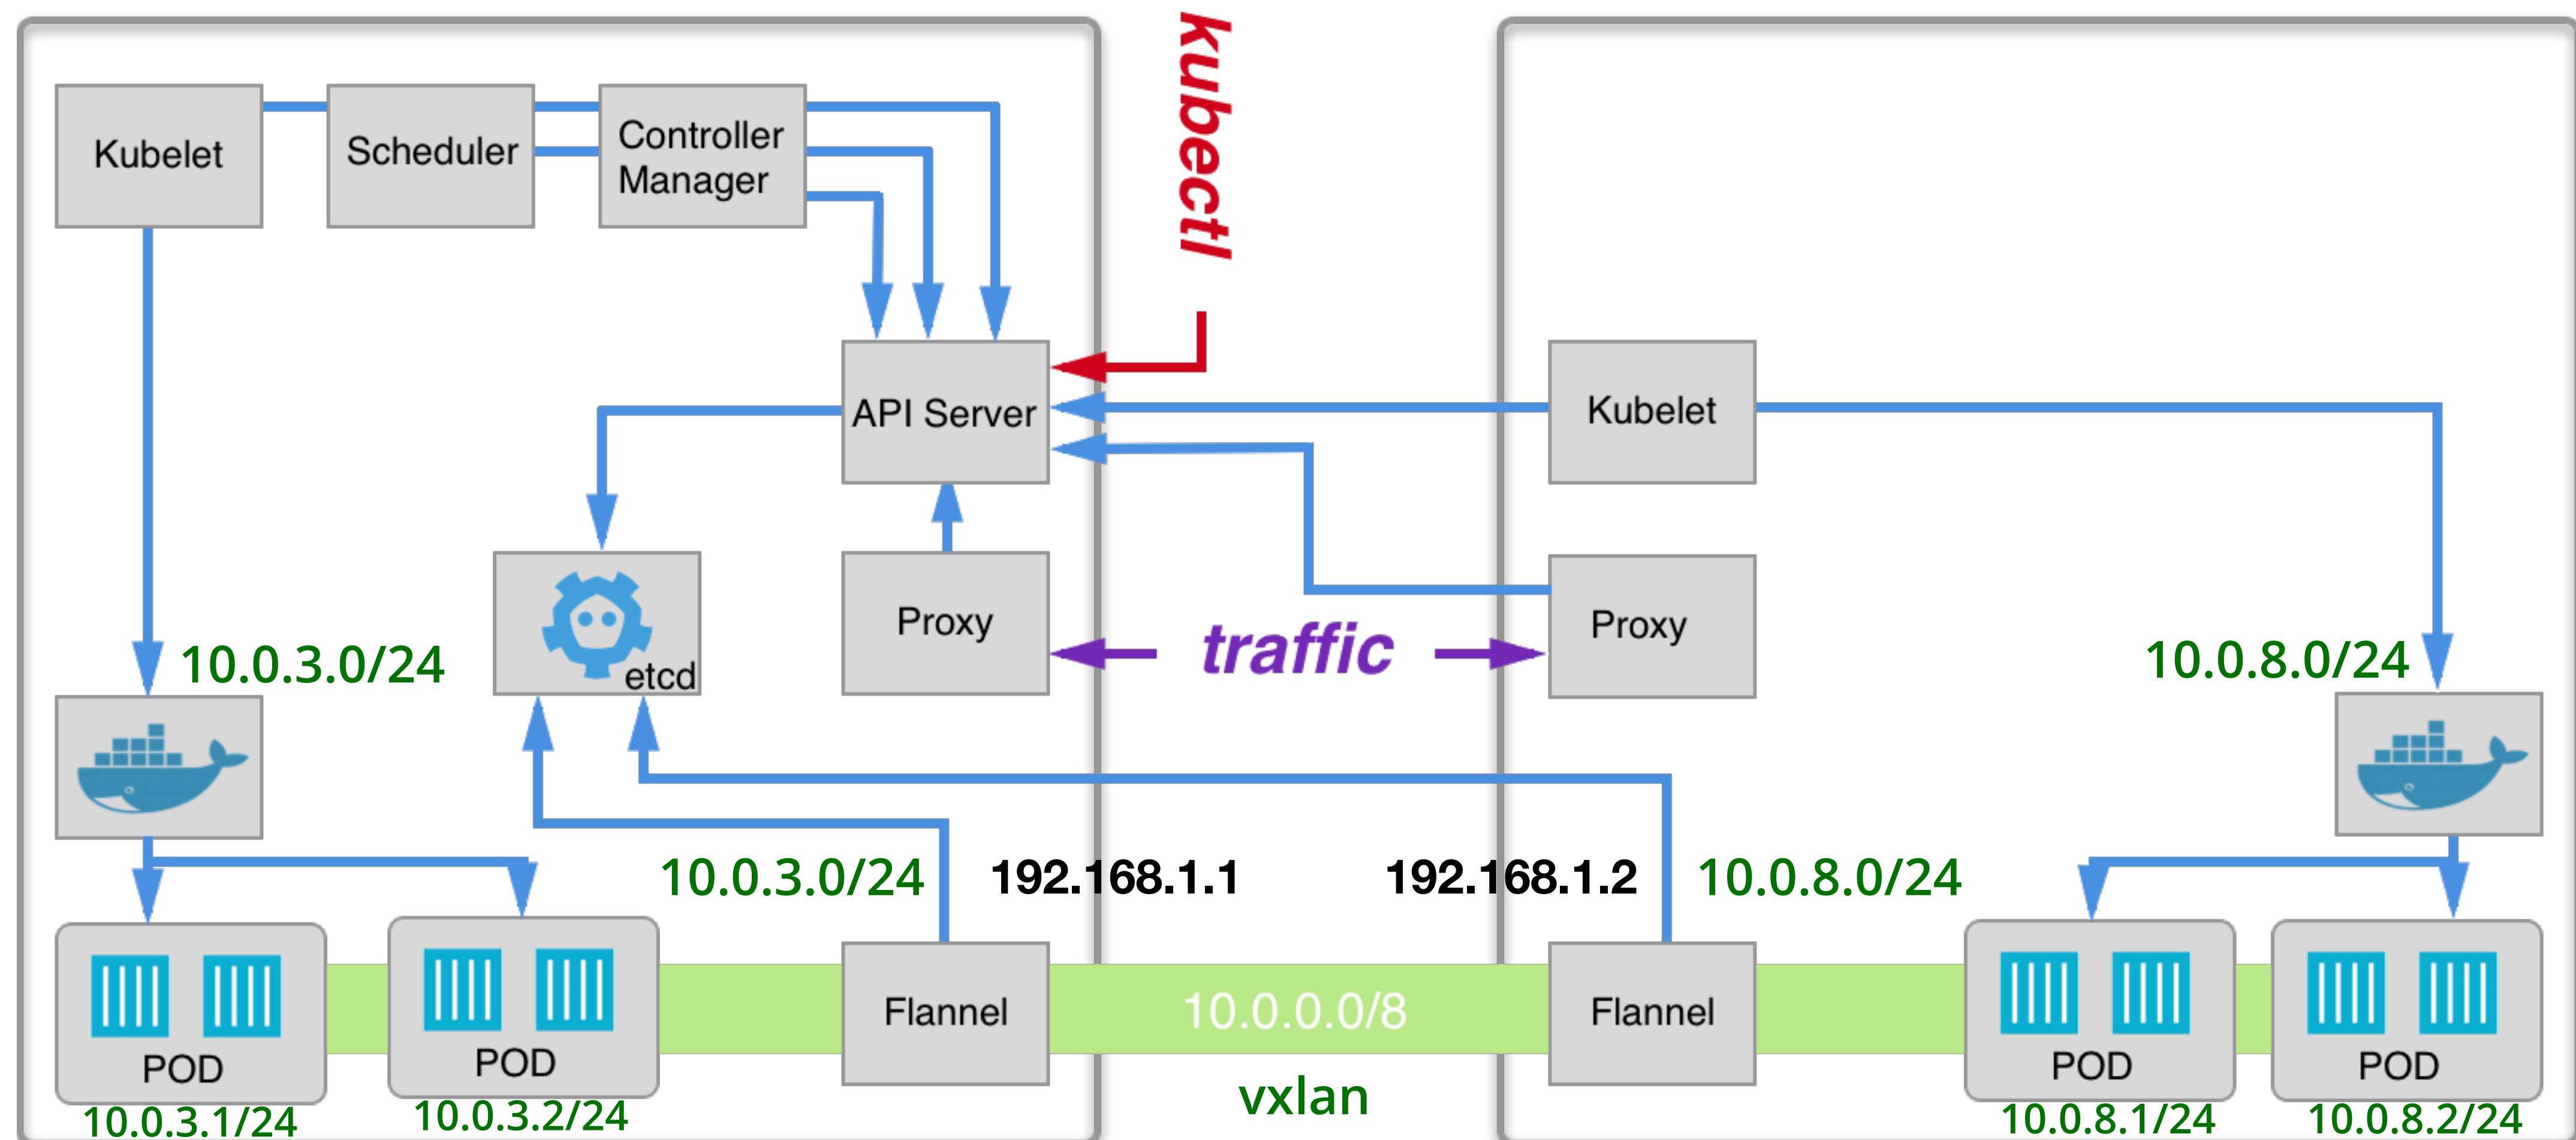
# NETWORKING



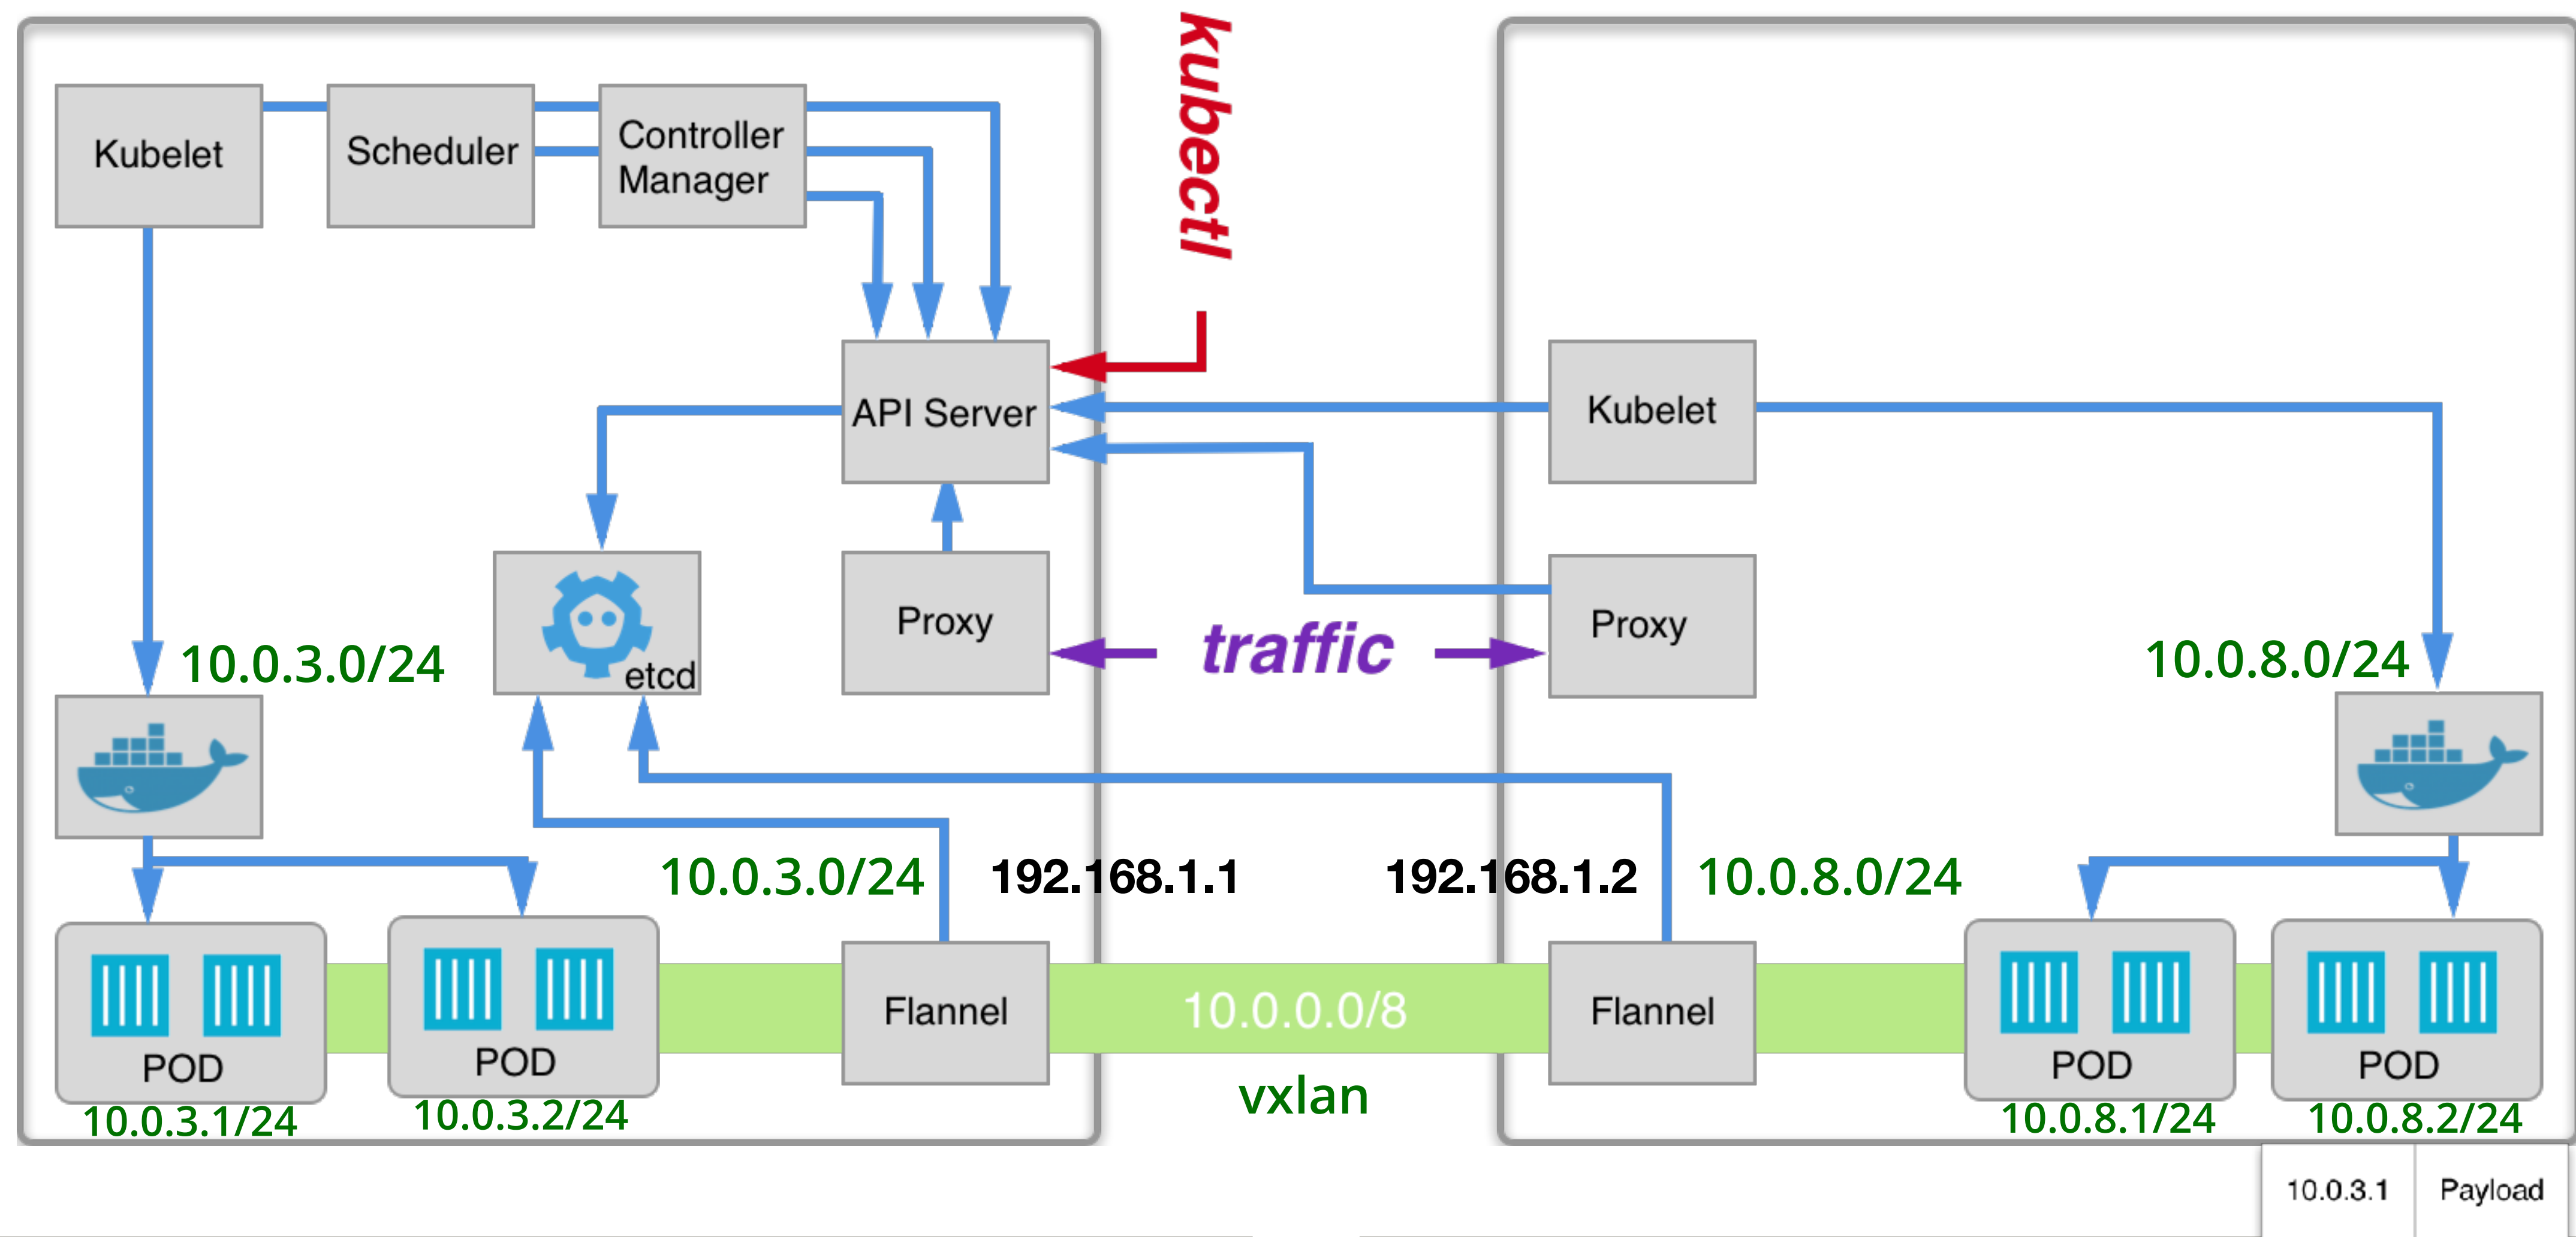
# NETWORKING



# NETWORKING

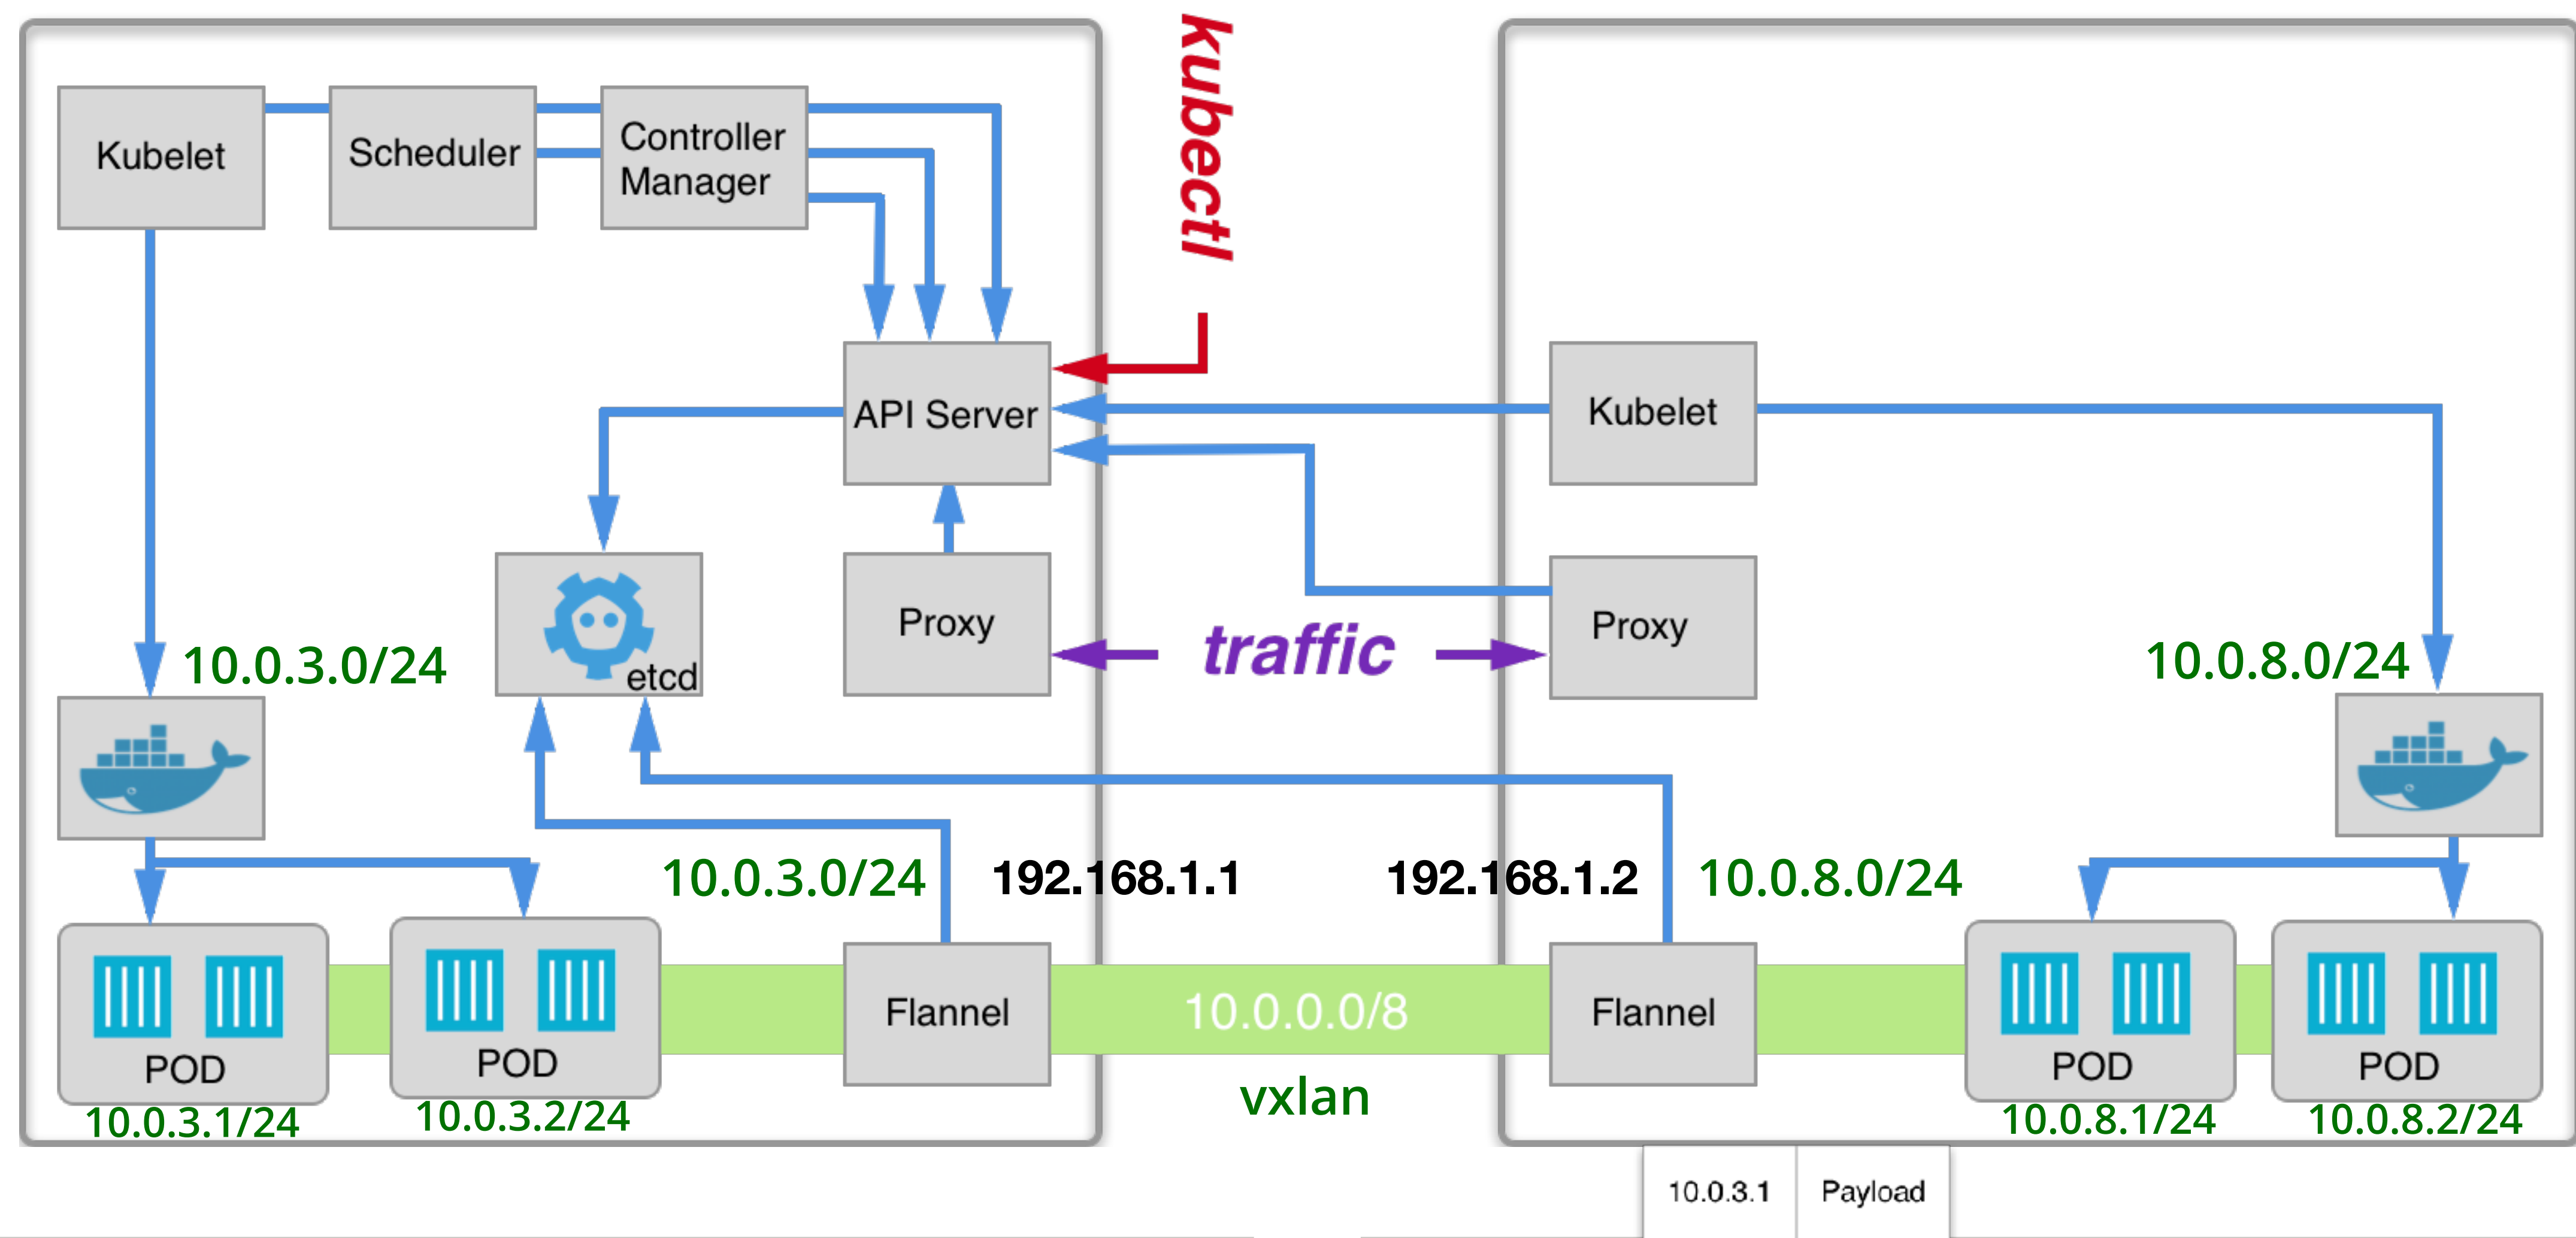


# NETWORKING



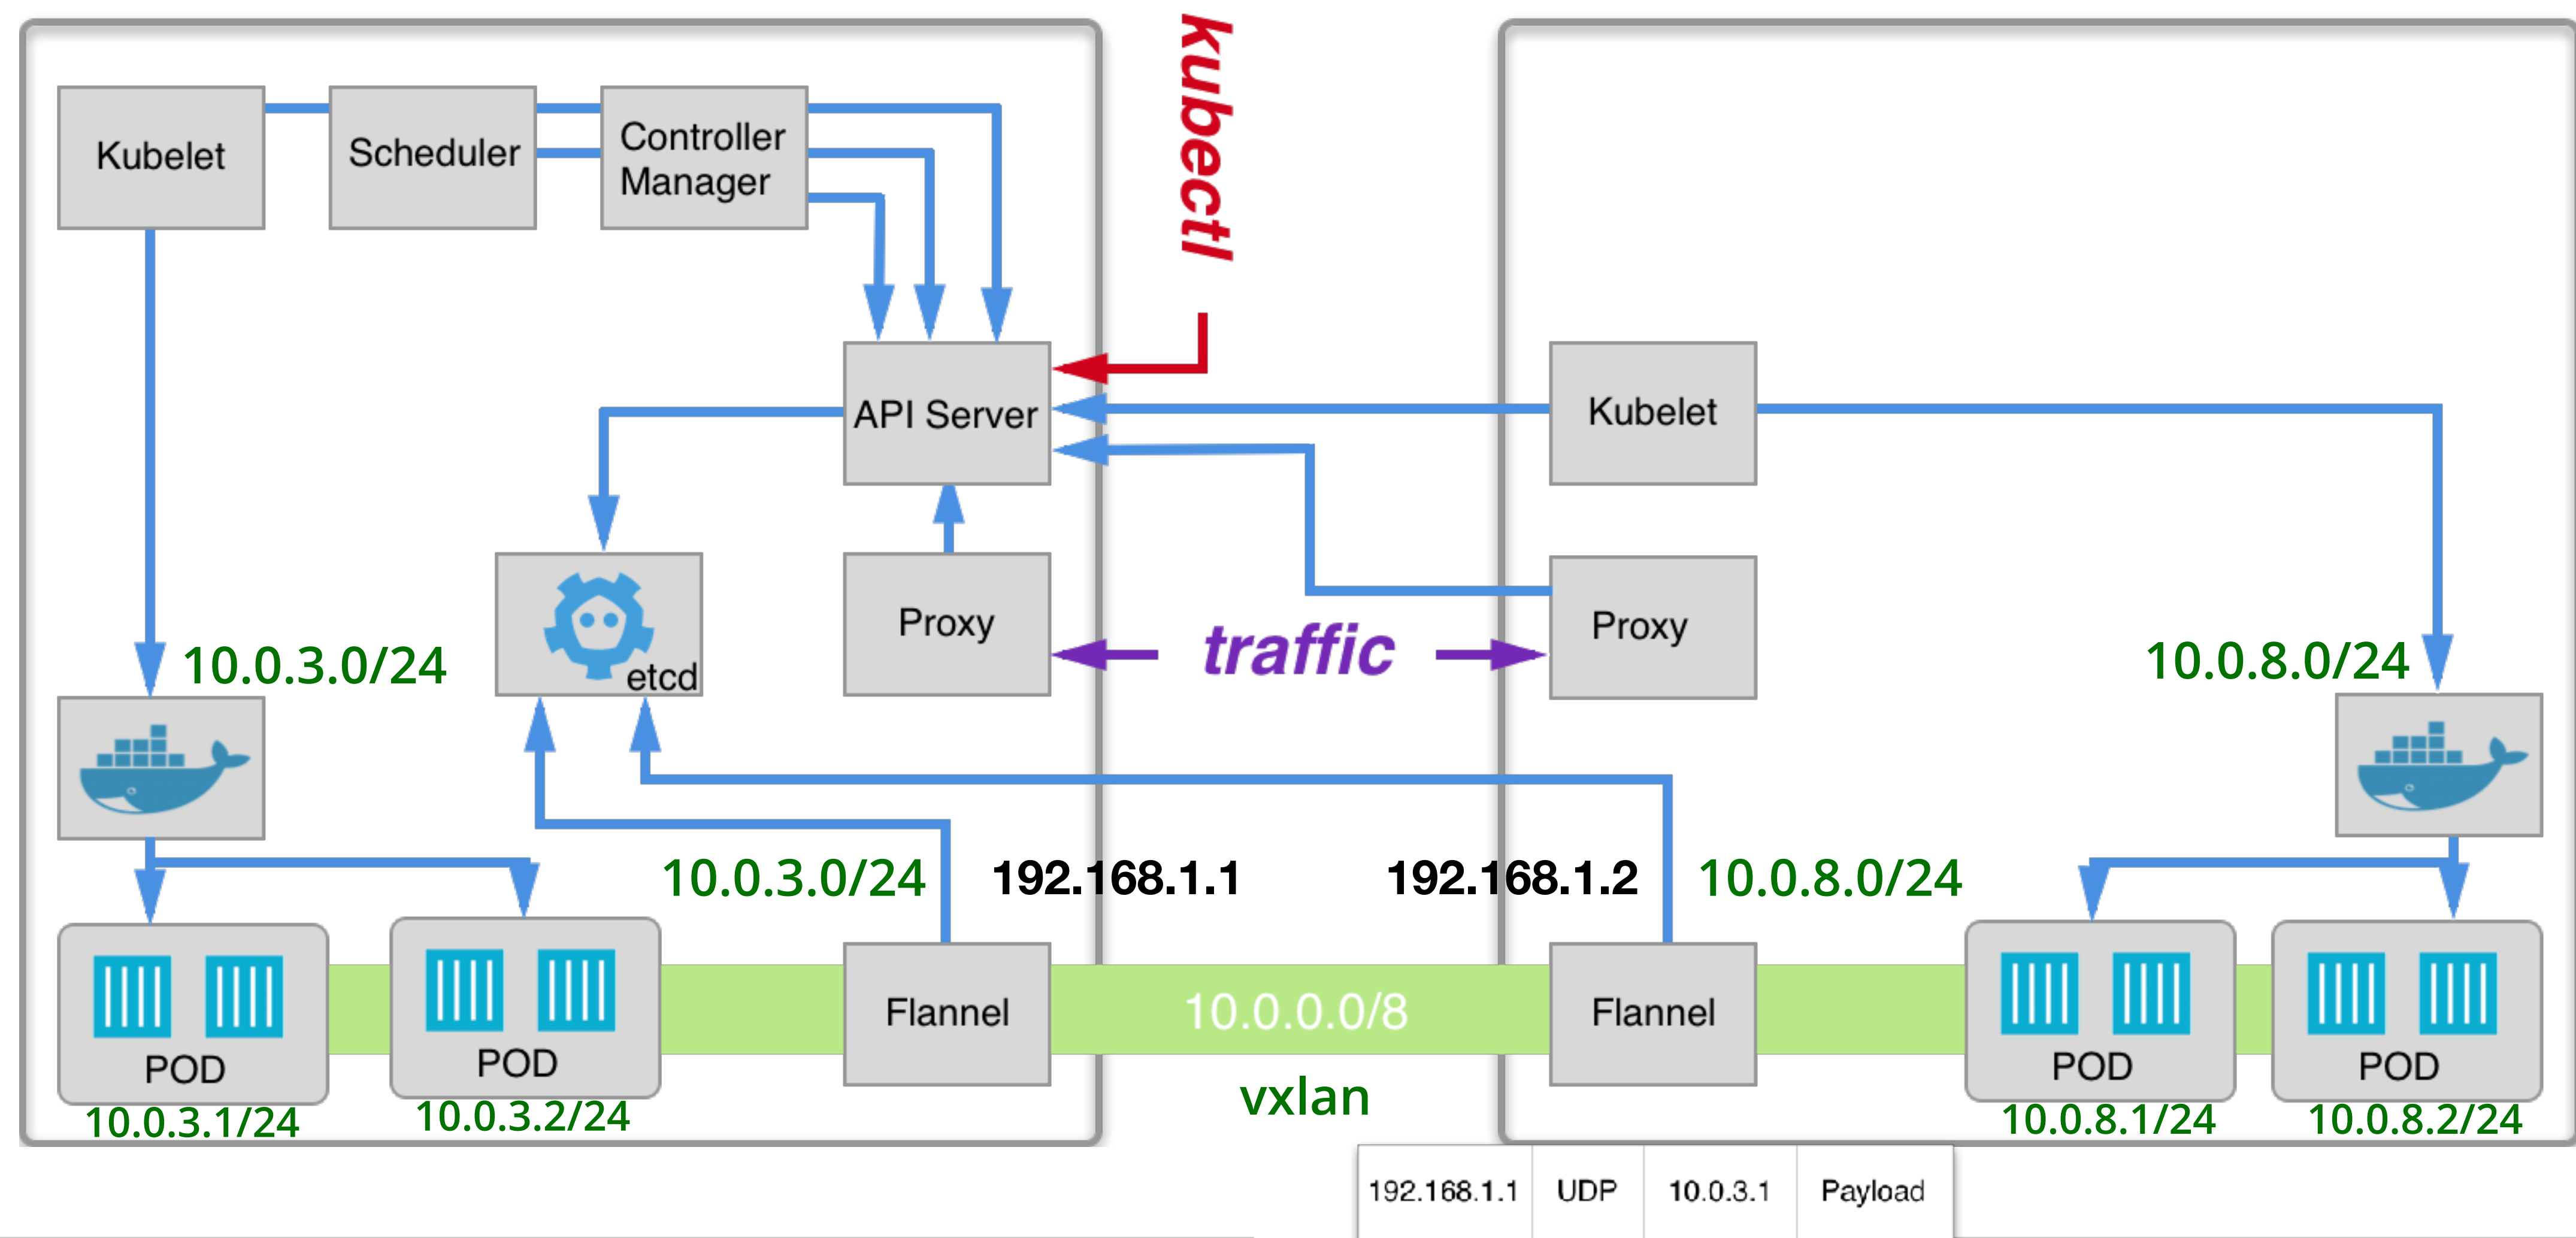


# NETWORKING

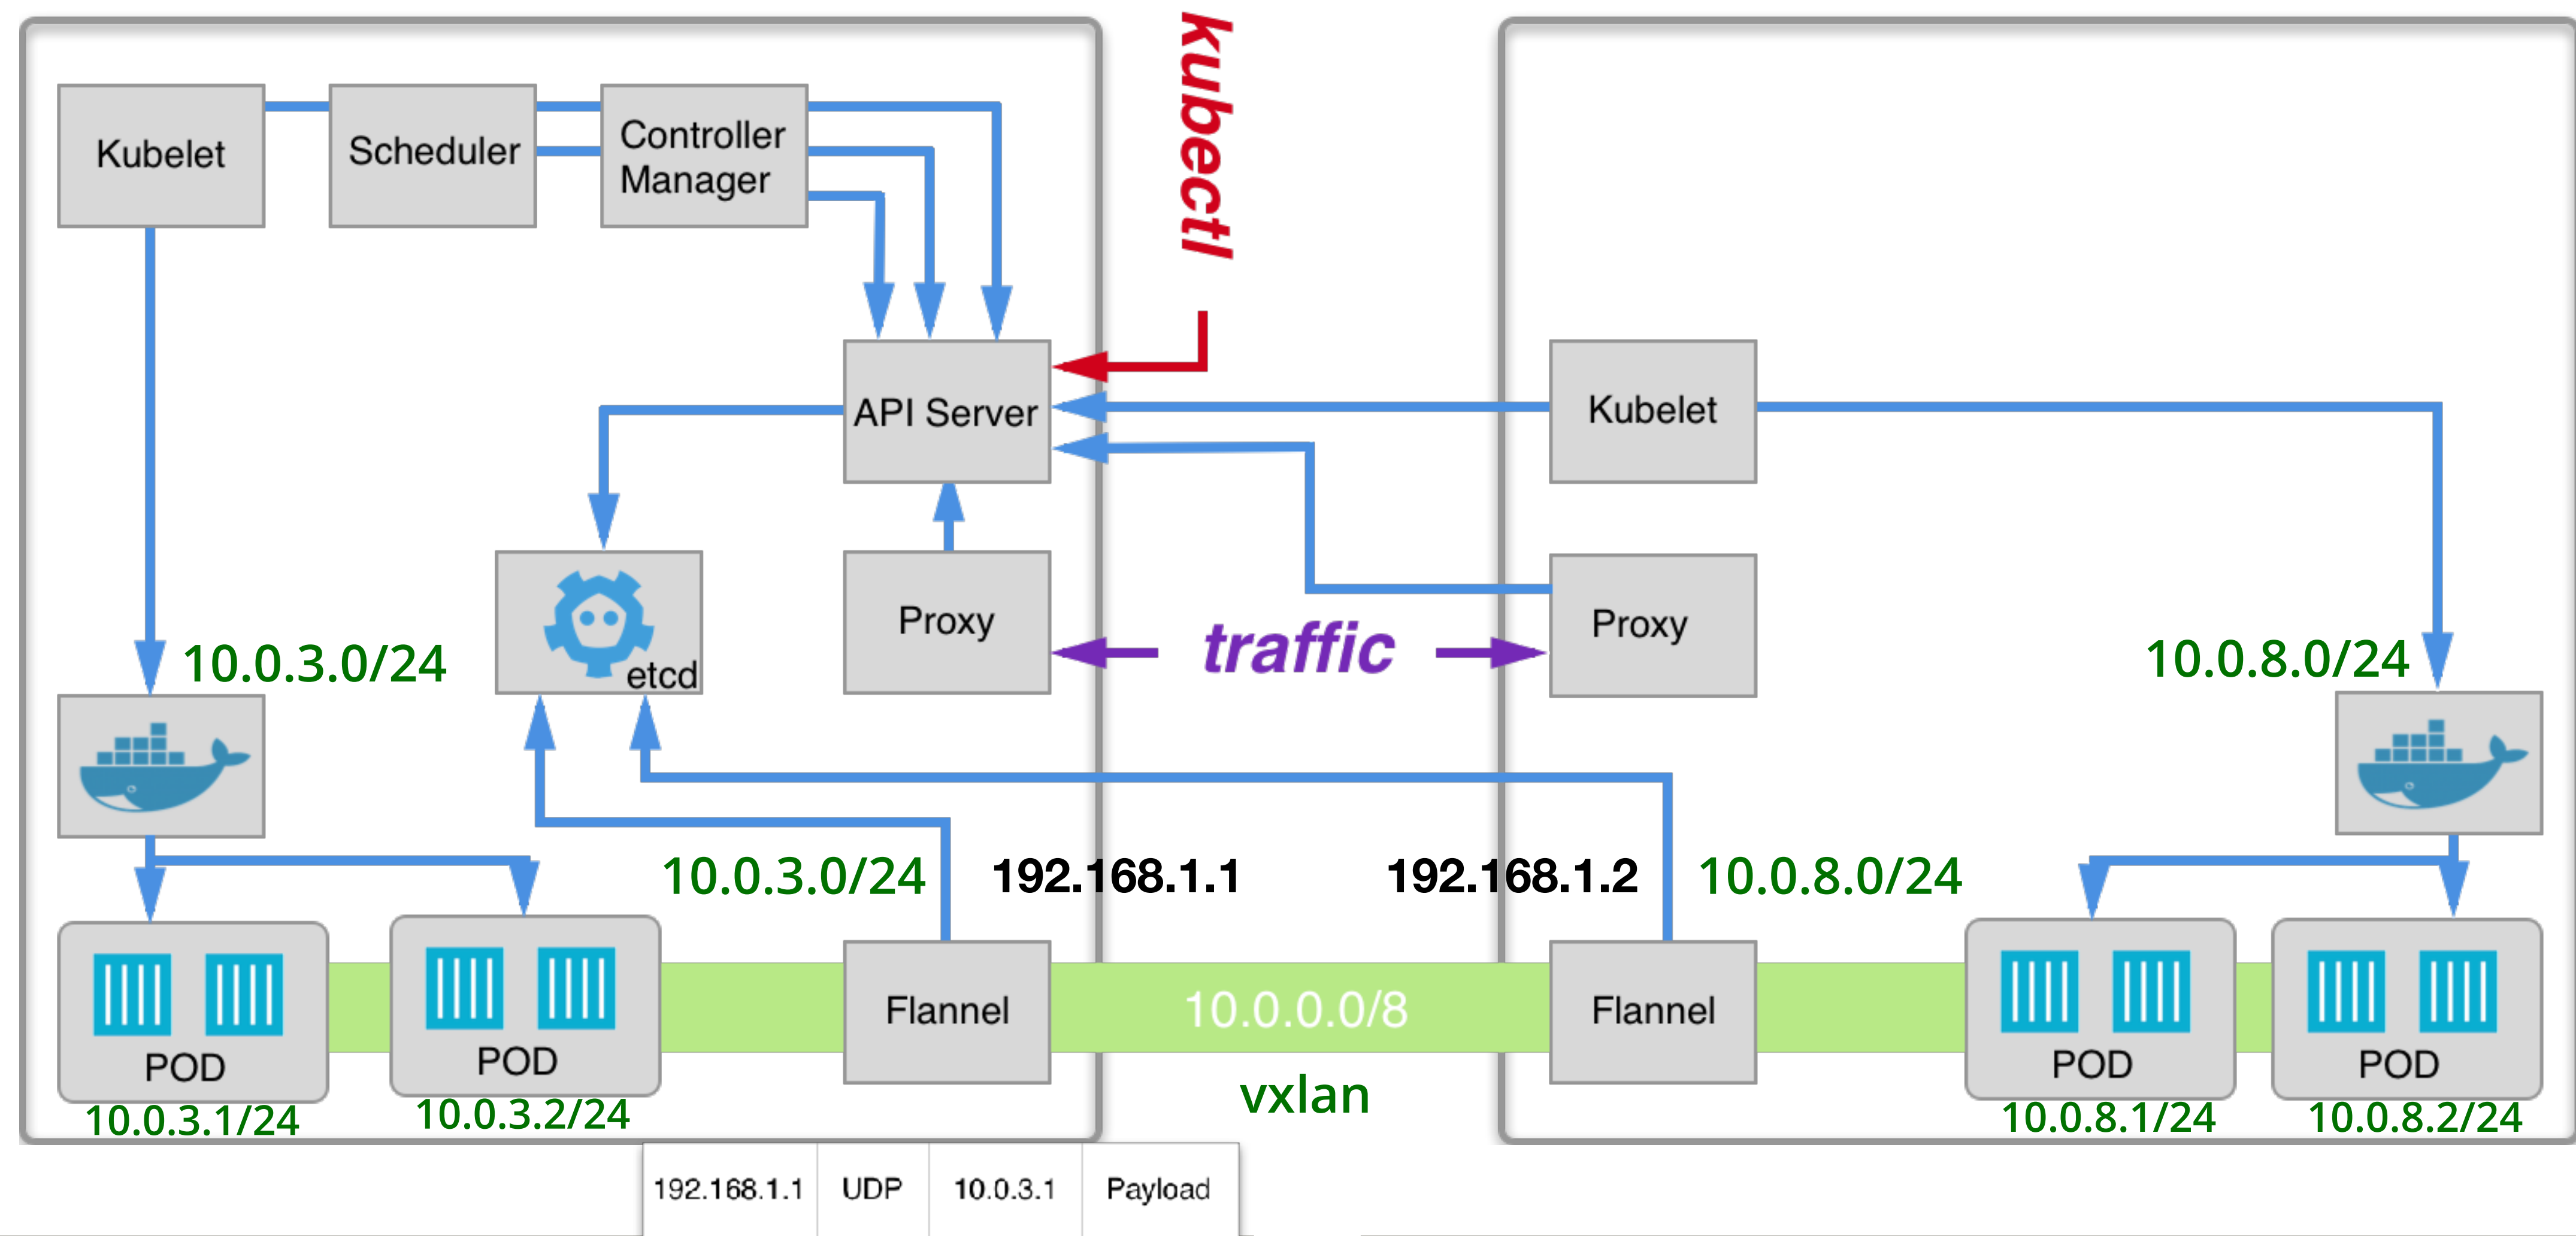




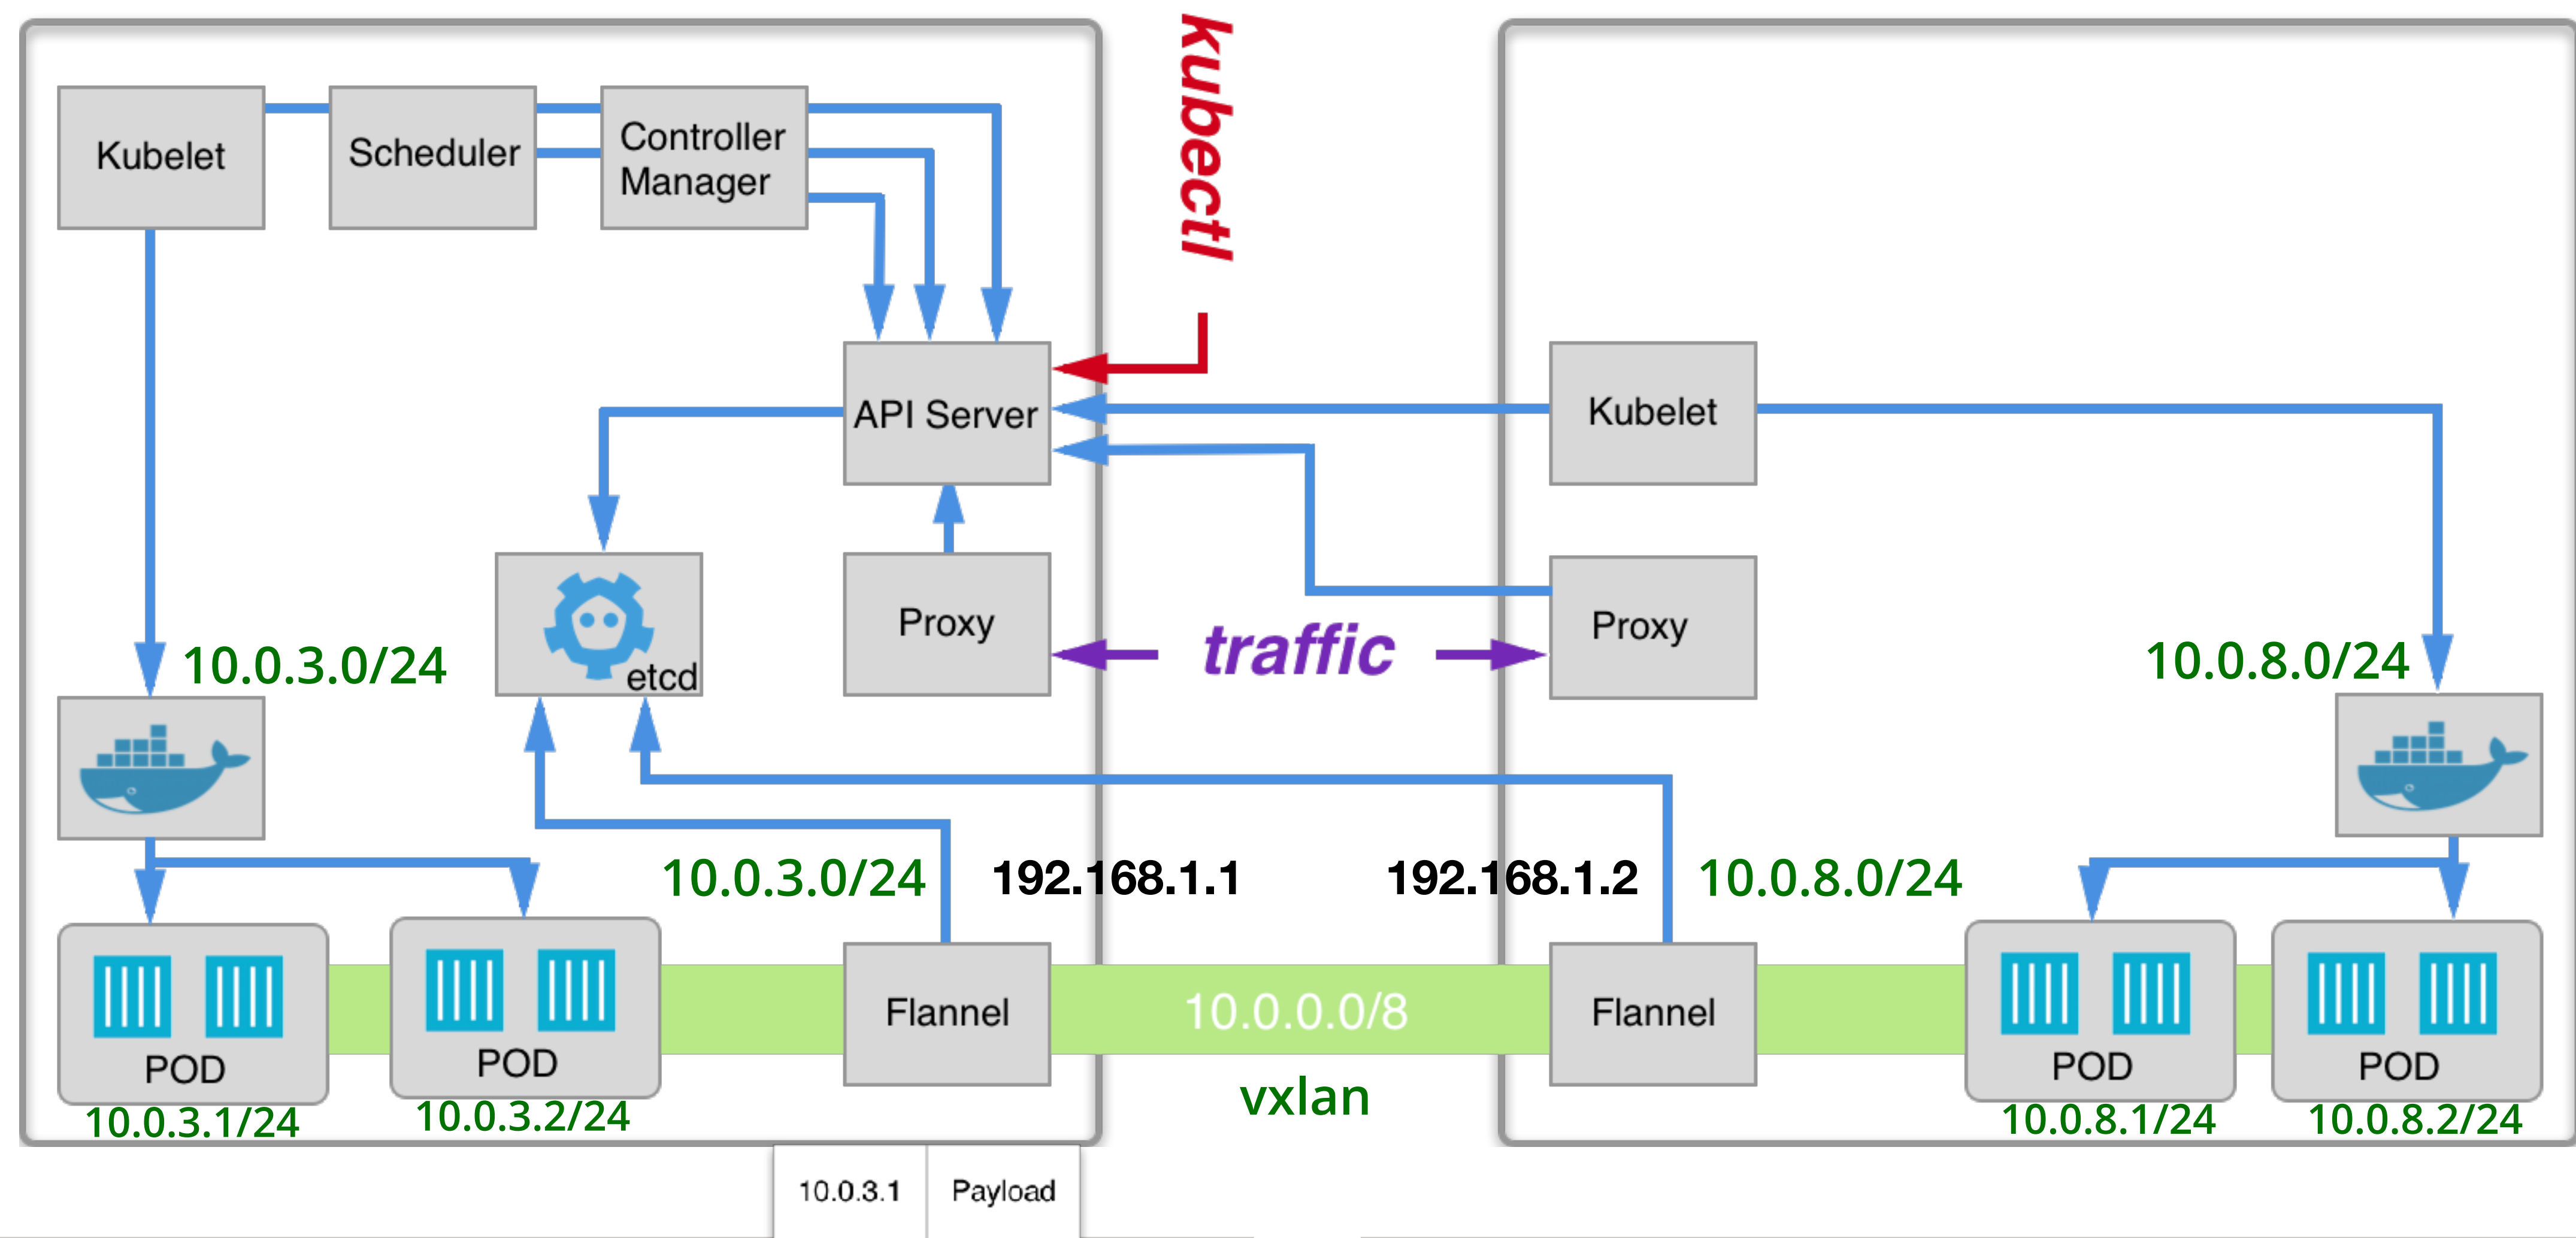
# NETWORKING



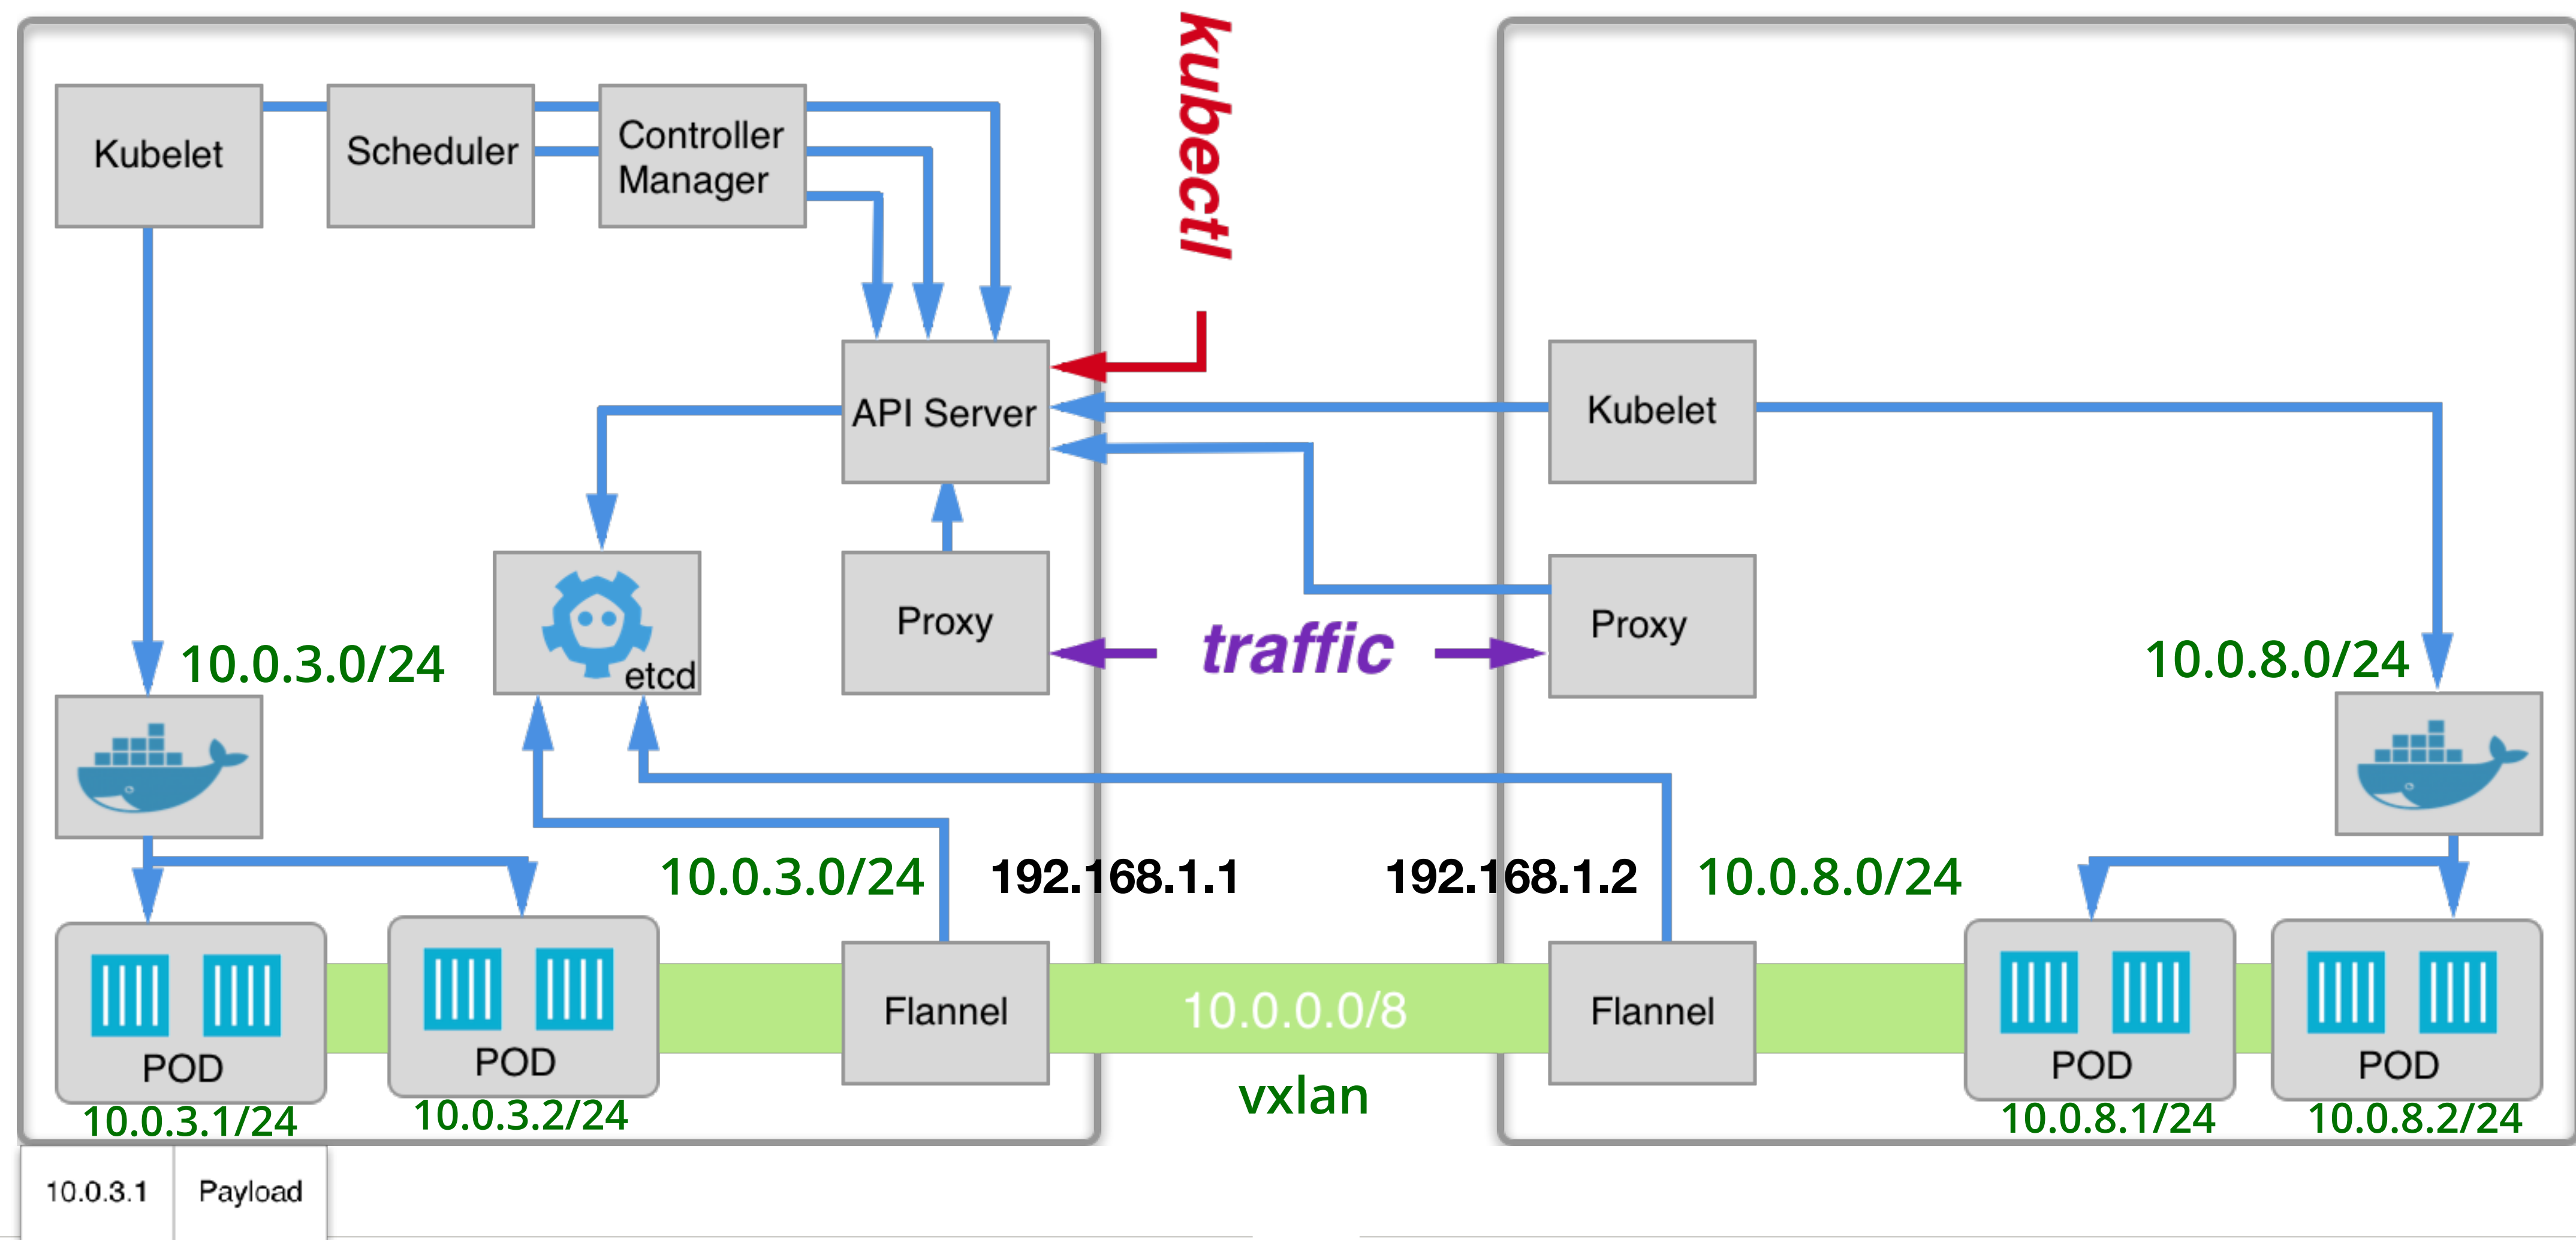
# NETWORKING



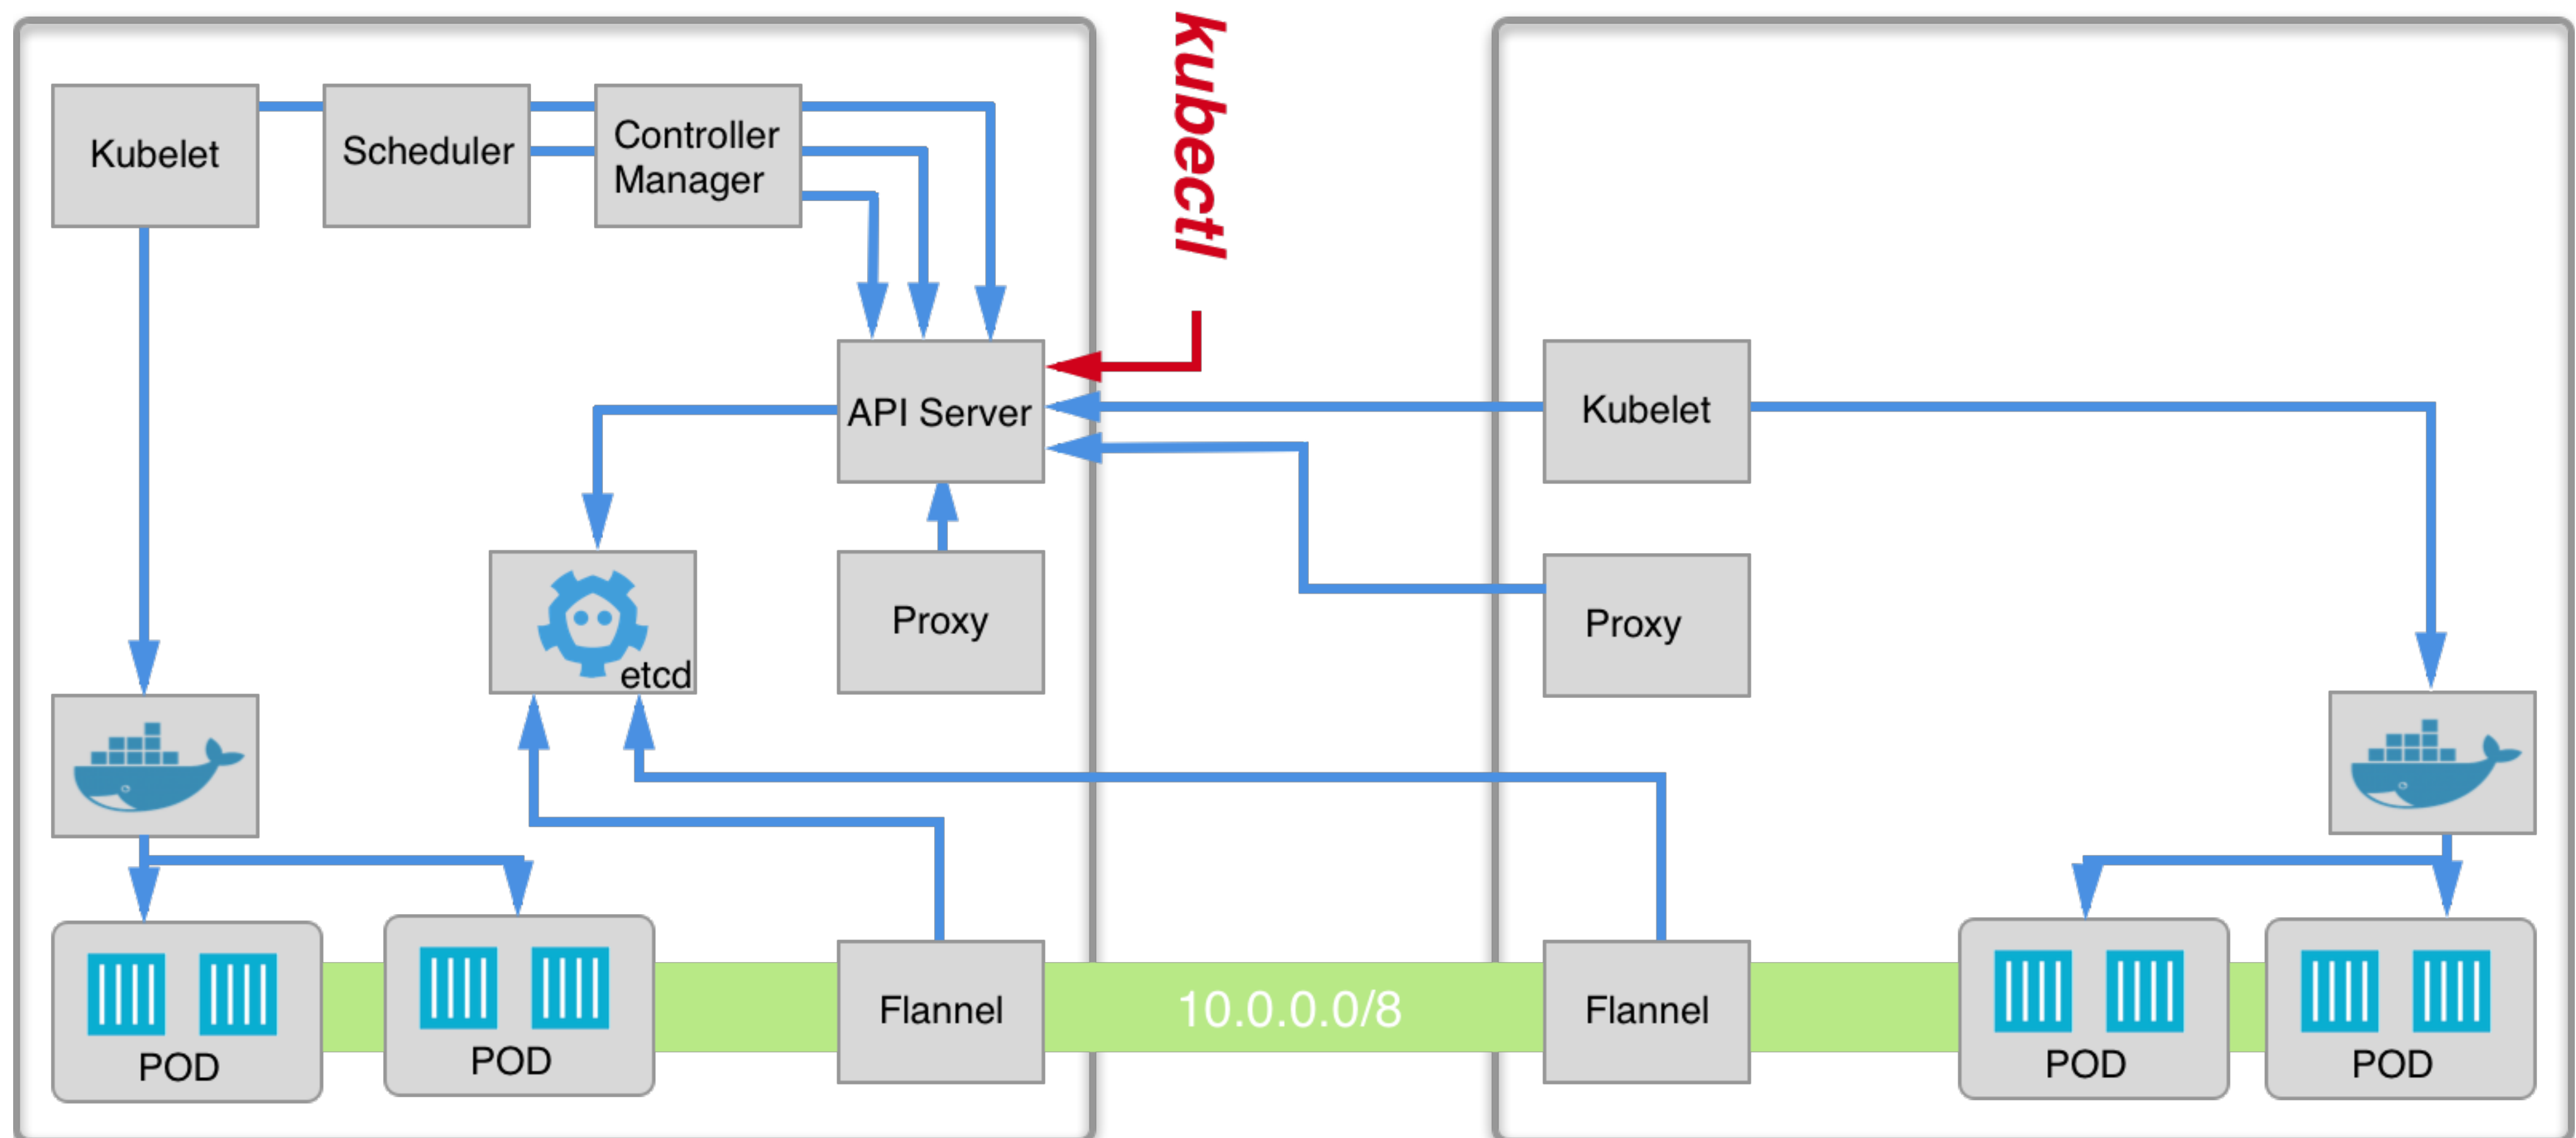
# NETWORKING



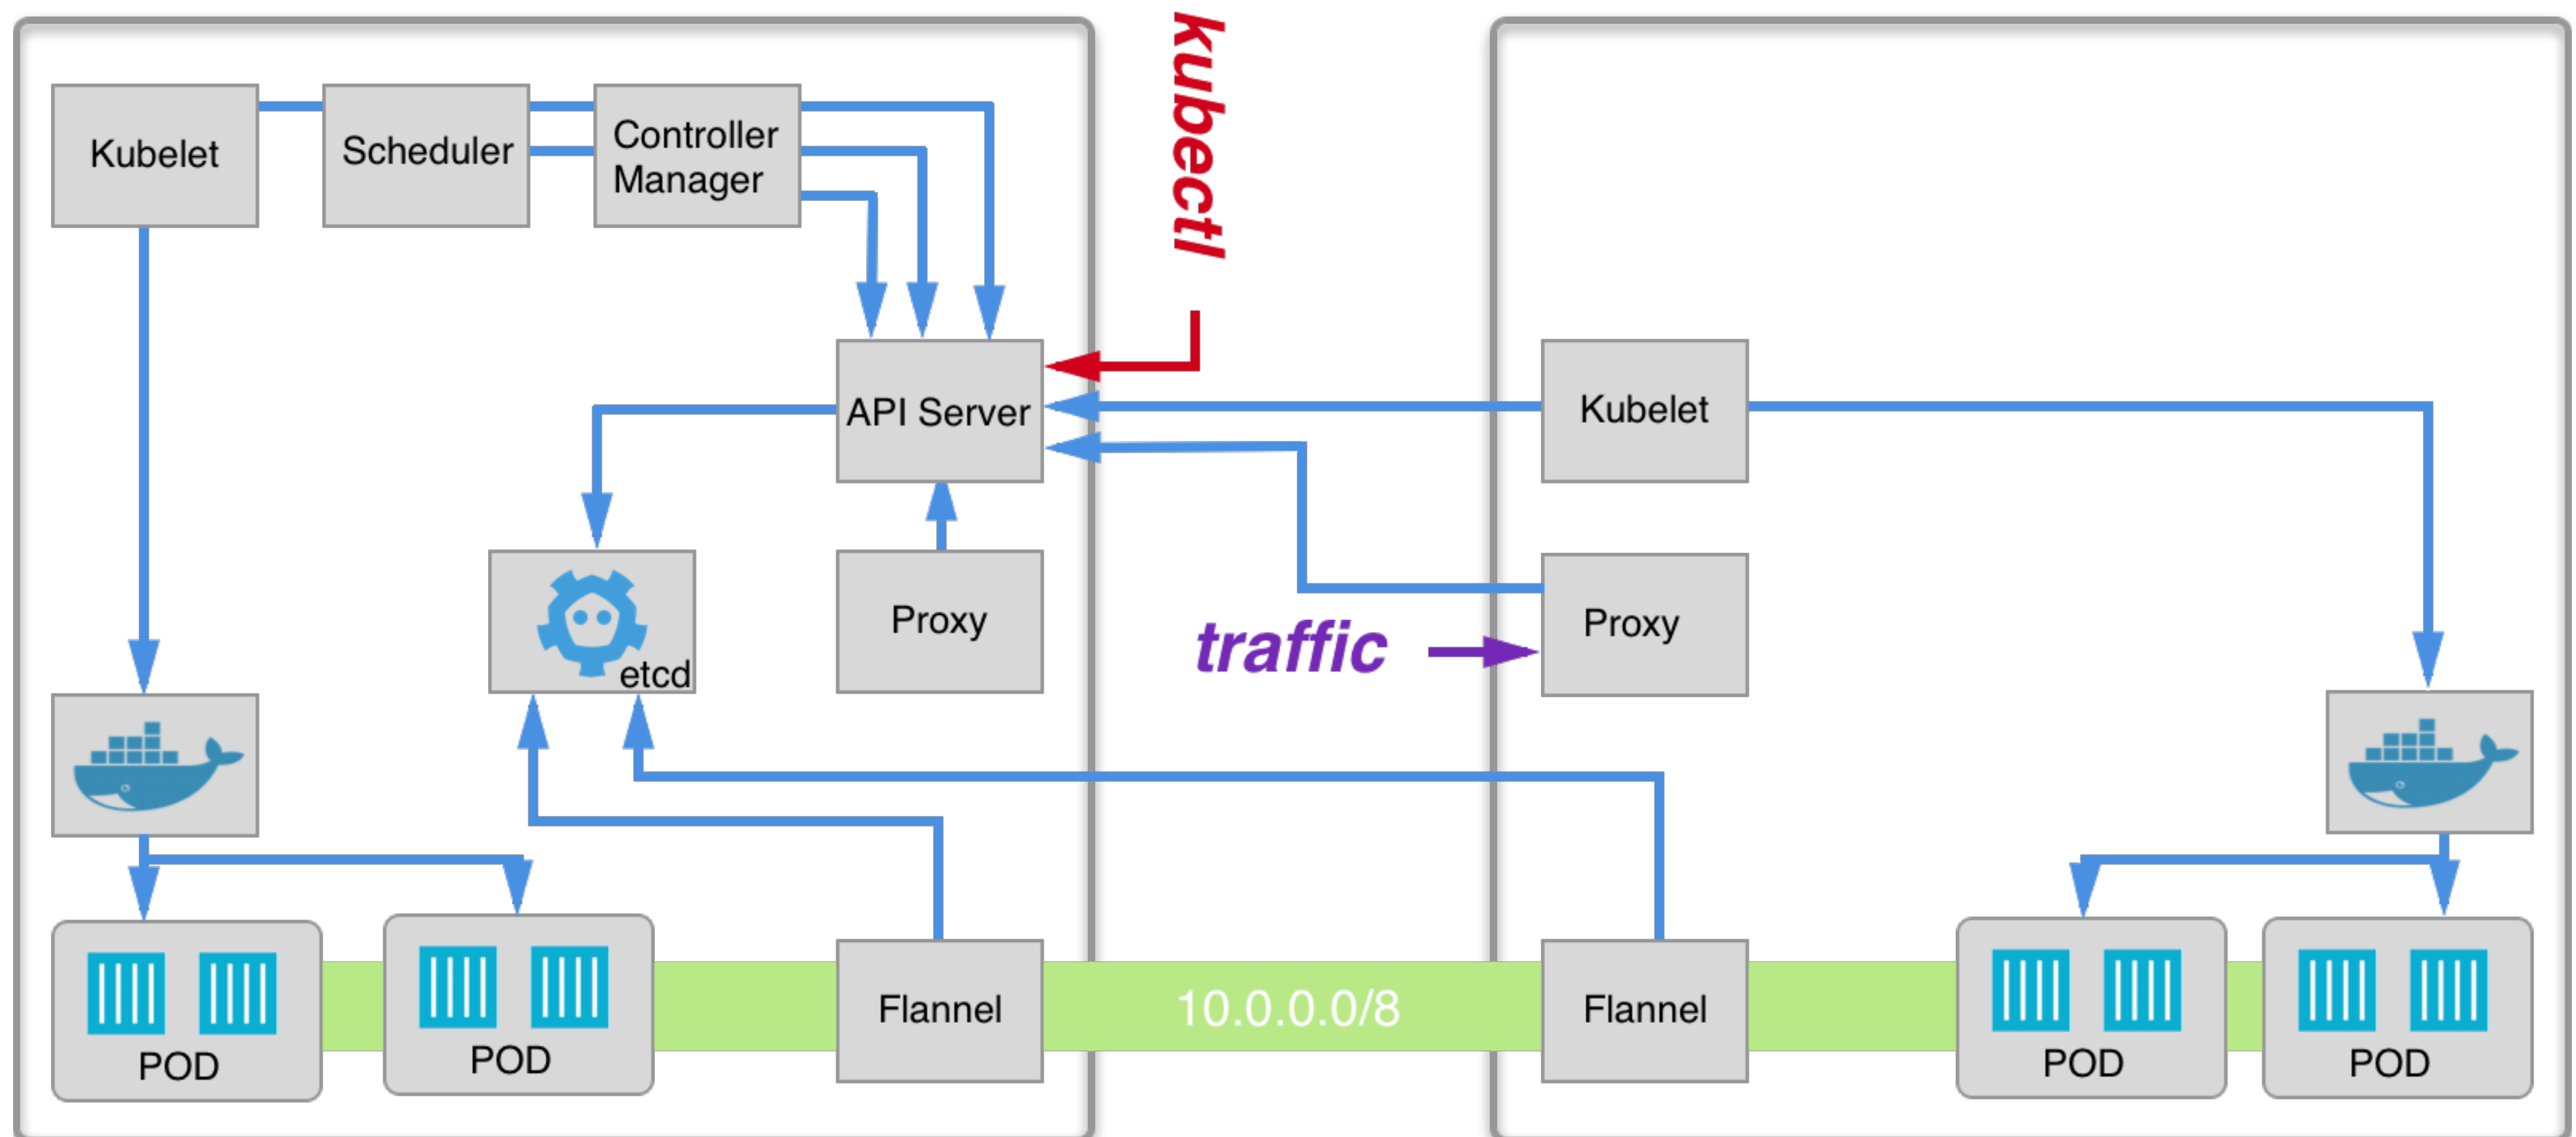
# NETWORKING



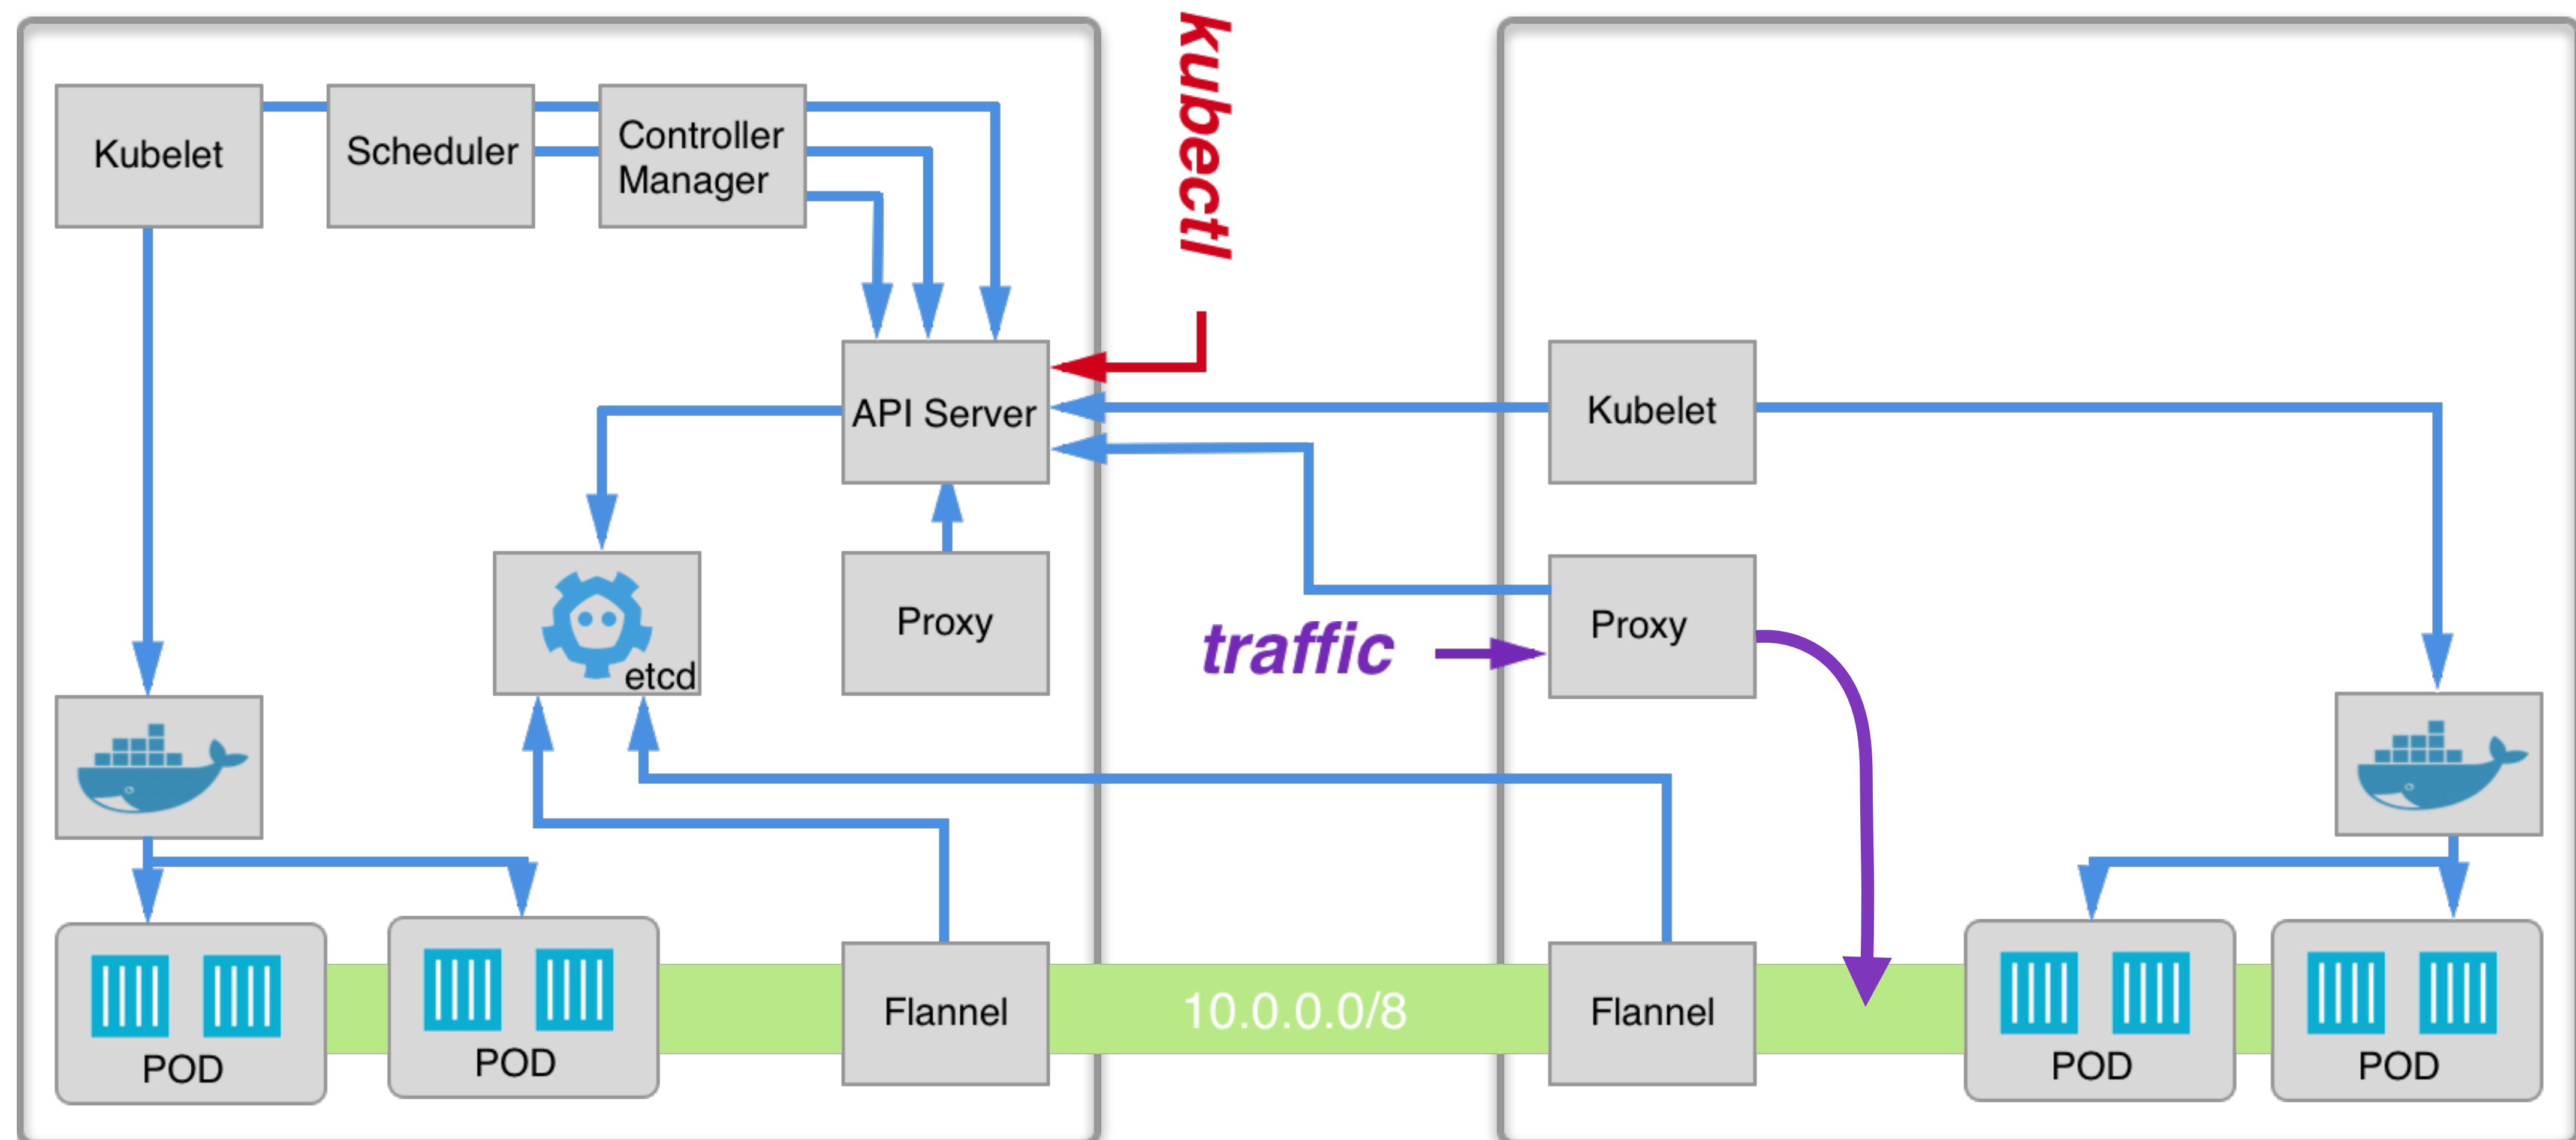
# NETWORKING



# NETWORKING

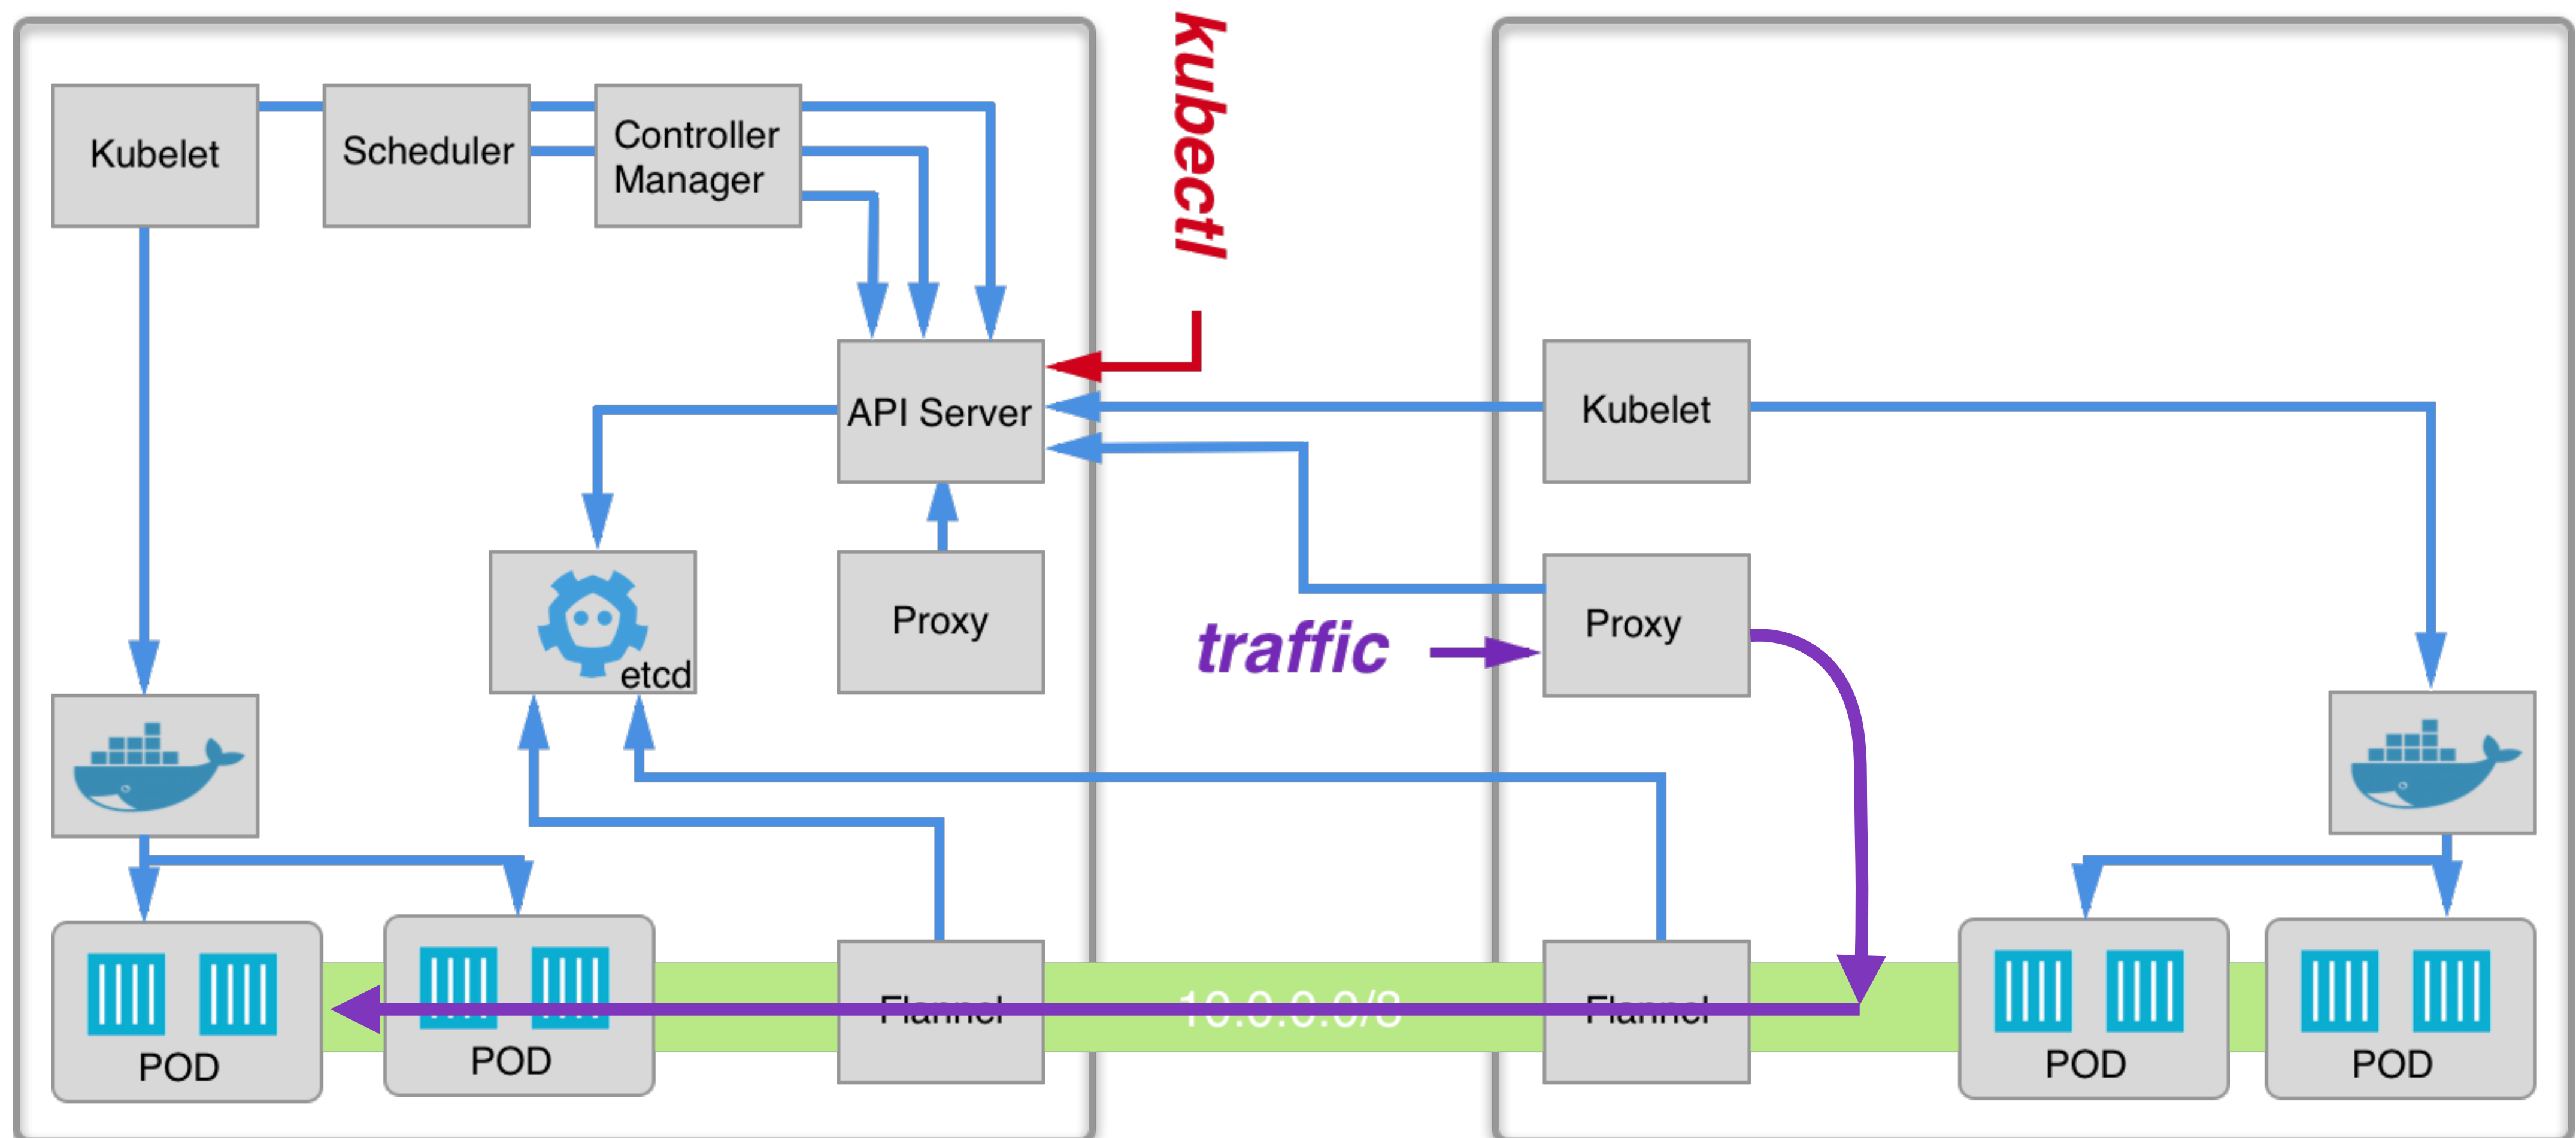


# NETWORKING



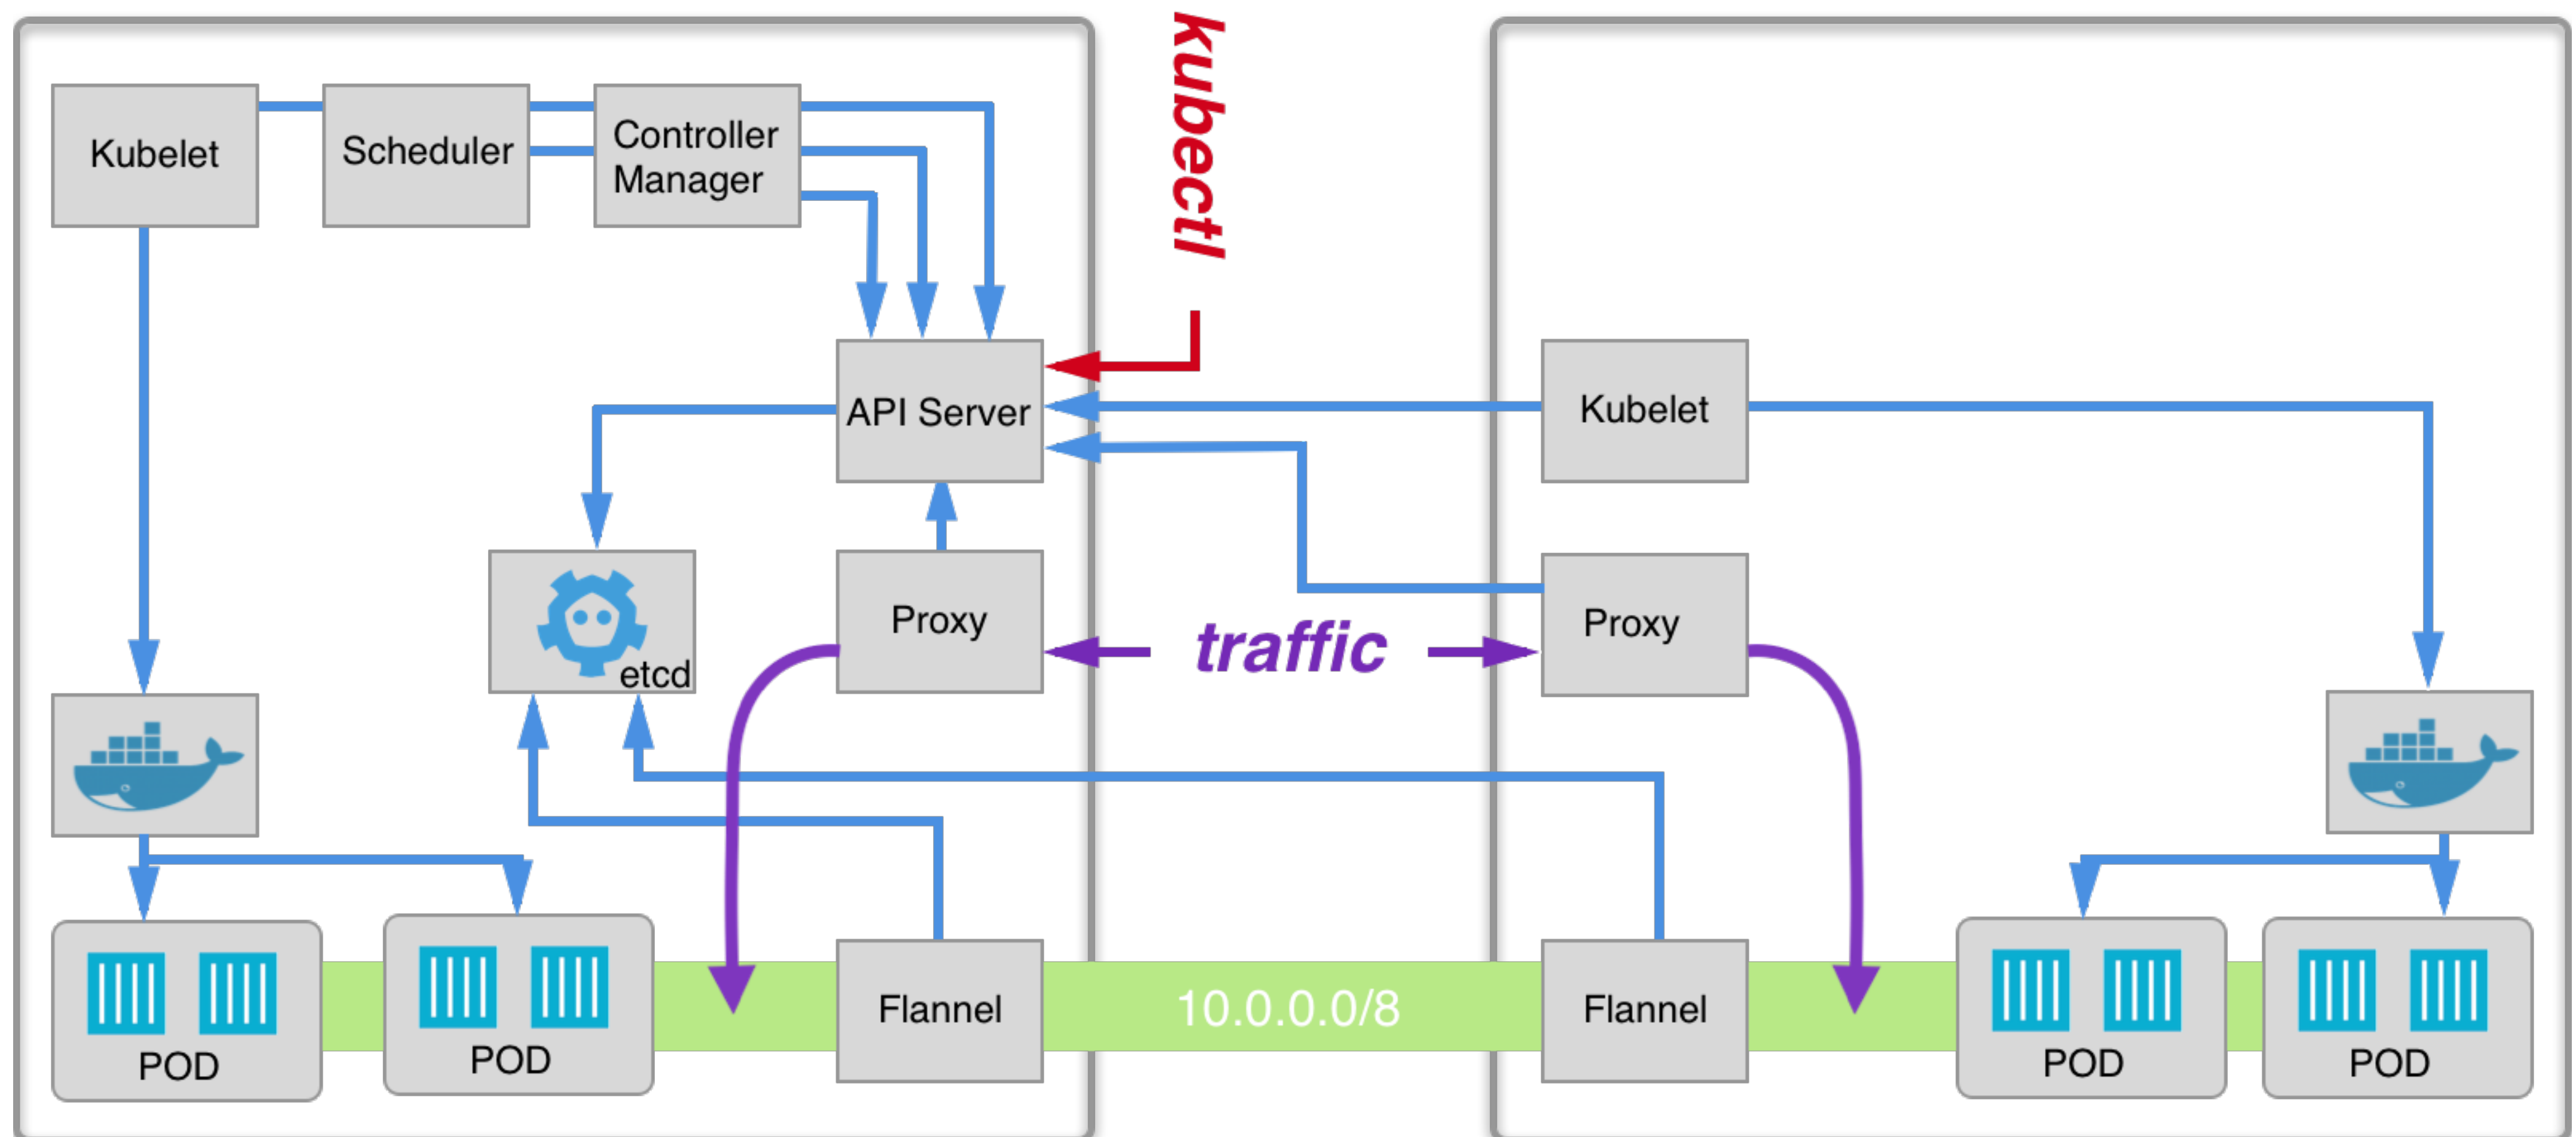


# NETWORKING

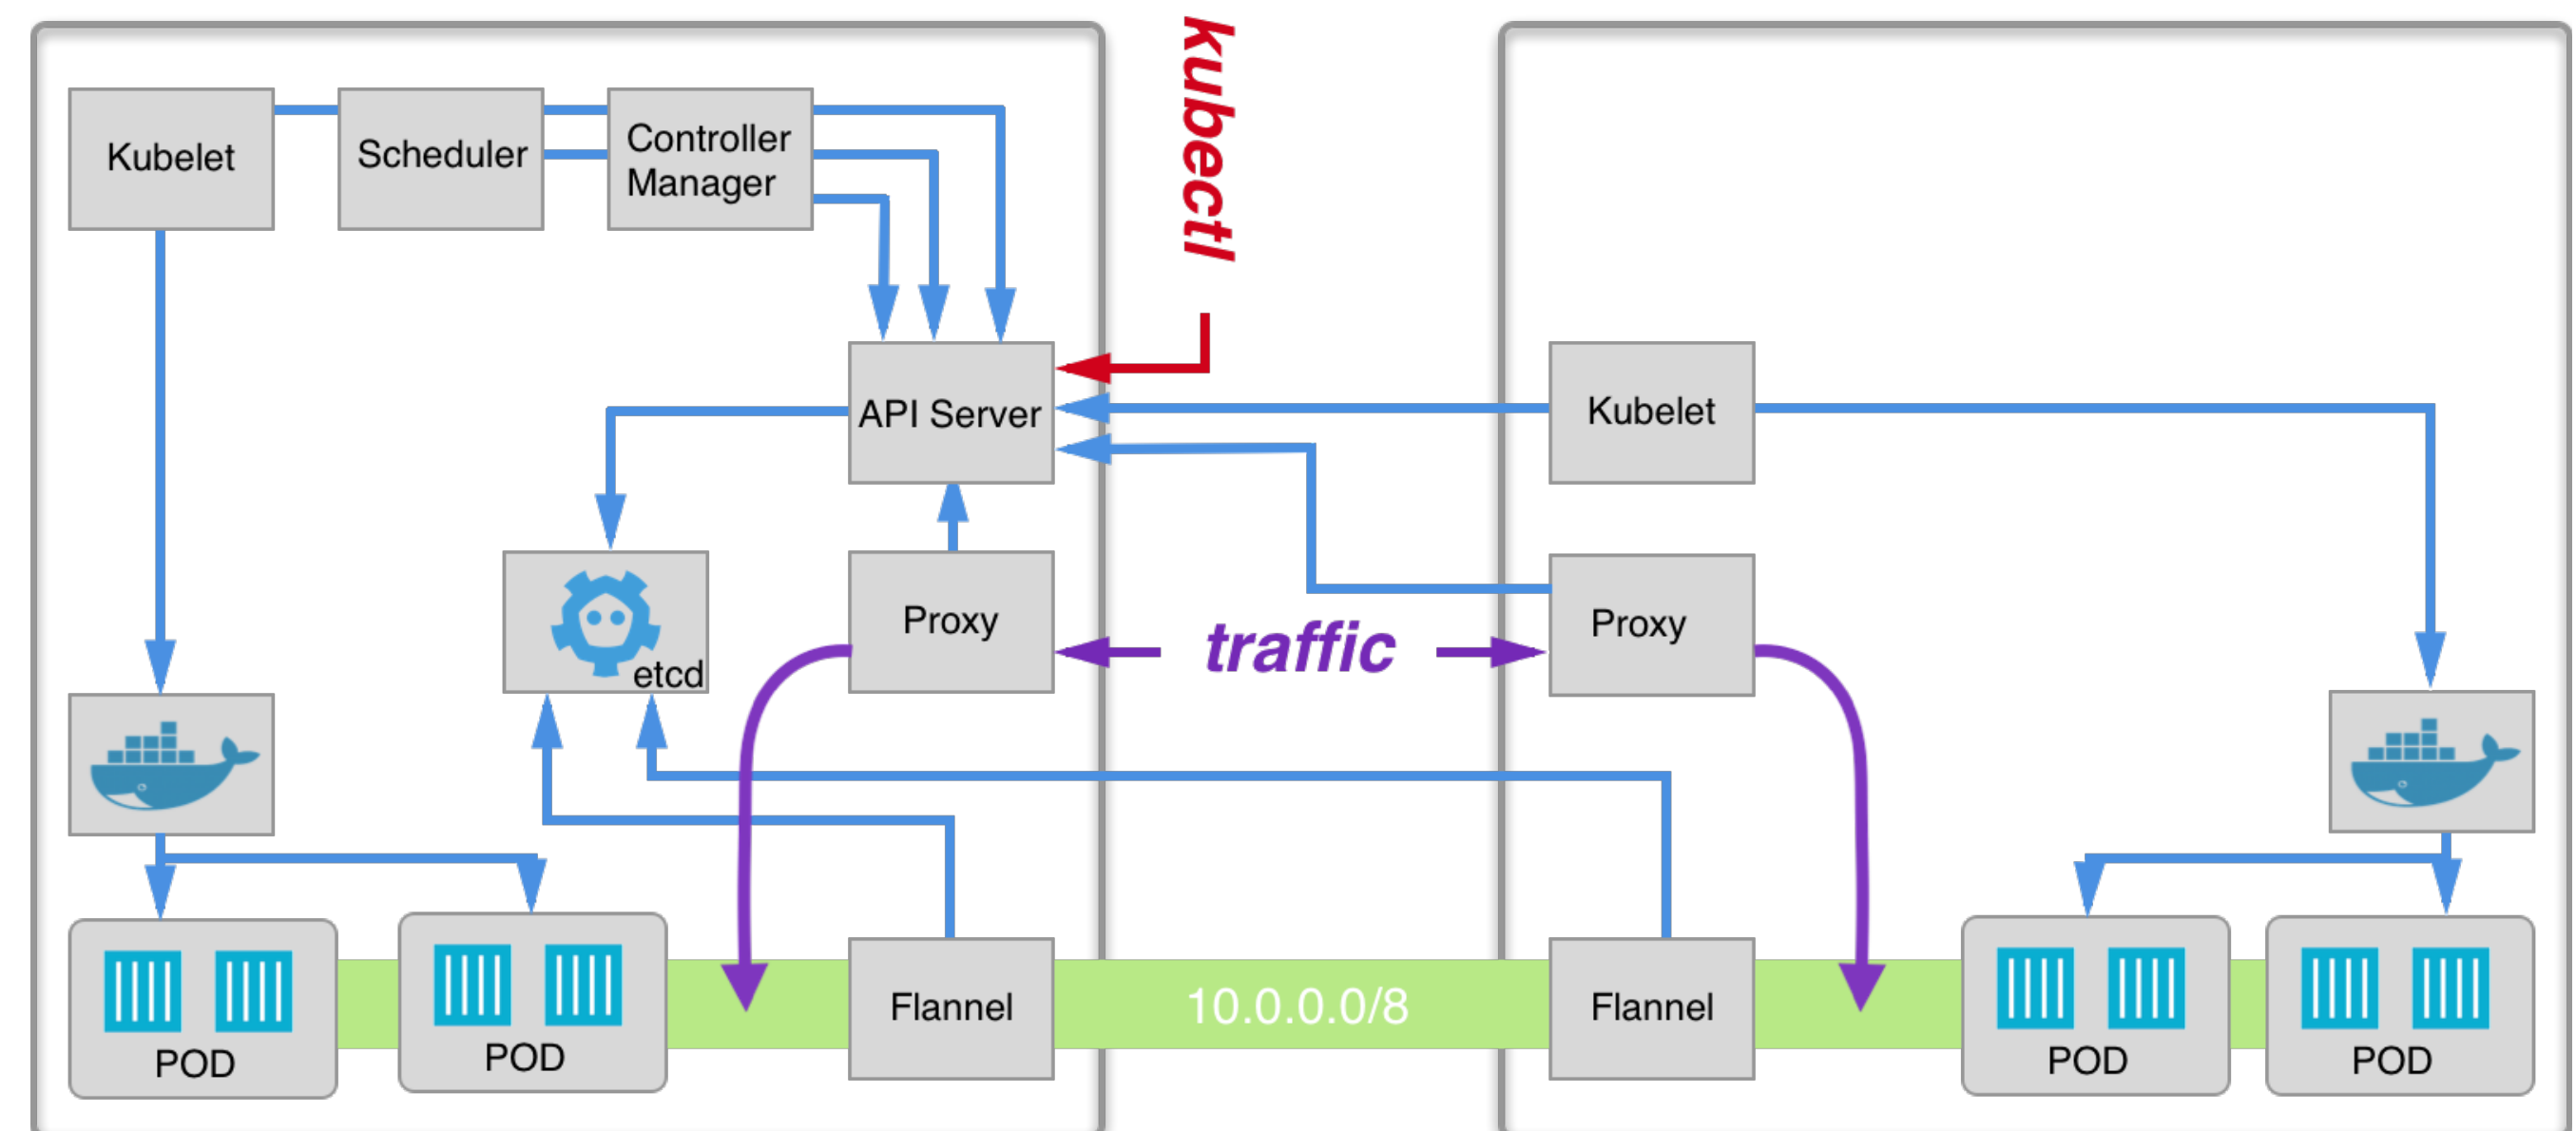




# NETWORKING

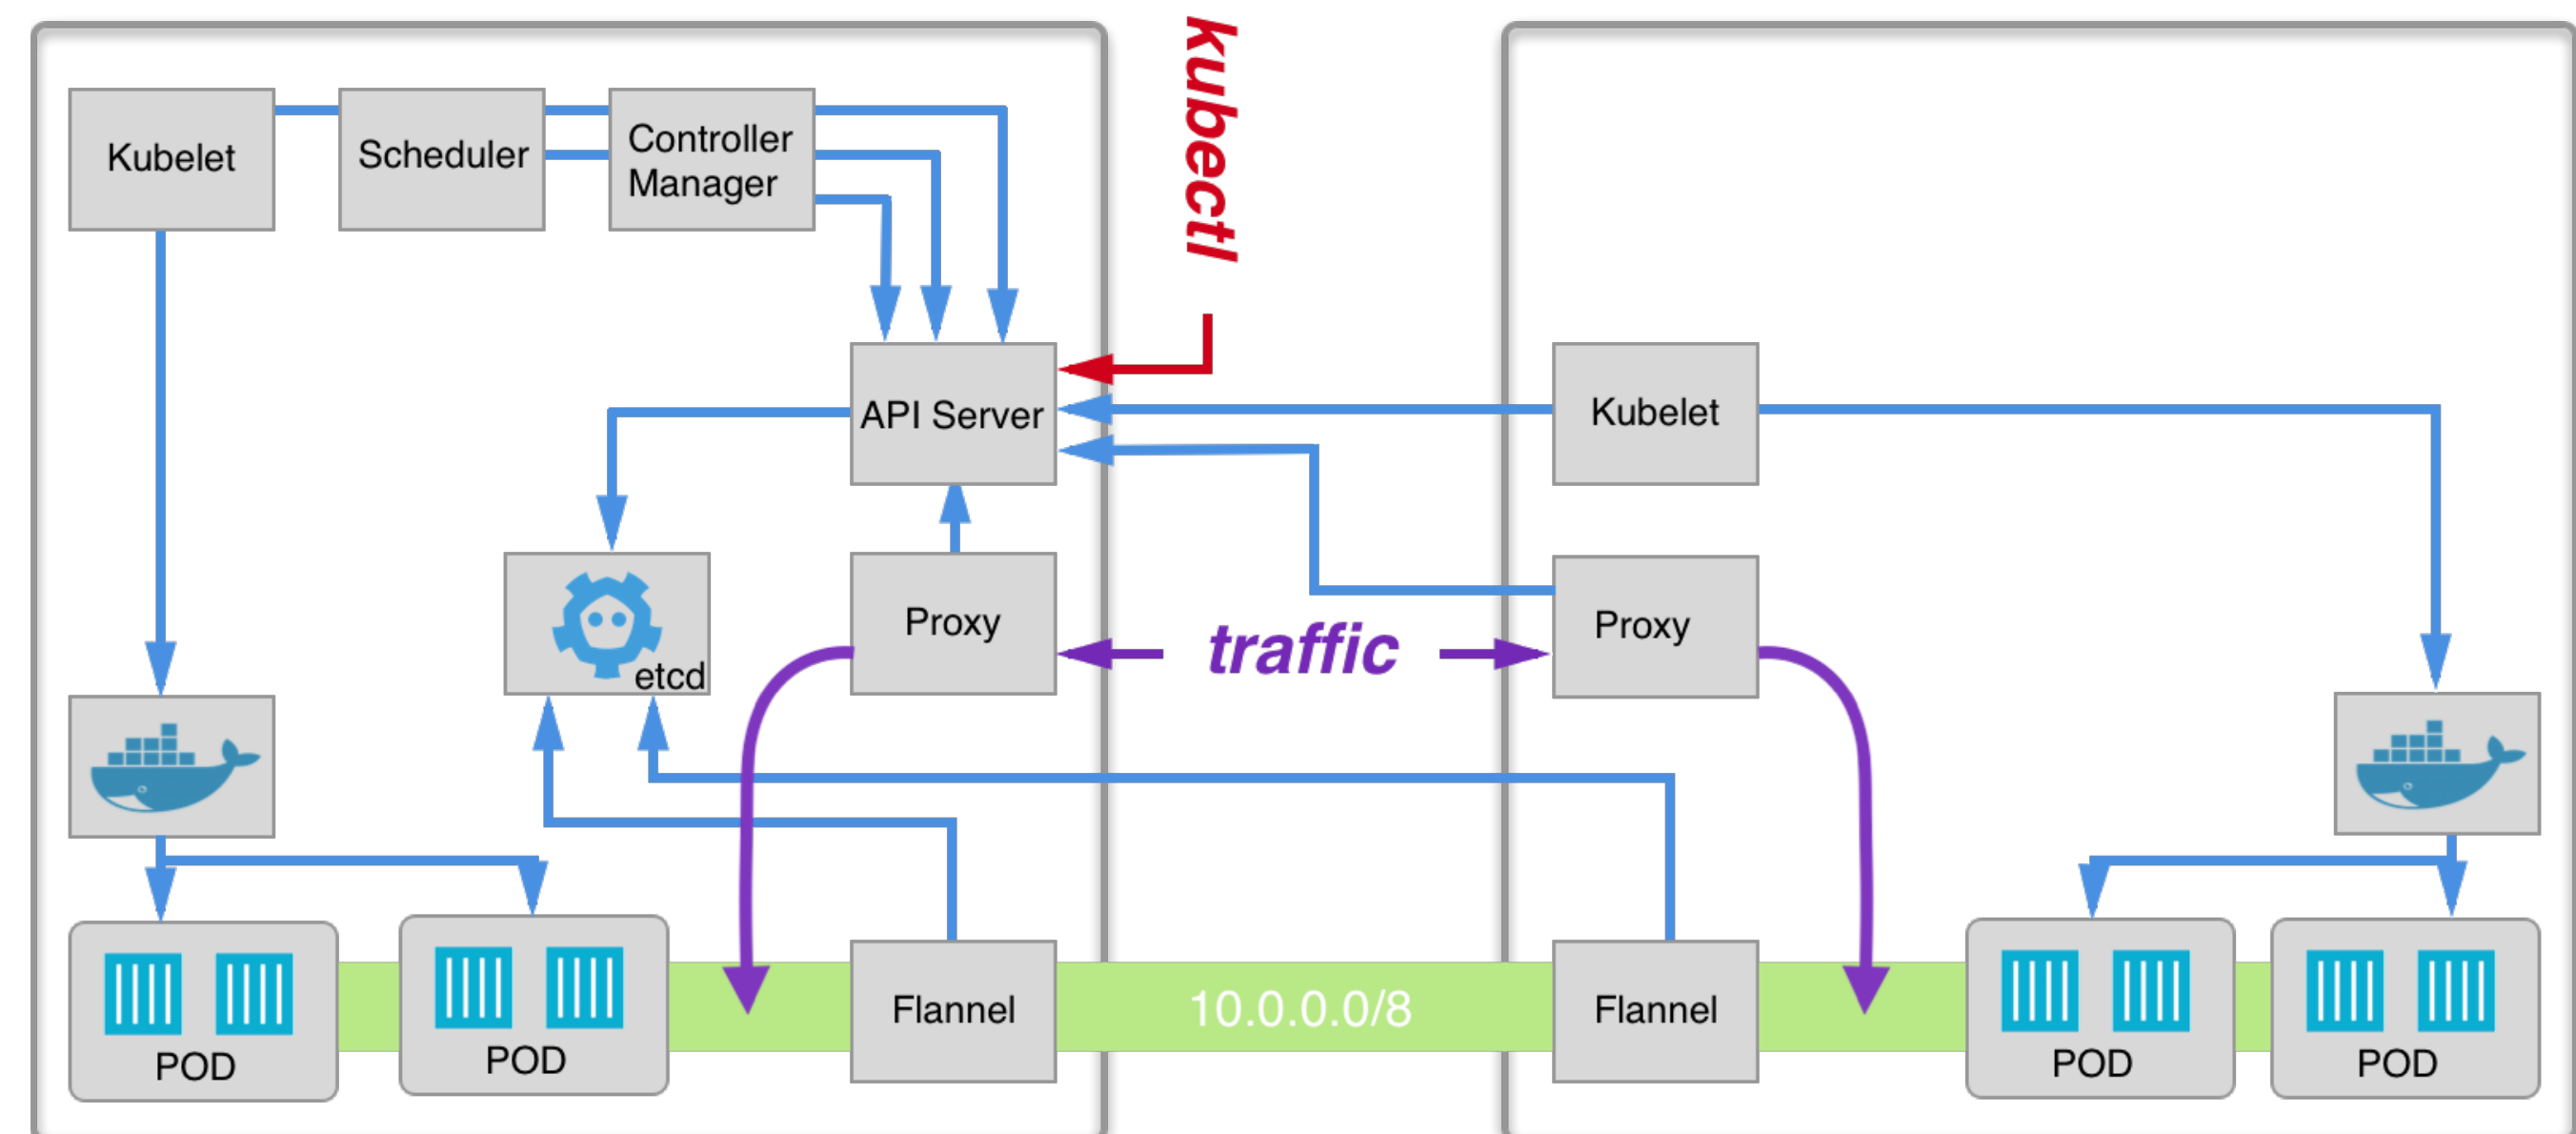


# NETWORKING



# NETWORKING

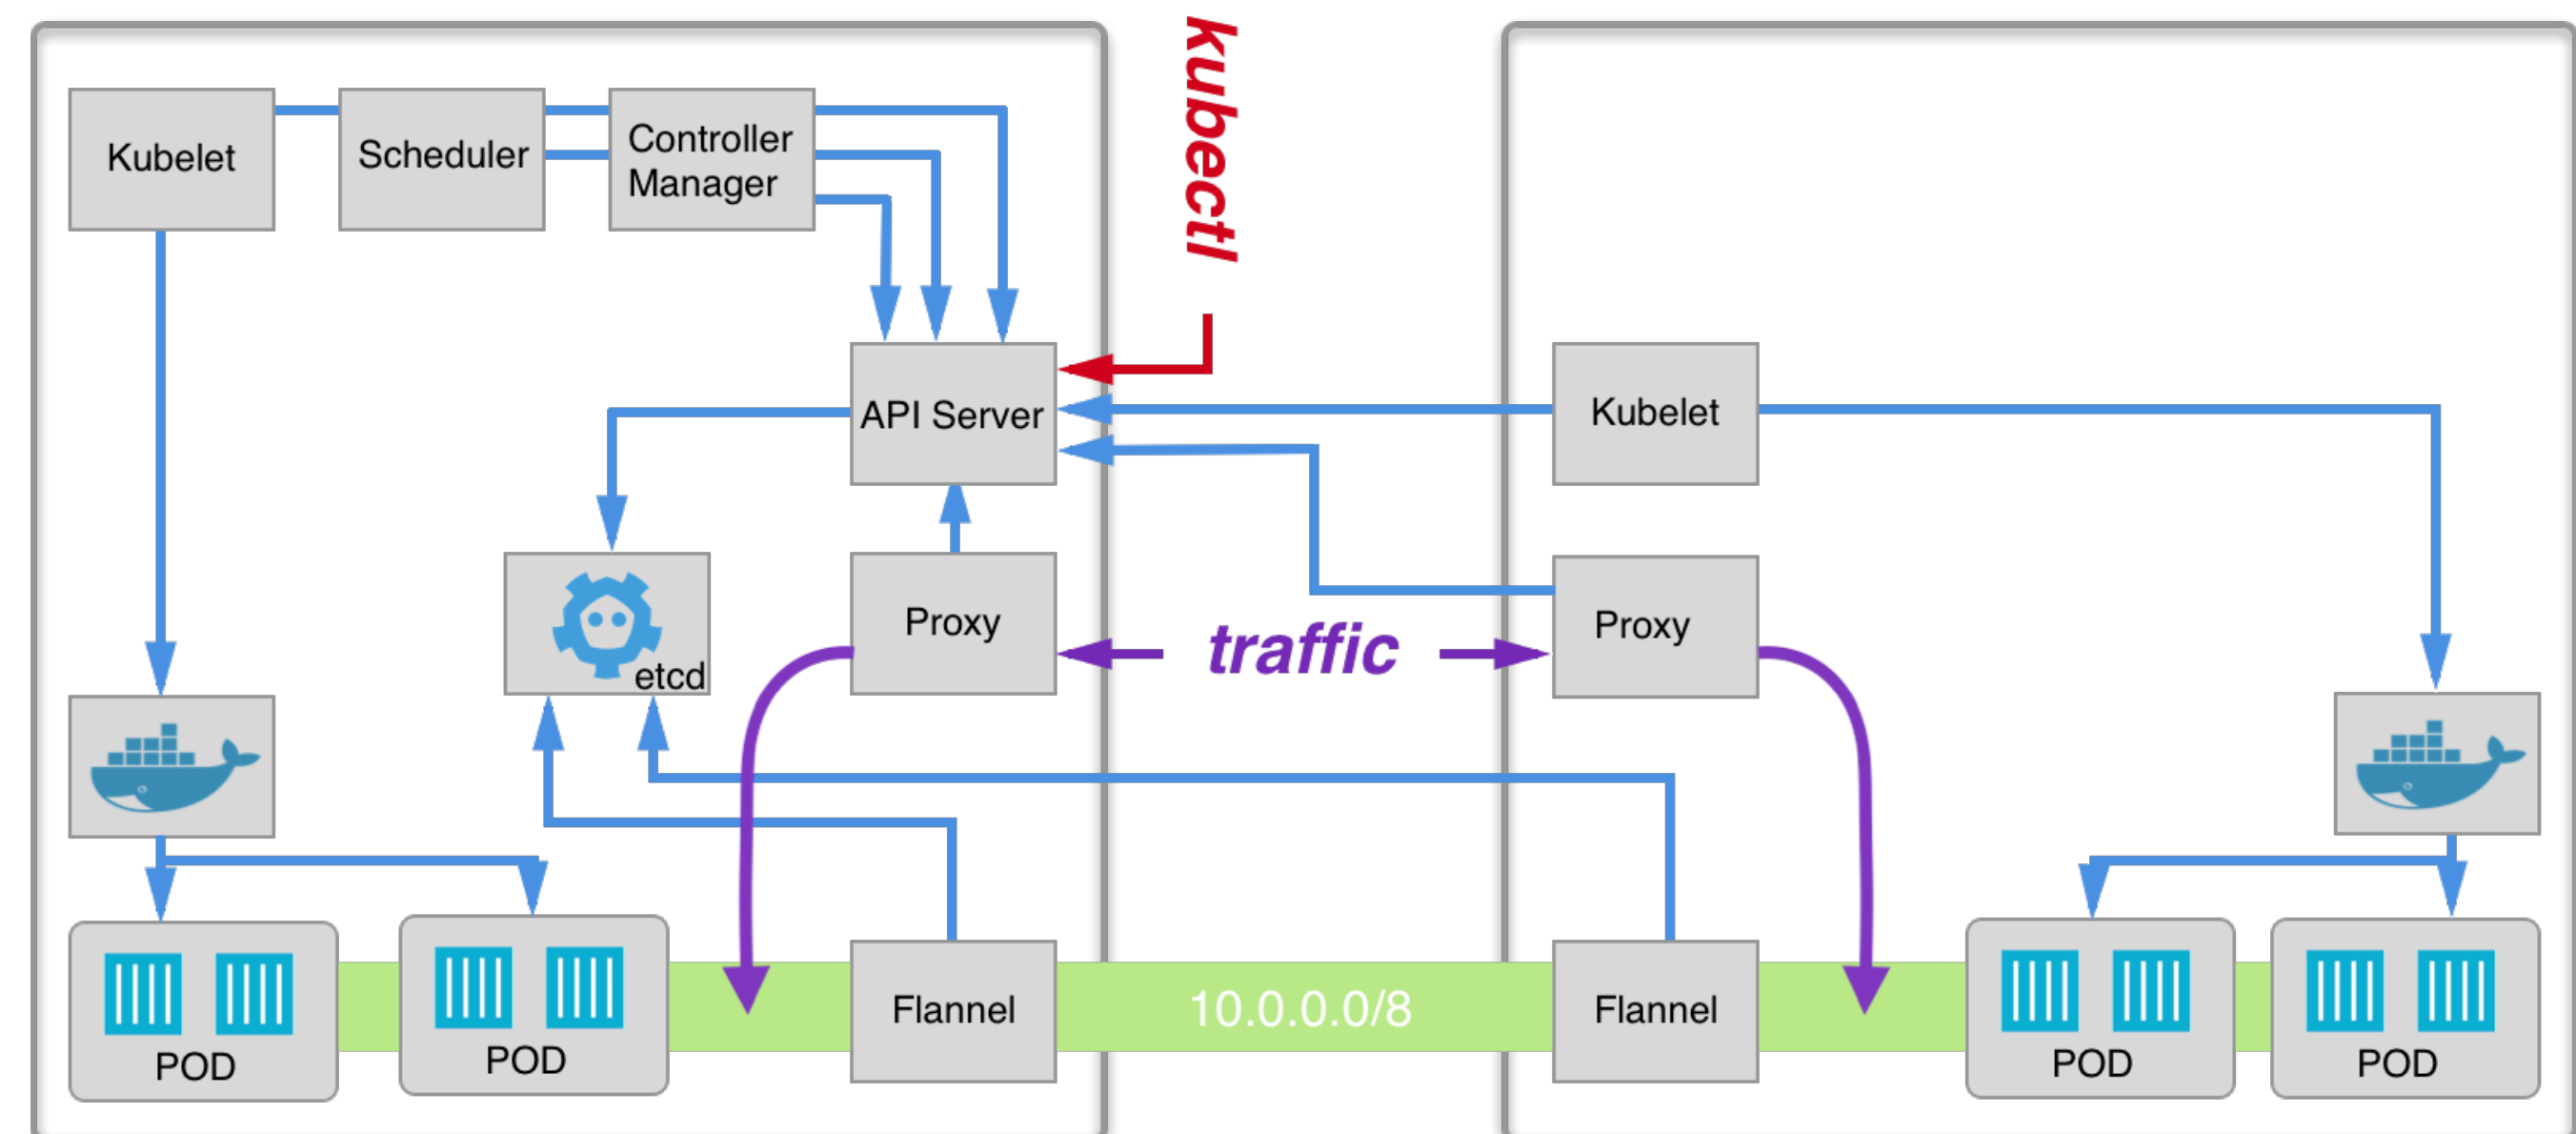
Flannel:



# NETWORKING

## Flannel:

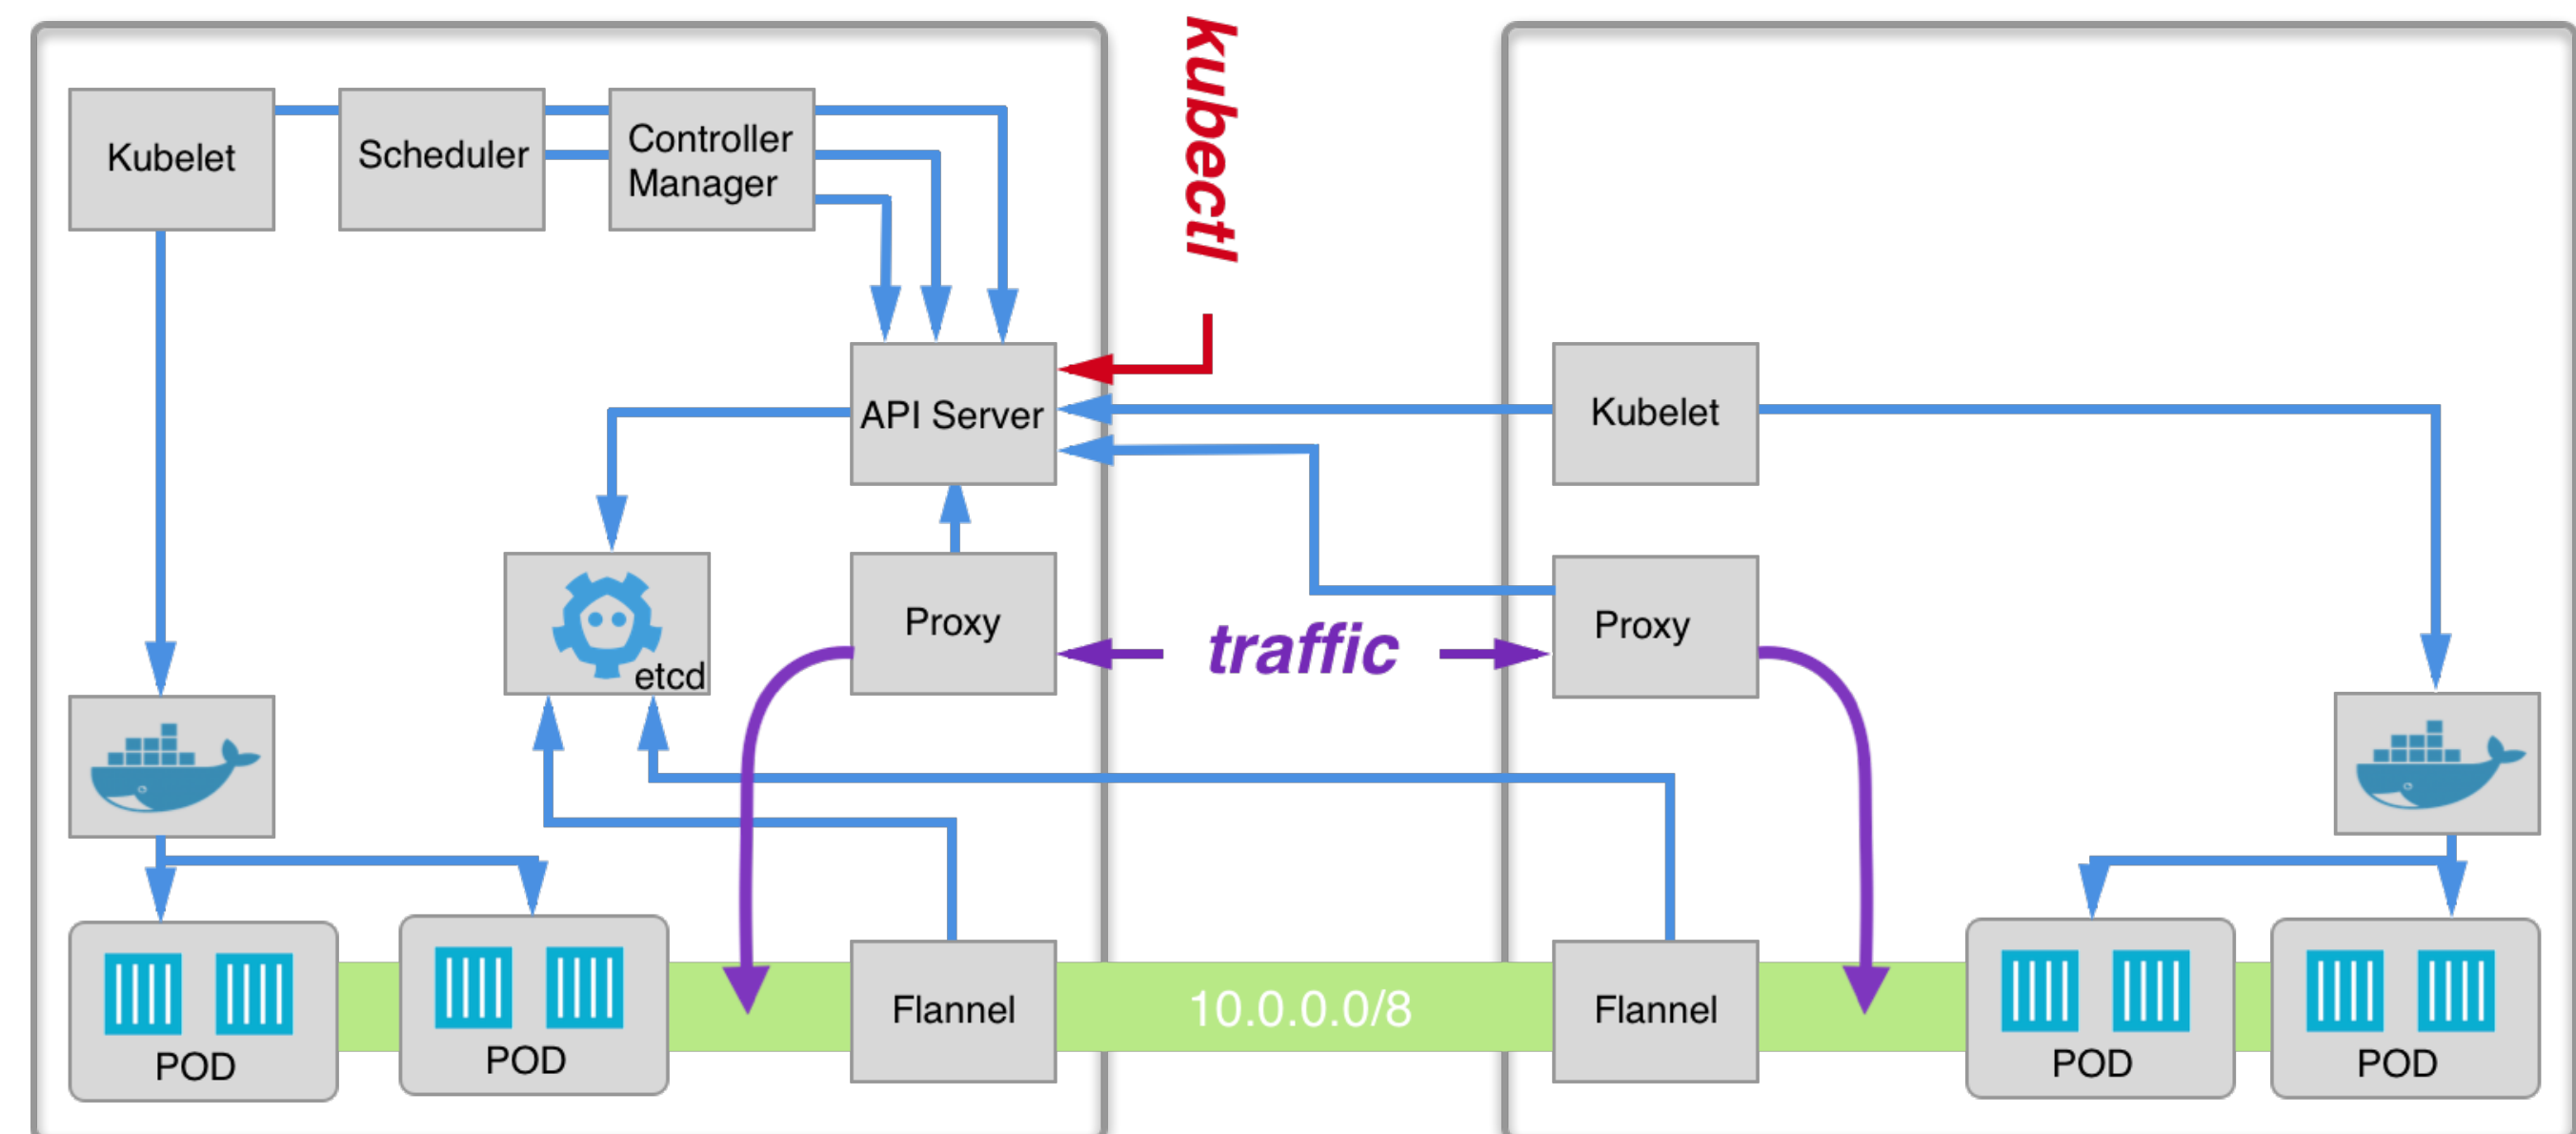
- uses etcd as a database



# NETWORKING

## Flannel:

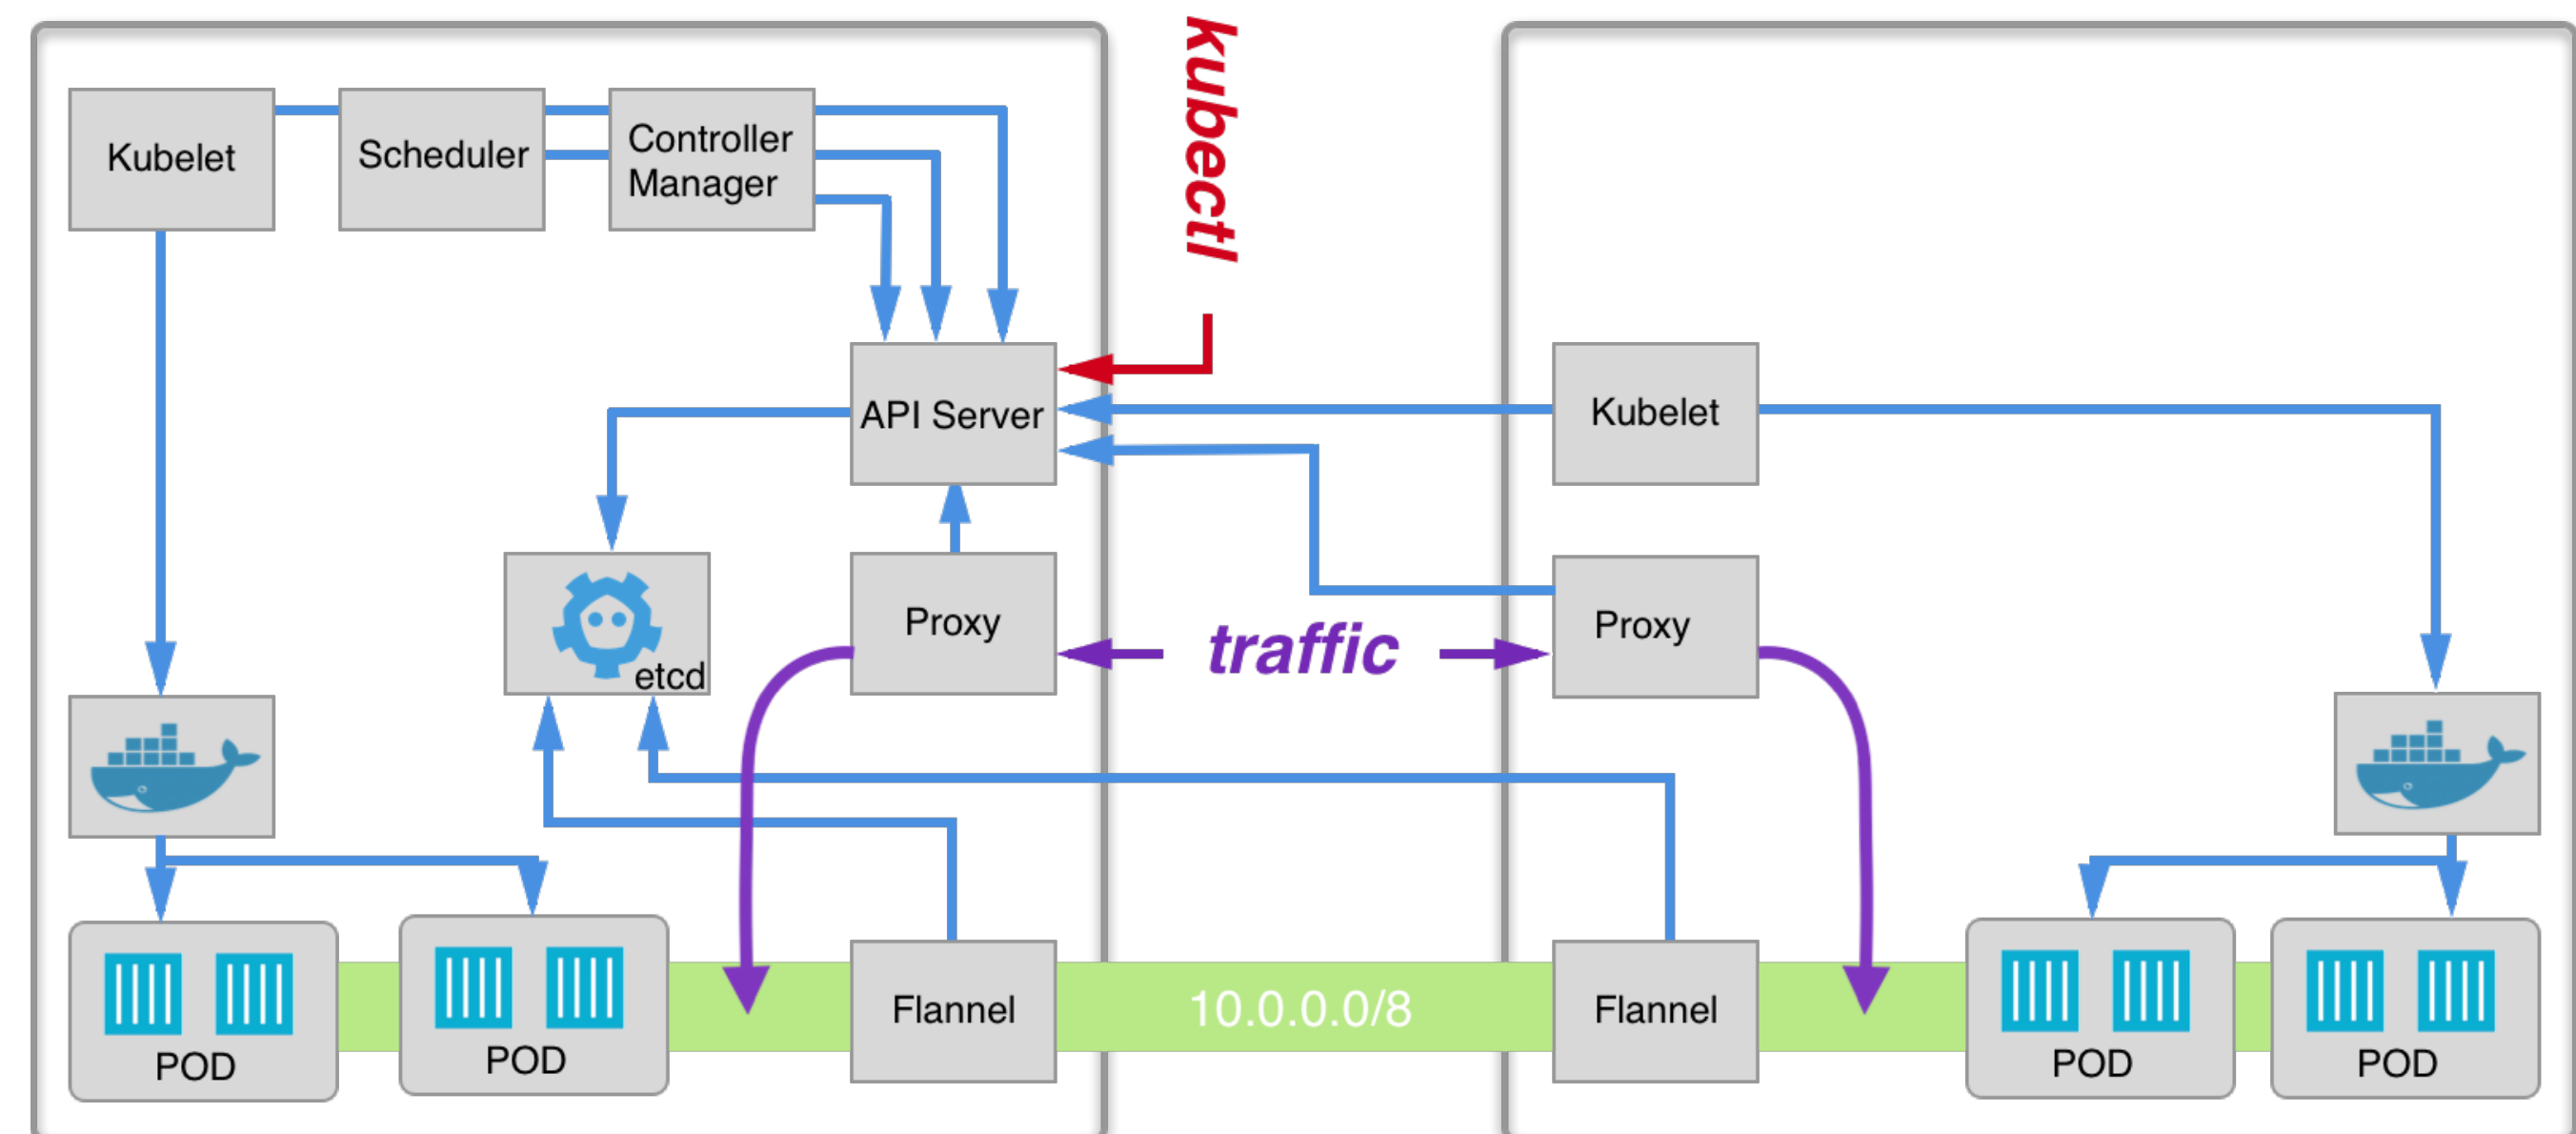
- uses etcd as a database
- agent allocates a subnet for a node out of larger address space



# NETWORKING

## Flannel:

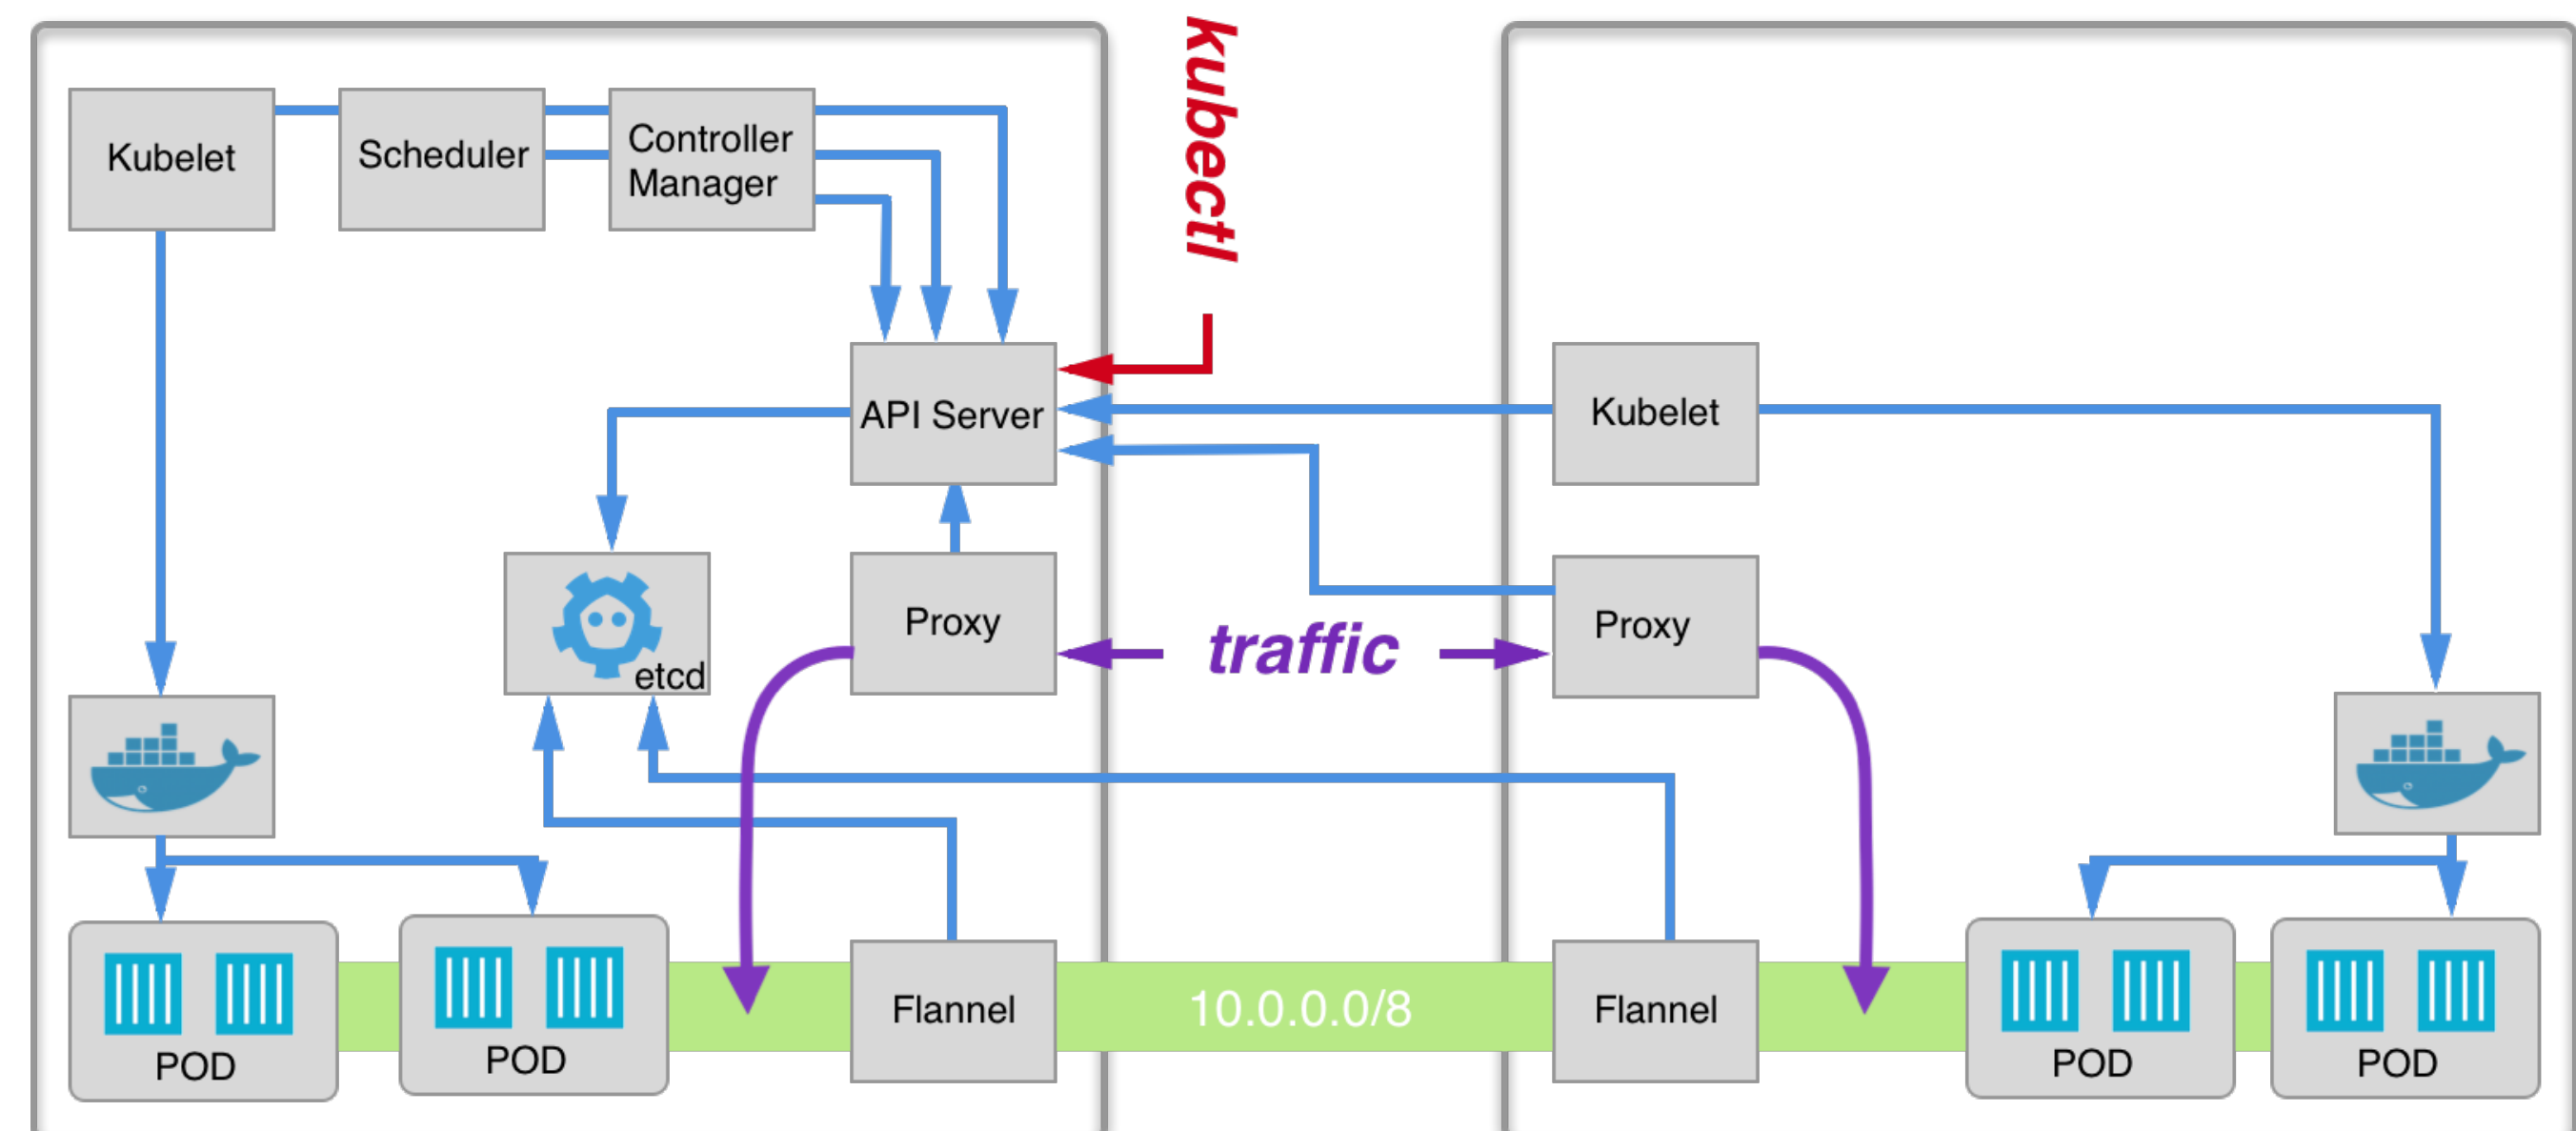
- uses etcd as a database
- agent allocates a subnet for a node out of larger address space
- tells docker to use that subnet



# NETWORKING

## Flannel:

- uses etcd as a database
- agent allocates a subnet for a node out of larger address space
- tells docker to use that subnet
- uses vxlan to encapsulate and forward packages between nodes



# DEMO



# Q&A

# THANK YOU